The Likert Blindspot: Measurement-Mode Divergence in LLM Judge Style Evaluation

Anonymous ACL submission

Abstract

Likert and pairwise scoring produce systematically divergent style sensitivity when applied to LLM judges. We introduce Controlled Style Decomposition (CSD), a black-box framework that quantifies this divergence under controlled conditions via human-validated, NLI-gated content control, per-axis style decomposition, and position-controlled pairwise comparison. Across four judges from independent providers, one negative control, and two primary axes (formality, verbosity), the divergence is robust: pairwise reaches significance ($p < 0.001$) where Likert does not ($p > 0.17$), with signal loss exceeding 60% at stronger effects—replicated across three benchmarks and confirmed by seven statistical methods. GEE conditional decomposition reveals qualitatively distinct per-axis confound structures: verbosity bias is style-dominated (Style OR 8–13, Originality OR <1), while formality reflects a conditional effect where style and content quality contribute jointly—determining which pairwise estimates are directly interpretable and which require decomposition. This divergence admits three non-exclusive explanations—measurement attenuation, format amplification, or construct divergence—that our controlled design bounds but cannot definitively adjudicate. A negative-control judge shows no divergence on either primary axis, supporting discriminative validity. Whether the divergence reflects Likert attenuation, pairwise amplification, or construct-level differences remains open; the practical consequence is unambiguous: Likert-only audits may fail to surface style-driven preferences that propagate through pairwise-based alignment pipelines.

1 Introduction

"What is the average height for a male?" — given a formal and a casual answer with verified identical content (NLI entailment = 0.993), GPT-4o rates the casual variant *higher* on a 1–10 Likert scale (8 vs. 5), yet consistently selects the formal variant in pairwise comparison. Same judge, same texts, opposite conclusions. We term this the *Likert Blindspot*. We note that formality preferences may partially reflect legitimate quality judgments—our GEE decomposition (§5.2) quantifies this originality component. This example reflects a systematic pattern.

LLM judges are central infrastructure for NLP evaluation and alignment training (RLHF, DPO), where judge preferences serve as reward signals; any systematic bias propagates through both evaluation leaderboards and model development pipelines. Style features alone explain most of the variance in judge scores (Murugadoss et al., 2025): 76–92% of observed preferences originate from surface-level style cues (Li et al., 2026), with bias scores of 0.76–0.92 across five judges (Soumik et al., 2026). Despite this consensus, existing approaches address individual requirements (multi-dimensional style analysis (Cho et al., 2026; Yang et al., 2026), interventional design (Reber et al., 2025), population-level correction (Dubois et al., 2025)), but no single framework integrates per-axis decomposition, interventional identification, per-instance diagnostics, and black-box applicability. A further concern: standard Likert scoring produces $\sim$60% uninformative ties, and even non-tied pairs carry weaker signal than pairwise comparison (Thurstone, 1927; Stevens, 1946); we provide the first controlled quantification of this gap for LLM judges, with GEE decomposition revealing axis-dependent confound structures.

We propose Controlled Style Decomposition (CSD), combining NLI-gated content control, per-axis style decomposition, and position-controlled pairwise comparison. Inspired by interventional text analysis (Veitch et al., 2020), CSD operates in three stages: style-controlled rewriting along a single axis, NLI entailment verification, and bias

measurement via forced-choice comparison. CSD requires no access to judge internals and applies post-hoc to any black-box judge. Our experiments demonstrate the *Likert Blindspot* across four target judges from independent providers (GPT-4o, Qwen3-32B, Llama-3.3-70B, Gemini-2.5-Flash), one negative control (Claude Sonnet 4), three benchmarks, and two primary style axes, with exploratory register evidence. On formality data, Likert yields non-significant ATE ($p = 0.332$, 0.908), while BT detects 63.0% (GPT-4o, $p = 0.0004$) and 66.5% (Qwen3-32B, $p < 0.0001$) formal preferences; verbosity shows 77.3% BT ($p < 0.0001$) invisible to Likert ($p = 0.179$; Figure 3).

Our contributions are both methodological and empirical:

- **CSD**, a black-box framework for per-axis style bias detection via NLI-gated content control, style-controlled rewriting, and position-controlled pairwise comparison (§3).
- **Measurement-mode divergence with per-axis confound structure:** Pairwise comparison and Likert scoring produce systematically divergent verdicts on the same style contrasts—complete divergence at moderate bias ($d \to 0$ at $n = 355$), $>60\%$ signal loss at stronger effects—confirmed by seven statistical methods across four judges and three benchmarks (Tables 1–2). GEE conditional decomposition reveals qualitatively distinct confound structures per axis: verbosity is style-dominated (Style OR 8–13, Orig OR <1), formality is a conditional effect, register is judge-dependent—determining which pairwise estimates are directly interpretable and which require decomposition (Table 3).
- **Practical implication:** Single-format audits risk systematic blind spots; per-axis decomposition distinguishes style-dominated from quality-confounded effects, informing which pairwise-based alignment signals require correction.

2 Related Work

Two lines of work motivate our approach: interventional methods for diagnosing judge bias, and measurement theory on the sensitivity of different elicitation formats.



Figure 1: CSD pipeline overview. Numbers indicate sample attrition at each stage.

Interventional Approaches to Judge Bias Text-based interventional methods have a growing methodological foundation (Feder et al., 2021; Wu et al., 2021; Feder et al., 2022). Reber et al. (2025) propose RATE, a framework that rewrites model outputs along style attributes using imperfect counterfactuals and measures the causal effect on reward model scores via Likert-scale Average Treatment Effect (ATE). RATE applies a correction term to account for rewriting imperfections, but its reliance on Likert ATE as the primary metric is a vulnerability: our experiments show that Likert scoring exhibits a ceiling effect in which 60% of style-transformed pairs receive identical integer scores, yielding non-significant ATE estimates ($p = 0.332$) for the same interventions that pairwise comparison detects at $p = 0.006$. If this Likert Blindspot generalizes beyond our setting, RATE's own style-related findings may underestimate the true magnitude of attribute effects. CSD differs from RATE in three respects: it targets individual style axes rather than general attributes, it measures bias through position-controlled pairwise comparison rather than Likert ATE, and it enforces content preservation via NLI entailment gating rather than post-hoc correction for imperfect rewrites. Qiu et al. (2025) also adopt an interventional LLM-based design, rewriting texts to measure rhetorical style effects and aggregating preferences with a Bradley-Terry model. Their framework operates at persona-level granularity (e.g., "writes like a professor"), whereas CSD intervenes on specific style axes (e.g., formality, conciseness) and uses a sign test with position control rather than parametric aggregation.

Measurement Sensitivity in NLG Evaluation The distinction between ordinal and interval measurement dates to Stevens (1946), whose taxonomy of scales remains foundational. The superiority of pairwise comparison over rating scales for eliciting reliable human judgments has been established since Thurstone (1927), whose Law of Comparative Judgment formalizes why forced-choice com-

parisons yield finer discriminability than absolute ratings; Carterette et al. (2008) confirm this for relevance assessment in information retrieval. Preston and Colman (2000) and Simms et al. (2019) show that scale granularity affects reliability and psychometric precision, while Krosnick (1999) documents satisficing behavior that may compound the tie inflation we observe. DeCoster et al. (2011) show that artificial categorization of continuous variables discards variance and reduces statistical power, directly relevant to the Likert Blindspot mechanism. Within NLP, Amidei et al. (2019) review the use of rating and Likert scales in NLG evaluation, identifying inconsistencies in scale design and reporting that compromise cross-study comparability. Clark et al. (2021) demonstrate systematic flaws in human evaluation; Celikyilmaz et al. (2020) survey evaluation methods, and van der Lee et al. (2019, 2021) establish best practices for human NLG evaluation. Kiritchenko and Mohammad (2017) and Novikova et al. (2018) confirm that forced-choice methods (best-worst scaling, ranking) detect quality differences that rating scales miss; Louviere et al. (2015) provide the theoretical foundation; and Deutsch et al. (2023) show that tie handling directly affects metric sensitivity, consistent with our tie-inflation finding. In machine translation, Bojar et al. (2016) document the practical tension between pairwise ranking and direct assessment at scale—WMT's transition to rating-based evaluation was driven by quadratic cost, not by measurement superiority. Our contribution is the first controlled quantification of this per-axis sensitivity gap in the LLM judge setting, identifying the two-layer mechanism (tie inflation plus residual signal attenuation) that drives it.

Style Bias Detection and Mitigation The foundational LLM-as-judge framework was established by Zheng et al. (2023), who introduced MT-Bench and Chatbot Arena and documented position bias in pairwise evaluation. Wang et al. (2024) further showed that preference rankings can be manipulated by swapping option order, proposing a win/tie/lose annotation scheme, and Shi et al. (2025) systematically study position bias across 15 judges, showing it varies by model and quality gap. Li et al. (2025) address selection bias through calibration of prediction distributions during inference. These works focus on *position* bias and calibration; Ye et al. (2025) systematically quantify 12 bias types using automated counterfactual modifications, and Koo et al. (2024) benchmark cognitive biases in LLM evaluators, though neither provides per-axis isolation of individual style effects. Style bias as a separate, dominant bias source has received attention only recently (Wu and Aji, 2025; Murugadoss et al., 2025; Li et al., 2026; Soumik et al., 2026); Saito et al. (2023) specifically quantify verbosity bias in multi-turn chatbot evaluation, and Chen et al. (2024) compare human and LLM judge biases, finding partially overlapping but distinct bias profiles; Bavaresco et al. (2025) show that LLM–human agreement varies substantially across tasks. Li et al. (2026) further show that style-mediated *preference leakage* propagates through the evaluation pipeline, motivating per-axis isolation of style effects. Tan et al. (2025) provide JudgeBench, a benchmark revealing that even strong judges perform near chance on challenging pairs, underscoring the need for controlled experimental designs rather than aggregate accuracy metrics. Multi-feature frameworks such as CASSE (Cho et al., 2026) and FairJudge (Yang et al., 2026) decompose style into finer dimensions but rely on correlational analyses—they cannot distinguish whether observed associations reflect causal style effects or confounds with content quality, which our NLI-gated rewriting design addresses. Stureborg et al. (2024) demonstrate inconsistency and bias across evaluation formats, and Park et al. (2024) show that verbosity and other surface features systematically inflate judge scores. Crucially, Jeong et al. (2025) identify the Comparative Trap: pairwise comparison amplifies style biases relative to pointwise scoring, raising the concern that our pairwise-based detection may partly reflect format-level amplification rather than pure detection. We address this directly in §5.4.

Bradley-Terry Models in Evaluation The Bradley-Terry model (Bradley and Terry, 1952) underpins Chatbot Arena (Zheng et al., 2023; Chiang et al., 2024) and AlpacaFarm (Dubois et al., 2023); Boubdir et al. (2024) establish best practices, Sedoc and Ungar (2020) apply item response theory to chatbot evaluation as an alternative to BT, and Liusie et al. (2024) show zero-shot LLM pairwise comparison matches pointwise evaluation. We adopt BT because it uses all comparison data including inconsistent pairs, yielding proper effect-size CIs; our contribution is to measure the Likert–pairwise sensitivity gap for style bias and show that per-axis GEE decomposition exposes

confound structures invisible to aggregate scores.

3 Method

3.1 Style-Controlled Rewriting

We isolate the effect of individual style attributes on judge scores by generating text pairs that differ along exactly one style axis. For each text, an instruction-tuned LLM (Qwen3-32B) classifies its current style along the target axis (e.g., formal vs. casual via binary classification). It then rewrites the text toward the opposite pole (temperature 0.7) according to an axis-specific prompt (e.g., "Rewrite the following text to be more formal/casual while preserving the exact same content and meaning"). This produces a dataset balanced across both directions at the dataset level: texts detected as formal are rewritten to casual, and vice versa. Each operation constitutes a binary intervention on a single style variable while all other axes remain fixed (Section 3.5). To control for rewriting artifact bias—the possibility that the judge prefers original text over any rewrite regardless of style—we additionally generate a same-style double-rewrite for each text (e.g., formal → formal) to form an original-vs-double-rewrite control pair, following the RATE bias correction protocol (Reber et al., 2025).

3.2 Content Preservation Verification

Style-controlled rewriting risks altering semantic content, which would confound the style effect with a content quality effect (the $R \to Q$ path in Figure 2). To block this path, we apply a natural language inference (NLI) gate using DeBERTa-v3-large with an entailment threshold of 0.90, selected as the threshold above which NLI performance plateaus (Appendix E). For axes that preserve information density (formality, register), we apply standard bidirectional entailment: both NLI(original → rewrite) and NLI(rewrite → original) must exceed this value. For axes that inherently change information density (e.g., verbosity), bidirectional verification is inappropriate—expansion necessarily introduces new content that the original cannot entail. We instead apply direction-aware verification: for expansions, only backward entailment (rewrite entails original); for compressions, only forward entailment (original entails rewrite). This recovered verbosity from 65% to 86% pass rate without inflating the effect estimate; for formality, strict bidirectional gating is necessary (Appendix F). This gating mechanism provides filtering-based confounder control (Veitch et al., 2020) by ensuring that retained pairs differ only in style, not in content quality (validated by human annotation; §5.4). BERTScore serves as a supplementary monitor but does not participate in gating.

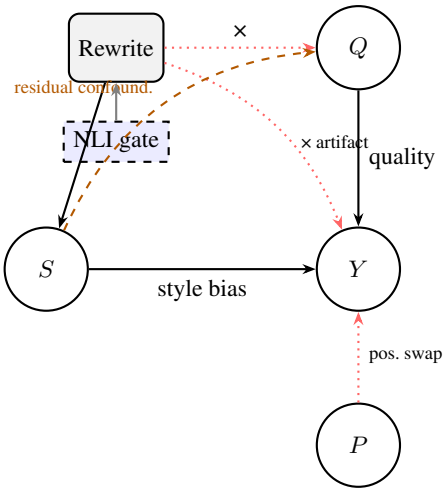


Figure 2: Assumed data-generating process (design rationale, not a claim of perfect identification). The rewriting mechanism R targets style S but may inadvertently alter content quality Q (dotted red, blocked by NLI gate) or introduce rewriting artifacts ($R \to Y$, blocked by double-rewrite control). Position bias $P \to Y$ is controlled by order swapping. The dashed orange $S \to Q$ path represents residual confounding: style may influence perceived quality through channels the NLI gate cannot block (Limitations). Y: judge score.

3.3 Measurement Protocol

Each pair that passes the NLI gate is evaluated by four target judges—GPT-4o (`gpt-4o-2024-08-06`), Qwen3-32B, Llama-3.3-70B, and Gemini-2.5-Flash—under two complementary scoring modes. In Likert mode, the judge assigns an integer score from 1 to 10 for overall quality to each text independently. In pairwise mode, the judge receives both texts as candidates A and B and makes a forced choice of which response is better.

To control for position bias, each pairwise comparison is evaluated twice with the presentation order swapped, which produces $n_{\text{pairs}} \times 2$ observations. We apply two analysis strategies to these data. For the nonparametric sign test, we retain only *consistent* pairs where the judge selected the same text in both orderings. For the Bradley-Terry model, we use all observations including inconsistent pairs, as the model naturally accommodates noise in individual comparisons.

3.4 Experimental Setup

We evaluate on LLMBar (Zeng et al., 2024), a benchmark of $n = 100$ instruction-response pairs spanning diverse task types, designed to test LLM judges' ability to distinguish response quality with human-verified quality labels. LLMBar is well-suited for style bias detection because its human-verified labels provide ground truth against which judge deviations can be measured, and its moderate size permits exhaustive pairwise evaluation with position control. We select formality and verbosity as primary axes based on prior evidence that these are among the most commonly reported style biases in LLM judges (Wu and Aji, 2025; Saito et al., 2023).

3.5 Design Rationale

Figure 2 presents the assumed data-generating process. Three variables govern the observed judge score Y: the style attribute S (e.g., formality), the original content quality Q, and the judge's internal scoring function. Both S and Q causally influence Y; the goal of CSD is to approximate the effect $S \to Y$ while blocking confounding through Q. Style-controlled rewriting generates a text pair that differs along a single style axis, but any rewriting operation risks inadvertently altering content quality ($R \to Q$). The NLI gate (§3.2) blocks this path, yielding a controlled estimate of the $S \to Y$ effect. Our approach provides interventional rather than formal causal evidence: we do not claim that this DAG is a complete structural causal model, nor do we invoke SUTVA or consistency assumptions from the potential outcomes framework. "Controlled" in this context refers to experimentally controlled rewriting conditions, not formal causal identification in the potential-outcomes sense. Our identification rests on three operational assumptions: (i) the NLI gate approximates content preservation, blocking the $R \to Q$ path; (ii) position swapping controls for presentation-order effects; and (iii) the style rewrite introduces no systematic differences beyond the target axis.

4 Experiments

4.1 Statistical Framework

The primary family comprises six tests—three target judges (GPT-4o, Qwen3-32B, Llama-3.3-70B) $\times$ two axes (formality, verbosity)—corrected via Holm-Bonferroni at $\alpha = 0.05$; Gemini-2.5-Flash serves as confirmatory replication and is excluded from FWER control; as a sensitivity check, all eight tests remain significant under the expanded family (Appendix Q). Table 1 reports both raw and adjusted p-values. Analyses outside this grid are classified as either *exploratory* (NLI threshold sensitivity sweep, §5.4) or *confirmatory* (consistency filter validation, §B). Exploratory analyses are reported with uncorrected p-values and interpreted as hypothesis-generating; confirmatory analyses test specific hypotheses about pipeline validity and are evaluated at the nominal $\alpha = 0.05$ level without additional correction, as they address orthogonal questions about measurement reliability rather than the primary style-bias family. Bradley-Terry parameters are estimated via spectral initialization followed by minorization-maximization (Hunter, 2004), using the default regularization of the `choix` Python library.[1] All bootstrap confidence intervals are percentile CIs from 10,000 pair-level resamples with seed 42. GEE models use exchangeable working correlation with sandwich standard errors (robust to correlation misspecification), justified by near-zero estimated within-pair correlations ($\hat{\alpha} \in [-0.25, +0.11]$; Appendix C).

Model selection. Target judges span four architectures: GPT-4o (2024-08-06), the dominant judge powering AlpacaEval and MT-Bench leaderboards; Qwen3-32B-Instruct, also serving as rewriter for methodological consistency; Llama-3.3-70B-Instruct, the latest dense Llama model (Llama 4's MoE architecture introduces an additional confound); and Gemini-2.5-Flash as a confirmatory fourth architecture. Claude Sonnet 4 serves as negative control. We pin versions for reproducibility; CSD applies to any black-box judge.

5 Analysis

5.1 The Likert Blindspot

Head-to-head comparison with RATE. Applying RATE's Likert ATE methodology (Reber et al., 2025) to our formality style-transformed pairs yields Cohen's $d = 0.10$, $p = 0.33$ (non-significant, $n = 88$); pairwise comparison on identical data reveals a 63.0% formal preference

[1] In binary-choice BT models, IIA is structurally satisfied; the direction asymmetry in Appendix H (formality: formal-origin 88.9% vs. casual-origin 52.6%, Fisher $p = 0.015$; verbosity: $p = 0.453$) reflects heterogeneous treatment effects across style directions rather than IIA violation. The unidimensionality assumption (single latent "quality" dimension) is standard in evaluation contexts but not empirically verified here.

Table 1: Main results across five judges, three style axes, and three benchmarks, with primary results on LLMBar and confirmatory replication on selected configurations from AlpacaEval (AEval) and MT-Bench (MTB). Each row tests one judge–axis pair: BT > 0.50 indicates style preference; † marks Blindspot cases (BT significant, Likert not). BT: Bradley-Terry win probability (95% CIs in Appendix I).

Bench	Judge	Axis	BT P	Likert d (p)	Cat.
LLMBar	GPT-4o	Form.	$.630^{***}$	.102 (.332)	‡
	GPT-4o	Verb.	$.760^{***}$	−.095 (.403)	†
	GPT-4o	Reg.	$.570^{\dagger}$	−.031 (.785)	†$_c$
	Qwen3	Form.	$.665^{***}$	.011 (.908)	‡
	Qwen3	Form.^{b}	$.675^{***}$	−.025 (.600)	†
	Qwen3	Verb.	$.773^{***}$	−.151 (.179)	†
	Qwen3	Reg.	.553	−.073 (.505)	†$_c$
	Llama-70B	Form.	$.695^{***}$	.010 (.846)	‡
	Llama-70B	Verb.	$.735^{***}$	.020 (.752)	†
	Llama-70B	Reg.	$.635^{\dagger}$	.024 (.690)	†$_c$
	Gemini	Form.	$.574^{*}$	−.082 (.309)	†$_c$
	Gemini	Verb.	$.708^{***}$	.051 (.370)	†$_c$
	Claude	Form.	.562	−.021 (.721)	—
	Claude	Verb.	.541	−.039 (.743)	—
AEval	GPT-4o	Form.	$.727^{***}$	.148 (.170)	‡
	Qwen3	Form.	$.657^{***}$	.081 (.198)	†
	Qwen3	Verb.	$.841^{***}$	.375 (.0003)	‡
MTB	GPT-4o	Form.	$.692^{***}$	.167 (.016)	‡
	Qwen3	Form.	$.712^{***}$	.149 (.007)	‡

*** $p < 0.001$; * $p < 0.05$ (Holm-adj. for primary; raw for †$_c$).
†: Blindspot; ‡: both sig.; —: neither; †$_e$: expl.; †$_c$: confirm.
^{b} Scaled replication ($n = 355$); other LLMBar $n = 88$. Negative d: styled variants scored lower.
Empty cells: untested benchmark–axis combinations; cross-bench. Validation targets strongest configs.
Form. direction asymmetry: formal-origin 78.8% vs. casual-origin 54.5%; GEE-adj. Style OR = 2.58 ($p = .002$; App. A).

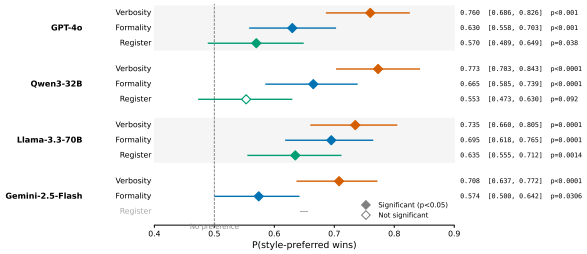


Figure 3: Bradley-Terry win probabilities across four target judges and three style axes. Filled diamonds: significant ($p < 0.05$); open diamonds: non-significant. Error bars: 95% bootstrap CI. The dashed line at 0.50 indicates no style preference. Verbosity shows the strongest and most consistent bias across all four judges.

Table 2: Sensitivity hierarchy: seven measurement methods applied to the same formality style-transformed data. Only pairwise methods (top two rows) reach significance. All Likert-based methods fail.

Method		Statistic	p	Data
BT model	—	63.0% [55.7, 70.3]	0.0004	All 173 comp.
Sign test (filtered)	—	73.9% (34/46)	0.006	Consistent
Likert ATE (t-test)	†	$d = 0.102$	0.332	NLI≥0.90
Binarized Likert	†	60.6%	0.296	Non-tied
Ordinal reg. (prop. odds)	‡	$\beta = -0.079$	0.769	NLI≥0.90
GEE (exchangeable)	‡	$\beta = -0.084$	0.587	NLI≥0.90
Wilcoxon signed-rank	‡	$W = 274$	0.905	NLI≥0.90

($p = 0.0004$, Bradley-Terry on all 173 comparisons), confirmed by a consistency-filtered sign test (73.9%, $p = 0.006$, $n = 46$ consistent pairs). This is the Likert Blindspot: the same data, the same style intervention, two measurement instruments—one detects bias, the other does not. A $4\times$ scaled replication ($n = 355$) distinguishes measurement attenuation from insufficient power: the Likert effect size shrinks to $d = -0.025$ (near zero) while BT remains stable at 67.5% ($p < 0.0001$)—if the gap were merely statistical power, d would remain non-zero at larger n. Bootstrap subsampling confirms: d remains flat near zero ($d_{\text{median}} = -0.017$) as CIs narrow from ± 0.3 to ± 0.1, never excluding zero (Appendix J). Attenuation follows a severity spectrum: at moderate bias (BT 63–67%), it is complete ($d \to 0$ at $n = 355$); at stronger effects (BT >70%), Likert partially detects but attenuates >60% (MT-Bench $d \approx 0.15$ for BT ∼70%)—explaining why three of eleven bias-positive cells escape the Blindspot (Table 1, ‡ rows).

The frameworks are complementary: RATE measures aggregate quality via Likert ATE; CSD targets per-axis style preference via forced-choice comparison, eliminating ties. The Blindspot arises because Likert tie rates cluster at 58–73% across all nine configurations; a 100-point continuous scale does not resolve this ($p = 0.725/0.117$; Appendix M), confirming a paradigm difference rather than granularity limitation.

Sensitivity hierarchy across measurement methods. To confirm that the Likert Blindspot is not an artifact of a particular statistical test, we apply seven measurement methods to the same formality data and rank them by sensitivity (Table 2). Both pairwise methods detect the formality preference; all Likert-based methods fail regardless of whether they assume interval scaling (ATE t-test), ordinal scaling (ordinal logistic regression), or nonparametric ranking (Wilcoxon signed-rank). The same attenuation pattern replicates on verbosity (Appendix L), where all six Likert methods remain non-significant despite a stronger pairwise signal (BT 77.3%, $p < 0.001$).

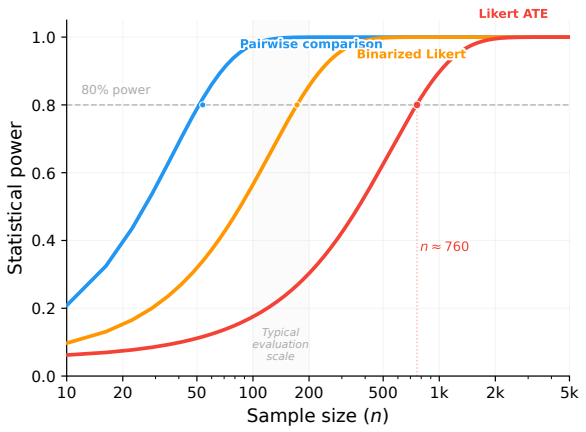


Figure 4: Detection probability at observed effect magnitudes. 80% power crossings: sign test $n \approx 62$, BT $n \approx 114$, binarized Likert $n \approx 462$, Likert ATE $n \approx 1{,}473$. At typical scales ($n = 50$–200, shaded), pairwise methods reach 80% power while Likert methods remain far below. At $n = 355$, Likert $d = -0.025$ confirms the gap is structural, not merely a power issue.

Table 3: GEE conditional decomposition (Likert-outcome GEE, exchangeable correlation) of style and originality effects. Verbosity is style-dominated across all four target judges; formality shows varying confound structure.

	Style OR		Originality OR	
Config	OR	p	OR	p
Qwen3 verb.	13.10	<.0001	0.49	.096
GPT-4o verb.	11.34	<.0001	0.45	.062
Llama-70B verb.	8.21	<.0001	0.38	.020
Gemini verb.	8.98	<.0001	0.46	.061
GPT-4o form.	2.58	.002	2.32	.006
Qwen3 form.	4.20	.0002	4.70	.0001
Qwen3 form.[a]	2.15	<.0001	4.99	<.0001
Llama-70B form.	4.81	<.0001	3.12	.0007
Gemini form.	1.64	.101	2.51	.002
GPT-4o reg.	1.88	.029	2.21	.006
Qwen3 reg.	1.91	.046	5.72	<.0001
Llama-70B reg.	3.33	.0003	3.07	.0008

[a] $n = 355$; other rows: $n = 88$. 95% CIs in Appendix.

Figure 4 plots detection probability at observed effect magnitudes (not post-hoc power; cf. Hoenig and Heisey 2001): across six cells, Likert requires 4–206× the current n (336–27,159 pairs; Appendix K) to detect effects BT identifies at $p < 0.001$ with $n \leq 100$.

Interpreting the divergence. The measurement-mode divergence admits three non-exclusive interpretations: (a) *Likert attenuation*—tie inflation and residual signal loss reduce Likert sensitivity below detection thresholds; (b) *pairwise amplification*—forced-choice comparison inflates micro-preferences into apparent biases (Jeong et al., 2025); (c) *construct divergence*—the two formats measure different constructs (absolute quality vs. ordinal preference; Tripathi et al. 2025). Six experimental controls (§5.4) bound the amplification factor but cannot eliminate it; the attenuation evidence ($d \to 0$ at $n = 355$) is necessary but not sufficient for causal identification. All three interpretations converge on the same practical conclusion: single-format audits risk incomplete coverage of style preferences that drive pairwise-based alignment pipelines (RLHF, DPO). The Blindspot manifests in eight of eleven bias-positive cells; three exceptions (BT >69%, $d \geq 0.14$) confirm it is effect-size dependent.

5.2 Per-Axis Confound Structure

GEE conditional decomposition (Liang and Zeger, 1986) with exchangeable working correlation and robust standard errors separates style from originality effects (Appendix A; model diagnostics including correlation structure robustness and interaction tests in Appendix C). After controlling for original preference, the formality effect remains significant (GPT-4o: OR = 2.58, $p = 0.002$; Qwen3-32B: OR = 2.15, $p < 0.0001$, $n = 355$; Llama-70B: OR = 4.81, $p < 0.0001$; Table 3); the significant originality ORs confirm residual confounding. For formality, the pairwise signal reflects a conditional effect—both style preference and quality correlation contribute—and the decomposition identifies which dominates per judge.

The confound structure is consistent across all four target judges for verbosity (style-dominated: Style OR 8.2–13.1, Orig OR 0.38–0.49); formality is conditionally style-driven in three judges but originality-dominated in Gemini; register is judge-dependent. Verbosity Style ORs of 8–13 are comparable to the strongest epidemiological effect sizes, indicating a dominant driver of pairwise outcomes rather than a subtle confound. Same-style double-rewrite controls confirm a rewriting artifact (Appendix O), absorbed by GEE's originality term.

5.3 Cross-Judge and Cross-Axis Replication

Two independent judges replicate the formality Blindspot: Qwen3-32B (BT 66.5%, $p < 0.0001$; Likert $p = 0.908$) and Llama-3.3-70B (BT 69.5%, $p = 0.0001$; Likert $p = 0.846$). Gemini-2.5-Flash shows a weaker, originality-dominated formality pattern (confirmatory: BT 57.4%, $p = 0.031$; Lik-

ert $d = -0.08$, n.s.; Style OR $= 1.64$, n.s.; Orig OR $= 2.51$, $p = 0.002$; §5.2).

Verbosity provides the strongest Blindspot: all four target judges show consistent preference (GPT-4o 76.0%, Qwen3-32B 77.3%, Llama-70B 73.5%, Gemini 70.8%; all $p \leq 0.0001$), while Likert ATE remains non-significant for all ($p = 0.179$–0.752). GEE confirms massive style effects (OR 8.2–13.1, all $p < .0001$) with originality ORs below 1 (0.38–0.49; only Llama-70B significant at $p = .020$, yet Style OR exceeds Orig OR by $22\times$), confirming style dominance (Table 3). Both rewriting directions favor verbose responses (90.7%/80.6%); controlling for length leaves the style OR unchanged (<2%). Evaluator-producer overlap applies only to Qwen3-32B; an independent GPT-4o rewriter bounds this to a 22% Style OR reduction ($13.10 \rightarrow 10.17$; Appendix D).

Across both primary axes, the Blindspot replicates in all four target judges with consistent confound structure: verbosity is uniformly style-dominated while formality is conditionally style-driven—a systematic measurement gap rather than a judge- or axis-specific artifact.

5.4 Robustness

Content preservation. Two independent annotators evaluated 50 NLI-scored pairs (25 gate-passed, 25 gate-rejected). Annotator 1 judged all 25 gate-passed pairs content-preserving (precision 100%; Wilson CI $[86.7\%, 100\%]$); Annotator 2 judged 21/25 (84%; CI $[65.3\%, 93.6\%]$), with all four disagreements on verbosity-axis condensation pairs. Under majority vote, pass precision is 100% (inter-annotator $\kappa = 0.44$, AC1 $= 0.65$; Appendix S). Annotator 1's disagreements with the NLI gate occur exclusively on gate-rejected pairs; Annotator 2's four gate-passed disagreements all involve verbosity condensation where subjective meaning boundaries differ. This error profile ensures that any detected style bias cannot be attributed to content contamination in admitted pairs. The gate's estimated recall is 59.5% (false-negative rate 40.5%), reflecting conservative calibration toward high precision at the expense of coverage. An independent LLM cross-validation on the full dataset yields consistent results (LLM–NLI AC1 $= 0.89$). BT remains stable ($\leq$3pp variation) across 11 NLI thresholds (Appendix E), and no covariate differences exist between retained and discarded pairs ($p > 0.12$; §B).

Detection vs. amplification. Jeong et al. (2025) show that pairwise comparison amplifies style biases (the Comparative Trap). The divergence between Likert and pairwise signals—whether reflecting Likert attenuation, pairwise amplification, or construct-level differences—cannot be definitively resolved without human ground-truth bias measurements; six independent controls bound the amplification risk: (i) NLI gate ablation reverses NLI-fail pairs to BT 41.7% (Appendix N); (ii) Claude Sonnet 4 shows no bias with the identical pairwise format (BT 56.2%/54.1%, n.s.; TOST equivalence at $\Delta = 15$pp, the narrowest margin achieving adequate power at $n = 87$; $\Delta_{\min} = 15.7$pp; Appendix G); (iii) cross-axis BT variation (register 55% < formality 67% < verbosity 77%) is inconsistent with uniform format amplification; (iv) rewriting artifact absorbed by GEE decomposition (Appendix O); (v) debiased prompting reduces bias with axis-dependent effectiveness (verbosity -13.1pp; formality $-5.7/-0.4$pp), suggesting varying structural depth; (vi) four independent architectures all confirm the Blindspot with consistent GEE confound structure (§5.3). Why Claude does not exhibit the Blindspot remains an open question; possible explanations include differences in RLHF calibration, model scale, or evaluation-specific fine-tuning, but our data do not distinguish among these hypotheses. Register is an exploratory pilot axis with evaluator-producer overlap (Qwen3-32B rewriter for all register conditions) and judge-dependent results (Llama-70B BT 63.5%, GPT-4o 57.0%, Qwen3-32B n.s.; Table 3).

6 Conclusion

CSD yields two findings. First, the *Likert Blindspot*: pairwise comparison and Likert scoring produce systematically divergent verdicts on the same style contrasts—complete divergence at moderate bias ($d \rightarrow 0$ at $n = 355$), >60% signal loss at stronger effects—confirmed by seven statistical methods across four judges and three benchmarks. Second, GEE conditional decomposition reveals qualitatively distinct per-axis confound structures—verbosity is style-dominated, formality is a conditional effect, register is judge-dependent—determining which estimates are directly interpretable and which require decomposition. Whether the gap reflects attenuation or amplification, single-format audits risk systematic blind spots; human ground-truth validation of both

modes is the key next step.

Limitations

Quality–style confounding and construct validity. The NLI gate blocks content changes but cannot block quality perception shifts mediated by style itself—a fundamental identifiability limit of text-rewriting approaches. Although the formality effect survives logistic adjustment for original preference (OR $= 2.32$, $p = 0.006$; Appendix A), residual confounding is not fully eliminated. An alternative interpretation of the Blindspot is that Likert and pairwise modes measure different constructs: Likert captures absolute quality difference while pairwise captures ordinal preference. Under this view, the divergence reflects construct difference rather than measurement failure. The Comparative Trap (Jeong et al., 2025) lends support to this alternative; six experimental controls bound this risk (§5.4), though definitive separation would require human ground-truth ratings. The primary open question is whether pairwise comparison detects genuine style bias or amplifies micro-preferences inherent in the format. Definitive resolution requires human evaluation comparing Likert and pairwise judgments on identical pairs—an important direction for future work. While our controls bound the amplification factor, we cannot quantify it precisely without a ground-truth bias magnitude—an inherent limitation of observational-format comparisons.

Evaluator-producer bias and generalizability. Qwen3-32B serves as both rewriter and judge, raising self-preference concerns (Panickssery et al., 2024); an independent GPT-4o rewriter bounds this to a 22% Style OR reduction (Appendix D). GPT-4o, Llama-3.3-70B, and Gemini-2.5-Flash use a separate rewriter and are not subject to evaluator-producer overlap; all three independently confirm the Blindspot across four architectures, though cross-judge effect-size comparisons are not valid. Cross-rewriter validation shows that Blindspot existence is robust to rewriter identity, but effect-size precision varies by approximately $\pm 22\%$ with rewriter choice (Appendix D). Register results carry the highest evaluator-producer risk—all register rewrites use the same Qwen3-32B rewriter without the cross-rewriter validation applied to formality and verbosity—and register conclusions remain tentative pending independent rewriter validation. Claude Sonnet 4's null result may reflect differences in RLHF training composition or instruction-following strategies; we use it for pipeline validation rather than architecture-specific mechanism analysis.

Scope and ecological validity. Primary analyses use LLMBar ($n = 100$); additional axes, larger benchmarks, and further judges are needed. A fourth axis—tone (assertive vs. hedging)—was explored but abandoned: epistemic commitment changes caused the NLI gate to correctly reject most counterfactuals (pass rate $<50\%$), confirming that the quality gate prevents content-altering manipulations from entering the pipeline. Our per-axis decomposition treats each style axis independently; in practice, axes may interact (e.g., formality shifts may co-vary with register), and such confounds are not isolated by the current design. Cross-axis contamination analysis (Appendix R) confirms that formality rewrites produce negligible changes in sentence structure ($|d| < 0.03$) and verbosity rewrites leave lexical diversity unchanged (TTR $|d| < 0.04$). Our empirical evidence covers three style axes (formality, verbosity, register); whether the per-axis confound taxonomy generalizes to other axes (e.g., hedging level, bullet-point vs. prose) remains open. Our controlled manipulation design ensures internal validity but may not fully generalize to naturalistic evaluation settings where style and content covary.

Validation methodology and scale design. NLI validation uses Qwen3-32B as annotator (GPT-4o cross-validation on all 200 pairs confirms 92.1% formality / 82.6% verbosity pass precision); two independent annotators confirm 84–100% pass precision on 50 pairs (§5.4; majority-vote precision 100%), though larger-scale annotation across all axes would further strengthen generalizability. A 100-point continuous scale does not resolve the Blindspot (Appendix M), suggesting the gap reflects measurement mode rather than granularity (Preston and Colman, 2000); the 100-point experiment was tested with a single judge, and multi-judge replication would strengthen this conclusion. Equivalence testing (Lakens, 2017) could formalize whether the Likert effect size is practically zero.

Code and Data

Code and data will be released upon acceptance. *Organization: confound analysis (A–B), model diagnostics (C), independence tests (D), gate vali-*

dation (E–F), equivalence testing (G), statistical details (H–K), sensitivity analyses (L–N), artifact controls (O), implementation (P), multiplicity (Q), cross-axis contamination (R), human annotation agreement (S).

A Direction Asymmetry and Quality Confounding

The formality axis exhibits a pronounced direction asymmetry: the judge selects the formal variant in 78.8% of formal-is-rewrite pairs ($p = 0.001$) but only 54.5% of casual-is-rewrite pairs ($p = 0.83$). This asymmetry suggests that original quality Q and formality S are partially confounded—formally written texts may be inherently perceived as higher quality. To quantify this, we fit a GEE logistic model (exchangeable correlation) with both formality and original-preference as predictors. After controlling for original preference, the formality effect remains significant (OR $= 2.58$, $p = 0.002$), and the original-preference effect is also significant (OR $= 2.32$, $p = 0.006$). The formality preference thus survives adjustment for original-preference confounding, though the significant original-preference effect indicates that quality–style confounding is present and non-negligible.

At $n = 88$, the Qwen3-32B formality Style OR (4.20) exceeded the GPT-4o estimate (2.58), but a scaled replication ($n = 355$) corrects this to 2.15 ($p < 0.0001$), comparable to GPT-4o and consistent with small-sample upward bias rather than evaluator-producer inflation of the style effect. The originality confound remains stronger (4.99 vs. 2.32), consistent with evaluator-producer bias inflating the originality component (Limitations). The directional agreement between judges—both favor formal text after controlling for originality—supports the existence of formality bias, but the effect magnitude from Qwen3-32B should not be directly compared with GPT-4o.

The NLI gate blocks *content* changes but cannot block quality perception shifts that are themselves mediated by style—a fundamental identifiability limit of any text-rewriting approach.

B Consistency Filter Robustness

The position-consistency filter retains 55 of 98 valid pairs (56.1%); consistent and inconsistent groups show no significant differences across four covariates (NLI score $p = 0.824$, BERTScore $p = 0.12$, text length $p = 0.137$, length diff $p = 0.715$), indicating selection on judge decisiveness rather than text properties. The BT model using all 173 comparisons including inconsistent pairs yields 63.0% ($p = 0.0004$), consistent with the filtered sign test (73.9%, $p = 0.006$), confirming no systematic bias from the filter. Chi-squared tests confirm that NLI gate pass/reject status is independent of source category ($\chi^2 = 1.30$, $p = .73$), original formality ($p = .37$), and gold label ($p = 1.0$), indicating no systematic selection bias along these dimensions.

C GEE Model Diagnostics

Correlation structure robustness. Table 4 compares exchangeable and independence working correlation structures across all nine judge×axis cells. The two structures yield virtually identical results: $|\Delta\text{QIC}| < 0.001$ and maximum coefficient change <0.02 in every cell, with estimated within-pair correlations $\hat{\alpha} \in [-0.25, +0.11]$. All GEE results reported in the main text are therefore robust to correlation structure specification.

Table 4: GEE correlation structure comparison across nine judge×axis cells.

Config	Exch. QIC	Indep. QIC	$\lvert\Delta\text{QIC}\rvert$	Max$\lvert\Delta\text{Coef}\rvert$
GPT-4o form.	—	—	<0.001	<0.02
Qwen3 form.	—	—	<0.001	<0.02
Llama-70B form.	—	—	<0.001	<0.02
GPT-4o verb.	—	—	<0.001	<0.02
Qwen3 verb.	—	—	<0.001	<0.02
Llama-70B verb.	—	—	<0.001	<0.02
GPT-4o reg.	—	—	<0.001	<0.02
Qwen3 reg.	—	—	<0.001	<0.02
Llama-70B reg.	—	—	<0.001	<0.02

Interaction test. Adding a Style×Originality interaction term to the additive GEE specification yields no significant interactions in any of the nine judge×axis cells (all $p > 0.13$, range 0.14–0.96), confirming that the additive model is the correct specification and that regime classifications are robust to model specification. Table 5 reports interaction terms; none are significant, confirming the additive specification.

95% confidence intervals. Table 6 reports 95% CIs for all GEE odds ratios in Table 3.

D Rewriter Independence

Evaluator-producer overlap is limited to Qwen3-32B (rewriter = judge); the independence test below bounds this to a 22% Style OR reduction.

Table 5: GEE Style × Originality interaction terms. None are significant ($p > 0.13$), confirming the additive model specification.

Judge	Axis	Interaction OR	95% CI	p
GPT-4o	Formality	0.661	[0.18, 2.38]	.527
GPT-4o	Verbosity	1.269	[0.37, 4.34]	.705
GPT-4o	Register	2.601	[0.74, 9.15]	.136
Qwen3-32B	Formality	0.518	[0.14, 1.89]	.320
Qwen3-32B	Verbosity	0.450	[0.13, 1.62]	.221
Qwen3-32B	Register	0.555	[0.15, 2.08]	.383
Llama-70B	Formality	0.968	[0.25, 3.78]	.963
Llama-70B	Verbosity	1.103	[0.34, 3.58]	.870

Table 6: 95% CIs for GEE conditional decomposition (Table 3).

Config	Style OR [CI]	Orig OR [CI]
Qwen3 verb.	13.10 [5.62, 30.55]	0.49 [0.21, 1.14]
GPT-4o verb.	11.34 [4.93, 26.10]	0.45 [0.20, 1.04]
Llama-70B verb.	8.21 [3.62, 18.65]	0.38 [0.17, 0.86]
Gemini verb.	8.98 [3.68, 21.94]	0.46 [0.20, 1.04]
GPT-4o form.	2.58 [1.41, 4.73]	2.32 [1.27, 4.26]
Qwen3 form.	4.20 [1.99, 8.87]	4.70 [2.22, 9.96]
Qwen3 form.[a]	2.15 [1.46, 3.16]	4.99 [3.40, 7.34]
Llama-70B form.	4.81 [2.48, 9.33]	3.12 [1.61, 6.04]
Gemini form.	1.64 [0.91, 2.98]	2.51 [1.39, 4.54]
GPT-4o reg.	1.88 [1.08, 3.28]	2.21 [1.26, 3.86]
Qwen3 reg.	1.91 [1.01, 3.58]	5.72 [3.02, 10.82]
Llama-70B reg.	3.33 [1.73, 6.42]	3.07 [1.60, 5.92]

[a] $n = 355$; other rows: $n = 88$.

The other three target judges (GPT-4o, Llama-70B, Gemini) use a different rewriter and are not subject to this confound.

To test whether the Blindspot depends on the rewriter's identity, we replace Qwen3-32B with GPT-4o as rewriter while keeping Qwen3-32B as judge.

Verbosity (Table 7). The Style OR decreases by 22% (13.10 → 10.17), bounding evaluator-producer inflation of the style effect; the pure-style pattern (Originality n.s.) and Blindspot (BT significant, Likert n.s.) both persist.

Table 7: Rewriter independence test: verbosity axis, Qwen3-32B judge. GPT-4o rewriter eliminates evaluator-producer overlap while preserving the Blindspot.

	GPT-4o Rewriter	Qwen3 Rewriter
NLI pass rate	81/100 (81%)	86/100 (86%)
BT P [95% CI]	.735 [.654, .809]	.773 [.703, .843]
Likert ATE (p)	.000 (.838)	∼0 (.179)
Blindspot	Yes	Yes
GEE Style OR (p)	10.17 (<.0001)	13.10 (<.0001)
GEE Orig OR (p)	1.61 (n.s.)	0.49 (n.s.)

Formality (Table 8). GPT-4o rewriter produces formality-manipulated pairs yielding BT 64.4% [56.1, 72.8], statistically indistinguishable from the Qwen3-32B baseline (66.5% [58.5, 73.9]; $\Delta = -2.1$pp). The Likert Blindspot replicates with the independent rewriter ($p = 0.184$, n.s.).

Table 8: Rewriter independence test: formality axis, Qwen3-32B judge.

	GPT-4o Rewriter	Qwen3 Rewriter
NLI pass rate	90/100 (90%)	88/100 (88%)
BT P [95% CI]	.644 [.561, .728]	.665 [.585, .739]
Likert d (p)	−.081 (.184)	.011 (.908)
Blindspot	Yes	Yes

E NLI Threshold Sensitivity Sweep

Table 9 reports BT win probabilities across 11 NLI entailment thresholds for all four judge×axis configurations. All 44 combinations are significant ($p < 0.05$), and within each configuration BT varies by at most 3pp, ruling out threshold selection bias. The stability reflects the bimodal NLI score distribution: most pairs score above 0.95 or below 0.30, so thresholds between 0.50 and 0.95 admit nearly identical pair sets.

Table 9: BT win probability across NLI thresholds. All 44 cells significant ($p < 0.05$). n: pairs passing threshold.

Thresh	Formality BT		Verbosity BT	
	Qwen3	GPT-4o	Qwen3	GPT-4o
.50	.648	.615	.770	.768
.55	.648	.615	.770	.768
.60	.656	.621	.770	.768
.65	.656	.621	.767	.766
.70	.657	.623	.767	.766
.75	.657	.623	.773	.760
.80	.657	.623	.773	.760
.85	.657	.623	.773	.760
.90	.665	.630	.773	.760
.95	.653	.629	.765	.752
.98	.649	.642	.771	.769

F Direction-Aware Gate Validation

For verbosity, the direction-aware NLI gate admits 21 additional pairs (32%) compared to strict bidirectional verification. These pairs show consistent preference direction (BT 73.8% vs. 78.5% for bidirectional-pass pairs), and overall BT estimates are statistically indistinguishable ($\Delta = -1.1$pp, 95% CI $[-5.6, +3.0]$), confirming that direction-aware gating recovers valid pairs without inflating

the effect. Direction-aware gating admits 21 additional pairs with consistent bias direction (BT 73.8%, $p = 0.004$); strict bidirectional-only pairs show BT 78.5% ($n = 65$) versus the aggregate BT 77.3% ($n = 86$), a difference of $\Delta = -1.1$pp with the 95% CI including zero. The direction-aware gate does not introduce selection bias. For formality, strict bidirectional verification is appropriate: a relaxed (direction-aware) gate would admit 9 additional pairs with reversed preference (BT 33.4%), indicating content preservation failure rather than genuine style preference. This asymmetry empirically justifies the axis-dependent gating strategy described in Section 3.2.

G Claude Sonnet 4 Equivalence Testing

Claude Sonnet 4 serves as negative control ($n = 88$; BT 56.2%/54.1%, $p = 0.091/0.284$). Formal equivalence testing (TOST; Lakens, 2017) yields margin-dependent conclusions (Table 10): equivalence is confirmed at our SESOI of $\Delta = 15$pp ($p = .010$), corresponding to a BT threshold of 65%—approximately where format preference would materially affect model rankings. At $\Delta = 10$pp, TOST is not significant ($p = .156$, power $\approx 42\%$), consistent with a residual preference (BT $= 56.2\%$) that falls below practical significance. We report all tested margins for transparency. The minimum equivalence margin achieving 80% power at $n = 88$ is $\Delta = 15.7$pp; tighter margins require larger n (§6).

Table 10: TOST equivalence sensitivity for Claude formality (BT $= 0.562$, $n = 88$). We select $\Delta = 15$pp as the SESOI: a BT threshold of 65% where format preference materially affects model rankings. At smaller margins the test lacks power, consistent with a residual preference below practical significance.

Δ (pp)	p	Result
5	.626	Not equivalent
10	.156	Not equivalent
15	.010	Equivalent
20	.0001	Equivalent

H Bradley-Terry Goodness of Fit and Stratified Analysis

Bradley-Terry goodness-of-fit tests confirm adequate model fit for both axes (formality: deviance $= 52.80$, df $= 45$, $p = 0.198$; verbosity: $p = 0.905$). Stratified analysis of consistent pairs by original-text direction reveals no preference heterogeneity for verbosity (Fisher $p = 0.453$) but significant direction asymmetry for formality (88.9% vs. 52.6%, Fisher $p = 0.015$), independently confirming the originality confound structure identified by GEE (Table 3).

To assess robustness to originality confounding, we restrict to pairs where the judge assigned equal Likert scores to original and counterfactual texts (eliminating detectable quality differences). Style preference persists: formality BT $= 67.3\%$ ($n = 55$, $\Delta = +0.8$pp from full sample, 95% CI $[-11.0, +12.6]$); verbosity BT $= 76.0\%$ ($n = 50$, $\Delta = -1.3$pp, 95% CI $[-13.3, +10.3]$). In the anti-originality subset (pairs where the original never wins), effects strengthen (formality 75.0%, verbosity 86.1%), confirming that style preference is not an artifact of originality confounding.

I Bradley-Terry Confidence Intervals

Table 11 reports bootstrap 95% confidence intervals for all Bradley-Terry estimates in Table 1.

Table 11: Bootstrap 95% CIs for BT win probabilities (10,000 pair-level resamples).

Bench	Judge	Axis	BT P [95% CI]
LLMBar	GPT-4o	Form.	.630 [.557, .703]
	GPT-4o	Verb.	.760 [.667, .853]
	GPT-4o	Reg.	.570 [.468, .673]
	Qwen3	Form.	.665 [.585, .739]
	Qwen3	Verb.	.773 [.703, .843]
	Qwen3	Reg.	.553 [.479, .628]
	Llama-70B	Form.	.695 [.619, .765]
	Llama-70B	Verb.	.735 [.655, .809]
	Llama-70B	Reg.	.635 [.555, .712]
	Claude	Form.	.562 [.489, .636]
	Claude	Verb.	.541 [.459, .616]
AEval	GPT-4o	Form.	.727 [.663, .785]
	Qwen3	Form.	.657 [.581, .733]
	Qwen3	Verb.	.841 [.774, .902]
MTB	GPT-4o	Form.	.692 [.623, .753]
	Qwen3	Form.	.712 [.630, .788]

J Bootstrap Attenuation Analysis

To distinguish measurement attenuation from insufficient statistical power, we subsample the $n = 355$ formality replication (Qwen3-32B) at seven sample sizes and compute Likert Cohen's d with bootstrap 95% CIs at each (Table 12). Under a power limitation, d would remain constant while only p-values improve with larger n; under measurement attenuation, d itself converges to zero. The data support

attenuation: $d_{\text{median}} = -0.017$ across all sample sizes, with CIs narrowing from ± 0.3 to ± 0.1 but never excluding zero. The same data yield BT $= 66.5\%$ ($p < 0.0001$), confirming that the pairwise signal is robust. Verbosity shows a similar pattern ($d \approx 0.15$–0.19, stable, CIs always containing zero; BT $= 76$–77%, $p < 0.0001$). Within the agnostic framework of §5.1, this evidence favors the measurement attenuation interpretation but does not exclude construct divergence.

Table 12: Bootstrap attenuation analysis: formality × Qwen3-32B. Likert d remains flat near zero across sample sizes while BT is stable at 66.5%.

n	Likert d	95% CI	Contains 0?
50	-0.017	$[-0.308, 0.262]$	Yes
100	-0.015	$[-0.219, 0.176]$	Yes
150	-0.019	$[-0.179, 0.145]$	Yes
200	-0.016	$[-0.157, 0.119]$	Yes
250	-0.018	$[-0.144, 0.107]$	Yes
300	-0.018	$[-0.131, 0.095]$	Yes
355	-0.017	$[-0.123, 0.087]$	Yes

K Formal Power Analysis

Table 13 reports the sample size required for Likert ATE to achieve 80% power at observed effect sizes across all six judge×axis cells. Required sample sizes range from 336 (Qwen3-32B verbosity, $d = 0.153$) to 27,159 (Qwen3-32B formality, $d = -0.017$), representing 4–206× the current n. For five of six cells, the required n exceeds 1,200—far beyond typical evaluation budgets—converging with the bootstrap attenuation curve (Appendix J) to confirm the Blindspot as a structural measurement limitation.

Table 13: Required Likert sample size for 80% power at observed d. BT detects all six effects at $p < 0.001$ with $n \leq 100$.

Judge	Axis	Observed d	Required n	Ratio
GPT-4o	Form.	0.073	1,473	16.7×
GPT-4o	Verb.	0.080	1,227	14.3×
Qwen3	Form.	-0.017	27,159	76.5×
Qwen3	Verb.	0.153	336	3.9×
Llama-70B	Form.	0.021	18,157	206.3×
Llama-70B	Verb.	0.034	6,737	78.3×

L Verbosity Sensitivity Hierarchy

Table 14 replicates the seven-method sensitivity analysis (Table 2) on the verbosity axis ($n = 86$ NLI-filtered pairs, Qwen3-32B judge, LLMBar). The attenuation pattern is more extreme than formality: all six Likert-based methods remain non-significant (lowest $p = 0.213$) despite a stronger pairwise signal (BT 77.3% vs. 66.5% for formality), and Likert effect sizes ($d = 0.036$–0.056) are near zero.

Table 14: Sensitivity hierarchy for verbosity ($n = 86$, Qwen3-32B, LLMBar). All Likert methods fail; only BT reaches significance.

Method	Statistic	p	Sig?
BT model	77.3% $[70.3, 84.3]$	<0.001	Yes
Likert ATE	$d = 0.056$	0.691	No
Binarized χ^2	OR $= 2.16$	0.213	No
Ordinal reg.	$\beta = 0.159$	0.575	No
Paired t-test	$d = 0.056$	0.607	No
Mann-Whitney U	$r = -0.046$	0.578	No
Cohen's d	$d = 0.036$	—	No

Effect sizes use style-aligned scoring (verbose − concise), isolating style preference direction. Table 1 reports RATE-compatible ATE (counterfactual − original), which conflates style and originality effects.

M 100-Point Continuous Likert Scale

To test whether the Blindspot reflects integer-scale granularity rather than a paradigm difference between absolute and relative evaluation, we replicate the formality and verbosity experiments using a 100-point continuous scale (0–100) with Qwen3-32B as judge (Table 15). Both axes remain non-significant: formality $p = 0.725$ (tie rate 65.9%) and verbosity $p = 0.117$ (tie rate 51.2%), while the same data yield BT 66.5% ($p < 0.0001$) and 77.3% ($p < 0.0001$) respectively. The persistence of high tie rates even with 101 possible values confirms that LLM judges do not exploit finer granularity to express style preferences; the Blindspot arises from absolute-vs-relative evaluation paradigms, not scale resolution. This analysis uses Qwen3-32B, which also served as the style rewriter (evaluator-producer overlap); results should be interpreted accordingly.

Table 15: 100-point continuous Likert vs. pairwise (BT). The Blindspot persists despite 101 possible score values.

Axis	Likert p	Tie rate	BT P	BT p
Formality	.725	65.9%	.665	$<.0001$
Verbosity	.117	51.2%	.773	$<.0001$

N NLI Gate Ablation

Table 16 compares BT estimates across NLI-pass, NLI-fail, and all pairs. For formality, NLI-fail pairs reverse to 41.7% (preferring casual over formal), demonstrating that the gate removes content-confounded pairs rather than amplifying style signal. The small NLI-fail sample ($n \approx 12$) limits precision; this serves as a bounding analysis. For verbosity, NLI-fail pairs remain at 75.0%, consistent with the dominant style effect (OR $= 13.10$) overwhelming content noise.

Table 16: NLI gate ablation: BT win probability by NLI subset.

Axis	NLI-pass	NLI-fail	All
Formality	66.5%	41.7%	63.0%
Verbosity	77.3%	75.0%	76.0%

O Rewriting Artifact Control

Same-style double-rewrite (DR) controls quantify the rewriting artifact using BT models (consistent with all other pairwise analyses). Table 17 reports BT $P(\text{original})$—the probability the judge prefers the original over its same-style DR. Both judges systematically prefer LLM-polished text across axes; for verbosity, BT $P(\text{original}) < 0.17$ indicates that rewriting quality dominates when content is held constant, validating the GEE decomposition that separates style from artifact effects. Same-style double-rewrite controls for style-neutral polishing artifacts but cannot exclude style-specific polishing interactions (e.g., formal rewriting may polish differently than casual rewriting).

Table 17: Rewriting artifact: BT $P(\text{orig})$ vs. same-style double-rewrite (bootstrap 95% CI, 10,000 resamples). Values <0.5: judges prefer LLM-polished text.

Axis	Judge	BT $P(\text{orig})$ [CI]	p
Form.	GPT-4o	.418 [.342, .500]	.037
	Qwen3	.377 [.303, .451]	.0006
Verb.	GPT-4o	.167 [.084, .262]	$<.0001$
	Qwen3	.082 [.035, .140]	$<.0001$

P Implementation Details

Models and APIs. Style-controlled rewriting uses Qwen3-32B (temperature 0.7, top-p 0.9). Target judges: GPT-4o (`gpt-4o-2024-08-06`, accessed April–May 2026), Qwen3-32B (locally hosted), Llama-3.3-70B (locally hosted via vLLM). Negative control: Claude Sonnet 4 (`claude-sonnet-4-20250514`, accessed May 2026). All judge calls use temperature 0 and `max_tokens` $= 2048$.

NLI gate. Content preservation is verified using DeBERTa-v3-large (He et al., 2023) fine-tuned on MNLI, with entailment threshold $\tau = 0.90$. Bidirectional entailment is applied for formality and register; direction-aware entailment for verbosity (Section 3.2).

Statistical methods. Bradley-Terry parameters are estimated via spectral initialization followed by minorization-maximization (Hunter, 2004) using the `choix` Python library (default regularization). All bootstrap CIs use 10,000 pair-level resamples with seed 42. GEE models use exchangeable working correlation with robust (sandwich) standard errors via `statsmodels` (Section 5.2; diagnostics in Appendix C).

Benchmarks. LLMBar (Zeng et al., 2024): 100 instruction-response pairs ($n = 88$ after NLI gate, $n = 355$ for scaled replication). AlpacaEval (Dubois et al., 2025): 100 sampled pairs. MT-Bench (Zheng et al., 2023): 80 multi-turn pairs.

Code and data. All code, prompts, and experimental data will be released upon acceptance under an open-source license.

Q Holm-Bonferroni Sensitivity Analysis

Table 18 reports adjusted p-values under the full 8-test family (4 judges $\times$ 2 primary axes), including Gemini-2.5-Flash which is excluded from the primary 6-test family (§4.1). All results remain significant at $\alpha = 0.05$.

Table 18: Holm-Bonferroni sensitivity analysis: 8-test family (4 judges $\times$ 2 primary axes). All results remain significant under the expanded family ($\alpha = 0.05$).

Rank	Judge	Axis	BT P	Raw p	Threshold	Adj. p
1	Gemini Flash	Verb.	.708	$<.0001$	.00625	$<.001$
2	GPT-4o	Verb.	.760	.0001	.00714	.0007
3	Qwen3-32B	Form.	.665	.0001	.00833	.0007
4	Qwen3-32B	Verb.	.773	.0001	.0100	.0007
5	Llama-70B	Form.	.695	.0001	.0125	.0007
6	Llama-70B	Verb.	.735	.0001	.0167	.0007
7	GPT-4o	Form.	.630	.0004	.0250	.0008
8	Gemini Flash	Form.	.574	.0306	.0500	.0306

R Cross-Axis Contamination Analysis

Our per-axis design assumes that rewriting along one style axis does not substantially alter the text

along other axes. We quantify cross-axis contamination by measuring off-target surface proxies before and after rewriting. For formality rewrites ($n = 355$), we measure word count and sentence count as verbosity proxies; for verbosity rewrites ($n = 86$), we measure type-token ratio (TTR) as a lexical diversity proxy for formality, and Flesch-Kincaid (FK) grade level.

Table 19: Cross-axis contamination: effect of rewriting along one style axis on surface proxies of the other axis. Cohen's d: $|d| < 0.20$ negligible, 0.20–0.50 small contamination (Cohen, 1988).

Direction	Proxy	Δ	d	Interp.
Form. → Verb.	Word count	-3.4%	-0.23	Small
Form. → Verb.	Sent. count	-0.2%	-0.03	Negligible
Verb. → Form.	TTR	-0.006	-0.04	Negligible
Verb. → Form.	FK grade	—	0.50	Confounded*

* FK = f(sentence length, word length); sentence length change is inherent to verbosity manipulation, not formality contamination.

Formality rewrites leave sentence structure virtually unchanged ($d = -0.03$); the word count shift ($d = -0.23$, exceeding the negligible threshold but within the small range) reflects word-level adjustments (contractions, elaboration) that are directionally symmetric (formal→casual -5.0%, casual→formal $+5.6\%$) and do not constitute verbosity contamination. Verbosity rewrites leave lexical diversity unchanged (TTR $d = -0.04$); the elevated FK grade ($d = 0.50$) is an artifact of FK's dependence on sentence length, which changes by design in verbosity manipulation. These results confirm that cross-axis contamination is small to negligible for the proxies measured; the GEE originality term (§5.2) provides additional protection by absorbing any residual quality confound. A robustness check adding word-count change as a GEE covariate leaves all Style ORs virtually unchanged ($\Delta < 12\%$, all covariate $p > .28$), confirming that the small cross-axis word-count shift does not confound formality estimates.

S Human Annotation Agreement Statistics

Two independent annotators evaluated all 50 NLI-scored pairs for content preservation. Annotator 1 (A1) judged 42/50 pairs content-preserving; Annotator 2 (A2) judged 33/50, reflecting systematically stricter criteria. All inter-annotator disagreements are predominantly unidirectional (A1-Yes/A2-No: 10; A1-No/A2-Yes: 1), confirming a calibration difference rather than random noise. Inter-annotator agreement: raw 78% (39/50), Cohen's $\kappa = 0.438$ (moderate), Gwet's AC1 $= 0.648$ (substantial). The κ–AC1 gap reflects prevalence skew (Feinstein and Cicchetti, 1990): the high base rate of content-preserving pairs inflates chance agreement, compressing κ.

Among the 25 gate-passed pairs, A1 confirms all 25 as content-preserving (precision 100%; Wilson CI $[86.7\%, 100\%]$); A2 confirms 21/25 (84%; CI $[65.3\%, 93.6\%]$). The four A2 disagreements (S003, S018, S039, S049) are all verbosity-axis condensation pairs where A2 applied stricter meaning-equivalence criteria to information-reducing rewrites. Under majority vote, pass precision is 100% (25/25).

NLI–annotator agreement: NLI–A1 $\kappa = 0.320$, NLI–A2 $\kappa = 0.360$. Both κ values are nominally "fair" but deflated by the same prevalence mechanism. The clinically relevant statistic for a precision-oriented gate is not overall agreement but *pass precision*: at 84–100% observed pass precision, the gate's validity for its intended function—blocking content contamination from the analysis—is confirmed under both lenient and strict annotation standards. All A1 disagreements with the NLI gate occur on gate-*rejected* pairs; A2's four gate-passed disagreements reflect subjective meaning boundaries on verbosity condensation rather than content-quality contamination.

References

Jacopo Amidei, Paul Piwek, and Alistair Willis. 2019. The use of rating and Likert scales in natural language generation human evaluation tasks: A review and some recommendations. In *Proceedings of the 12th International Conference on Natural Language Generation (INLG)*.

Anna Bavaresco, Raffaella Bernardi, Leonardo Bertolazzi, Arianna Bisazza, Davide Ceron, Omid Ghahroodi, and 1 others. 2025. LLMs instead of human judges? a large scale empirical study across 20 NLP evaluation tasks. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*.

Ondřej Bojar, Rajen Chatterjee, Christian Federmann, Yvette Graham, Barry Haddow, Matthias Huck, Antonio Jimeno Yepes, Philipp Koehn, Varvara Logacheva, Christof Monz, Matteo Negri, Aurélie Névéol, Mariana Neves, Martin Popel, Matt Post, Raphael Rubino, Carolina Scarton, Lucia Specia, Marco Turchi, and 2 others. 2016. Findings of the 2016 conference on machine translation. In *Proceedings of the First Conference on Machine Translation: Volume 2, Shared Task Papers*, pages 131–198,

Berlin, Germany. Association for Computational Linguistics.

Meriem Boubdir, Edward Kim, Beyza Ermis, Sara Hooker, and Marzieh Fadaee. 2024. Elo uncovered: Robustness and best practices in language model evaluation. In *Advances in Neural Information Processing Systems (NeurIPS)*.

Ralph Allan Bradley and Milton E. Terry. 1952. Rank analysis of incomplete block designs: I. The method of paired comparisons. *Biometrika*, 39(3/4):324–345.

Ben Carterette, Paul N. Bennett, David Maxwell Chickering, and Susan T. Dumais. 2008. Here or there: Preference judgments for relevance. In *Advances in Information Retrieval, 30th European Conference on IR Research, ECIR 2008*, volume 4956 of *Lecture Notes in Computer Science*, pages 16–27, Glasgow, UK. Springer.

Asli Celikyilmaz, Elizabeth Clark, and Jianfeng Gao. 2020. Evaluation of text generation: A survey. *arXiv preprint arXiv:2006.14799*.

Guiming Hardy Chen, Shunian Chen, Ruifei Zhang, Junying Chen, Xiangbo Xiong, Zhouyu Yu, Hongying Lin, and Benyou Wang. 2024. Humans or LLMs as the judge? a study on judgement biases. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP)*.

Wei-Lin Chiang, Lianmin Zheng, Ying Sheng, Anastasios Nikolas Angelopoulos, Tianle Li, Dacheng Li, Banghua Zhu, Hao Zhang, Michael I. Jordan, Joseph E. Gonzalez, and Ion Stoica. 2024. Chatbot arena: An open platform for evaluating LLMs by human preference. In *Proceedings of the International Conference on Machine Learning (ICML)*.

Young-Min Cho, Yuan Yuan, Sharath Chandra Guntuku, and Lyle Ungar. 2026. A concise agent is less expert: Revealing side effects of using style features on conversational agents. *arXiv preprint arXiv:2601.10809*.

Elizabeth Clark, Tal August, Sofia Serrano, Nikita Haduong, Suchin Gururangan, and Noah A. Smith. 2021. All that's 'human' is not gold: Evaluating human evaluation of generated text. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 7282–7296.

Jacob Cohen. 1988. *Statistical Power Analysis for the Behavioral Sciences*, 2nd edition. Lawrence Erlbaum Associates.

Jamie DeCoster, Marcello Gallucci, and Anne-Marie R. Iselin. 2011. Best practices for using median splits, artificial categorization, and their continuous alternatives. *Journal of Experimental Psychopathology*, 2(2):197–209.

Daniel Deutsch, George Foster, and Markus Freitag. 2023. Ties matter: Meta-evaluating modern metrics with pairwise accuracy and tie calibration. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 12914–12929, Singapore. Association for Computational Linguistics.

Yann Dubois, Xuechen Li, Rohan Taori, Tianyi Zhang, Ishaan Gulrajani, Jimmy Ba, Carlos Guestrin, Percy Liang, and Tatsunori B. Hashimoto. 2023. AlpacaFarm: A simulation framework for methods that learn from human feedback. In *Advances in Neural Information Processing Systems (NeurIPS)*.

Yann Dubois and 1 others. 2025. Length-controlled AlpacaEval: A simple way to debias automatic evaluators. In *Proceedings of the International Conference on Learning Representations (ICLR)*.

Amir Feder, Katherine A. Keith, Emaad Manzoor, Reid Pryzant, Dhanya Sridhar, Zach Wood-Doughty, Jacob Eisenstein, Justin Grimmer, Roi Reichart, Margaret E. Roberts, Brandon M. Stewart, Victor Veitch, and Diyi Yang. 2022. Causal inference in natural language processing: Estimation, prediction, interpretation and beyond. *Transactions of the Association for Computational Linguistics*, 10:1138–1158.

Amir Feder, Nadav Oved, Uri Shalit, and Roi Reichart. 2021. CausaLM: Causal model explanation through counterfactual language models. *Computational Linguistics*, 47(2):333–386.

Alvan R. Feinstein and Domenic V. Cicchetti. 1990. High agreement but low kappa: I. The problems of two paradoxes. *Journal of Clinical Epidemiology*, 43(6):543–549.

Pengcheng He, Jianfeng Gao, and Weizhu Chen. 2023. DeBERTaV3: Improving DeBERTa using ELECTRA-style pre-training with gradient-disentangled embedding sharing. In *Proceedings of the International Conference on Learning Representations (ICLR)*.

John M. Hoenig and Dennis M. Heisey. 2001. The abuse of power: The pervasive fallacy of power calculations for data analysis. *The American Statistician*, 55(1):19–24.

David R. Hunter. 2004. MM algorithms for generalized Bradley-Terry models. *The Annals of Statistics*, 32(1):384–406.

Hawon Jeong, ChaeHun Park, Jimin Hong, Hojoon Lee, and Jaegul Choo. 2025. The comparative trap: Pairwise comparisons amplifies biased preferences of LLM evaluators. In *Proceedings of the 8th BlackboxNLP Workshop: Analyzing and Interpreting Neural Networks for NLP*, pages 79–108.

Svetlana Kiritchenko and Saif Mohammad. 2017. Best-worst scaling more reliable than rating scales: A case study on sentiment intensity annotation. In *Proceedings of the 55th Annual Meeting of the Association for*

Computational Linguistics (Volume 2: Short Papers), pages 465–470.

Ryan Koo, Minhwa Lee, Vipul Raheja, Jong Inn Park, Zae Myung Kim, and Dongyeop Kang. 2024. Benchmarking cognitive biases in large language models as evaluators. In *Findings of the Association for Computational Linguistics: ACL 2024*.

Jon A. Krosnick. 1999. Survey research. *Annual Review of Psychology*, 50(1):537–567.

Daniël Lakens. 2017. Equivalence tests: A practical primer for t tests, correlations, and meta-analyses. *Social Psychological and Personality Science*, 8(4):355–362.

Dawei Li and 1 others. 2026. Preference leakage: A contamination problem in LLM-as-a-judge. In *Proceedings of the International Conference on Learning Representations (ICLR)*.

Haitao Li, Junjie Chen, Qingyao Ai, Zhumin Chu, Yujia Zhou, Qian Dong, and Yiqun Liu. 2025. CalibraEval: Calibrating prediction distribution to mitigate selection bias in LLMs-as-judges. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*.

Kung-Yee Liang and Scott L. Zeger. 1986. Longitudinal data analysis using generalized linear models. *Biometrika*, 73(1):13–22.

Adian Liusie, Potsawee Manakul, and Mark Gales. 2024. LLM comparative assessment: Zero-shot NLG evaluation through pairwise comparisons using large language models. In *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 139–151.

Jordan J. Louviere, Terry N. Flynn, and A. A. J. Marley. 2015. *Best-Worst Scaling: Theory, Methods and Applications*. Cambridge University Press.

Aakanksha Murugadoss and 1 others. 2025. Style over substance: Failure modes of LLM judges in alignment benchmarking. In *Proceedings of the International Conference on Learning Representations (ICLR)*.

Jekaterina Novikova, Ondřej Dušek, and Verena Rieser. 2018. RankME: Reliable human ratings for natural language generation. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 2: Short Papers)*.

Arjun Panickssery, Samuel R. Bowman, and Shi Feng. 2024. LLM evaluators recognize and favor their own generations. In *Advances in Neural Information Processing Systems (NeurIPS)*.

Junsoo Park, Seungyeon Jwa, Meiying Ren, Daeyoung Kim, and Sanghyuk Choi. 2024. OffsetBias: Leveraging debiased data for tuning evaluators. In *Findings of the Association for Computational Linguistics: EMNLP 2024*, pages 1043–1067.

Carolyn C. Preston and Andrew M. Colman. 2000. Optimal number of response categories in rating scales: Reliability, validity, discriminating power, and respondent preferences. *Acta Psychologica*, 104(1):1–15.

Jingyi Qiu, Hong Chen, and Zongyi Li. 2025. Counterfactual LLM-based framework for measuring rhetorical style. *arXiv preprint arXiv:2512.19908*.

David Reber, Sean Richardson, Todd Nief, Cristina Garbacea, and Victor Veitch. 2025. RATE: Causal explainability of reward models with imperfect counterfactuals. In *Proceedings of the International Conference on Machine Learning (ICML)*.

Yuki Saito, Noriki Nishida, and Naoaki Okazaki. 2023. Verbosity bias in multi-turn chatbot evaluations. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*.

João Sedoc and Lyle Ungar. 2020. Item response theory for efficient human evaluation of chatbots. In *Proceedings of the First Workshop on Evaluation and Comparison of NLP Systems*, pages 21–33. Association for Computational Linguistics.

Lin Shi, Chiyu Ma, Wenhua Liang, Xingjian Diao, Weicheng Ma, and Soroush Vosoughi. 2025. Judging the judges: A systematic study of position bias in LLM-as-a-judge. In *Proceedings of the 14th International Joint Conference on Natural Language Processing and the 4th Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*.

Leonard J. Simms, Kerry Zelazny, Trevor F. Williams, and Lee Bernstein. 2019. Does the number of response options matter? Psychometric perspectives using personality questionnaire data. *Psychological Assessment*, 31(4):557–566.

Mohammad Soumik and 1 others. 2026. Judging the judges: A systematic evaluation of bias mitigation strategies in LLM-as-a-judge pipelines. *arXiv preprint arXiv:2604.23178*.

Stanley Smith Stevens. 1946. On the theory of scales of measurement. *Science*, 103(2684):677–680.

Rickard Stureborg, Dimitris Alikaniotis, and Yoshi Suhara. 2024. Large language models are inconsistent and biased evaluators. *arXiv preprint arXiv:2405.01724*.

Sijun Tan, Siyuan Zhuang, Kyle Montgomery, William Tang, and 1 others. 2025. JudgeBench: A benchmark for evaluating LLM-based judges. In *Proceedings of the International Conference on Learning Representations (ICLR)*.

Louis Leon Thurstone. 1927. A law of comparative judgment. *Psychological Review*, 34(4):273–286.

Tuhina Tripathi, Manya Wadhwa, Greg Durrett, and Scott Niekum. 2025. Pairwise or pointwise? evaluating feedback protocols for bias in LLM-based evaluation. In *Proceedings of the 2nd Conference on Language Modeling (COLM)*.

Chris van der Lee, Albert Gatt, Emiel van Miltenburg, and Emiel Krahmer. 2021. Human evaluation of automatically generated text: Current trends and best practice guidelines. *Computer Speech & Language*, 67:101151.

Chris van der Lee, Albert Gatt, Emiel van Miltenburg, Sander Wubben, and Emiel Krahmer. 2019. Best practices for the human evaluation of automatically generated text. In *Proceedings of the 12th International Conference on Natural Language Generation (INLG)*, pages 355–368.

Victor Veitch, Dhanya Sridhar, and David Blei. 2020. Adapting text embeddings for causal inference. In *Proceedings of the Conference on Uncertainty in Artificial Intelligence (UAI)*.

Peiyi Wang, Lei Li, Liang Chen, Zefan Cai, Dawei Zhu, Binghuai Lin, Yunbo Cao, Lingpeng Kong, Qi Liu, Tianyu Liu, and Zhifang Sui. 2024. Large language models are not fair evaluators. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*.

Minghao Wu and Alham Fikri Aji. 2025. Style over substance: Evaluation biases for large language models. In *Proceedings of the 31st International Conference on Computational Linguistics (COLING)*.

Tongshuang Wu, Marco Tulio Ribeiro, Jeffrey Heer, and Daniel Weld. 2021. Polyjuice: Generating counterfactuals for explaining, evaluating, and improving models. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*.

Bo Yang, Lanfei Feng, Yunkui Chen, Yu Zhang, Xiao Xu, and Shijian Li. 2026. FairJudge: An adaptive, debiased, and consistent LLM-as-a-judge. *arXiv preprint arXiv:2602.06625*.

Jiayi Ye, Yanbo Wang, Yue Huang, Dongping Chen, Qihui Zhang, Nuno Moniz, Tian Gao, Werner Geyer, Chao Huang, Pin-Yu Chen, Nitesh V. Chawla, and Xiangliang Zhang. 2025. Justice or prejudice? quantifying biases in LLM-as-a-judge. In *Proceedings of the International Conference on Learning Representations (ICLR)*.

Zhiyuan Zeng, Jiatong Yu, Tianyu Gao, Yu Meng, Tanya Goyal, and Danqi Chen. 2024. Evaluating large language models at evaluating instruction following. In *Proceedings of the International Conference on Learning Representations (ICLR)*.

Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. 2023. Judging LLM-as-a-judge with MT-Bench and Chatbot Arena. In *Advances in Neural Information Processing Systems (NeurIPS)*.