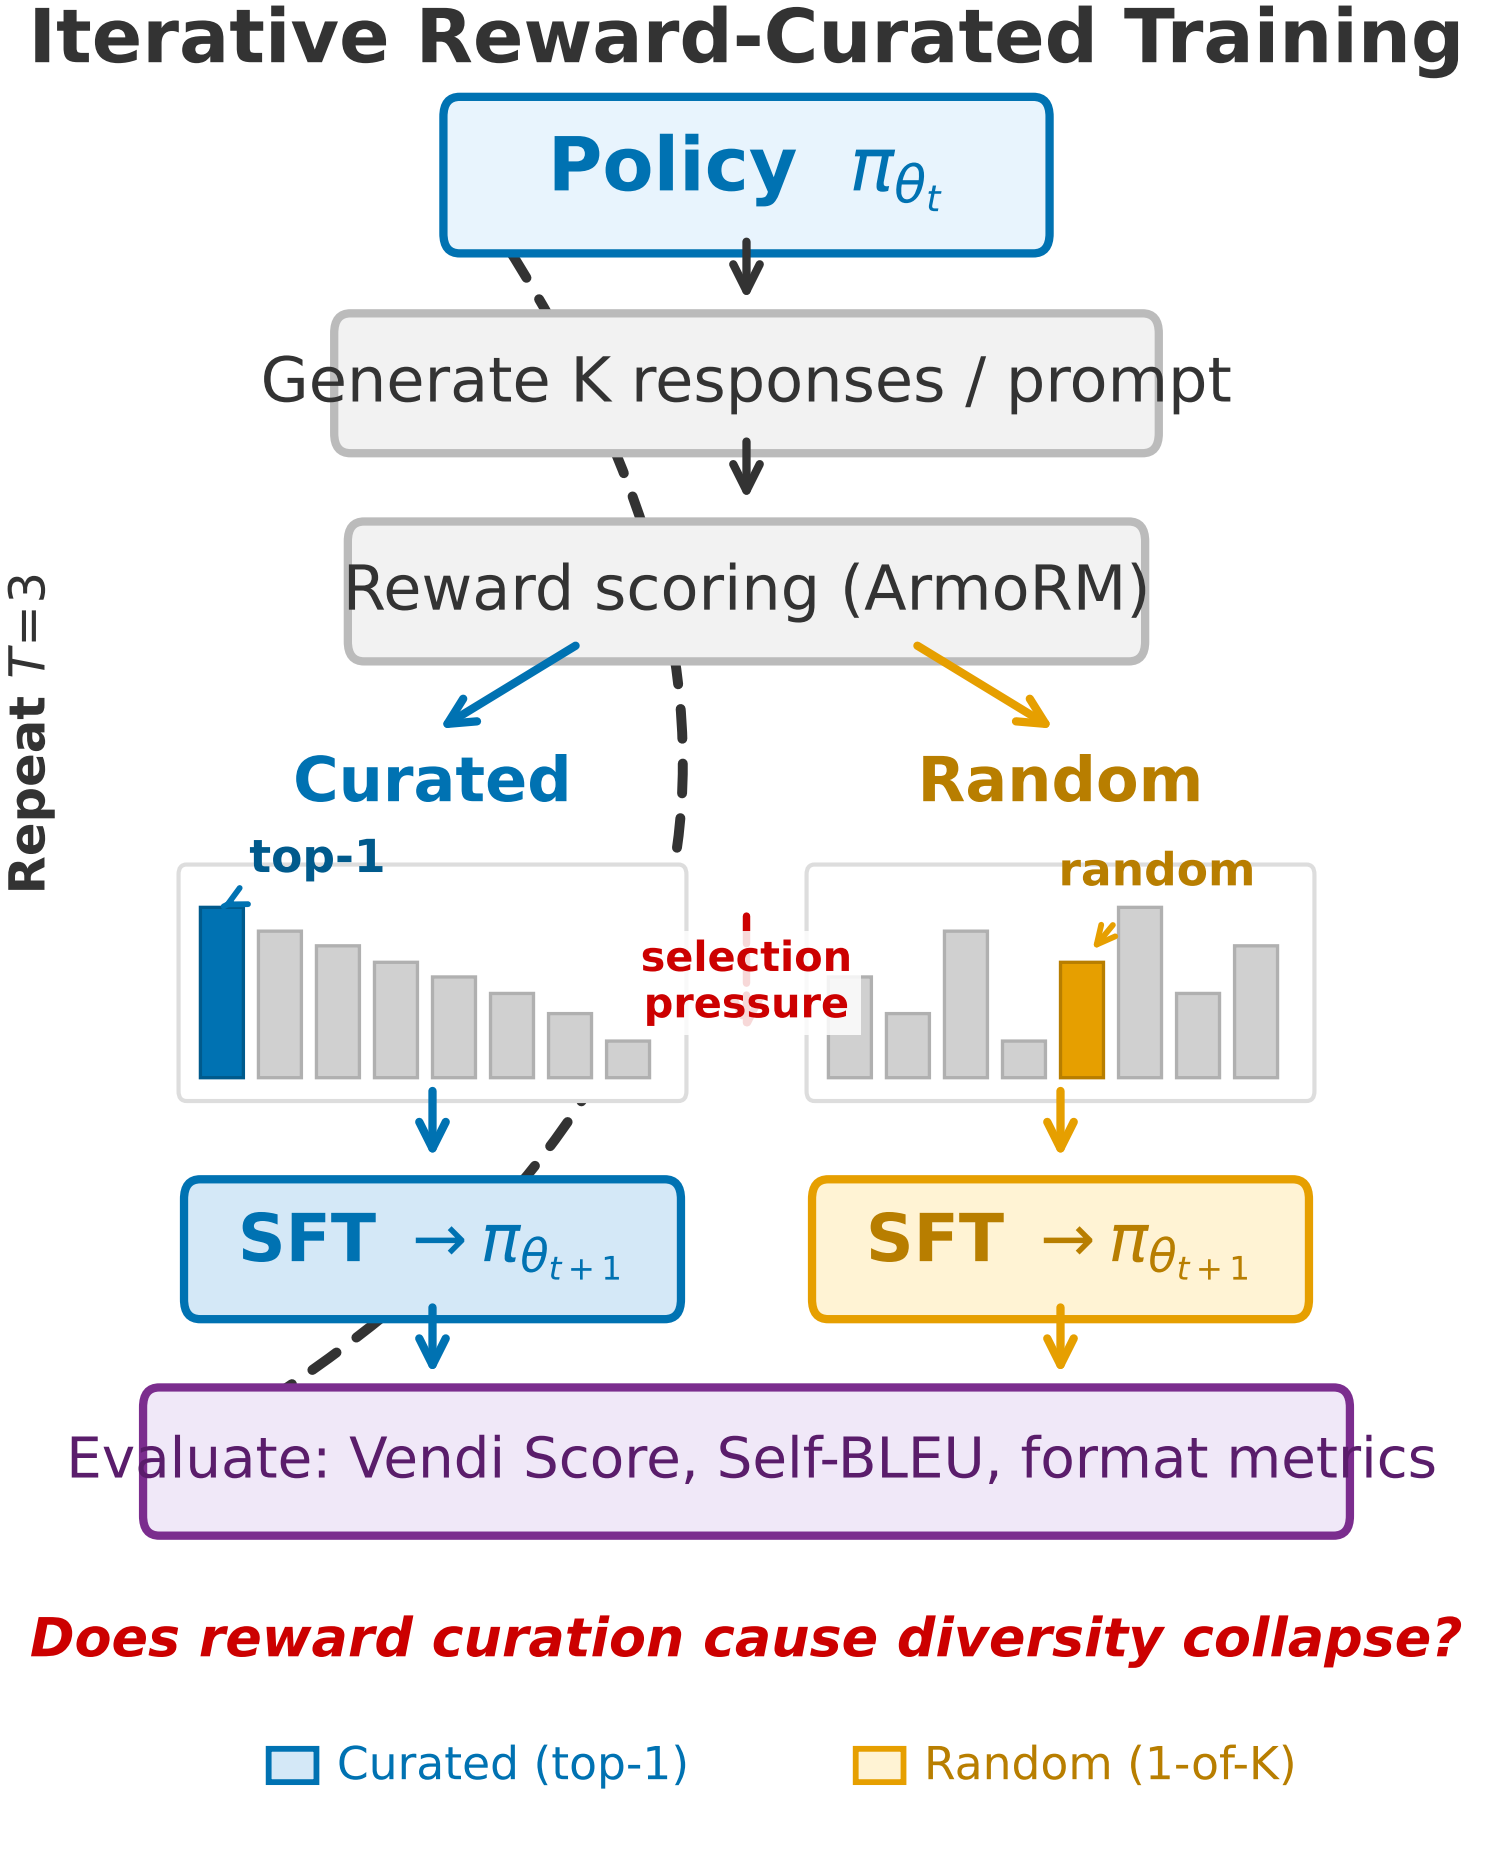
Iterative Reward-Curated Training
Policy π_{θ_t}
Generate K responses / prompt
Reward scoring (ArmoRM)
Repeat $T=3$
Curated
Random
top-1
random
selection pressure
SFT $\rightarrow \pi_{\theta_{t+1}}$
SFT $\rightarrow \pi_{\theta_{t+1}}$
Evaluate: Vendi Score, Self-BLEU, format metrics
Does reward curation cause diversity collapse?
Curated (top-1)
Random (1-of-K)