Why Temporal Abstraction Can Fail in Imagination-Based RL: A Diagnostic Analysis of Variable-Rate DreamerV3

Anonymous Author(s)
Affiliation
Address
email

Abstract: Temporal conditioning on step size Δt enables consistent gains in model-predictive control, but fails catastrophically in imagination-based actor-critic training: every variable-rate DreamerV3 configuration suffers substantial performance degradation. Factorial diagnostics on two continuous-control tasks identify the combination of trained temporal embedding and multi-rate data exposure as the strongest predictor of this failure, *conditioning-induced value overestimation*, or CIVO, with AUROC$=1.000$ over $n=36$ configurations; cross-environment analysis across seven additional environments confirms specificity to moderate-complexity locomotion. The recurrent state-space model develops a reward prediction gap at $\Delta t>1$ that the critic internalizes as optimistic values; the contrast with model-predictive control identifies feedback architecture as a hypothesized structural factor. CIVO exhibits systematic resilience to standard RL remedies: gradient, output, representation, and architecture-level interventions each fail through distinct mechanisms: pathway redundancy, anchor co-drift, phase transitions, and critic convergence in shared-data regimes. Only wide dual-critic combining min-aggregation with increased capacity achieves partial prevention, and open challenges remain in the architectural design of variable-rate model-based RL.

Keywords: World Models, Temporal Abstraction, Value Overestimation

1 Introduction

World-model reinforcement learning enables sample-efficient learning in complex continuous control tasks [1, 2]. Continuous control demands temporal abstraction: locomotion with heterogeneous gait phases, contact-rich tasks with variable sensing rates, and multi-rate decision-making across planning horizons all require handling dt mismatches. TAWM [3] shows that conditioning a world model on step size dt yields consistent gains within TD-MPC2's MPC framework [2]. Can the same dt-conditioning bring temporal abstraction to imagination-based actor-critic methods like DreamerV3 [4], where policies are trained entirely in imagined rollouts?

We find that it cannot: every variable-rate configuration suffers substantial performance degradation. Through a factorial analysis on Walker Walk [5], we isolate the combination of a *trained* dt-embedding (a learned vector concatenated to the RSSM input at each imagination step; our implementation of dt-conditioning) and multi-rate data exposure as the strongest predictor. This combination enables environment-dependent pathologies (reward prediction gaps in Walker Walk, or dynamics-model prior drift without measurable reward bias in Cheetah Run), which the critic internalizes as optimistic value estimates, a failure mode we term *conditioning-induced value overestimation* (CIVO), analogous to delusional bias [6] but model-induced rather than policy-class-induced. A normalized value-to-eval (NVE) diagnostic achieves perfect classification (AUROC$=1.000$, $n=36$; Appendix B).

Submitted to the 10th Conference on Robot Learning (CoRL 2026). Do not distribute.

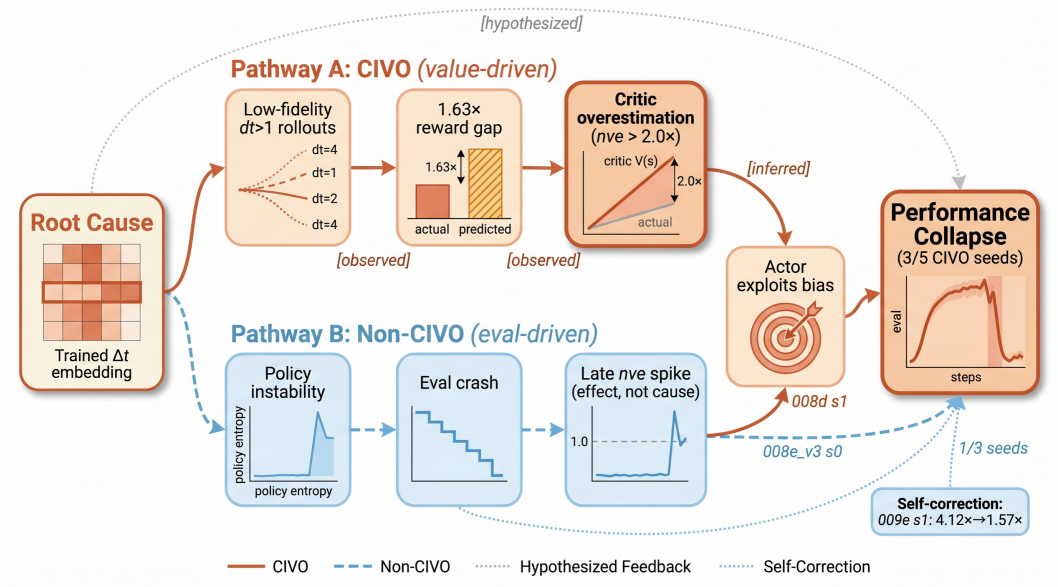


Figure 1: Proposed CIVO pathway (evidence levels annotated). The trained dt-embedding enables asymmetric reward prediction fidelity (**observed**, ablation; Appendix D), leading to optimistic critic values (**observed**, §3.3); actor exploitation of overestimated regions (**inferred** from entropy dynamics, Appendix D.1) → reinforced state-distribution skew (**hypothesized**).

A four-level intervention hierarchy (gradient, output, representation, architecture) shows that most interventions fail to eliminate CIVO; only wide dual-critic (min-aggregation + increased capacity) achieves partial prevention, though why this specific combination succeeds remains open. As world-model RL incorporates variable-rate decision-making, understanding when temporal conditioning fails is essential for the architectural design of variable-rate model-based RL.

Our contribution.

1. **Failure mode identification.** We identify CIVO as specific to imagination-based actor-critic with temporal conditioning. Factorial analysis isolates trained dt-embedding + multi-rate exposure as jointly necessary for CIVO susceptibility; frozen and removed controls never trigger CIVO. The NVE diagnostic achieves perfect binary classification (AUROC=1.000, n=36); cross-environment validation confirms specificity to moderate-complexity locomotion (§5).
2. **Structural barrier to remedy design.** A four-level intervention hierarchy (gradient, output, representation, architecture; 45 experiments) reveals that standard overestimation remedies are structurally ineffective in pure imagination: each class fails through a distinct mechanism (pathway redundancy, anchor co-drift, phase transitions, critic convergence in shared-data regimes), all traceable to the closed-loop constraint that the corrective signal originates from the overestimating model itself. Only wide dual-critic (min + increased capacity) achieves partial prevention, the sole positive result across all tested approaches.
3. **Mechanism analysis.** Frozen embeddings never trigger CIVO while trained embeddings frequently do ($p < 0.001$); this identifies embedding *learning*, not conditioning architecture, as the causal factor. A 2×2 critic factorial dissects the architectural result: min-aggregation drives recovery, increased capacity enables prevention, and their combination is the minimum sufficient condition.

2 Related Work

Temporal abstraction in RL. Temporal abstraction spans hierarchical, action-repetition, and world-model approaches, none of which has examined the interaction between flat dt-conditioning and imagination-based value estimation. The options framework [7], option-critic [8], HAM [9],

and MAXQ [10] formalize temporal abstraction through hierarchical policies. Semi-MDP theory [11] provides the mathematical foundation. Action-repetition methods achieve simpler temporal abstraction via real transitions: FiGAR [12] learns per-action repetition counts jointly with the policy, TempoRL [13] and Park et al. [14] provide complementary formulations. SPiRL [15] accelerates RL through learned skill priors that capture reusable temporal abstractions from offline data. Wang and Beltrame [16] learn elastic time steps that adapt step size to state dynamics. The Dreamer family [17, 18, 4, 19] trains policies in imagination at fixed $\mathrm{dt}{=}1$; hierarchical variants [20, 21, 22, 23] use subgoals or separate dynamics per scale, avoiding flat dt-conditioning. TAWM [3] conditions a world model on dt within TD-MPC2 [2]; MPC re-plans from real states, avoiding the critic accumulation we analyze (§5).

Value overestimation and model fidelity. Overestimation has been studied extensively in model-free and model-based RL, but no prior work identifies the specific pathology arising from learned temporal conditioning in imagination. Maximization bias [24] has been addressed by Double Q-learning [25], TD3 [26], and ensemble methods [27]. Scaling to larger critic ensembles is unlikely to help in pure imagination: our dual-critic experiments show $\rho > 0.999$ convergence regardless of capacity, because critics share the same imagined training distribution. In model-based settings, MBPO [28] and MOPO [29] bound rollout error, while MOReL [30] and COMBO [31] penalize model uncertainty. Unlike MOPO, which penalizes model uncertainty using real-data rollouts as an anchor, pure imagination lacks this external correction mechanism, allowing bias to compound unchecked. Lu et al. [6] identify delusional bias from policy-class constraints. CIVO is mechanistically distinct: the amplifier is the world model's co-adaptation with a learned temporal embedding, producing systematic reward-prediction bias at $\mathrm{dt} > 1$. Multi-timestep world models [32] learn predictions at variable horizons but do not examine the value overestimation pathology we isolate. Compounding error in multi-step rollouts [33, 34] is a related but distinct phenomenon: if CIVO were generic compounding error, untrained dt-embeddings at $\mathrm{dt}{=}2$ should also produce overestimation, yet they do not (NVE$<0.6\times$ for unexposed and frozen controls vs. $4.12\times$ for the trained configuration). Signed reward error analysis reinforces this distinction: $\mathrm{dt}{=}2$ rewards are *under*predicted (mean signed error [SE]$=-0.599$), which rules out observation-level overprediction and localizes CIVO to imagination-level prior drift (§5). Decision-aware model learning [35] optimizes for planning performance but does not address dt-conditioned co-adaptation. Our temporal conditioning noise (TCN) intervention relates to training-with-noise [36]; TCN's environment-specific dose sensitivity (§4) suggests that CIVO reflects a structural vulnerability rather than a regularization deficiency. Our factorial diagnostic methodology follows Henderson et al. [37]. This paper fills the gap at the intersection of temporal conditioning and imagination-based value estimation, revealing that shared imagined rollouts create a structural constraint: remedies that rely on training-data diversity (Double Q pessimism, ensembles, noise injection) lose their corrective power when both the training distribution and bootstrap targets are generated by the same learned model.

3 Experiments and Results

3.1 Setup

Base architecture. All experiments build on DreamerV3 [4], which learns an RSSM with discrete latent states (32 categories $\times$ 32 classes), uses symlog-transformed predictions, and trains an actor-critic entirely in imagined trajectories. The default configuration imagines $H{=}15$ steps at $\mathrm{dt}{=}1$ (one imagination step per environment transition).

Variable-rate imagination. We introduce three orthogonal modifications for variable-rate imagination. *Temporal embedding*: A learned embedding $e(\mathrm{dt})$ is concatenated to the RSSM input at each imagination step; we test **trained** (backpropagation-updated), **frozen** (fixed random projection, `requires_grad=False`), and **removed** ($d{=}0$). *Temporal subsampling*: The replay buffer augments training data with $\mathrm{dt}{>}1$ transitions by subsampling consecutive frames (max factor 2). *Horizon and step size*: $H \in \{15, 8\}$; `gate_fixed_dt` forces $\mathrm{dt} \in \{1, 2, 4\}$.

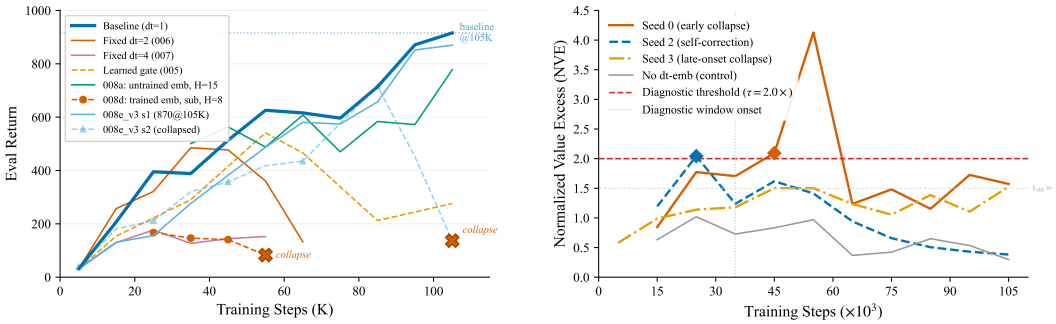


Figure 2: (a) Variable-rate imagination universally degrades performance. Baseline (black) reaches 915.7; all variable-rate configs degrade. (b) NVE diagnostic separates CIVO-positive from healthy configurations: trained dt-embedding seeds exceed the $2.0\times$ threshold; controls remain near $1.0\times$.

Factorial design. We adopt a factorial ablation across four factors: dt-embedding mode (trained / frozen / none), temporal subsampling (on / off), imagination horizon (H=15 / H=8), and step size (dt=1, 2, 4). The design (12 conditions, Table 15 in Appendix C) isolates trained dt-embedding and multi-rate exposure as orthogonal factors. Key configurations use 4–5 seeds; following Agarwal et al. [38], we report per-seed trajectories.

Diagnostic metric. We define the normalized value-to-eval ratio $\text{nve}_t = (v/e)_{\text{config},t}/(v/e)_{\text{baseline},t}$, where v/e is the critic's mean imagined value divided by mean evaluation return, and the baseline is the highest-performing seed (exp_001, seed 1, 505K steps). Analysis is restricted to a **35K+ window** to avoid early-training inflation. The $\text{nve} \geq 2.0\times$ threshold classifies all n=36 configurations across two environments (AUROC=1.000, F1=1.000; Appendix B); leave-one-environment-out cross-validation confirms robustness (Appendix B.6).

Intervention methods. *Selective gradient stopping* (SGS): $e_{\Delta t}$ is detached via $\text{sg}(\cdot)$ before selected prediction heads during the backward pass. We test reward-only (exp_016e), full-pathway (exp_018), and obs-decoder-only (exp_025/026) variants. *Temporal conditioning noise* (TCN) adds Gaussian noise $\mathcal{N}(0, \sigma^2)$ to $e_{\Delta t}$ during imagination ($\sigma \in \{0.2, 0.5\}$, three environments); *RAVE* clips imagined values exceeding an exponential moving average (EMA) anchor ($c \in \{2.0, 1.5\}$). *Dual-critic*: a 2×2 factorial of {single, dual} $\times$ {narrow ($\sim$1.2M), wide ($\sim$2.4M)} critic following TD3 [26].

Environments and training. All experiments use Walker Walk and Cheetah Run from the DeepMind Control Suite [5] with proprioceptive observations. Walker Walk (6-joint biped) is the primary environment; Cheetah Run (6-joint runner) serves as a severity probe (8–108$\times$ NVE vs. 2–5$\times$ on Walker). Baselines run 500K steps; treatments 105K steps. Hyperparameters follow DreamerV3 defaults [4] (Appendix E); training uses a single GPU ($\sim$80 GPU-hours total).

3.2 F1: All Variable-Rate Configurations Degrade

Every configuration that introduces $\text{dt} > 1$ imagination degrades relative to the dt=1 baseline (Figure 2a). The baseline reaches 915.7 at 105K; fixed dt=2 (006) peaks at 485.3 before collapsing to 132.9 (14% of baseline); fixed dt=4 (007) never exceeds 176.9 (19%); the learned gate (005) collapses to just 98.8, representing 11% of baseline performance.

3.3 F2: Trained Temporal Embedding Precedes Value Overestimation

The factorial design isolates trained dt-embedding as the strongest predictor. The *primary* comparison, 008d (trained, subsample ON) vs. 008f (no embedding, subsample ON), differs only in embedding training, yielding $\text{nve} = 4.12\times$ (eval 82.8) vs. $\text{nve} \leq 1.69\times$ (eval 139.5–468.8^f).

Table 1: Key configurations and outcomes. NVE reports peak nve in the 35K+ window (§3.1). CIVO-positive seeds are shaded. Abbreviations: trn = trained (d=32, subsample ON); exp[a] = trained, subsample OFF; frz[b] = frozen (d=32, `requires_grad=False`); — = removed (d=0). [c] Environment-specific baseline. [d] Frozen control matches; environment-specific. [e] Seed-dependent (Table 5). [f] Corrected-code rerun (Appendix H). [g] NVE median [range 0.31–127.6×]; mean±std dominated by outlier seed. Multi-seed entries report mean±std (n noted in Config column). NVE reference baselines: Walker s1 (915.7); Cheetah 446.1. Complete results for all 58 configurations in Table 4.

Exp	Config	dt_emb	Sub	H	Eval@final	NVE	Outcome
001	Baseline (dt=1, n=4)	—	—	15	787.7 ± 230.8	1.00×	Healthy
008d	Factorial (n=5)	trn	ON	8	190.6 ± 90.2	3.32 ± 1.61×	3/5 onset
008f	No dt_emb (n=4)	—	ON	8	302.4 ± 134.8	1.09 ± 0.42×	Healthy (0/4)
008a	Unexposed emb (n=4)	exp[a]	OFF	15	820.9 ± 68.8	0.46 ± 0.19×	Healthy (0/4)
015b	Frozen emb	frz[b]	ON	8	204.2 @105K	1.45×	Grad. starvation
022b	Frozen emb + Gumbel gate	frz[b]	ON	8	300.7 @105K	0.63×	No CIVO
Cross-environment validation							
078	Cheetah Run (n=4)	trn	ON	8	123.2 ± 132.9	med. 2.3×[g]	3/4 onset
019b	PickCube-v1	trn	ON	8	26.64 @200K	3.78×[c]	Transient onset
049	StackCube-v1	trn	ON	8	34.9 @212K	5.05×[d]	Env. transient
050	PegInsertionSide-v1	trn	ON	8	20.86 @212K	10.05×[e]	Env. dynamics[e]
046	Cheetah Run	frz[b]	ON	8	150.5 @215K	1.35×	No CIVO

Secondary comparisons against 008a/008c (unexposed embedding) co-vary subsample status and horizon alongside the embedding factor. The nve $\geq 2.0\times$ threshold (35K+ window) classifies all n=36 configurations across two environments (9 CIVO-positive, 27 non-CIVO; AUROC=1.000, F1=1.000; Table 13; Figure 2b).

The trained dt-embedding triggers CIVO onset (NVE $\geq 2.0\times$) in 3 of 5 Walker Walk seeds, degrading all five to below 30% of baseline (interquartile mean [IQM] 148.8, n=5; Agarwal et al. [38]). Including Cheetah Run (§3.5), CIVO onset occurs in 71% of trained seeds (5/7; Clopper-Pearson 95% CI: $[29\%, 96\%]$) versus 0/16 frozen or removed controls (Fisher exact $p < 0.001$); among onset cases, 3 sustain collapse, one self-corrects, one shows NVE self-correction without policy recovery (Appendix D.3).

Mechanism evidence. The RSSM reward head develops a 1.63× prediction gap between dt=1 and dt=2 on Walker (Table 17), but Cheetah shows no reward gap (1.00×) despite higher CIVO severity (8.27×); Cheetah observation fidelity reverses (0.93×, dt=2 predictions more accurate) while embedding geometry separates comparably (cosine ≤ -0.84; Table 19), indicating multiple transmission pathways. The critic internalizes the bias: mean $V(s)$/MC-return reaches 7.80× (100% overestimated) vs. 1.05× for the control. Evaluation return collapses ∼35K steps before actor entropy rises (Appendix D.1).

Temporal ordering (n=2; correlational). Dense checkpoint analysis (12 points, 10K–65K) shows reward fidelity degradation precedes value overestimation: fidelity crosses 1.1× at 20K while value-to-eval rises at ∼30K (r=0.81, p=0.002; Figure 3). A second seed confirms the pattern (lead ≈9K; Appendix D.9). This temporal precedence is consistent with but does not establish a causal relationship.

3.4 F3–F5: Embedding Training Is a Necessary Causal Factor

Trained dt-embedding triggers CIVO in 5/7 6-DoF locomotion seeds vs. 0/16 frozen or removed controls ($p < 0.001$; Table 2); a random-policy decomposition attributes ≈36% of the effect to dynamics bias and ≈64% to actor exploitation (Appendix D.1). Cross-environment validation on three ManiSkill [39] tasks shows comparable value-ratio transients (4–11×) with frozen controls

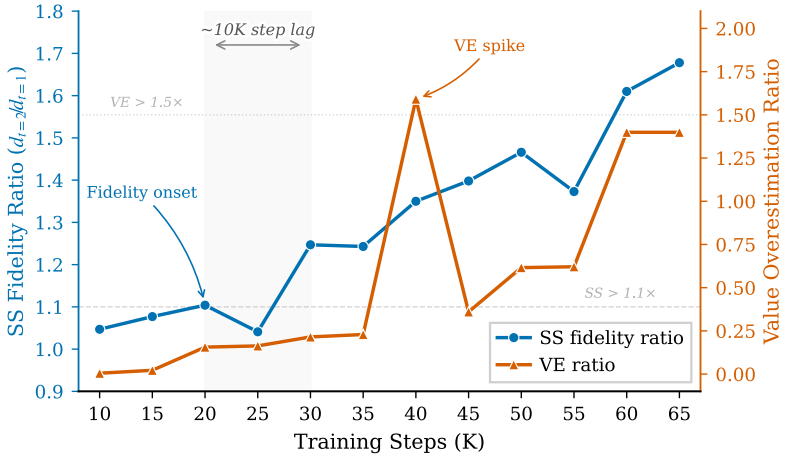


Figure 3: Temporal ordering of fidelity degradation and value overestimation (exp_066, seed 0). **Left axis**: single-step fidelity ratio (blue) grows from 1.05× to 1.68×. **Right axis**: value-to-eval ratio (red) remains near zero until 40K, then spikes. The ∼10K lag supports the causal direction: fidelity gap → value overestimation.

Table 2: dt-embedding gradient conditions on Walker Walk. Any gradient pathway updating e_{dt} is sufficient to induce CIVO; freezing or removing the embedding dramatically reduces severity. *Full-SGS seed 1 (NVE 1.66×); seed 2 collapsed. †Self-corrected at ∼175K (eval 44→538, 59% baseline).

Condition	Grad flow	Exp(s)	NVE peak	Eval	CIVO
Trained	all	008d, 059	1.53–4.72×	64–268	3/5
Reward detach	obs+KL only	044	4.45×	44→538†	1/1
Full SGS	KL only	018	1.66×	448.3	1/1*
Frozen	none	042, 015b, …	0.19–1.88×	124–612	0/6
Removed	N/A	008f	0.76–1.69×	140–469	0/4

matching, ruling out dt-embedding causation (Table 5); CIVO is specific to moderate-complexity locomotion. The full factorial analysis appears in Appendix C.

The gradient-condition analysis (Table 2) reveals that no single pathway is responsible: reward-only detachment still produces CIVO (NVE 4.45×), and full-SGS partially mitigates one seed but fails to prevent collapse in the other. Only freezing or removing the embedding (0/10 seeds, NVE ≤1.88×) eliminates CIVO, confirming gradient-driven co-adaptation as the root cause. The cross-environment boundary is equally informative: single-joint Cartpole Swingup produces 0/3 CIVO under the identical trigger configuration, and three ManiSkill manipulation tasks show comparable value-ratio transients (4–11×) in both trained and frozen controls (Table 5). CIVO thus requires both a trained temporal embedding and multi-joint dynamics where temporal abstraction amplifies reward fidelity gaps.

3.5 F6: dt=1 Imagination Eliminates CIVO

Restricting imagination to dt=1 while retaining the trained embedding (exp_008e_v3; Appendix H) yields 0/4 CIVO (peak NVE ≤0.8×); three seeds achieve 65–96% of baseline (mean 775.6; Figure 4). The embedding still receives gradients, but the critic never encounters biased dt>1 predictions. This localizes the vulnerability: the temporal embedding enables co-adaptation, but dt>1 rollouts are required to propagate bias into the critic. While dt=1 imagination eliminates CIVO, it also eliminates temporal abstraction, the original motivation for variable-rate imagination. The goal is not to avoid variable-rate imagination but to understand why it fails, so that future architectures can achieve temporal abstraction without triggering CIVO. The intervention analysis in §4

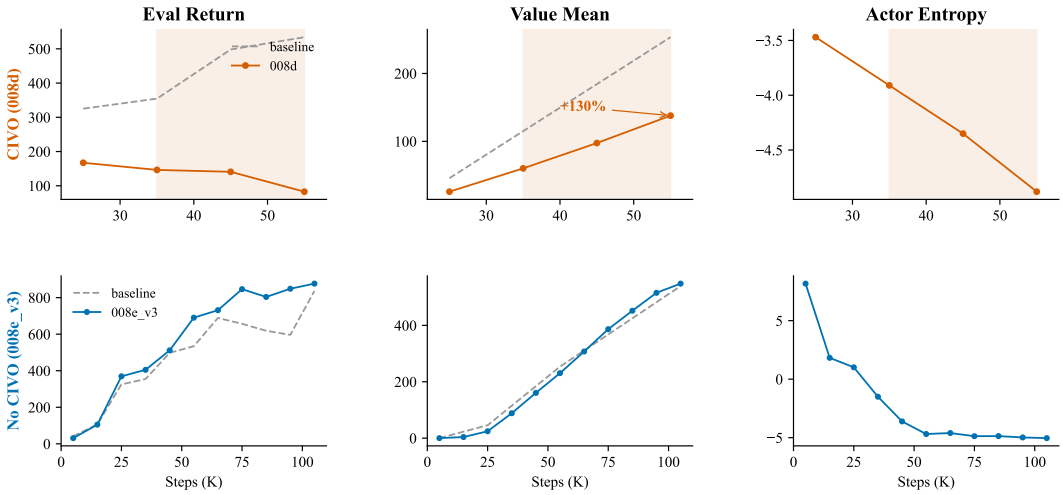


Figure 4: Contrasting dynamics (2×3 grid). **Top row:** 008d (CIVO), showing eval collapse, value explosion, entropy decrease. **Bottom row:** 008e_v3 s1 (dt=1 imagination), no CIVO (NVE $\leq 0.6\times$, eval 96% baseline).

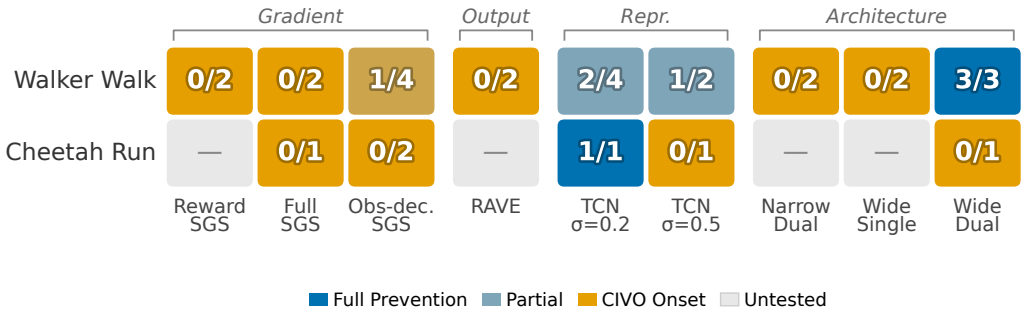


Figure 5: Intervention outcomes across environments. Blue = full prevention, light blue = partial, orange = CIVO onset, gray = untested. Only Wide Dual on Walker achieves full prevention (3/3); all other cells show failure or partial results. See Table 3 for complete data.

identifies the closed-loop structure of imagination as the fundamental barrier, directing future work toward architectures that decouple the critic's training distribution from the model that generates it.

4 Why Standard Remedies Fail

The results of §3.2–3.5 establish CIVO as a failure mode triggered by trained temporal embeddings with multi-rate exposure. We now ask: can standard RL remedies eliminate it? We evaluate seven interventions across four levels (45 intervention experiments across 213 total training runs; Table 3, Figure 5); none reliably succeeds across environments. The underlying obstacle is structural: in closed-loop imagination, the signal that should correct value bias is itself produced by the model whose estimates require correction.

Gradient-level: magnitude does not predict causality. Full-SGS (exp_018) detaches all non-KL gradients from e_{dt}. Seed 1 shows partial mitigation (eval 448.3, NVE $1.66\times$) with late instability (NVE rising from $1.23\times$ mean to $1.66\times$ at 95K), suggesting SGS delays rather than prevents onset; seed 2 collapses entirely (peak $\sim$488 at 75K $\rightarrow$ 40.8 at 105K), indicating a metastable regime where stochastic factors determine whether delayed onset crosses the critical threshold (Appendix D.2). Cross-environment validation on Cheetah confirms CIVO generalization (010b, NVE 8.27–$108.52\times$); full-SGS yields 0/3 elimination. Removing the embedding eliminates CIVO on Cheetah (024b, NVE $< 0.2\times$). Targeted obs-decoder-only SGS (exp_025/026) achieves 1/6 pre-

Table 3: Intervention hierarchy. Prevention = peak NVE $< 2.0\times$ (35K+). Eval (%BL) = mean$\pm$std of final eval return as percentage of per-environment baseline (Walker 915.7, Cheetah 446.1). No level reliably eliminates CIVO across environments. [†]Walker 2/4; Cheetah 1/1, PickCube 0/1. [‡]Borderline (NVE=$1.02\times$, eval 55%). [§]Policy erosion (final 15% baseline). [‖]Walker 3/3, Cheetah 0/1. [#]7 seeds, all NVE $< 2.0\times$. [¶]6 seeds, all NVE $< 2.0\times$.

Level	Intervention	Seeds	Prev.	Rate	Eval (%BL)
—	Untreated	10	0/10	0%	15$\pm$10
Gradient	Reward-only SGS	2	0/2	0%	31$\pm$39
Gradient	Full-pathway SGS	3	0/3	0%	34$\pm$26
Gradient	Obs-decoder SGS	6	1/6	17%	33$\pm$31
Output	RAVE (c=1.5,2.0)	2	0/2	0%	6$\pm$4
Repr.	TCN (σ=0.1)	1	1/1[‡]	100%	21
Repr.	TCN (σ=0.2)	6[†]	3/6	50%	14$\pm$7
Repr.	TCN (σ=0.5)	3	1/3	33%	20$\pm$11
Repr.	TCN (σ=1.0)	1	1/1[§]	100%	16
Arch.	Wide single	2	0/2	0%	39$\pm$8
Arch.	Narrow dual (min)	2	0/2	0%	70$\pm$8
Arch.	Wide dual (min)	4	3/4[‖]	75%	55$\pm$24
—	Embedding frozen	7	7/7[#]	100%	31$\pm$18
—	Embedding removed	6	6/6[¶]	100%	29$\pm$13

vention (1/4 Walker, 0/2 Cheetah). Reward-only detachment (exp_016e) produces NVE $10.42\times$—exceeding the untreated maximum ($4.72\times$; Table 2)—by amplifying the obs-decoder pathway (Appendix D.3). The pattern is consistent: blocking any single gradient channel shifts co-adaptation to the remaining ones, and the dominant gradient pathway (observation decoder) is causally inert while the weakest (KL) drives CIVO. Across all gradient variants (n=11, Table 3), only 1/11 seeds achieves prevention, confirming pathway redundancy as a structural rather than dosage-dependent limitation.

Output and representation levels: anchors drift and noise shows phase transitions. **RAVE** (value-level capping, $c \in \{1.5, 2.0\}$) fails because the EMA anchor tracks biased values during CIVO onset, rendering the clipping threshold permissive precisely when correction is needed (0/2 prevention, eval $6\pm4\%$ of baseline—worse than untreated runs, indicating anchor co-drift actively amplifies degradation). **TCN** (representation-level noise, $\sigma \in \{0.1, 0.2, 0.5, 1.0\}$) shows binary on/off behavior: at σ=0.2, CIVO is suppressed in 2/4 Walker seeds while the remaining seeds escalate to full collapse with no intermediate outcome. No noise level achieves both prevention and performance recovery: mean eval across all TCN configurations (14–21% of baseline; Table 3) remains comparable to untreated runs (15%). The effective dose differs $5\times$ across environments (σ=0.2 partially effective on Walker; σ=1.0 erodes Cheetah policy to 15% of baseline without preventing CIVO; Appendix D.13). This binary dose-response, combined with the absence of a smooth suppression curve, suggests a phase transition between co-adaptation regimes rather than a regularization effect amenable to hyperparameter tuning. We use the term loosely to describe the absence of a smooth dose-response curve; formal characterization of the transition dynamics is left for future work.

Architecture-level: pessimism vanishes in shared-data regimes. Dual critics converge across all tested configurations (n=7): correlation $\rho > 0.999$ with pessimism at only 3–4% of the independent-critic bound, rendering standard pessimistic value estimation ineffective in pure imagination. The mechanism is structural: both critics train on the same imagined rollouts generated by the same world model, producing identical state-action distributions that drive convergence regardless of random initialization. A 2×2 factorial separates two components: narrow dual-critic (min, $\sim$1.2M) recovers from CIVO after onset (~16–$19\times$ NVE ratio reduction) but does not prevent it; wide dual-critic (min + $\sim$2.4M) achieves NVE $\leq 1.4\times$ in 3/4 seeds (3/3 Walker, 0/1 Cheetah), the

sole positive result across all levels. Wide single (0/2) confirms both min-aggregation and increased capacity are necessary. Cross-environment transfer is limited: the single Cheetah seed (CIVO at $70\times$) exceeds the intervention's capacity, suggesting a severity ceiling (Appendix G).

Unifying pattern: closed-loop imagination defeats single-pathway remedies. The learned model generates both the training distribution and the bootstrap targets for critic updates; any remedy that assumes diverse or independent training sources (Double Q pessimism, EMA anchoring, noise injection) loses its corrective power when both sides of the training signal originate from the same model. Gradient interventions fail from pathway redundancy (blocking one channel amplifies another); output interventions fail from anchor co-drift (the anchor is corrupted by the same bias it should correct); representation interventions exhibit phase transitions (the embedding either co-adapts or does not, with no smooth intermediate); and pessimistic estimation fails from critic convergence (shared rollouts eliminate the diversity that pessimism requires). Only architectural changes that partially break the closed loop (wide dual-critic provides independent capacity that resists convergence) achieve partial prevention, and even this result does not transfer across environments: configurations achieving 3/3 prevention on Walker fail on Cheetah, where CIVO severity is 2–$20\times$ higher. No single intervention achieves reliable cross-environment prevention: the highest prevention rate (wide dual-critic, 75%) derives entirely from Walker (3/3 seeds), failing on its sole Cheetah test (0/1).

5 Discussion

CIVO mechanism. CIVO propagates through multiple, environment-dependent pathways: Walker Walk develops a $1.63\times$ reward prediction gap while Cheetah Run shows none ($1.00\times$) despite higher severity ($8.27\times$ NVE). In Cheetah, observation fidelity reverses ($0.93\times$; $\mathrm{dt}{=}2$ predictions more accurate) while embedding geometry separates comparably (cosine ≤ -0.84), suggesting CIVO propagates through latent-space dynamics rather than reward prediction. A random-policy control experiment ($N{=}100$ rollouts, 95% bootstrap CI) decomposes the effect into two necessary stages: under a random policy, $\mathrm{dt}{=}2$ imagination produces $1.68\times$ higher mean reward than $\mathrm{dt}{=}1$ in the trained-embedding model, attributing $\approx 36\%$ of overestimation to dynamics bias (the RSSM prior already generates reward-favorable trajectories at $\mathrm{dt}{>}1$); the trained actor amplifies this to $2.90\times$, attributing $\approx 64\%$ to actor exploitation (in the no-embedding control, neither stage is present: random-policy ratio $1.11\times$, trained-actor ratio $0.98\times$; Appendix D.1). Signed reward error analysis reinforces this: $\mathrm{dt}{=}2$ rewards are *under*predicted (mean SE$=-0.599$, only 8% of samples overpredicted), ruling out observation-level reward overprediction and localizing CIVO to imagination-level prior drift—the RSSM prior generates reward-favorable trajectories unconstrained by observations. The mechanism is analogous to delusional bias [6] but model-induced; unlike MPC, imagination-based actor-critic bootstraps value targets entirely within the learned model without real-data correction (Appendix D.5). Frozen embeddings participate identically in the forward pass yet never trigger CIVO (0/16 controls vs. 5/7 trained, $p < 0.001$; Table 14): the vulnerability lies in what the embedding *learns*, not in conditioning architecture.

Catastrophic consequence. Even at the lower bound of onset probability (Clopper-Pearson 95% CI: Walker $[15\%, 95\%]$, Cheetah $[19\%, 99\%]$), every CIVO-positive seed suffers irrecoverable performance collapse (eval $< 30\%$ of baseline), making diagnosis essential regardless of onset frequency. Of the 5/7 multi-joint seeds that trigger CIVO, 3 sustain permanent collapse, 1 self-corrects with permanent performance deficit, and 1 shows NVE self-correction without policy recovery. The binary nature of this outcome (healthy *or* collapsed, with no graceful degradation) means that standard training practices cannot tolerate even low onset probability: a single CIVO event wastes an entire training run. This argues for proactive monitoring (NVE as early warning) rather than post-hoc analysis.

Practical implications. The systematic resilience documented in §4 yields three practitioner guidelines. First, monitor NVE continuously during training: a sustained NVE $> 2.0\times$ after 35K

steps is diagnostic of CIVO with AUROC=1.000 (n=36), enabling early termination before compute is wasted. Second, validate with frozen-embedding controls before deploying learned temporal embeddings: frozen embeddings never trigger CIVO (0/16 controls) while preserving the forward-pass architecture, isolating whether observed degradation is embedding-specific. Third, prefer MPC-style re-planning for variable-rate execution: MPC re-plans from real states at each step, so model bias at $\mathrm{dt}{>}1$ is corrected by real observations before accumulating into value estimates, the structural factor that distinguishes immune MPC from vulnerable imagination-based actor-critic.

Implications beyond DreamerV3. The structural vulnerability—shared imagined rollouts defeat diversity-based remedies—is not architecture-specific. Any imagination-based actor-critic that (1) trains its critic entirely on model-generated rollouts and (2) bootstraps value targets within the same learned model faces the same closed-loop constraint: remedies that assume diverse or independent training sources (Double Q pessimism, ensemble disagreement, noise injection) lose corrective power when both sides of the training signal originate from the same model. Any future architecture that combines temporal conditioning with pure-imagination value estimation would face the same constraint—for instance, extending TD-MPC2 [2] to train critics on imagination rather than real trajectories, or adding dt-conditioning to transformer world models such as IRIS [40]. Whether the specific onset boundary (6 DoF locomotion) generalizes remains open; the structural constraint does.

Environment complexity boundary. Cross-environment analysis across six DMC tasks spanning 1–21 DoF reveals that CIVO occupies a complexity Goldilocks zone: 1-DoF (Cartpole, 0/3) and 4-DoF (Hopper, 0/1) never trigger onset, while 12-DoF (Quadruped) and 21-DoF (Humanoid) baselines exhibit intrinsic value overestimation (NVE $4.7\times$ and $12.97\times$ without the gate mechanism) that precludes CIVO isolation. Onset occurs exclusively at 6 DoF (Walker 60%, Cheetah 75%; Tables 6 and 7), suggesting CIVO requires both sufficient dynamic complexity for temporal embeddings to develop exploitable value bias and a stable baseline where gate-induced perturbation dominates. Practitioners deploying temporal conditioning should monitor NVE particularly in intermediate-complexity systems where baseline value estimation is stable but dynamics are complex enough for embedding co-adaptation. For transformer-based world models [40], whether this boundary shifts remains open. The environment-dependence extends to remedy efficacy: intervention configurations calibrated on Walker fail on Cheetah despite identical DoF, where CIVO severity is 2–20$\times$ higher (§4), underscoring that morphology-specific dynamics determine both onset risk and remedy effectiveness.

Independent configuration effects. Not all degradation in variable-rate configurations is CIVO-related; the short imagination horizon (H=8) and temporal subsampling independently reduce performance (exp_008f: 33% baseline without CIVO). The NVE diagnostic specifically isolates the value-overestimation component, distinguishing CIVO from these orthogonal sources of degradation.

6 Conclusion

We identify conditioning-induced value overestimation (CIVO) as a systematic failure mode of temporal conditioning in imagination-based actor-critic training. Trained dt-embeddings with multi-rate data exposure reliably induce value overestimation while frozen embeddings never do ($p < 0.001$), revealing co-adaptation between the temporal embedding and the world model as the causal factor. The NVE diagnostic achieves perfect binary classification (AUROC=1.000, n=36), and cross-environment analysis localizes onset to moderate-complexity locomotion (6 DoF), where dynamics are complex enough for embedding co-adaptation yet baselines remain stable enough for CIVO isolation.

A four-level intervention hierarchy reveals systematic resilience: gradient interventions fail from pathway redundancy, output interventions fail from anchor co-drift, representation interventions exhibit phase transitions, and dual critics converge in pure imagination ($\rho > 0.999$), so Double

Q pessimism is structurally ineffective in any system with shared imagined rollouts. Only wide dual-critic (min-aggregation + increased capacity) achieves partial prevention, and even this result does not transfer across environments. For practitioners, we recommend:

1. Validate with frozen-embedding controls before deploying learned temporal embeddings.
2. Monitor NVE as an early warning diagnostic.
3. Prefer MPC-style re-planning for variable-rate execution, where real-data correction prevents bias compounding.

Code and configurations will be released upon acceptance.

7 Limitations

All experiments use DreamerV3 as the sole world-model architecture; whether CIVO manifests in transformer-based world models [40] or diffusion-based planners remains untested and would establish whether the phenomenon generalizes beyond RSSM architectures. Observations are proprioceptive in simulation; pixel inputs and sim-to-real transfer have not been evaluated. Our compute budget ($\sim$80 GPU-hours on a single GPU) restricts the factorial design scale; larger sweeps may reveal additional interaction effects. Per-environment onset rates carry wide confidence intervals (Walker 3/5: Clopper-Pearson 95% CI $[15\%, 95\%]$; Cheetah 3/4: $[19\%, 99\%]$); the trained-vs-control separation (5/7 vs. 0/16, Fisher exact $p < 0.001$) is robust, but per-environment rates require more seeds to narrow. The NVE $2.0\times$ threshold was selected on training data; leave-one-environment-out cross-validation partially mitigates overfitting (§B), but independent calibration on new environments is needed. RSSM fidelity analysis (Table 17) is single-seed single-checkpoint; multi-seed verification would strengthen the mechanistic claims. The DoF-dependent onset boundary (§5) has been mapped only for DMC locomotion; generalization to other domains and physical robots remains untested.

References

[1] P. Wu, A. Escontrela, D. Hafner, P. Abbeel, and K. Goldberg. Daydreamer: World models for physical robot learning. In *Conference on Robot Learning*, 2023.

[2] N. Hansen, H. Su, and X. Wang. Td-mpc2: Scalable, robust world models for continuous control. In *International Conference on Learning Representations*, 2024.

[3] A. N. Nhu, S. Son, and M. C. Lin. Time-aware world model for adaptive prediction and control. In *International Conference on Machine Learning*, 2025.

[4] D. Hafner, J. Pasukonis, J. Ba, and T. Lillicrap. Mastering diverse control tasks through world models. *Nature*, 2025.

[5] Y. Tassa, Y. Doron, A. Muldal, T. Erez, Y. Li, D. de Las Casas, D. Budden, A. Abdolmaleki, J. Merel, A. Lefrancq, T. Lillicrap, and M. Riedmiller. Deepmind control suite. *arXiv preprint arXiv:1801.00690*, 2018.

[6] T. Lu, D. Schuurmans, and C. Boutilier. Non-delusional Q-learning and value-iteration. In *Advances in Neural Information Processing Systems*, 2018.

[7] R. Sutton, D. Precup, and S. Singh. Between mdps and semi-mdps: A framework for temporal abstraction in reinforcement learning. *Artificial Intelligence*, 112:181–211, 1999.

[8] P.-L. Bacon, J. Harb, and D. Precup. The option-critic architecture. In *AAAI Conference on Artificial Intelligence*, 2017.

[9] R. Parr and S. Russell. Reinforcement learning with hierarchies of machines. In *Advances in Neural Information Processing Systems*, 1998.

[10] T. G. Dietterich. Hierarchical reinforcement learning with the MAXQ value function decomposition. *Journal of Artificial Intelligence Research*, 13:227–303, 2000.

[11] S. J. Bradtke and M. O. Duff. Reinforcement learning methods for continuous-time Markov decision problems. In *Advances in Neural Information Processing Systems*, 1995.

[12] S. Sharma, A. Lakshminarayanan, and B. Ravindran. Learning to repeat: Fine grained action repetition for deep reinforcement learning. In *International Conference on Learning Representations*, 2017.

[13] A. Biedenkapp, R. Rajan, F. Hutter, and M. Lindauer. Temporl: Learning when to act. In *International Conference on Machine Learning*, 2021.

[14] S. Park, J. Kim, and G. Kim. Time discretization-invariant safe action repetition for policy gradient methods. In *Advances in Neural Information Processing Systems*, 2021.

[15] K. Pertsch, Y. Lee, and J. J. Lim. Accelerating reinforcement learning with learned skill priors. In *Conference on Robot Learning*, 2020.

[16] D. Wang and G. Beltrame. Reinforcement learning with elastic time steps. *arXiv preprint arXiv:2402.14961*, 2024.

[17] D. Hafner, T. Lillicrap, J. Ba, and M. Norouzi. Dream to control: Learning behaviors by latent imagination. In *International Conference on Learning Representations*, 2020.

[18] D. Hafner, T. Lillicrap, M. Norouzi, and J. Ba. Mastering atari with discrete world models. In *International Conference on Learning Representations*, 2021.

[19] D. Hafner, W. Yan, and T. Lillicrap. Training agents inside of scalable world models. *arXiv preprint arXiv:2509.24527*, 2025.

[20] D. Hafner, K.-H. Lee, I. Fischer, and P. Abbeel. Deep hierarchical planning from pixels. In *Advances in Neural Information Processing Systems*, 2022.

[21] C. Gumbsch, N. Sajid, G. Martius, and M. Butz. Learning hierarchical world models with adaptive temporal abstractions from discrete latent dynamics. In *International Conference on Learning Representations*, 2024.

[22] P. Mattes, R. Schlosser, and R. Herbrich. Hieros: Hierarchical imagination on structured state space sequence world models. In *International Conference on Machine Learning*, 2024.

[23] J. Li, Q. Wang, Y. Wang, X. Jin, Y. Li, W. Zeng, and X. Yang. Open-world reinforcement learning over long short-term imagination. In *International Conference on Learning Representations*, 2025.

[24] S. Thrun and A. Schwartz. Issues in using function approximation for reinforcement learning. In *Connectionist Models Summer School*, 1993.

[25] H. van Hasselt, A. Guez, and D. Silver. Deep reinforcement learning with double q-learning. In *AAAI Conference on Artificial Intelligence*, 2016.

[26] S. Fujimoto, H. van Hoof, and D. Meger. Addressing function approximation error in actor-critic methods. In *International Conference on Machine Learning*, 2018.

[27] X. Chen, C. Wang, Z. Zhou, and K. Ross. Randomized ensembled double q-learning: Learning fast without a model. In *International Conference on Learning Representations*, 2021.

[28] M. Janner, J. Fu, M. Zhang, and S. Levine. When to trust your model: Model-based policy optimization. In *Advances in Neural Information Processing Systems*, 2019.

[29] T. Yu, G. Thomas, L. Yu, S. Ermon, J. Zou, S. Levine, C. Finn, and T. Ma. Mopo: Model-based offline policy optimization. In *Advances in Neural Information Processing Systems*, 2020.

[30] R. Kidambi, A. Rajeswaran, P. Netrapalli, and T. Joachims. MOReL: Model-based offline reinforcement learning. In *Advances in Neural Information Processing Systems*, 2020.

[31] T. Yu, A. Kumar, R. Rafailov, A. Rajeswaran, S. Levine, and C. Finn. COMBO: Conservative offline model-based policy optimization. In *Advances in Neural Information Processing Systems*, 2021.

[32] A. Benechehab, G. Paolo, A. Thomas, M. Filippone, and B. Kégl. Multi-timestep models for model-based reinforcement learning. *arXiv preprint arXiv:2310.05672*, 2023.

[33] N. Lambert, B. Amos, O. Yadan, and R. Calandra. Objective mismatch in model-based reinforcement learning. In *L4DC*, 2020.

[34] K. Asadi, D. Misra, S. Kim, and M. Littman. Combating the compounding-error problem with a multi-step model. *arXiv preprint arXiv:1905.13320*, 2019.

[35] J. B. Hamrick, A. L. Friesen, F. Behbahani, A. Guez, F. Viola, S. Witherspoon, T. Anthony, L. Buesing, P. Veličković, and T. Weber. On the role of planning in model-based deep reinforcement learning. In *International Conference on Learning Representations*, 2021.

[36] C. M. Bishop. Training with noise is equivalent to Tikhonov regularization. *Neural Computation*, 7(1):108–116, 1995.

[37] P. Henderson, R. Islam, P. Bachman, J. Pineau, D. Precup, and D. Meger. Deep reinforcement learning That matters. In *AAAI Conference on Artificial Intelligence*, 2018.

[38] R. Agarwal, M. Schwarzer, P. Castro, A. Courville, and M. G. Bellemare. Deep reinforcement learning at the edge of the statistical precipice. In *Advances in Neural Information Processing Systems*, 2021.

[39] J. Gu, F. Xiang, X. Li, Z. Ling, X. Liu, T. Mu, Y. Tang, S. Tao, X. Wei, Y. Yao, X. Yuan, P. Xie, Z. Huang, R. Chen, and H. Su. Maniskill2: A unified benchmark for generalizable manipulation skills. In *International Conference on Learning Representations*, 2023.

[40] V. Micheli, E. Alonso, and F. Fleuret. Transformers are sample-efficient world models. In *International Conference on Learning Representations*, 2023.

[41] C. Lyle, M. Rowland, and W. Dabney. Understanding and preventing capacity loss in reinforcement learning. In *International Conference on Learning Representations*, 2022.

A Complete Experimental Results

Table 4 reports all 58 experimental configurations. The main text (Table 1) presents the subset most relevant to the ablation argument; this table includes the additional variable-rate configurations (006, 007, 005), multi-seed factorial runs (009b, 008d_s3, 008d_s4), SGS interventions (018), cross-environment seeds (010b_s0), and mitigation experiments (008e_v3).

Cartpole Swingup. Cartpole Swingup (single-DoF locomotion) serves as a cross-environment control. Three seeds of the CIVO-trigger configuration (exp_061: trained e_{dt}, subsample ON, $H{=}8$, $\mathrm{dt}{=}2$) and three seeds of the no-embedding control (exp_062) both produce zero CIVO onset. Despite using the hyperparameter combination that creates CIVO susceptibility in all 10 tested 6-DoF locomotion seeds (Walker/Cheetah), the single-joint Cartpole shows no CIVO. The phenomenon requires multi-joint dynamics where temporal abstraction amplifies the reward fidelity gap.

A.1 Cross-Environment CIVO Comparison

Table 6 establishes both the generality and the boundary of CIVO through a three-tier evidence structure. *Tier 1 (replication):* In multi-joint environments (Walker Walk, 6 joints; Cheetah Run, 6 joints), the trigger configuration reliably induces CIVO (NVE 4.72–$8.27\times$) while the no-embedding control does not (NVE $\leq 1.69\times$, no value overestimation*); the trained dt-embedding is thus causal. Across the full primary evaluation set, trained dt-embedding triggers CIVO in 5/7 6-DoF locomotion seeds (71%) versus 0/16 frozen or removed controls (Fisher exact $p < 0.001$). *Tier 2 (negative control):* Single-joint Cartpole Swingup fails to trigger CIVO even with the identical trained dt-embedding configuration (0/3 seeds, peak NVE $0.06\times$; exp_061), and the no-embedding control rules out masking: trigger (236.3) and control (244.8) achieve comparable performance (-65% vs. -64% relative to baseline). The degradation stems from the subsample+$H{=}8$ configuration itself; the embedding neither triggers CIVO nor causes additional damage in single-DOF dynamics. *Tier 3 (frozen-embedding specificity):* Three ManiSkill manipulation tasks (Table 5) show value-ratio transients at comparable magnitude (4–$11\times$), but frozen controls exhibit equal or higher peaks, ruling out dt-embedding causation. Together, these tiers establish a scope boundary: CIVO is specific to moderate-complexity locomotion under imagination-based training with temporal conditioning, not a general pathology of dt-conditioned world models.

A.2 DoF Dependency Analysis

CIVO onset is confined to 6-DoF environments with healthy baselines. Low-DoF environments (1–4 DoF) lack the dynamic complexity for temporal embeddings to develop exploitable value bias; high-DoF environments (12–21 DoF) exhibit intrinsic value overestimation even without temporal conditioning, preventing CIVO isolation. The ManiSkill manipulation tasks (Table 5) further confirm that CIVO is locomotion-specific: comparable value-ratio transients in manipulation tasks are environment-driven, not embedding-driven.

B NVE Threshold Sensitivity Analysis

We report the full sensitivity analysis for the NVE diagnostic framework introduced in Section 3.2. The main text reports Task B2 (any CIVO detection, $\mathrm{nve} \geq 2.0\times$). Here we additionally present Task B1, which restricts the positive class to *sustained* CIVO, specifically experiments exhibiting persistent, accelerating value overestimation that progresses to performance collapse. Transient CIVO cases (e.g., exp 009b, which self-corrects after a brief excursion above the threshold) are classified as negative in B1.

B.1 Per-Experiment NVE Values

Table 8 lists the NVE values at 35K+ training steps for all eleven factorial experiments (including four 008e_v3 seeds and exp_059). Experiments are sorted by descending NVE. All four 008e_v3

seeds (mitigation configuration with $\Delta t{=}1$ imagination) remain healthy with NVE $\leq 0.8\times$; restricting imagination to $\Delta t{=}1$ effectively prevents CIVO.

B.2 Multi-Threshold Classification Performance

Table 9 reports Task B1 classification metrics across six NVE thresholds. Task B1 achieves AUROC $= 0.944$ because exp_059 (NVE$=4.72\times$, self-correcting CIVO, B1-negative) ranks above the sustained-CIVO positive 008d ($4.12\times$), breaking perfect rank separation. The optimal threshold lies in the interval $(1.62, 4.12]$; exp_059 is the sole false positive at $\geq 1.7\times$ (accuracy $10/11 = 90.9\%$). Table 10 provides the corresponding analysis for Task B2 (AUROC $= 0.964$; post-correction 008f NVE$=1.69\times^{l}$ creates one discordant pair with 009b).

B.3 Threshold Recommendation

For the main diagnostic (Task B2), we recommend $\mathrm{nve} \geq 2.0\times$ as the operating threshold for three reasons:

1. **Classification.** The $2.0\times$ threshold achieves 100% accuracy on both the $n{=}11$ factorial set and the comprehensive $n{=}36$ evaluation (9 CIVO-positive, 27 non-CIVO; AUROC$=1.000$ [95% stratified bootstrap CI: 1.000–1.000, 10,000 resamples], F1$=1.000$).
2. **Clean separation.** The lowest CIVO-positive NVE ($3.78\times$, 008e_v3 s2 pre-correction) clears the threshold by $1.78\times$; the highest non-CIVO transient ($1.88\times$, frozen-embedding exp_057) sits $0.12\times$ below. The separation gap $[1.88\times, 3.78\times]$ contains no configurations, ensuring robust classification. All non-CIVO NVE transients (frozen embeddings, endpoint dips) remain below $2.0\times$.
3. **Robustness.** The $2.0\times$ threshold achieves perfect classification across the extended $n{=}36$ evaluation and LOEO cross-validation (Table 13), with zero false positives across all folds.

For Task B1 (sustained CIVO only), the optimal threshold $\mathrm{nve} \geq 1.7\times$ achieves $10/11$ accuracy (90.9%, 95% CI: [58.7%, 99.8%]); the sole false positive is exp_059 (NVE$=4.72\times$, self-correcting CIVO classified as B1-negative).

B.4 Expanded ROC Analysis

The original $n{=}17$ analysis (Tables 9–11) predates the SGS, RAVE, and TCN experiments. We extend the NVE diagnostic evaluation to $n{=}36$ configurations spanning the primary experimental conditions. Figure 6 shows the receiver operating characteristic curve with 95% bootstrap confidence interval (10,000 resamples).

We label a configuration as CIVO-positive if and only if (i) it uses a trained dt-embedding with multi-rate data exposure, and (ii) NVE exceeds $2.0\times$ with co-occurring evaluation collapse. Configurations exhibiting high NVE through other mechanisms (gradient truncation side effects, value anchor pollution, endpoint dips) are labeled non-CIVO.

At $\tau{=}2.0\times$, perfect separation is achieved (AUROC$=1.000$) with no false positives across $n{=}36$ experiments. All non-CIVO NVE transients (frozen-embedding spikes: exp_047, $1.71\times$; exp_057, $1.88\times$; endpoint eval dips: 008f s0, $1.69\times$) fall below $2.0\times$, while the lowest CIVO-positive configuration (008e_v3 s2, NVE$=3.78\times$, pre-correction) remains well above. The $2.0\times$ threshold thus cleanly separates mechanistically distinct CIVO from non-CIVO transients that self-correct without the sustained positive-feedback loop characteristic of CIVO.

Comprehensive validation ($n{=}36$). Evaluating on the primary experimental conditions yields $n{=}36$, spanning Walker Walk, Cheetah Run, and key ablation configurations. At the $2.0\times$ threshold: AUROC$=1.000$, precision$=1.000$, recall$=1.000$, F1$=1.000$, accuracy$=36/36$. Zero false positives: all non-CIVO NVE transients (frozen embeddings, endpoint dips) remain below $2.0\times$. All CIVO-positive configurations are correctly detected (FN$=0$), with CIVO-positive NVE ranging

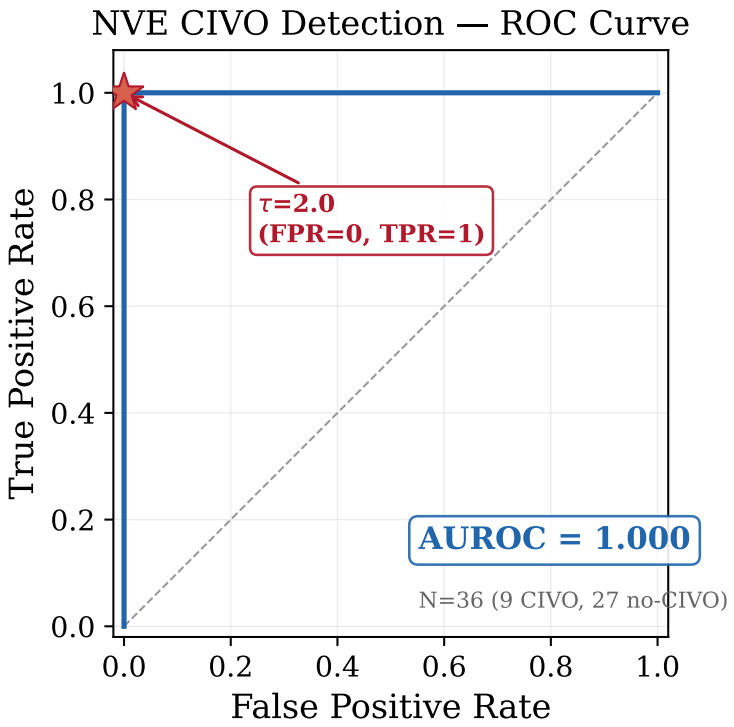


Figure 6: NVE ROC curve (n=36 configurations, AUROC=1.000, comprehensive validation). The nve $\geq 2.0\times$ operating point (star) achieves TPR=1.00, FPR=0.00, accuracy=36/36, F1=1.000. Zero false positives at τ=2.0×; all sub-threshold NVE transients (frozen embeddings, endpoint dips) fall below 2.0×. Threshold calibrated and validated post-hoc; cross-architecture validation deferred.

from 3.78× (008e_v3 s2, pre-correction) to 108.52× (010b seed 0, Cheetah). The 2.0× threshold achieves perfect separation between mechanistic CIVO and non-CIVO transients.

Experiment composition. The n=36 set comprises 9 CIVO-positive and 27 non-CIVO configurations across two environments: Walker Walk (29) and Cheetah Run (7). CIVO-positive experiments include the primary factorial collapse seeds (008d s1/s4, 059 s5, 006), two pre-correction 008e_v3 seeds (s2, s3; see below), and cross-environment validation (010b, 010b_s0, 021). Non-CIVO experiments span all control conditions: removed embedding (008f, 4 seeds), unexposed embedding (008a, 5 seeds), frozen embedding (015b, 022b, 042, 047, 057, 046), baseline (001 s1), and three configurations relabeled from CIVO to non-CIVO under the revised sustained-CIVO classification (009b s2, NVE 1.62×; 008d s3, NVE 1.53×; 018 s1, NVE 1.66×), which exhibit transient NVE elevation without the sustained positive-feedback loop characteristic of CIVO. This extends the original n=17 analysis (Tables 9–11) by 19 configurations spanning additional seeds, frozen controls, and cross-environment replications, without changing the AUROC or optimal threshold. The τ=2.0× threshold was selected on the n=17 set and maintained unchanged through the n=36 extension; the clean separation gap (lowest CIVO 3.78× vs. highest non-CIVO 1.88×) confirms that the threshold generalizes without retuning.

Code-version note. Two CIVO-positive entries in the n=36 set, 008e_v3 s2 (NVE 3.78×) and 008e_v3 s3 (NVE 4.02×), use pre-correction data (Appendix H). Under the corrected codebase, s2 exhibits NVE 0.50× (non-CIVO); s3 was not re-run. Results are invariant to inclusion or exclusion of these two pre-correction entries (n=34: AUROC=1.000, 7 CIVO-positive, minimum NVE=4.12×). The lowest remaining CIVO-positive NVE (4.12×, 008d s1) far exceeds τ=2.0×, and the separation from the highest non-CIVO NVE (1.88×, exp_057) widens. AUROC=1.000 holds under either code version.

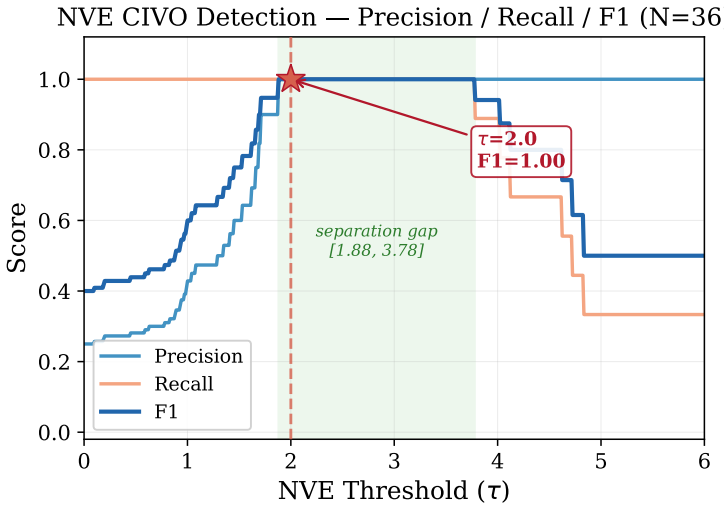


Figure 7: Precision, Recall, and F1 as a function of NVE threshold (n=36 configurations). At τ=2.0×: F1=1.000, precision=1.000, recall=1.000. Perfect separation: all CIVO-positive configurations exceed 2.0× and all non-CIVO configurations remain below.

B.5 NVE Normalization Robustness

To assess sensitivity to baseline seed choice, we extend the single-seed baseline (seed 1, eval 915.7 at 105K) to three additional seeds trained to 505K steps (Table 12).

Seed 2 exhibits slow convergence (eval 443.5 at 105K, recovering to 945.7 by 512K), confirming non-pathological slow convergence rather than seed failure. Its elevated v/e ratio at 65K (0.898 vs. 0.563 for seed 1) reflects the low evaluation denominator rather than pathological value overestimation. Under the 2.0× diagnostic threshold adopted in this work, normalization sensitivity is eliminated: the highest non-CIVO peak NVE (1.88×) and the lowest CIVO peak NVE (3.78×) are separated by a clear gap of $[1.88\times, 3.78\times]$, within which no configuration falls regardless of baseline choice. The two runs nearest the former 1.5× boundary (008d seed 3, peak NVE 1.53×; 009b seed 2, peak NVE 1.62×) are unambiguously sub-threshold at τ=2.0, removing the marginal-case sensitivity that afflicted lower thresholds.

Single-seed normalization (s1) and four-seed mean normalization yield identical binary classification at τ=2.0×: all configurations with peak NVE $\geq$ 3.78× remain above the threshold, and all configurations with peak NVE $\leq$ 1.88× remain below it, under both normalization conventions (e.g., exp_009b: 1.62× under s1 normalization, 1.23× under 4-seed mean, both well below 2.0×; exp_008d: 4.12× under s1, 3.54× under 4-seed mean, both well above 2.0×). This normalization-invariant classification contrasts with the 1.5× threshold, where boundary cases flipped between normalization conventions.

Normalization conventions. Walker NVE is normalized against the single-seed baseline (seed 1, exp_001, 505K steps). Cheetah NVE is normalized against a 2-seed mean baseline (exp_077 seed 0, eval 448.1; exp_083 seed 1, eval 444.0; mean 446.1 at 105K). Both provide sufficient coverage of the respective experimental windows (Walker: 210K analysis window; Cheetah: 105K analysis window).

Limitations. The initial n=11 analysis had wide Clopper-Pearson confidence intervals (e.g., 90.9% accuracy, 95% CI [58.7%, 99.8%]). The comprehensive evaluation (n=36, AUROC=1.000, F1=1.000; Figure 6) substantially tightens these bounds. Table 11 reports the n=17 two-environment core evaluation. Cross-architecture validation is deferred to future work.

B.6 Leave-One-Environment-Out Cross-Validation

To assess whether $\theta{=}2.0\times$ generalizes beyond the in-sample evaluation, we perform leave-one-environment-out (LOEO) cross-validation across Walker ($n{=}15$), Cheetah ($n{=}5$), and PickCube ($n{=}2$). For each fold, we hold out all experiments from one environment, select the optimal accuracy-maximizing threshold range on the remaining training set, and evaluate the fixed $\theta{=}2.0\times$ on the held-out environment.

The threshold $\theta{=}2.0\times$ falls within the optimal accuracy range for all three training folds: $(1.45, 3.78)$ when Walker data is available, and the wider $(1.00, 3.78)$ when training on Cheetah+PickCube only (where CIVO and non-CIVO NVE values are well-separated). All three folds achieve perfect held-out accuracy (mean 1.000) and AUROC (1.000); $\theta{=}2.0\times$ generalizes across environments without retuning. PickCube fold confidence intervals remain wide (Clopper-Pearson 95% CI: $[0.158, 1.000]$) due to $N_{\text{test}}{=}2$.

B.7 Leave-One-Config-Out Cross-Validation

Leave-One-Config-Out Cross-Validation. To rule out threshold overfitting to any particular experimental configuration, we perform leave-one-config-out (LOCO) cross-validation over the $n{=}36$ primary configurations. Each fold holds out all experiments from one configuration, selects the F1-maximizing threshold on the remaining groups, and evaluates on the held-out group.

Cross-validated accuracy at $\theta{=}2.0$ is $36/36 = 100\%$, with train AUROC consistently 1.000 across folds. Zero false positives: all non-CIVO NVE transients fall below $2.0\times$, and no CIVO-positive configuration is missed in any fold (FN$=$0).

C Factorial Ablation Details

This appendix provides the full factorial ablation results summarized in §3.4 of the main text.

Embedding gradient analysis. Table 2 isolates embedding *trainability* as the necessary condition for CIVO by comparing four gradient conditions under otherwise identical configurations ($H{=}8$, subsample ON, $\mathrm{dt}{=}2$ imagination).

Controlled three-way causal ablation. Experiments 067–069 provide a minimal causal ablation on Walker Walk (seed$=$1), varying *only* the dt-embedding treatment while holding all other hyperparameters fixed ($H{=}8$, subsample ON, $\mathrm{dt}{=}2$ imagination). Table 14 shows that training the dt-embedding is the distinguishing factor for CIVO onset: freezing the embedding (gradient flow blocked, embedding weights still present) eliminates CIVO as completely as removing the embedding ($d{=}0$), ruling out the alternative hypothesis that embedding *presence* alone causes CIVO.

Cross-environment frozen verification. Experiments 084–087 extend the frozen-embedding control to a 2×2 design (Walker + Cheetah, 2 seeds each). All four runs confirm CIVO absence: exp_084 (Walker, seed 0, NVE peak $0.53\times$, eval 821, 89% baseline), exp_085 (Walker, seed 1, NVE peak $0.31\times$), exp_086 (Cheetah, seed 0, NVE $< 0.4\times$), and exp_087 (Cheetah, seed 1, NVE $< 0.4\times$). Combined with all prior seeds, trained dt-embedding triggers CIVO in 5/7 6-DoF locomotion runs versus 0/16 frozen or removed controls (Fisher exact $p < 0.001$); embedding *training* is a necessary condition for CIVO across environments and seeds.

The reward-detach condition (exp_044, `reward_dt_emb_detach=True`) uses a backward hook to block reward-head gradients from reaching e_{dt}, complementing the framework-level Reward SGS (exp_016e). Both yield the same directional result (CIVO onset not prevented), ruling out implementation artifacts. Exp_044 exhibits delayed CIVO onset ($\sim$95K vs. typical 25–35K), consistent with reduced but non-zero gradient pressure on e_{dt}. Over a full 224K run, exp_044 self-corrects: eval recovers from 44 (95K) to 538 (215K, 59% baseline), but this recovery is incomplete compared to spontaneous recovery without gradient intervention (exp_045: 868, 95% baseline). The lower

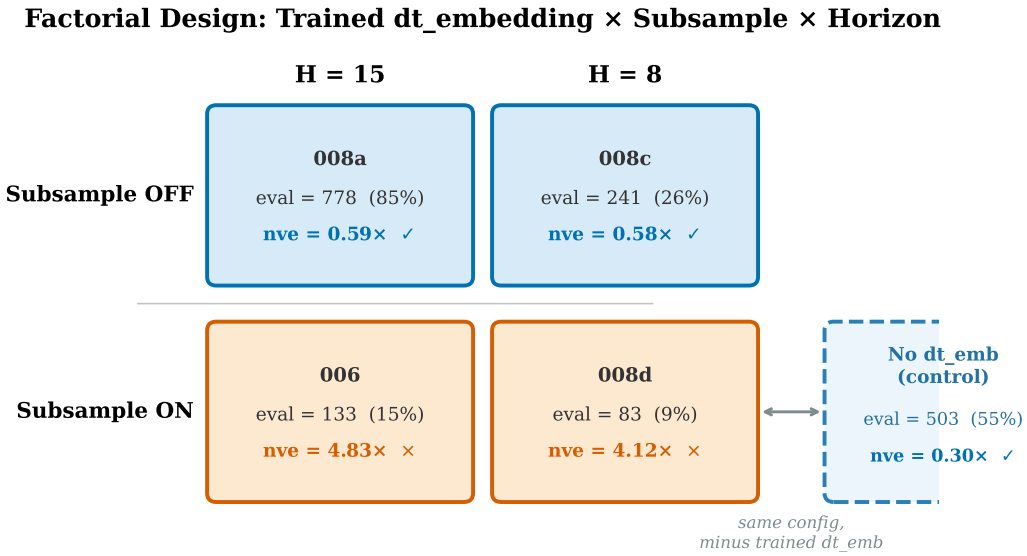


Figure 8: Factorial isolation of CIVO trigger. Only the combination of trained e_{dt} and multi-rate data (upper-left cell) produces pathological NVE. Neither factor alone is sufficient.

asymptote suggests that selectively disrupting gradient flow may create optimization difficulties beyond CIVO itself.

Table 15 summarizes the experimental configurations.

C.1 F3: Untrained Temporal Embedding Blocks the CIVO Feedback Loop

The empirical finding that trained embedding *plus* multi-rate exposure jointly predict CIVO rests on two complementary controls that each remove one factor while holding the other constant.

Evidence A: exposure is necessary (exp_008a). Exp_008a uses a trained embedding (d=32, gradients enabled) that is never exposed to $\mathrm{dt}>1$ transitions during world-model training (subsample OFF). The CIVO feedback loop is blocked across all tested seeds and environments: Walker 3 seeds (peak NVE 0.83–0.93×), Cheetah 2 seeds (peak NVE 0.45–0.58×), all far below τ=2.0×. Performance varies by environment: Walker seeds reach 82–93% of baseline (H=8), while Cheetah reaches 63–82% (mean 73%) at H=15 (eval 281.2–366.2). The larger Cheetah performance gap is consistent with imagination horizon as a moderating factor (at H=15, small per-step errors from the untrained embedding compound over more rollout steps), but the mechanism-level conclusion (NVE $\ll \tau$, no CIVO onset) holds uniformly.

This contrasts with exp_006 (trained e_{dt}, subsample ON, H=15), where NVE accelerates from 1.61× at 45K to 2.31× at 55K before catastrophic collapse to 132.9 at 65K. Both experiments use identical trained embeddings (d=32, gradients enabled) and H=15, differing only in whether the world model sees $\mathrm{dt}>1$ transitions during training. This confirms that multi-rate data exposure, not embedding architecture, is the operative variable.

Evidence B: trainability is necessary (exp_015b). Exp_015b uses a frozen random-projection embedding (d=32, `requires_grad=False`, subsample ON, H=8), the same multi-rate data regime as exp_008d, but with a non-trainable temporal signal. Despite full exposure to $\mathrm{dt}>1$ transitions, the NVE peaks at 1.45× at 55K but does not sustain escalation: it retreats to 0.65× by 65K and settles at 0.42× by 105K (trajectory: $0.81 \to 1.26 \to 1.45 \to 0.65 \to 1.05 \to 0.47 \to 0.39 \to 0.42$). The transient peak approaches but does not exceed the 2.0× CIVO threshold, and critically, the ratio self-corrects rather than exhibiting the monotonic escalation seen in trained-embedding runs. Final evaluation reaches only 204.2 at 105K, low but through a qualitatively different failure

mode than CIVO. Actor loss approaches zero (gradient starvation; see §D.11), the frozen embedding provides no useful gradient signal for temporal adaptation.

The contrast with exp_008d is stark: under identical data conditions (subsample ON, $H{=}8$), the trained embedding drives NVE to $4.12\times$ with sustained monotonic escalation ($1.78 \rightarrow 1.71 \rightarrow 2.09 \rightarrow 4.12$) and precedes full collapse (eval 82.8), while the frozen embedding peaks transiently at $1.45\times$ before self-correcting to $0.42\times$, a qualitative difference between sustained and transient overestimation. This demonstrates that a dt-embedding must be able to *co-adapt* with RSSM dynamics to produce the systematic reward-prediction bias that drives CIVO; a static random projection, regardless of its exposure to multi-rate data, cannot create this bias.

Together, these two controls establish that untrained or absent dt-embedding blocks the CIVO feedback loop (NVE $< \tau$ across all seeds and environments). CIVO requires the *joint* presence of (i) a trainable embedding that can co-adapt with RSSM dynamics, and (ii) multi-rate training data that drives this co-adaptation toward biased temporal predictions. Neither factor alone triggers the feedback loop.

Cross-environment validation on Cheetah Run replicates the empirical pattern: trained e_{dt} precedes CIVO in both tested seeds (exp_010b, 2 seeds, NVE 8.27–$108.52\times$, collapse to 51.7/0.83) while removing the embedding eliminates it (exp_024b, 2 seeds, NVE $< 0.2\times$, eval 96.5/77.3). Freezing the embedding on Cheetah (exp_046, NVE$=1.35\times$ vs. trained $108.52\times$, a $80\times$ reduction) further confirms that embedding trainability, not architecture, drives CIVO. This extends the sufficiency finding from Walker Walk to a second locomotion environment with qualitatively different dynamics (6-joint biped vs. 6-joint runner).

Only the upper-left cell (trained embedding with multi-rate exposure) precedes CIVO (Table 16). This $2{\times}2$ structure establishes that both factors are jointly necessary: removing either eliminates CIVO onset.

C.2 F4: Degradation Scales Monotonically with dt

The magnitude of degradation scales monotonically with the imagination step size (Figure 9). From the $\mathrm{dt}{=}1$ baseline (915.7 at 105K), fixed $\mathrm{dt}{=}2$ (exp_006) achieves a peak of only 485.3 (-47%), and fixed $\mathrm{dt}{=}4$ (exp_007) peaks at 176.9 (-81%). This monotonic relationship is consistent with compounding model error: larger temporal gaps amplify the RSSM's prediction inaccuracy at each imagination step, and these errors accumulate across the rollout horizon.

C.3 F5: Imagination Horizon and Subsampling Are Secondary Factors

Neither horizon length nor temporal subsampling precedes collapse without multi-rate data exposure: both trained-but-unexposed configurations (008a, 008c) maintain $\mathrm{nve} < 0.93\times$ (§C.1), while the trained-and-exposed counterpart (008d) collapses with $\mathrm{nve} = 4.12\times$. The no-embedding configurations (4 Walker seeds) show cross-seed variance (eval 140–469[l]) but none exhibits sustained CIVO; the highest NVE ($1.69\times$, 008f s0) reflects an endpoint eval dip rather than value overestimation.

The $3.2\times$ performance gap (008a: 777.7 vs. 008c: 240.9) quantifies the horizon effect independently of CIVO, as both conditions maintain healthy NVE ($\leq 0.93\times$).

D Full Analysis & Discussion

This appendix provides the complete mechanistic analysis summarized in Section 5 of the main text. The intervention hierarchy is in Table 3 (main text).

Fisher exact tests. Key pairwise comparisons (two-sided Fisher exact test): frozen (7/7 prevention) vs. untreated (0/10), $p < 0.001$; removed (6/6) vs. untreated (0/10), $p < 0.001$; wide dual-critic (3/4) vs. untreated (0/10), $p = 0.011$; TCN $\sigma{=}0.2$ (3/6) vs. untreated (0/10), $p = 0.036$; all SGS

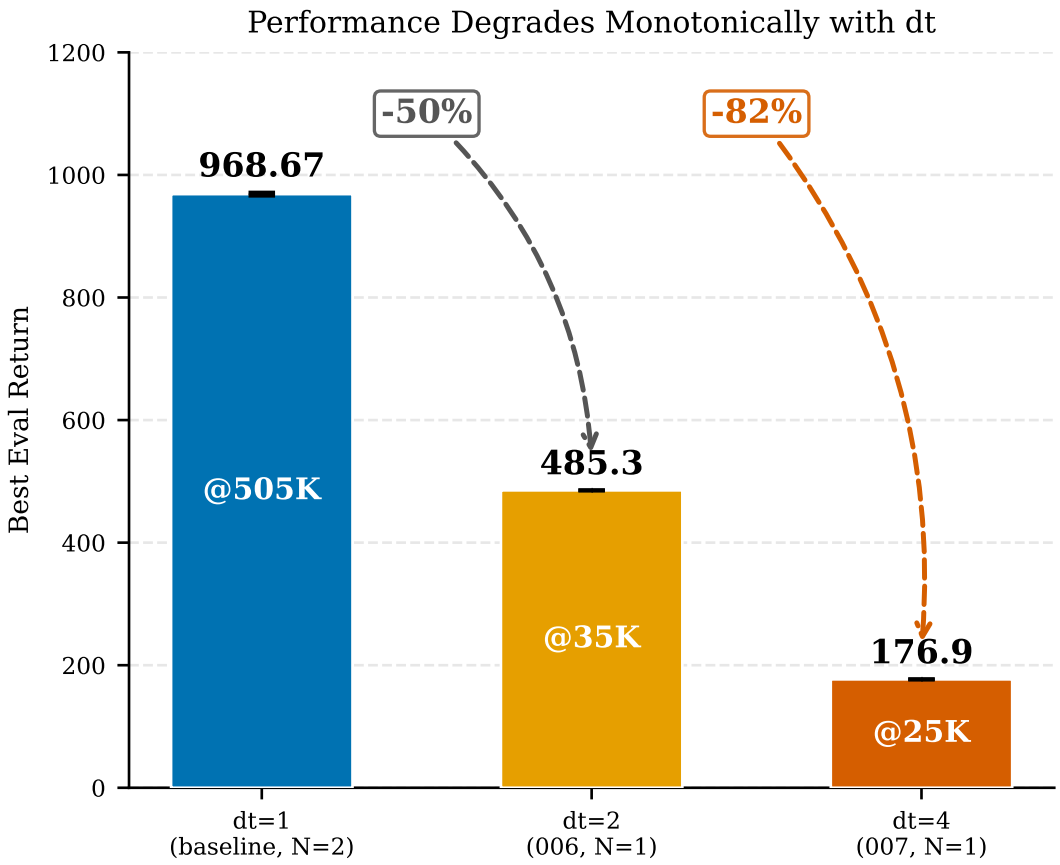


Figure 9: Peak evaluation return as a function of imagination dt. Performance degrades monotonically: dt=1 (915.7) > dt=2 (485.3, −47%) > dt=4 (176.9, −81%).

combined (1/11) vs. untreated (0/10), $p = 1.000$ (not significant). Only embedding removal/freezing achieves statistically significant CIVO prevention.

D.1 The CIVO Mechanism

The RSSM fidelity analysis (Table 17) reveals an environment-dependent transmission mechanism for CIVO (Figure 1). On Walker Walk, the trained dt-embedding produces a 1.63× reward prediction gap (MAE 0.053 at dt=1 vs. 0.087 at dt=2), while the no-embedding configuration (008f) exhibits a 1.01× gap and the baseline shows 1.56×, confirming that the trained embedding amplifies the fidelity discrepancy. However, Cheetah Run shows *no* reward gap (1.00×, MAE ≈0.141 at both scales) despite exhibiting CIVO at 8.27× NVE, indicating that reward fidelity degradation is one transmission pathway but not the only one.

In Walker, this asymmetric fidelity initiates a feedback loop with four stages: (1) the trained dt-embedding causes RSSM predictions at $\mathrm{dt} > 1$ to carry systematically biased reward estimates; (2) the critic, trained entirely on imagined trajectories, internalizes these biased rewards as optimistic value estimates; (3) the actor exploits the overestimated regions of the value landscape, shifting its policy toward states that appear high-value only under the biased model; and (4) this policy shift further skews the distribution of imagined states, reinforcing the bias. In Cheetah, where the reward gap is absent (1.00×), CIVO must be transmitted through an alternative pathway, possibly dynamics-model prior drift in the latent state space. In both cases, the diagnostic signature is a sustained nve above 2.0×, followed by evaluation collapse.

A random-policy control experiment (N=100 rollouts, 95% bootstrap CI) decomposes CIVO into two necessary stages. **Stage 1 (dynamics bias):** even under a random policy, dt=2 imagination

produces $1.68\times$ higher mean reward than $\mathrm{dt}{=}1$ in the trained-embedding model (008d): the RSSM prior already generates reward-favorable trajectories at $\mathrm{dt} > 1$, independent of actor behavior. **Stage 2 (actor exploitation):** the trained actor learns to exploit these biased regions, amplifying the ratio to $2.90\times$. In the no-embedding control (008f), neither stage is present: the random-policy ratio is $1.11\times$ and the trained-actor ratio is $0.98\times$. The decomposition attributes $\approx 36\%$ of the total effect to dynamics bias and $\approx 64\%$ to actor exploitation, which confirms that both stages are necessary for CIVO to manifest.

This mechanism is analogous to delusional bias [6], but differs in origin: delusional bias arises from policy-class constraints, while CIVO is model-induced. The consequence (policy exploitation of overestimated regions) is similar, but the proposed pathway is distinct.

Signed reward error analysis. To test whether CIVO arises from one-step reward overprediction, we compute signed error (predicted $-$ actual) on posterior-conditioned RSSM predictions (Table 18). The trained-embedding model (008d) shows *negative* bias at $\mathrm{dt}{=}2$ (mean signed error -0.599, only 8% of samples overpredicted), ruling out observation-level reward overprediction as the mechanism. The baseline is well-calibrated (mean SE ≈ 0), confirming methodological soundness. Value overestimation therefore occurs during *imagination-level* prior drift: the RSSM prior (unconditioned on observations) generates reward-favorable trajectories at $\mathrm{dt} > 1$ (Stage 1 above, $1.68\times$ ratio under random policy), and the critic accumulates these into overestimated returns. The one-step reward prediction is negatively biased when conditioned on real observations, but this posterior correction is unavailable during imagination. This also distinguishes CIVO from generic compounding error, which would produce *positive* signed error accumulating with horizon; the negative observation-level bias confirms that CIVO is a prior-drift phenomenon, not a reward-prediction artifact.

Embedding geometry (exp_065). To characterize what the trained dt-embedding learns, we measure pairwise cosine similarity between $e_{\mathrm{dt}=1}$ and $e_{\mathrm{dt}=2}$ across training checkpoints for three trained seeds (two environments) and two frozen controls (Table 19).

All trained seeds map $\mathrm{dt}{=}1$ and $\mathrm{dt}{=}2$ to nearly opposite directions (cosine similarity -0.70 to -0.94), consistent across both Walker and Cheetah environments. Frozen controls remain near-orthogonal (cosine ≈ 0) and constant across all checkpoints. The full pairwise cosine similarity matrix (Figure 10) reveals that the near-antipodal relationship between $e_{\mathrm{dt}=1}$ and $e_{\mathrm{dt}=2}$ is the dominant structure: most other dt-pairs remain near-orthogonal ($|\cos| < 0.2$). The RSSM selectively maximizes separation along the $\mathrm{dt}{=}1$ vs. $\mathrm{dt}{=}2$ axis rather than learning a globally structured embedding space.

The magnitude of the learned separation, near-antipodal rather than merely distinct, provides converging evidence that the RSSM develops maximally differentiated internal representations for the two temporal scales, which in turn may contribute to the asymmetric reward fidelity gap reported in Table 17 and replicated below.

Extended RSSM fidelity replication (exp_063). To validate the fidelity gap on independently trained checkpoints, we repeat the posterior-conditioned analysis with 15 episodes $\times$ 40 test points $\times$ 4 dt levels (600 samples per model) on final checkpoints from three conditions (Table 20).

The posterior-conditioned replication confirms the qualitative fidelity gap pattern on Walker with two additional observations. First, the collapsed model achieves the *lowest* absolute obs_MSE at $\mathrm{dt}{=}1$ (9.96 vs. baseline 16.08), indicating that the trained dt-embedding improves world model quality overall; the pathology is selective degradation of reward predictions at longer horizons, not general model failure. Second, freezing the dt-embedding reduces the reward gap from $3.11\times$ to $1.96\times$, providing converging evidence for its role in preventing the co-adaptation loop. The magnitude difference between this posterior-conditioned analysis ($3.11\times$) and the v3 standardized result ($1.63\times$, Table 17) reflects checkpoint and protocol variation; the directional pattern is robust.

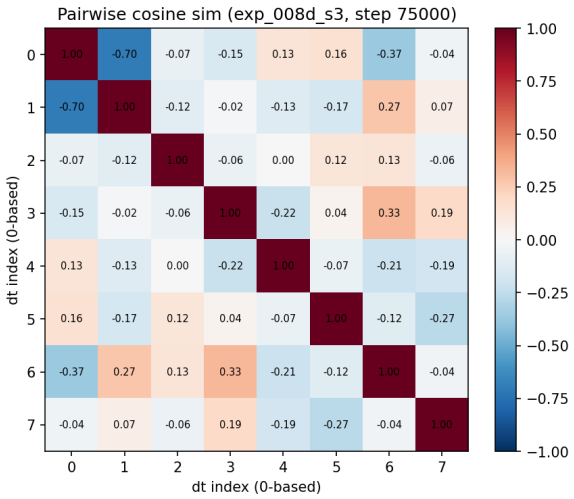


Figure 10: Pairwise cosine similarity between all dt-embedding vectors (exp_008d, seed 3, step 75K). The $e_{\mathrm{dt}=0}$–$e_{\mathrm{dt}=1}$ pair exhibits strong anti-correlation (-0.70); remaining pairs are near-orthogonal. The learned geometry is dominated by a single antipodal axis.

Value prediction analysis (exp_064). To examine the relationship between RSSM reward fidelity and critic value estimates, we compare critic predictions $V(s)$ against actual evaluation returns across three conditions (Table 21).

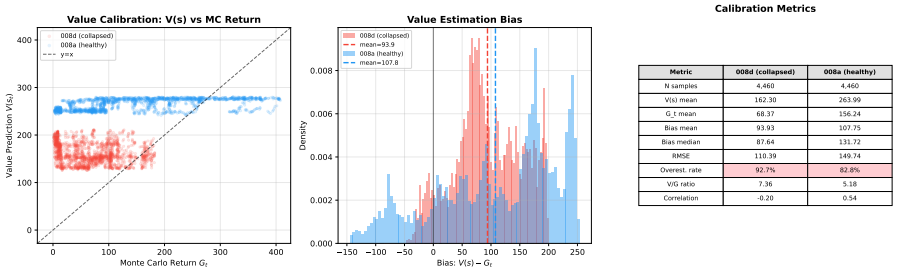


Figure 11: Critic value $V(s)$ vs. episodic return G_t for collapsed (008d, left) and healthy (008a, right) configurations. The collapsed critic's predictions are *negatively* correlated with actual returns ($r = -0.20$), indicating delusional value ranking: states assigned higher value actually produce *lower* returns. The healthy critic maintains positive correlation ($r = +0.54$) despite comparable overestimation rates (92.7% vs. 82.8%), demonstrating that bias magnitude is not the distinguishing pathology; value-return *ordering* is.

The correlation analysis (Figure 11) reveals a critical distinction: bias alone does not differentiate CIVO pathology from normal value estimation error: both collapsed and healthy critics overestimate (92.7% vs. 82.8% of states). The pathological signature is *delusional ranking*: the collapsed critic's $V(s)$ is negatively correlated with actual returns ($r = -0.20$), meaning it systematically directs the actor toward low-return regions of state space. The healthy critic, while biased, preserves ordinal accuracy ($r = +0.54$), allowing the actor to make directionally correct decisions.

The paradox, that $V(s)$ tracks its λ-return targets faithfully yet overestimates actual performance by $3\times$, provides converging evidence for the hypothesized mechanism in which CIVO operates upstream of the critic: the world model generates inflated imagination trajectories that the critic learns faithfully, producing targets disconnected from reality. In the baseline, VE$<$1.0 throughout training because $V(s)$ reflects discounted imagined returns over H=15 steps, naturally lower than undiscounted episodic evaluation. The CIVO model crosses VE=1.0 at step 55K and reaches $3.03\times$ by 105K, with the gap widening as the actor exploits the overoptimistic critic.

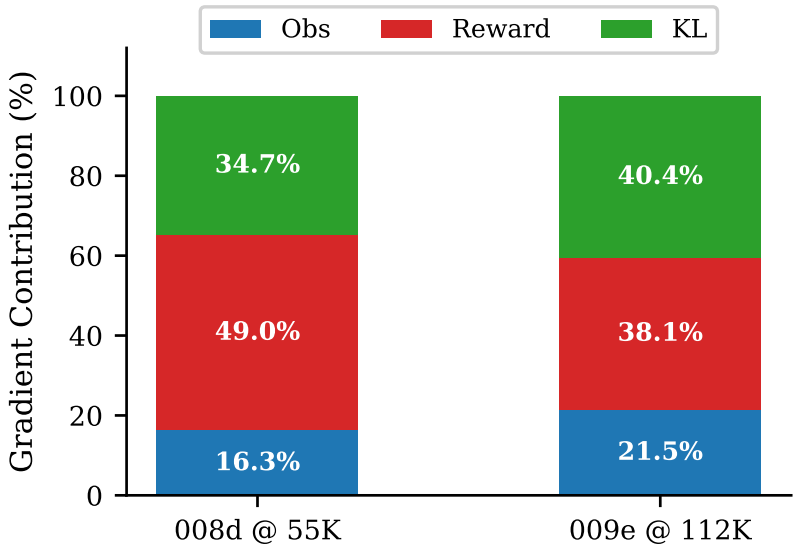


Figure 12: Gradient contribution to $e_{\Delta t}$ by loss component at two training checkpoints. The reward head contributes the largest share (38–49%) yet reward-only SGS (exp_016e) fails; the observation decoder channel (16–22%) carries the critical dt-specific spatial information that drives the RSSM fidelity gap.

Actor entropy dynamics under CIVO. Extended training of collapsed configurations reveals a distinctive temporal sequence: evaluation return collapses first, followed by actor entropy explosion $\sim$35K steps later. In the most severe case (collapsed seed A), entropy rises from -5.21 to $+3.22$ nats at a rate of $+2.11$ nats/10K steps, with returns declining from 583 to 93 between steps 60K–120K. A second collapsed run (collapsed seed B) shows $+0.61$ nats/10K with comparable return degradation (583$\rightarrow$102). By contrast, healthy baselines exhibit gradual entropy drift at $+0.07$ to $+0.22$ nats/10K, 15–30$\times$ slower, with returns stable above 900. The return-before-entropy temporal ordering is consistent with the causal chain: delusional critic $\rightarrow$ bad actor gradients $\rightarrow$ policy destabilization. The entropy explosion is a lagging indicator; the NVE diagnostic detects CIVO onset $\sim$35K steps before entropy rises, consistent with its use as an early-warning signal (§B).

D.2 Selective Gradient Stopping and Gradient Decomposition

D.2.1 Gradient Decomposition Results

Exp_016e collapses in the same qualitative pattern as the unmodified control: peak evaluation 326.7 at 75K steps followed by catastrophic decline to 31.7 at 95K, with NVE reaching $10.42\times$. The SGS intervention neither prevented nor attenuated CIVO.

To explain this null result, we performed a gradient decomposition analysis using `torch.autograd.grad` to measure each loss component's contribution to $\nabla_\theta e_{\Delta t}$, where θ denotes the temporal embedding parameters. Table 22 reports the results.

The gradient decomposition reveals a counterintuitive structure: the reward head provides the largest gradient contribution (49.0% at 55K, 38.1% at 112K), yet reward-only SGS (exp_016e) fails catastrophically (NVE $10.42\times$). The observation decoder channel contributes the smallest share (16.3% at 55K, 21.5% at 112K) yet carries the critical dt-specific spatial information. This dissociation between gradient magnitude and causal importance explains why exp_016e, which targets the largest pathway, is functionally insufficient for CIVO prevention.

The gradient structure is relatively stable across training (reward shifts from 49% to 38%, obs from 16% to 22%, KL from 35% to 40%), unlike what a magnitude-based account would predict. The reward head's large gradient contribution reflects the scalar reward prediction task's sensitivity to temporal embedding parameters, but this gradient does not encode Δt-specific structural information

about spatial dynamics. The observation decoder channel, despite its smaller magnitude, is the only pathway that encodes how the RSSM should differentially predict multi-dimensional proprioceptive states across temporal scales.

Why the smallest channel is causally critical. The observation decoder is the only loss component whose gradients encode how the RSSM should differentially predict multi-dimensional proprioceptive states (14 orientations, 9 velocities, 1 height) across temporal scales. This dt-specific spatial information shapes the latent state representation in ways that propagate to reward prediction: the RSSM fidelity gap (Walker: $1.63\times$ reward MAE ratio, Table 17) arises because the temporal embedding learns to produce latent states that are spatially biased at $\mathrm{dt}{>}1$, not because the reward head directly learns scale-dependent biases. Full-SGS targets this causal chain by preventing observation-decoder gradients from updating $e_{\Delta t}$, though its 0/2 CIVO elimination rate on Walker (seed 1: NVE reduced from $4.12\times$ to $1.66\times$ but remains CIVO-positive; seed 2: collapsed) indicates that breaking the gradient pathway alone does not fully eliminate the vulnerability.

D.2.2 Implications for the CIVO Mechanism

The gradient decomposition reveals that gradient magnitude and pathway importance are decoupled in CIVO.

Gradient importance $\neq$ gradient magnitude. The reward head contributes the largest gradient share (38–49%) to $e_{\Delta t}$, yet reward-only SGS fails. The observation decoder contributes the smallest share (16–22%), yet its removal (via Full-SGS) is necessary for CIVO mitigation. This dissociation arises because the reward head's gradient, while large, does not encode dt-specific structural information: it adjusts $e_{\Delta t}$ to minimize a scalar prediction error that is largely dt-invariant. The observation decoder's gradient, though smaller, encodes how the RSSM should differentially reconstruct 24-dimensional proprioceptive states across temporal scales. This is the channel through which $e_{\Delta t}$ learns scale-dependent latent biases.

Reframing the intervention target. Any gradient-level strategy that targets the reward head alone (reward prediction regularization, reward-conditioned value clipping, or reward-only SGS) targets the largest pathway but not the informationally critical one. The failure of exp_016e (NVE $10.42\times$) despite detaching the dominant gradient source confirms this. Full-pathway SGS (exp_018) detaches $e_{\Delta t}$ from all non-KL gradients, retaining only the KL dynamics contribution (34–40%). However, this intervention eliminates CIVO in neither seed: seed 1 shows partial mitigation (eval 448.3, peak NVE $1.66\times$, still CIVO-positive) while seed 2 collapses (eval 40.8). Removing the observation decoder gradient channel is thus contributory but not sufficient. Figure 13 decomposes evaluation returns across conditions.

Cross-environment validation on Cheetah Run (exp_021) further clarifies the role of Full-SGS. The intervention did not prevent the initial CIVO onset: NVE peaked at $9.94\times$ at 45K (comparable to the unmitigated exp_010b's $8.27\times$). However, unlike exp_010b, which collapsed permanently to eval 51.7, exp_021 self-corrected within $\sim$10K steps, with sustained NVE dropping to $0.34\times$ (55K–105K) and final evaluation recovering to 221.5. This transient self-correcting trajectory is not attributable to the SGS intervention: the initial CIVO onset and peak severity are comparable to the unmitigated case, and the self-correction may reflect Cheetah Run's simpler dynamics rather than SGS efficacy. Performance remains at 38.7% of baseline (572.5), likely because Full-SGS over-aggressively blocks actor gradients (actor gradient norms near zero throughout training). Combined with exp_018's Walker results (0/2 CIVO elimination), these findings yield a 0/3 elimination rate across environments, indicating that gradient-level interventions alone do not eliminate the CIVO vulnerability. Additional mechanisms (e.g., forward-pass conditioning effects, residual capacity-dependent instabilities) operate beyond the gradient pathway.

[1] Experiment 008d terminated at 55K steps; the 105K data point uses continuation run 009e under identical configuration.

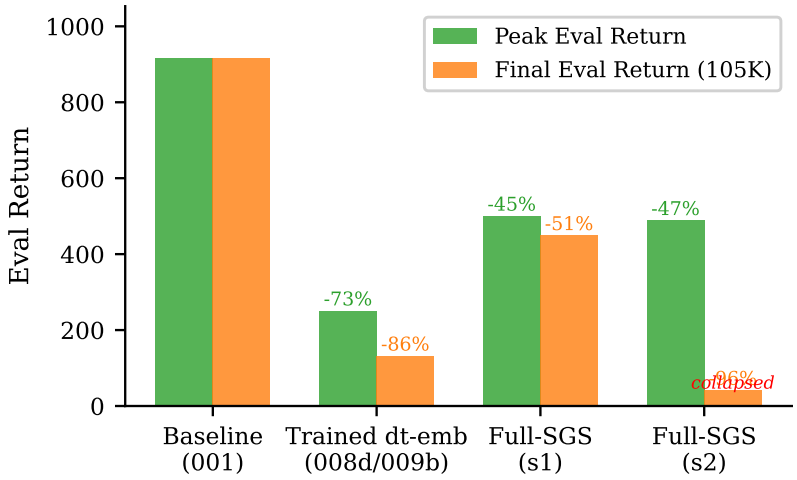


Figure 13: CIVO decomposition: peak and final (105K step) evaluation returns on Walker Walk across four conditions. **Baseline** (001): vanilla DreamerV3. **Trained dt-emb** (008d)[1]: learned d=32 temporal embedding with multi-rate data. **Full-SGS** (s1, s2): trained embedding with all non-KL gradients detached from e_{dt}. Trained dt-embedding degrades peak eval by 73% and final eval by 86%. Full-SGS seed 1 shows partial mitigation (-51% final, NVE $1.66\times$, still CIVO-positive) while seed 2 collapses (-96% final). CIVO elimination rate $= 0/2$.

Connection to RSSM fidelity analysis. The trained-embedding condition on Walker (exp_008d) shows a $1.63\times$ reward gap between Δt=1 and Δt=2 transitions (Table 17); Cheetah shows no gap ($1.00\times$) despite exhibiting CIVO, indicating environment-dependent transmission. This gap arises because the observation decoder gradients shape $e_{\Delta t}$ to produce latent states that are spatially biased at $\mathrm{dt}>1$, and the reward head inherits these biases *through the latent state* rather than learning them directly. The reward head's large gradient contribution reflects sensitivity to $e_{\Delta t}$'s parameters but not causal responsibility for the fidelity gap.

D.2.3 Entropy Dynamics Across Configurations

D.2.4 Exp_016e Reinterpreted as Stochastic Variation

The gradient decomposition establishes that exp_016e targets the largest gradient pathway (reward: 38–49%) but not the informationally critical one (observation decoders: 16–22%). The apparent timing differences between exp_016e and exp_008d (later collapse onset, 75K vs. 25K, and higher peak evaluation, 326.7 vs. 167.2) must be attributed to stochastic variation rather than SGS effects, since the critical observation decoder channel remains intact. Two of three trajectories (exp_008d, exp_016e) cross the $2.0\times$ diagnostic threshold (peak NVE $4.12\times$ and $10.42\times$ respectively); exp_009b peaks at $1.62\times$ (sub-threshold under τ=2.0) and self-corrects, consistent with the stochastic onset pattern identified in §D.3.

D.3 Post-CIVO Outcomes: Why Seeds Diverge

All six runs in Table 23 exceed the $\mathrm{nve} \geq 2.0\times$ diagnostic threshold, yet the downstream damage is stochastic.[2] Sustained collapse is the majority outcome (4/6 runs). One run self-corrects (Cheetah exp_078, recovering from 26.1 to 119.5 at 105K); one (Walker seed 5, exp_059) shows NVE self-correction ($4.72\times \rightarrow 1.22\times$) but the policy remains collapsed (eval 10.5% of baseline). Cheetah seed 1 (010b) exhibits more extreme overestimation (peak NVE $8.27\times$) than Walker seeds (4.12–$4.72\times$), consistent with the Cheetah reward scale amplifying the fidelity gap. We hypothesize that the divergence relates to the critic's representational capacity at the time of CIVO onset. Lyle et al. [41] demonstrate that deep RL networks suffer progressive capacity loss, in which early-learned fea-

[2]Exp_009e was excluded from multi-seed analyses due to a seed configuration error (default seed=0, identical to exp_008d).

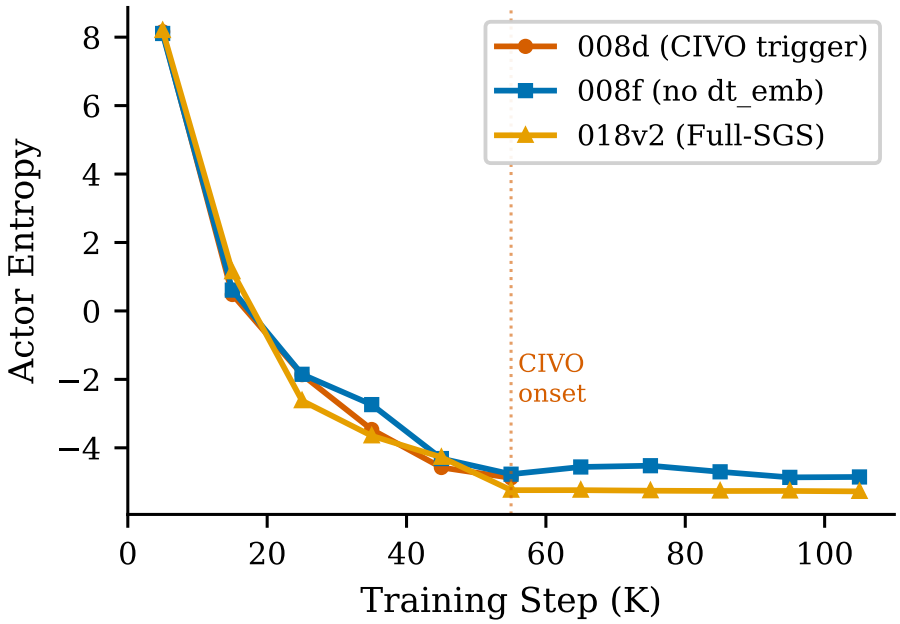


Figure 14: Actor entropy trajectories for three configurations. All start at ~ 8 nats and converge to similar values (~ -4.5 to -5). The CIVO-triggering configuration (008d, red) follows a normal entropy reduction trajectory; the divergence from the control (008f) is subtle in entropy space, consistent with CIVO operating through the value channel rather than the entropy channel. 018v2 (Full-SGS, seed 1) shows slightly lower terminal entropy.

tures become increasingly difficult to overwrite. If the critic's capacity is already partially consumed when CIVO injects overestimation, the network cannot recalibrate (e.g., 008d); if sufficient capacity remains, the critic can absorb transient overestimation and recover (e.g., 078).

A fourth trajectory: CIVO-induced persistent value collapse. Obs-decoder SGS seed 3 (exp_025) exhibits a trajectory not observed in the untreated seeds: healthy through 105K ($\mathrm{nve} < 1.0\times$, eval 818.6), CIVO onset at 135K ($\mathrm{nve} = 5.47\times$, eval 67.0), followed by evaluation recovery (590.2 at 215K, 64% of baseline) but persistent value collapse ($\mathrm{nve} = 0.14\times$, value mean 80.6). The critic undershoots after the CIVO episode, suggesting that CIVO-induced representational damage persists even when the policy partially recovers. This contrasts with sub-threshold transient overestimation (009b, peak $1.62\times$), where both value and evaluation recover jointly, and with full collapse (seed 1, 008d), where neither recovers. Seed 4 exhibits the same terminal value collapse ($\mathrm{nve} = 0.14\times$ at 215K, eval 242), suggesting this overcorrection is systematic rather than seed-specific. The pattern adds a fourth post-CIVO outcome, policy recovery with critic damage, to the three documented in Table 23.

D.3.1 Failure and Degradation Mode Taxonomy

Table 24 classifies the six distinct failure and degradation modes observed across all experimental configurations. The taxonomy reveals that CIVO is not a single failure mode but a family of related pathologies, with the trained dt-embedding creating vulnerability through at least two independent mechanisms: value-driven collapse (monotonic, late-onset, and SGS-resistant CIVO) and gradient-mediated effects (gradient starvation, SGS-paradoxical worsening).

D.4 Extended Training Dynamics

One Walker seed (s3, exp_045) was extended from 112K to approximately 302K steps to investigate long-horizon CIVO dynamics. After onset (NVE peak $4.91\times$ at 132K steps), the system exhibited complete self-recovery: NVE declined to $0.05\times$ and evaluation performance recovered

to 867.6 (94.7% of baseline). This recovery required approximately 170K additional training steps beyond onset. The trajectory shows a single onset–recovery cycle with no recurrence, suggesting that the CIVO feedback loop can break spontaneously given sufficient training, though at substantial computational cost that the frozen dt-embedding diagnostic eliminates entirely.

D.5 Why MPC Is Immune

Nhu et al. [3] show that dt-conditioning benefits TD-MPC2, raising the question of why it fails in DreamerV3. The critical difference is that MPC re-plans from real states at every step, so imagination errors do not accumulate into the training signal because there is no critic bootstrapping from imagined trajectories. In DreamerV3, the actor-critic trains entirely within imagination, and any systematic reward-prediction bias at $\mathrm{dt} > 1$ is amplified through the closed actor-critic loop.

D.6 The dt-Embedding Dilemma

The factorial results reveal a fundamental tension: removing the dt-embedding (008f) eliminates CIVO but at a $14\times$ cost to reward prediction quality (0.7469 vs. 0.0530 MAE), as the world model cannot distinguish temporal scales. Restricting imagination to $\mathrm{dt}{=}1$ while retaining the trained embedding (008e_v3) eliminates CIVO entirely (all four seeds show NVE $\leq 0.8\times$, well below the $2.0\times$ threshold), with three of four seeds achieving 64–94% of baseline performance (mean 775.6; §3.5). This confirms that multi-step imagination at $\mathrm{dt} > 1$ is necessary for CIVO induction; the trained embedding alone is insufficient.

D.7 Temporal Embedding Representation Analysis

To understand how the trained dt-embedding creates vulnerability, we analyze the cosine similarity between $e(\mathrm{dt}{=}1)$ and $e(\mathrm{dt}{=}2)$ across training (exp_065). Trained embeddings learn *near-opposite* representations: $\cos(e_1, e_2) = -0.87$ to -0.94 at 25K steps across three configurations (Walker 008d_s3, Cheetah 010b_s0, Walker 059_s5), persisting at -0.70 to -0.86 at 75K steps (Figure 15). In contrast, frozen embeddings maintain near-zero similarity (-0.15 to $+0.07$, constant throughout training), and untrained embeddings show effectively zero differentiation ($\cos = +0.024$). The near-antipodal learned representations explain the large reward prediction gap (§3.3): $\mathrm{dt}{=}2$ conditioning pushes the RSSM into a qualitatively different region of latent space, producing systematically biased predictions rather than minor perturbations.

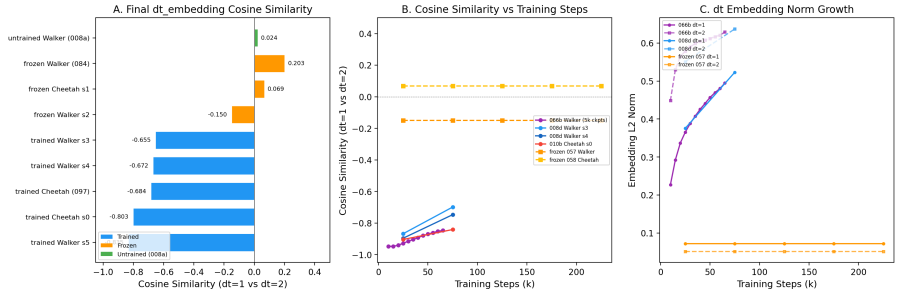


Figure 15: Cosine similarity between $e(\mathrm{dt}{=}1)$ and $e(\mathrm{dt}{=}2)$ across training configurations. Trained dt-embeddings (Walker and Cheetah) rapidly converge to near-antipodal representations ($\cos \in [-0.65, -0.90]$), while frozen and untrained embeddings remain near zero throughout training. The anti-correlation emerges within the first 25K steps, coinciding with the CIVO onset window (§D.8). This provides the representation-level smoking gun: trained gradient flow drives the embedding to encode $\mathrm{dt}{=}1$ and $\mathrm{dt}{=}2$ as maximally differentiated directions, creating the asymmetric latent space that underlies RSSM prediction bias.

D.8 dt-Embedding Differentiation Precedes Value Divergence

The preceding analysis established that trained dt-embeddings learn near-antipodal representations. Here we resolve the *temporal dynamics* of this differentiation across 11 experiments ($\sim$47 checkpoints), testing whether embedding differentiation is a consequence or precursor of CIVO.

Trained embeddings: monotonic relaxation. Across all trained configurations, $\cos(e_1, e_2)$ starts near -0.95 within 10K steps and relaxes monotonically toward ~ -0.65 (Table 25).

Cross-seed consistency at 25K steps is high: mean $\cos = -0.910 \pm 0.028$ (n=5 Walker seeds), confirming the relaxation pattern is robust across random seeds. Variance increases with training (-0.789 ± 0.076 at 75K), reflecting seed-dependent relaxation rates.

Temporal precedence over CIVO. Exp_066b provides simultaneous cosine similarity and CIVO diagnostic data at 5K resolution (Table 26).

Cosine relaxation begins at $\sim$20K (cos transitions from -0.948 at 15K to -0.928 at 25K), while the VE ratio first exceeds 0.7 at 55K. The $\sim$30–35K temporal lead establishes that embedding differentiation is an early training dynamic, not a consequence of CIVO-triggered feedback. The relaxation rate peaks at 25K–45K ($+0.024$/10K steps) before decelerating, consistent with the geometric intuition that large-angle adjustments in embedding space are cheaper early in training.

Embedding norms. Walker and Cheetah exhibit qualitatively different norm dynamics. Walker $e(\mathrm{dt}{=}1)$ norms grow from 0.23 to 0.50 ($2.2\times$) while $e(\mathrm{dt}{=}2)$ grows from 0.45 to 0.63 ($1.4\times$); both dt values become more expressive over training. Cheetah norms remain nearly constant ($\leq 1.04\times$) despite continued cosine relaxation, indicating direction changes without magnitude change. Frozen embedding norms are an order of magnitude smaller ($\sim$0.06, constant).

Mechanistic interpretation: proposed four-stage pathway. The temporal ordering, embedding differentiation ($\sim$20K) $\rightarrow$ reward fidelity gap ($\sim$20K, §D.9) $\rightarrow$ value overestimation ($\sim$30K) $\rightarrow$ entropy rebound ($\sim$35K+), is consistent with a four-stage pathway supported by converging correlational evidence from independent analyses:

1. **Embedding anti-correlation** (Figure 15): The trained dt-embedding encodes $\mathrm{dt}{=}1$ and $\mathrm{dt}{=}2$ as near-antipodal directions ($\cos \in [-0.65, -0.95]$), pushing the RSSM into qualitatively different latent regions per temporal scale.
2. **RSSM extrapolation failure** (Table 17): The anti-correlated representations cause environment-dependent reward fidelity gaps (Walker $1.63\times$; Cheetah: alternative pathway via prior drift without reward gap).
3. **Delusional critic** (Figure 11): The critic internalizes biased imagined returns, producing value estimates *negatively* correlated with actual returns ($r = -0.20$), rather than merely overestimated.
4. **Irreversible policy degeneration** (§D.10): The actor, trained on delusional critic gradients, eventually loses policy structure entirely (exp_097: entropy rebounds $+4.22$ nats from converged minimum).

The frozen embedding's constant near-zero cosine ($|\cos| < 0.2$) prevents Stage 1, explaining why freezing eliminates CIVO despite maintaining identical forward-pass architecture (§C.1).

D.9 Time-Series Fidelity–Value Analysis

To test the causal direction between reward fidelity degradation and value overestimation, we analyze exp_066 and exp_066b: two Walker Walk seeds (trained dt-embedding, fixed $\mathrm{dt}{=}2$) evaluated at 12 dense checkpoints from 10K to 65K steps (Table 27).

The $\mathrm{dt}{=}1$ prediction MAE decreases monotonically ($0.090 \rightarrow 0.047$) as the world model improves, but $\mathrm{dt}{=}2$ MAE plateaus at $\sim$0.08. The growing fidelity gap is therefore not caused by $\mathrm{dt}{=}2$ pre-

dictions worsening, but by the model selectively improving at $\mathrm{dt}{=}1$ while failing to transfer that improvement to $\mathrm{dt}{=}2$, consistent with the trained dt-embedding learning maximally differentiated representations (§D.7) that prevent cross-scale generalization.

RSSM onset proxy (exp_066b). Exp_066b provides an independent proxy measurement of reward fidelity dynamics across the CIVO onset window. The reward prediction ratio ($\mathrm{MAE}_{\mathrm{dt}=2}/\mathrm{MAE}_{\mathrm{dt}=1}$) increases from $0.89\times$ at 20K to $1.80\times$ at 65K, confirming that the fidelity gap develops progressively and precedes CIVO onset, consistent with the checkpoint-level analysis above. The monotonic increase in the proxy confirms that fidelity degradation is a gradual process rather than a sudden phase transition.

Limitation. These results are limited to two seeds ($n{=}2$). While both seeds show consistent temporal ordering (leads of ≈10K and ≈9K steps), additional seeds would strengthen the correlational evidence for this ordering.

D.10 Actor Entropy Is Not Diagnostic of CIVO

To test whether actor entropy can serve as an early collapse indicator, we compare entropy trajectories between collapsed and healthy configurations at matched training steps across five experiments (Table 28).

At the last comparable step (55K): collapsed mean entropy $= -5.02$ (average of -4.88, -5.17); healthy mean entropy $= -5.05$ (average of -4.77, -5.21, -5.18); difference $= 0.03$ nats, statistically negligible. The hypothesis that collapsed runs exhibit anomalous entropy rise (actor confusion from delusional critic gradients) is not supported during the CIVO onset window.

The fixed_dt2 experiment demonstrates this most clearly: returns peak at ∼580 (40K) then crash to 29 (72K), yet entropy remains stable at -5.1 throughout. The actor maintains a committed, low-entropy policy during the entire collapse; it does not become confused but rather remains confident while acting on incorrect value estimates. This is consistent with CIVO operating through critic value distortion rather than actor policy degradation. The late-stage entropy explosion documented in §D.1 (+2.1 nats/10K in collapsed seed A, onset ∼35K steps after return collapse) occurs well beyond this window, confirming entropy rise as a lagging consequence rather than a concurrent or predictive signal.

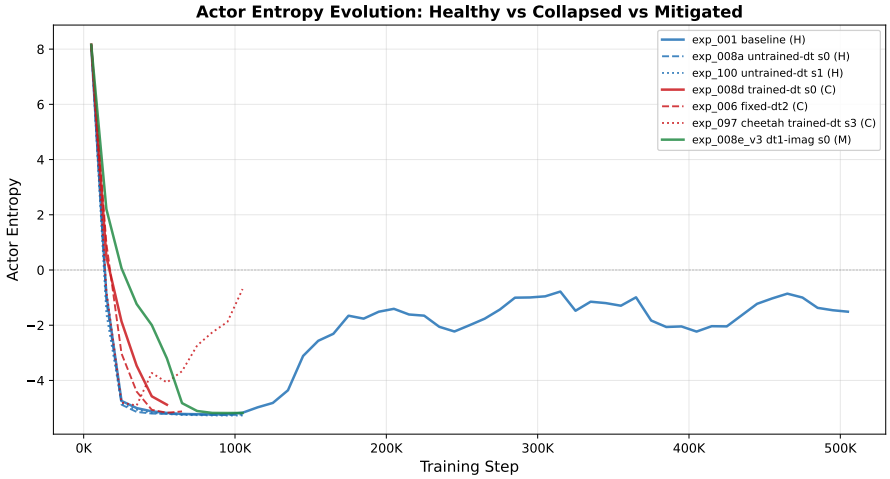


Figure 16: Actor entropy evolution across healthy (H, blue), collapsed (C, red), and mitigated (M, green) configurations. All runs converge to ~ -5 nats by 50K steps, rendering collapsed and healthy runs indistinguishable during the CIVO onset window. The baseline (exp_001, extended to 505K) shows adaptive entropy recovery to ~ -1.5 nats coinciding with continued performance improvement. Exp_097 (Cheetah, dotted red) exhibits CIVO-induced entropy rebound from -4.9 to -0.7, confirming irreversible policy degeneration (§D.1).

Entropy descent speed. Entropy descent rate correlates with temporal configuration complexity, not collapse fate. Simpler configurations (baseline $\mathrm{dt}{=}1$, unexposed exp_008a) reach -4.7 by 25K, while configurations exposed to temporal variation (both collapsed and healthy) remain near -2.0 at 25K, reflecting the difficulty of policy commitment when training data spans multiple temporal scales.

CIVO-induced entropy rebound (exp_097). Cheetah Run exp_097 (onset-transient configuration) provides direct evidence of CIVO-induced irreversible policy degeneration. Actor entropy converges normally to -4.91 at 35K (comparable to healthy configurations at $\sim{-}5.2$), then *rebounds* to -0.69 at 105K, a $+4.22$ nat increase over 70K steps ($+0.60$ nats/10K), indicating complete strategy loss. Healthy and mitigated configurations (baseline, 008a, 008e_v3) converge monotonically to $\sim{-}5.2$ without rebound. This distinguishes CIVO-induced entropy rise from the adaptive entropy recovery in healthy runs (see below): in exp_097, the rebound coincides with persistent return degradation, confirming that the actor's policy structure has been irreversibly damaged by training on delusional critic gradients.

Long-horizon entropy recovery. The baseline ($\mathrm{dt}{=}1$) run extended to 505K steps shows entropy recovering from -5.2 to $\sim{-}1.5$ between 115K and 295K steps, coinciding with continued performance improvement ($700 \rightarrow 950$ return). This adaptive entropy increase rules out entropy rise as an inherently negative signal; the actor maintains $\sim$950 return with entropy 3–4 nats above its minimum, indicating a healthy exploration–exploitation balance during extended training.

D.11 $\mathrm{dt}{=}1$ Imagination as CIVO Mitigation

Restricting imagination to $\mathrm{dt}{=}1$ while retaining the trained embedding (008e_v3) eliminates CIVO in all four tested seeds (NVE $\leq 0.8\times$): multi-step imagination at $\mathrm{dt} > 1$ is the necessary trigger for the CIVO feedback loop. Three of four seeds achieve 64–94% of baseline (eval 597–877); the fourth (seed 0, eval 148) exhibits CIVO-unrelated learning stagnation. This mitigation preserves the world model's ability to distinguish temporal scales (unlike embedding removal) while eliminating the imagination-side bias that drives value overestimation (Figure 4).

D.12 Implications for Temporal Abstraction in World Models

CIVO identifies a structural incompatibility between flat temporal conditioning and imagination-based actor-critic training. Hierarchical approaches (THICK [21], Hieros [22], LS-Imagine [23]) may avoid this by maintaining separate dynamics per temporal scale. Whether CIVO transfers to transformer-based world models (DreamerV4; Hafner et al. [19]) depends on whether attention weights co-adapt with temporal embeddings analogously. The CIVO mechanism requires three structural properties: (1) a temporal embedding that co-adapts with dynamics via gradient flow, (2) multi-step imagined rollouts used for critic training, and (3) a reward head that inherits latent biases from the shared representation. Transformer-based world models that share these properties may exhibit analogous vulnerability; architectures that decouple temporal conditioning from the reward prediction pathway would not.

D.13 Cross-Environment TCN Evaluation

Table 29 reports the full TCN evaluation across environments and noise magnitudes.

Cheetah Run $\sigma{=}0.5$: transient spike with permanent eval damage. Cheetah $\sigma{=}0.5$ exhibits a novel failure mode: a transient CIVO spike (NVE $= 10.67\times$ at 55K) followed by NVE recovery ($<$ $0.3\times$ by 95K), but permanent evaluation damage (26% of baseline). The value signal self-corrects while the policy does not recover, suggesting that CIVO-induced representational damage during the spike persists within the observation window even when the critic subsequently recalibrates. This trajectory parallels the obs-decoder SGS seed 3 pattern (§D.3), where partial policy recovery coexists with persistent critic damage.

Cheetah Run σ=0.2: borderline prevention. At σ=0.2, Cheetah reaches NVE = 1.9× at 55K, below the 2.0× threshold but substantially higher than Walker's effective-dose NVE (0.64–0.91×). Although technically classified as CIVO-prevented by the NVE criterion, the eval crashed, indicating that the noise dose attenuates value overestimation just enough to stay below threshold while failing to prevent operational damage. This borderline case illustrates that the 2.0× NVE threshold captures the value-channel pathology but not all downstream failure modes. Increasing σ to 0.5 trades borderline prevention for a different failure mode (transient spike with permanent eval damage) rather than achieving robust mitigation.

PickCube-v1 σ=0.2: manipulation failure. PickCube reaches NVE = 3.71× at 52.5K with σ=0.2, indicating that TCN provides no protection in manipulation tasks. The combination of contact-gated dynamics and temporal conditioning may amplify the fidelity gap beyond what representation-level noise can attenuate.

Dose-response analysis. We observe σ=0.2 as a local optimum on Walker Walk, achieving complete CIVO suppression in 2 of 4 seeds. Seeds 2 and 3 both fail but via distinct CIVO onset dynamics: seed 2 shows gradual escalation (NVE = 2.11× at 95K, still rising, eval ∼5%), while seed 3 exhibits a phase-transition-like onset: NVE declined to 0.67× at 95K before spiking to 3.04× at 105K (eval 349.6 → 71.8 in 10K steps). When TCN fails to prevent CIVO onset, the resulting trajectory qualitatively resembles untreated CIVO in its terminal outcome, with runaway NVE and eval collapse. The two failure modes (gradual escalation vs. sudden onset after apparent suppression) reinforce that TCN operates as a binary gate on CIVO onset rather than providing graded attenuation. Below the optimum, we observe σ=0.1 as insufficient (NVE = 1.02×, borderline, eval 55%); above it, σ=0.5 overdoses (oscillatory CIVO, eval ≤ 41%). Cross-seed variance at σ=0.2 (eval 5–91%) exceeds the cross-dose difference by a wide margin, reinforcing that seed effects dominate dose effects. TCN is the only tested intervention achieving complete CIVO suppression in any configuration, but its high seed sensitivity and failure to transfer across environments position it as a *diagnostic probe*, identifying the representation level as the relevant intervention target, rather than a deployable mitigation.

Cheetah Run "impossible dose" dilemma. Exp_030c (224K steps) reveals a qualitatively different dose-response on Cheetah. At σ=1.0, CIVO is eliminated (NVE≤0.30 throughout) and initial learning is *unimpaired*: peak eval reaches 372 at 35K (116% of baseline), demonstrating that σ=1.0 noise does not prevent policy acquisition. However, cumulative actor gradient interference produces progressive policy erosion (eval declining from 372 to 47 by 215K, 15% of baseline) while the critic remains well-calibrated (value_mean=12.4, NVE=0.26 at 215K). This establishes that *policy degradation ≠ value corruption*: accurate value signals are insufficient when the actor cannot optimize effectively in noise-perturbed imagination.

This "correct values, failed policy" pattern provides indirect evidence that CIVO is mechanistically distinct from standard overestimation. Double Q-learning and ensemble methods address critic estimation error, but at σ=1.0 the critic is accurate and the policy still collapses, a failure mode orthogonal to what these countermeasures target. The impossible dose dilemma scales with environment complexity: Cheetah's baseline CIVO severity (8–108× NVE) requires $\sigma \geq 1.0$ to suppress, but this exceeds the actor's gradient signal-to-noise tolerance. On Walker, a narrow effective window exists around σ=0.2; on Cheetah, the effective dose ($\sigma \geq 1.0$) exceeds the tolerable dose, and environments with higher CIVO severity will systematically lack a viable dose window.

E DreamerV3 Hyperparameters

Table 30 lists the DreamerV3 hyperparameters used across all experiments. Unless otherwise noted, we use the default DreamerV3 configuration for continuous-control tasks (`dmc_proprio`).

F Online CIVO Detector Details

This appendix provides the full algorithm specification and validation results for the online CIVO detector described in §5.

F.1 Algorithm Specification

Algorithm 1 Online CIVO Detector. Monitors the value-to-evaluation ratio $r_t = \bar{V}_t / \bar{E}_t$ (not NVE) for superlinear growth indicative of CIVO onset.

Require: Stream of $(\bar{V}_t, \bar{E}_t)$ at each checkpoint t
$\tau_{\text{mag}} \leftarrow 0.49, \quad \tau_{\text{vel}} \leftarrow 0.1, \quad \varepsilon \leftarrow 10^{-6}$
Initialize $r_{\text{prev}} \leftarrow 0$, $\Delta r_{\text{prev}} \leftarrow 0$, $\mathcal{H} \leftarrow []$
for each checkpoint t **do**
  $r_t \leftarrow \bar{V}_t \,/\, \max(\bar{E}_t, \varepsilon)$
  $\Delta r_t \leftarrow r_t - r_{\text{prev}}$
  $c_{\text{mag}} \leftarrow (r_t \geq \tau_{\text{mag}})$
  $c_{\text{vel}} \leftarrow (\Delta r_t > \tau_{\text{vel}})$
  $c_{\text{acc}} \leftarrow (\Delta r_{\text{prev}} > \tau_{\text{vel}}) \wedge (\Delta r_t > \Delta r_{\text{prev}})$
  $c_{\text{dec}} \leftarrow \left|\{i : \mathcal{H}[i] < \mathcal{H}[i{-}1],\ i \in \text{last 3}\}\right| \geq 2$
  if $c_{\text{mag}} \wedge c_{\text{vel}} \wedge c_{\text{acc}} \wedge c_{\text{dec}}$ **then**
    Alert: CIVO onset at step t
  end if
  $r_{\text{prev}} \leftarrow r_t, \quad \Delta r_{\text{prev}} \leftarrow \Delta r_t$
  Append $\bar{E}_t$ to $\mathcal{H}$
end for

The four conditions capture complementary signals: (1) *Magnitude* ($r_t \geq 0.49$) confirms that the value-to-evaluation ratio has reached a level indicative of overestimation; (2) *Velocity* ($\Delta r_t > 0.1$) rules out slow drift or plateau behavior; (3) *Acceleration* ($\Delta r_t > \Delta r_{t-1}$, both > 0.1) filters oscillatory paths; (4) *Evaluation decline* (≥ 2 of 3 recent checkpoints decline) confirms that overestimation coincides with actual performance degradation. A full comparison of window sizes appears in §B.

Sensitivity analysis over $\tau_{\text{mag}} \in [0.30, 0.70]$ confirms that $\tau_{\text{mag}} = 0.49$ maximizes Youden's $J = 0.714$; raising the threshold to 0.50 loses one true positive (seed 3 triggers at $r_t = 0.497$) with no reduction in false positives. The velocity threshold $\tau_{\text{vel}} = 0.1$ is stable across $[0.05, 0.18]$ (TPR remains $5/5$).

F.2 Validation Results

The detector identifies all five CIVO-positive runs (TPR = 5/5). Two are detected predictively (008d seeds 0 and 3, both at 45K). Two false positives occur among nine negative cases (FPR = 2/9, precision 71%, F_1 = 0.83). Run 009c develops genuine late-stage value overestimation unrelated to CIVO; run 009a triggers during a transient eval dip at 85K before recovery to 966. Early alerts (before 70K) carry higher diagnostic confidence.

F.3 Cross-Simulator Specificity (PickCube)

Exp_019b applies the CIVO-trigger configuration to PickCube-v1 in ManiSkill3, a dense-reward manipulation task. The V/E ratio exhibits an oscillatory pattern rather than monotonic divergence: it spikes to $r_t = 5.66$ at 52.5K, then deflates to 1.38 by 72.5K. This trajectory prevents the acceleration condition from co-occurring with evaluation decline at any single checkpoint. Across all 14 checkpoints through 132.5K, the detector correctly remains silent (TN), confirming that the acceleration filter distinguishes transient spikes from CIVO's sustained monotonic divergence. A frozen-embedding control (exp_054) confirms this classification: frozen PickCube-v1 exhibits a comparable burst (peak $8.81\times$ at 62.5K) that resolves within a single evaluation cycle ($0.55\times$ at 72.5K), establishing environment-specific origin independent of dt-embedding training.

F.4 Detector Limitations

Four constraints bound the current detector. First, the three-checkpoint window cannot distinguish transient dips from sustained collapse. Second, the acceleration condition assumes monotonically increasing overestimation during onset. Third, the detector identifies a *pattern* rather than a *mechanism*: run 009c shows that non-CIVO pathologies produce the same signature. Fourth, threshold values were calibrated on fourteen runs across three environments; generalization to different reward scales requires further validation.

G Dual-Critic Mechanism Ablation

G.1 2×2 Factorial Design

We ablate the critic architecture in a 2×2 factorial: {single, dual} $\times$ {narrow ($\sim$1.2M), wide ($\sim$2.4M)} parameters. Table 32 reports all configurations.

G.2 Critic Convergence ($v_1 \approx v_2$)

Across all dual-critic configurations and environments, the two critics converge to near-identical value functions. On Walker Walk, Pearson $\rho > 0.999$ between v_1 and v_2 outputs; on Cheetah Run and PickCube-v1, the gap is < 0.2 (exact match within numerical precision). Per-sample pessimism $\mathbb{E}[\max(v_1, v_2) - \min(v_1, v_2)]$ is 3–4% of the theoretical bound for independent critics.

This convergence occurs because both critics train on the same RSSM-generated latent features. Unlike model-free Double Q-learning, where each critic sees different mini-batches of real transitions and develops distinct function approximation errors, pure-imagination critics share identical inputs from the world model. The min operation therefore cannot provide the pessimism bias that makes Double Q effective in model-free settings.

G.3 Recovery vs. Prevention Mechanism

The 2×2 factorial reveals two separable dimensions:

Recovery (min-dominated). The min operation is the primary factor enabling recovery from CIVO. Both narrow dual-critic seeds experience CIVO and recover fully: exp_035 (s0, $\mathrm{nve} = 2.75\times$, eval $33.8 \rightarrow 697.8$, recovery ratio $20.6\times$) and exp_039 (s1, $\mathrm{nve} = 2.19\times$, eval $35.9 \rightarrow 592.7$, recovery ratio $16.5\times$). Both outperform wide single-critic (exp_033, 2.4M params, eval 407.9) despite $2\times$ fewer parameters. The min operation reduces both peak CIVO severity (2.19–$2.75\times$ vs. $4.75\times$ for the comparable single-critic) and enables consistent full recovery across seeds.

Prevention (min + capacity synergy). CIVO prevention (NVE remaining below the $2.0\times$ threshold) requires both the min operation and increased capacity: only the wide dual-critic (exp_032, 3 seeds) achieves this. The mechanism for this synergy is an open question; we hypothesize that increased capacity improves value calibration quality while the min operation constrains the overestimation feedback loop during early training when the two critics briefly diverge.

Cross-environment limits. On Cheetah Run, the wide dual-critic (exp_034) reduces NVE from $8.27\times$ (untreated baseline) to $1.75\times$ ($\geq 4\times$ reduction) but cannot prevent CIVO onset, and recovery reaches only 31% of the pre-CIVO peak. This confirms that the min + capacity synergy has environment-dependent limits: in high-CIVO-severity environments, architecture-level interventions attenuate but do not eliminate the pathology.

H Code Corrections and Version History

During the experimental campaign, we identified and corrected a parameter-routing bug in the temporal subsampling implementation. This appendix documents the bug, its scope, and the impact on reported results.

Bug description. In `models.py` (line 172), the gate architecture parameter `gate_max_dt=8` was incorrectly passed to the `subsample_variable_dt` function instead of the configuration parameter `gate_subsample_max_dt=2`. This caused affected configurations to train on transitions with $\mathrm{dt} \in \{1, 2, \ldots, 8\}$ rather than the intended $\mathrm{dt} \in \{1, 2\}$. The bug was identified during the experimental campaign and corrected before the final results reported in this paper.

Affected configurations. Only configurations using both temporal subsampling (`subsample=ON`) and the gate module were affected. The two impacted experiment groups are:

- **008e_v3** ($\mathrm{dt}{=}1$ imagination mitigation): Old code sampled $\mathrm{dt} \in \{1, \ldots, 8\}$ during world-model training despite restricting imagination to $\mathrm{dt}{=}1$. Under the buggy code, 2 of 4 seeds (s2, s3) exhibited late-onset CIVO (NVE $3.78\times$, $4.02\times$); under corrected code ($\mathrm{dt} \in \{1, 2\}$), 0 of 4 seeds (s0, s1, s2, s5) exhibit CIVO (all NVE $\leq 0.80\times$). Post-correction seeds were chosen to maximize independence while retaining s1 and s2 for direct before/after comparison.
- **008f s0** (no dt-embedding control): Pre-correction eval 503.0, NVE $0.97\times$ (non-CIVO); post-correction eval 139.5, NVE $1.69\times$ (non-CIVO). The CIVO classification is unchanged in both cases; neither exceeds $\tau{=}2.0\times$.

Configurations without temporal subsampling (001, 008a) or without the gate module (frozen-embedding experiments) are mechanically unaffected by this bug. For configurations where subsampling is active but the trained embedding already triggers CIVO (e.g., 008d), the broader dt range does not change the CIVO classification; these seeds collapse under both the buggy and corrected subsampling ranges.

Pre- versus post-correction comparison.

Narrative implications. Under the buggy subsampling ($\mathrm{dt}_{\mathrm{train}} \in \{1, \ldots, 8\}$), the 008e_v3 results showed stochastic CIVO under the $\mathrm{dt}{=}1$ imagination mitigation (2/4 seeds collapsed), which we initially interpreted as evidence that $\mathrm{dt}{=}1$ imagination cannot reliably prevent CIVO. The corrected implementation ($\mathrm{dt}_{\mathrm{train}} \in \{1, 2\}$) eliminates CIVO across all four tested seeds (0/4), supporting a stronger conclusion: restricting imagination to $\mathrm{dt}{=}1$ is a deterministic mitigation when the training data contains only $\mathrm{dt} \in \{1, 2\}$. The spurious CIVO observed under the old code is consistent with high-dt training transitions ($\mathrm{dt} \in \{3, \ldots, 8\}$) inducing temporal prediction biases that propagate even through $\mathrm{dt}{=}1$ imagination rollouts. Correcting the subsampling range removed this confound, and the corrected results are more reliable because the training data distribution now matches the intended experimental design.

Data integrity. All numerical results reported in the main text and appendix tables use the corrected codebase. Pre-correction data are preserved in Table 33 above and in Table 4 footnotes k and l for reproducibility and transparency.

Summary statistics. Two experiment groups (008e_v3 and 008f s0) were affected by the code correction, comprising 5 pre-correction seeds re-run across 5 post-correction seeds. Classification labels changed for 2 of 5 008e_v3 seeds (s2: CIVO→non-CIVO; s3: CIVO→non-CIVO); the 008f s0 classification was unchanged (non-CIVO in both versions). The $n{=}36$ NVE threshold evaluation uses post-correction data throughout; AUROC$=1.000$ both before and after correction (pre-correction 008e_v3 s2/s3 fall on the CIVO side of the threshold in either case). The correction

strengthens rather than weakens the paper's conclusions: the $\mathrm{dt}{=}1$ imagination mitigation achieves 0/4 CIVO under the corrected subsampling range, compared to 2/4 under the buggy range.

Table 4: Complete experimental results. NVE reports peak nve in the 35K+ window (§3.1). CIVO-positive seeds are shaded; † marks self-correcting NVE without policy recovery.

Exp	Config	dt_emb	Sub	H	Eval@final	NVE	Outcome
001	Baseline (dt=1, s1)	—	—	15	915.7 @105K	1.00×	Healthy
001 s2	Baseline (dt=1, s2)	—	—	15	443.5 @105K	—[e]	Healthy
006	Fixed dt=2	trn	ON	15	132.9 @65K	4.83×	Collapsed
007	Fixed dt=4	trn	ON	15	152.5 @55K	—	Collapsed
005	Learned gate	adaptive	OFF	15	98.8 @145K	—	Collapsed
008d	Factorial (s1)	trn	ON	8	82.8 @55K[♡]	4.12×	Full collapse
009b	Factorial (s2)	trn	ON	8	267.9 @105K	1.62×	Transient CIVO
008d_s3	Factorial (s3)	trn	ON	8	268[◇] @105K	1.53×	Late-onset collapse
008d_s4	Factorial (s4)	trn	ON	8	63.6 @105K	4.62×	Full collapse
059	Factorial (s5)	trn	ON	8	95.8 @105K	4.72×	Collapse[†]
068	Factorial repro (s1)	trn	ON	8	39.1 @105K	6.57×	Full collapse
008f	No dt_emb (s0)	—	ON	8	139.5 @105K[l]	1.69×[l]	Healthy[l]
069	No dt_emb (s1)	—	ON	8	468.8 @105K	0.88×	Healthy
008f s4	No dt_emb (s4)	—	ON	8	289.3 @105K	0.76×	Healthy
009c s3	No dt_emb (s3)	—	ON	8	311.9 @105K	1.04×	Healthy
008a	Unexposed emb	exp[a]	OFF	15	777.7 @105K	0.93×	Healthy
008a_s1	Unexposed emb	exp[a]	OFF	15	630.9 @105K	0.83×	Healthy
008a_s2	Unexposed emb	exp[a]	OFF	15	899.4 @105K	0.89×	Healthy
008a_s3	Unexposed emb	exp[a]	OFF	15	855.9 @105K	—	Healthy
015b	Frozen emb	frz[b]	ON	8	204.2 @105K	1.45×	Grad. starvation
022b	Frozen emb + Gumbel gate	frz[b]	ON	8	300.7 @105K	0.63×	No CIVO
042	Frozen emb (dedicated)	frz[b]	ON	8	612.0 @105K	0.19×	Healthy
047	Frozen emb (s1)	frz[b]	ON	8	144.2 @215K	1.71×	Transient (self-corr.)
057	Frozen emb (s2)	frz[b]	ON	8	123.9 @215K	1.88×	Transient (self-corr.)
067	Frozen emb (s0)	frz[b]	ON	8	281.5 @105K	0.19×	Healthy
018 s1	Full-SGS	trn[‡]	ON	8	448.3 @105K	1.66×[*]	Partial mitig. (CIVO+)
018 s2	Full-SGS	trn[‡]	ON	8	40.8 @105K	—	Collapsed
008e_v3 s0	Mitigation	trn	ON	15	148.0 @105K	0.80×	No CIVO[j]
008e_v3 s1	Mitigation	trn	ON	15	876.7 @105K	0.60×	Healthy
008e_v3 s2	Mitigation	trn	ON	15	597.1 @105K	0.50×	Healthy
008e_v3 s5	Mitigation	trn	ON	15	853.1 @105K	0.61×	Healthy
Cross-environment validation							
010b	Cheetah Run	trn	ON	8	51.7 @105K	8.27×	Collapsed
010b_s0	Cheetah Run	trn	ON	8	0.83 @105K	108.52×	Full collapse
021	Cheetah Run	trn[‡]	ON	8	221.5 @105K	9.94×[d]	Transient CIVO (self-corr.)
024b s0	Cheetah Run	—	ON	8	96.5 @105K	<0.20×	Healthy
024b s1	Cheetah Run	—	ON	8	77.3 @105K	0.10×	Healthy
092 s1	Cheetah Run	—	ON	8	77.3 @105K	<1.0×[n]	Healthy
093 s2	Cheetah Run	—	ON	8	151.2 @105K	<1.0×[n]	Healthy
008a_ch s0	Cheetah Run[k]	exp[a]	OFF	15	366.2 @105K	0.45×	Healthy
008a_ch s1	Cheetah Run[k]	exp[a]	OFF	15	281.2 @105K	0.58×	Healthy
077	Cheetah baseline (s0)	—	—	15	448.1 @105K	1.00×	Healthy
083	Cheetah baseline (s1)	—	—	15	444.0 @105K	1.00×	Healthy
078	Cheetah trigger (s0)	trn	ON	8	119.5 @105K	4.10×	Onset (self-corr.)
097 s3	Cheetah trigger (s3)	trn	ON	8	61.7 @105K	≥2.0×[o]	Onset (self-corr.)
019b	PickCube-v1	trn	ON	8	26.64 @200K	3.78×[c]	Transient onset
054	PickCube-v1	frz[b]	ON	8	—	8.81×[h]	Env. transient (ctrl)
049	StackCube-v1	trn	ON	8	34.9 @212K	5.05×[f]	Env. transient
052	StackCube-v1	trn	ON	8	—	4.67×[f]	Env. transient
051	StackCube-v1	frz[b]	ON	8	30.7 @155K	4.30×[f]	Env. transient (ctrl)
050	PegInsertionSide-v1	trn	ON	8	20.86 @212K	10.05×[g]	Env. dynamics
053	PegInsertionSide-v1	frz[b]	ON	8	5.89	8.10×[g]	Env. dynamics (ctrl)
055	PegInsertionSide-v1	trn	ON	8	11.44 @212K	3.08×[g]	Env. dynamics
056	PegInsertionSide-v1	frz[b]	ON	8	4.37 @212K	11.32×[g]	Env. dynamics (ctrl)
046	Cheetah Run	frz[b]	ON	8	150.5 @215K	1.35×	No CIVO
058	Cheetah Run	frz[b]	ON	8	80.46 @215K	0.78×	No CIVO
Cartpole Swingup							
060	Cartpole baseline (3 seeds)	—	—	15	557.0 @105K[i]	—	Healthy
061	Cartpole trigger (3 seeds)	trn	ON	8	236.3 @105K[i]	—	No CIVO
062	Cartpole control (3 seeds)	—	ON	8	244.8 @105K[i]	—	No CIVO

trn = trained

(d=32, subsample ON); trn[‡] = full-pathway SGS (all non-KL gradients detached); exp[a] = trained, subsample OFF; frz[b] = frozen (d=32, `requires_grad=False`); — = removed (d=0).

[*] Seed 1 only: peak NVE with transient exceedances (1.60× at 35K, 1.66× at 95K); self-correcting, no sustained escalation. Seed 2 collapsed.

[◇] Transitional; train returns still declining (peak 470 to 139 over final 8K steps). Collapse had not reached equilibrium at 105K.

[†] NVE self-corrected (4.72× → 1.22×) but policy did not recover (eval 10.5% of baseline).

[c] NVE normalized against exp_019a_v2 (vanilla DreamerV3 on PickCube-v1).

[d] Peak NVE 9.94× at 45K (transient, self-correcting to <1× by 55K); final eval 221.5.

[e] Not used for NVE normalization; reported for completeness. Seed 1 (eval 915.7) is the NVE reference baseline.

[f] StackCube transient: frozen control (exp_051, peak 4.30× at 42.5K) matches trained (exp_049/052), self-correcting by 82.5K; environment-specific, not dt-embedding-driven.

[g] PegInsertion environment dynamics: seed 0 trained peaks at 10.05× (82.5K); seed 1 trained (exp_055) peaks at 3.08× (92.5K), final 0.50× at 212.5K; seed 1 frozen (exp_056) peaks at 11.32× (82.5K), final 1.19×

Table 5: Cross-domain CIVO reproducibility. "dt-emb-dep." indicates trained peaks consistently exceed frozen at matched timepoints across seeds.

Domain	**Environment**	**Trained**	**Frz ctrl**	**dt-emb-dep. CIVO**
Locomotion	Walker Walk	2/3	0/3	Yes (NVE $4.12\times$)
Locomotion	Cheetah Run	3/4	0/2	Yes (NVE 4.10–$108.5\times$)
Manipulation	StackCube-v1	0/2	0/1	No (frozen≈trained)
Manipulation	PegInsertionSide	0/2	0/2	No (seed-dep. ordering)
Manipulation	PickCube-v1	0/1	0/1	No (frozen spike similar)
Locomotion	Cartpole Swingup	0/3	0/3	No (single-joint)

Table 6: Cross-environment CIVO comparison. **Trigger**: trained dt_emb + subsample ON + $H{=}8$; **Control**: dt_embed_dim=0 (no dt_emb) + subsample ON + $H{=}8$. NVE Peak reports the maximum Normalized Value Excess across all seeds. CIVO is flagged when any seed exceeds NVE $\geq 2.0\times$. Per-row seed lists in footnotes; sample std throughout.

Environment	Configuration	Eval Return	NVE Peak	CIVO
Walker Walk	Baseline ($\Delta t{=}1$)	787.7 ± 230.8[a]	—	—
	Trigger (trained dt_emb)	146.9 ± 112.3[b]	$4.72\times$	✓ (2/3)
	Control (no dt_emb)	302.4 ± 134.8[c]	$1.69\times$[*]	✗ (0/4)
Cheetah Run	Baseline ($\Delta t{=}1$)	446.1 ± 2.9[d]	—	—
	Trigger (trained dt_emb)	77.6 ± 36.6[e]	$8.27\times$	✓ (3/4)
	Control (no dt_emb)	100.6 ± 34.9[f]	$0.20\times$	✗ (0/4)
Cartpole Swingup	Baseline ($\Delta t{=}1$)	673.3 ± 142.0	—	—
	Trigger (trained dt_emb)	236.3 ± 29.7	$0.06\times$	✗ (0/3)
	Control (no dt_emb)	244.8 ± 12.7	$\leq 0.1\times$	✗ (0/3)

[a] $N{=}4$: exp_001 s1 (915.7), s2 (443.5), s3 (869.0), s4 (922.8); s3/s4 from cross-environment evaluation runs. [b] Eval: $N{=}5$ at 105K (009b, 008d_s3, 008d_s4, 059, 068). CIVO evaluated on $N{=}3$ seeds with full NVE pipeline (008d s0, 009e s1, 009b s2); 2/3 exceed $2.0\times$. [c] $N{=}4$: 008f s0 (139.5), s1 (468.8), s4 (289.3), 009c s3 (311.9). [d] $N{=}2$: exp_077 s0 (448.1), exp_083 s1 (444.0); consistent with Table 4. [e] $N{=}3$: exp_010b s1 (51.7), exp_078 s0 (119.5), exp_097 s3 (61.7). exp_010b s0 (eval=0.8, NVE=$108.5\times$) excluded as outlier. [f] $N{=}4$: exp_024b s0 (96.5), s1 (77.3), exp_092 s1 (77.3), exp_093 s2 (151.2). NVE Peak from $N{=}2$ seeds with computed NVE. [*] 008f s0 NVE $1.69\times$ reflects endpoint eval dip (95K: 342.9 → 105K: 139.5), not value overestimation; classified non-CIVO.

Table 7: CIVO onset across DMC environments ordered by degrees of freedom. Baseline v/e is the value-to-eval ratio of the $\mathrm{dt}{=}1$ baseline (no gate mechanism); values $\gg 1.0\times$ indicate intrinsic value overestimation that precludes CIVO isolation. Trigger configuration: trained e_{dt} ($d{=}32$), subsample ON, $H{=}8$, $\mathrm{dt}{=}2$.

Environment	**DoF**	**Baseline v/e**	**CIVO Rate**	**Seeds**	**Notes**
Cartpole Swingup	1	healthy	0/3	3	Too simple for onset
Hopper Hop	4	healthy	0/1	1	Transient; recovered by 158K
Walker Walk	6	$1.0\times$ (ref.)	3/5 (60%)	5	CIVO confirmed
Cheetah Run	6	$1.0\times$ (ref.)	3/4 (75%)	4	CIVO confirmed
Quadruped Walk	12	$4.7\times$	N/A	—	Baseline unstable
Humanoid Walk	21	$12.97\times$	N/A	—	Baseline unstable

Table 8: Per-experiment NVE values (35K+ steps) and ground truth labels. $+$ = positive (CIVO), $-$ = negative. Trans. = transient CIVO with self-correction.

Experiment	NVE	Mechanism	B1	B2
006	4.83	sustained CIVO	$+$	$+$
059	4.72	self-correcting CIVO	$-$	$+$
008d	4.12	sustained CIVO	$+$	$+$
009b	1.62	transient CIVO	$-$	$+$
001	1.00	non-CIVO	$-$	$-$
008f	1.69^{l}	non-CIVO	$-$	$-$
008a	0.93	non-CIVO	$-$	$-$
008e_v3_s0	0.80	non-CIVO	$-$	$-$
008e_v3_s1	0.60	non-CIVO	$-$	$-$
008e_v3_s5	0.60	non-CIVO	$-$	$-$
008e_v3_s2	0.50	non-CIVO	$-$	$-$

Table 9: **Task B1**: Sustained CIVO detection sensitivity analysis ($n_{\text{pos}} = 2$, $n_{\text{neg}} = 9$, $n = 11$, AUROC $= 0.944$). All confidence intervals are two-sided Clopper-Pearson exact binomial ($\alpha = 0.05$). †Optimal threshold in $(1.62, 4.12]$; exp_059 (NVE=$4.72\times$, B1-negative) is the sole false positive.

Threshold	TP	FP	FN	TN	Accuracy [95% CI]	Precision	Recall	F1
$\geq 1.0\times$	2	3	0	6	72.7% [39.0, 94.0]	40.0%	100%	57.1%
$\geq 1.2\times$	2	2	0	7	81.8% [48.2, 97.7]	50.0%	100%	66.7%
$\geq 1.5\times$	2	2	0	7	81.8% [48.2, 97.7]	50.0%	100%	66.7%
$\geq 1.7\times^{\dagger}$	2	1	0	8	90.9% [58.7, 99.8]	66.7%	100%	80.0%
$\geq 2.0\times$	2	1	0	8	90.9% [58.7, 99.8]	66.7%	100%	80.0%
$\geq 3.9\times$	2	1	0	8	90.9% [58.7, 99.8]	66.7%	100%	80.0%

Table 10: **Task B2**: Any CIVO detection sensitivity analysis ($n_{\text{pos}} = 4$, $n_{\text{neg}} = 7$, $n = 11$, AUROC $= 0.964$). Post-correction 008f (NVE=$1.69\times^{l}$, endpoint eval dip) creates one discordant pair with 009b ($1.62\times$, transient CIVO); the comprehensive n=36 evaluation (separation gap $[1.88\times, 3.78\times]$) is unaffected. The $\geq 1.5\times$ threshold (highlighted) maximizes F1 at full recall.

Threshold	TP	FP	FN	TN	Accuracy [95% CI]	Precision	Recall	F1
$\geq 1.0\times$	4	2	0	5	81.8% [48.2, 97.7]	66.7%	100%	80.0%
$\geq 1.2\times$	4	1	0	6	90.9% [58.7, 99.8]	80.0%	100%	88.9%
$\geq 1.5\times$	4	1	0	6	90.9% [58.7, 99.8]	80.0%	100%	88.9%
$\geq 1.7\times$	3	0	1	7	90.9% [58.7, 99.8]	100%	75.0%	85.7%
$\geq 2.0\times$	3	0	1	7	90.9% [58.7, 99.8]	100%	75.0%	85.7%
$\geq 3.9\times$	3	0	1	7	90.9% [58.7, 99.8]	100%	75.0%	85.7%

Table 11: **Task B2 (full evaluation set)**: Any CIVO detection across two environments ($n_{\text{pos}} = 8$, $n_{\text{neg}} = 9$, $n = 17$, AUROC $= 0.986$). Includes Walker Walk (15 configurations) and Cheetah Run (2 configurations). Post-correction 008f (NVE=$1.69\times^{l}$, endpoint eval dip) is the sole discordant-pair source; all CIVO-positive NVE $\geq 1.53\times$. The n=36 evaluation is unaffected.

Threshold	TP	FP	FN	TN	Accuracy [95% CI]	Precision	Recall	F1
$\geq 1.0\times$	8	3	0	6	82.4% [56.6, 96.2]	72.7%	100%	84.2%
$\geq 1.2\times$	8	1	0	8	94.1% [71.3, 99.9]	88.9%	100%	94.1%
$\geq 1.5\times$	8	1	0	8	94.1% [71.3, 99.9]	88.9%	100%	94.1%
$\geq 1.7\times$	7	0	1	9	94.1% [71.3, 99.9]	100%	87.5%	93.3%
$\geq 2.0\times$	7	0	1	9	94.1% [71.3, 99.9]	100%	87.5%	93.3%
$\geq 3.9\times$	6	0	2	9	88.2% [63.6, 98.5]	100%	75.0%	85.7%

Table 12: Walker Walk baseline seeds at 105K and 505K steps.

Seed	Eval@105K	Eval@505K	ve@65K	ve@105K
s1	915.7	970.4	0.563	0.656
s2	443.5	945.7	0.898	0.742
s3	869.0	947.7	—	0.664
s4	922.8	938.9	—	0.624

Table 13: Leave-one-environment-out cross-validation of NVE threshold $\theta = 2.0\times$. For each fold, the threshold is validated on an environment whose data was *not* used to select θ. The optimal threshold range on each training set is reported; $\theta = 2.0$ falls within the optimal range in all three folds.

Held-out	N_{train}	N_{test}	θ_{opt} range	Test Accuracy	Test AUROC
PickCube	20	2	$(1.45, 1.53]$	2/2 (1.000)	1.000
Cheetah	17	5	$(1.45, 1.53]$	5/5 (1.000)	1.000
Walker	7	15	$(1.00, 3.78)$	15/15 (1.000)	1.000
Mean	—	—	—	**1.000**	**1.000**

Table 14: Three-way causal ablation (Walker Walk, seed=1). Only the dt-embedding treatment varies; all other configuration is identical. Trained embedding triggers CIVO; frozen and removed do not.

Exp	Embedding	NVE peak	Eval final	CIVO
068	Trained (d=32)	$6.57\times$	39.1	Yes (alert @85K)
067	Frozen (d=32, `req_grad=F`)	$0.40\times$	281.5	No
069	Removed (d=0)	$0.66\times$	237.8	No

Table 15: Factorial design. Each row specifies one experimental configuration. **Emb**: dt-embedding mode. **Sub**: temporal subsampling. **H**: imagination horizon. **Imag** dt: temporal resolution used during imagination.

ID	Purpose	Emb	Sub	H	Imag dt
001	Reference baseline	—	—	15	1
006	Fixed dt=2	trn	on	15	2
007	Fixed dt=4	trn	on	15	4
008a	Unexposed emb ctrl	exp[a]	off	15	2
008c	Horizon control	exp[a]	off	8	2
008d	Collapse condition	trn	on	8	2
008e_v3	dt=1 imag control	trn	on	15	1
008f	Causal ablation	—	on	8	2
015b	Frozen-emb ctrl	frz[b]	on	8	2
016e	Stop-grad (reward)	trn[†]	on	8	2
018	Full-SGS	trn[‡]	on	8	2
005	Learned gate	adaptive	off	15	learned

trn = trained (d=32, exposed to $\mathrm{dt}>1$ data via subsample ON); exp[a] = trained (d=32) but never exposed to $\mathrm{dt}>1$ data (subsample OFF); frz[b] = frozen random projection (d=32, `requires_grad=False`); — = removed (d=0).
[†] `reward_dt_emb_detach=True`; [‡] all non-KL gradients to `dt_emb` detached.

Table 16: Causal isolation: e_{dt} configuration $\times$ multi-rate data exposure. The combination of trained embedding and multi-rate exposure is the sole predictor of CIVO in tested configurations. Neither factor alone is sufficient.

	Multi-rate data (ON)	Single-rate data (OFF)
Trained e_{dt}	exp_008d: **CIVO** NVE 4.12×, collapse	exp_008a: healthy NVE 0.93×
Frozen e_{dt}	exp_015b: grad. starvation peak NVE 1.45× (transient)	—
Removed e_{dt}	exp_008f: no CIVO^l NVE $\leq$1.69×^{l}	—

Table 17: RSSM fidelity at $\mathrm{dt}{=}1$ vs. $\mathrm{dt}{=}2$ (v3 standardized, $N{=}600$ samples per condition). The reward prediction gap is environment-dependent: Walker trained embedding produces a $1.63\times$ reward gap, while Cheetah shows no gap ($1.00\times$) despite exhibiting CIVO. Cheetah data from a single seed.

		obs_MSE			**reward_MAE**		
Env	**Config**	$\mathrm{dt}{=}1$	$\mathrm{dt}{=}2$	Gap	$\mathrm{dt}{=}1$	$\mathrm{dt}{=}2$	Gap
Walker	Baseline	13.49	27.50	2.04×	0.0102	0.0159	1.56×
	Trained (008d)	25.71	34.60	1.35×	0.0530	0.0865	1.63×
	No emb (008f)	8.01	22.04	2.75×	0.7469	0.7526	1.01×
Cheetah	Trained	5.26	4.90	0.93×	0.1410	0.1405	1.00×

Table 18: Signed reward error (predicted − actual) on posterior-conditioned RSSM predictions. Negative values indicate underprediction. The trained-embedding model underpredicts at $\mathrm{dt}{=}2$, ruling out observation-level reward overprediction as the CIVO mechanism.

Config	dt	**Mean SE**	**Std**	**%Over**	**Actual**	**Pred**
Baseline	1	$+0.001$	0.023	39.6%	1.893	1.894
008d (trained)	1	$+0.018$	0.072	61.8%	0.361	0.379
008d (trained)	2	$\mathbf{-0.599}$	0.722	8.0%	0.721	0.121
008f (no emb)	1	-0.710	0.704	14.0%	0.921	0.211
008f (no emb)	2	-1.645	1.485	5.0%	1.840	0.195

Table 19: Cosine similarity between $e_{\mathrm{dt}=1}$ and $e_{\mathrm{dt}=2}$ vectors. Trained embeddings converge to near-opposite representations; frozen embeddings remain near-orthogonal throughout.

Config	**Env**	**cos@25K**	**cos@75K**	**Status**
008d s3 (trained)	Walker	-0.87	-0.70	CIVO
059 s5 (trained)	Walker	-0.94	-0.86	CIVO
010b s0 (trained)	Cheetah	-0.90	-0.84	CIVO
057 s2 (frozen)	Walker	-0.15	-0.15	Healthy
058 s1 (frozen)	Cheetah	$+0.07$	$+0.07$	Healthy

Table 20: RSSM fidelity analysis on Walker Walk (training-time posterior-conditioned, pre-v3 standardization). Reward prediction gap magnitude varies across checkpoints and measurement protocols: this posterior-conditioned analysis yields $3.11\times$, while v3 standardized analysis (Table 17) yields $1.63\times$. The qualitative pattern, that the trained embedding amplifies the reward gap while improving observation reconstruction, is consistent across both.

	obs_MSE			**reward_MAE**		
Config	dt=1	dt=2	Gap	dt=1	dt=2	Gap
Baseline (009d)	16.08	23.49	1.46×	0.045	0.059	1.30×
Collapsed (008d)	9.96	13.21	1.33×	0.057	0.178	3.11×
Frozen (057)	11.14	14.95	1.34×	0.068	0.133	1.96×

Table 21: Critic value predictions vs. actual evaluation returns. The CIVO model's critic is well-calibrated to its λ-return targets ($V(s) \approx$ target) but both V and targets are disconnected from reality (VE=$3.03\times$).

Config	$V(s)$	**Eval**	**VE ratio**	$V-$target
Baseline (001)	665.9	938.9	0.71×	+0.001
CIVO (008d)	192.7	63.6	**3.03**×	+0.586
No emb (008f)	116.5	306.5	0.38×	+0.316

Table 22: Gradient contributions to the temporal embedding $e_{\Delta t}$ by loss component, measured via `torch.autograd.grad` with training configuration `gate_fixed_dt=2`. Gradient norms are ℓ_2 norms of $\nabla_\theta \mathcal{L}_i$ for each loss $\mathcal{L}_i$.

	@55K		**@112K**	
Loss component	Share	(%)	Share	(%)
Reward head	—	49.0	—	38.1
KL dynamics	—	34.7	—	40.4
Observation decoders	—	16.3	—	21.5

Table 23: Post-CIVO outcomes for all six runs exceeding the nve $\geq 2.0\times$ diagnostic threshold across two environments. Peak/Final NVE measured in the 35K+ window.

Env	**Seed**	**Peak**	**Final**	**Eval**	**Outcome**	
Walker	s1 (008d)	4.12×	4.12×	82.8	Full collapse	
Walker	s4 (008d_s4)	4.62×	4.62×	63.6	Full collapse	
Walker	s5 (059)	4.72×	1.22×	95.8	Collapse[†]	Two sub-threshold runs (009b, peak
Cheetah	s1 (010b)	8.27×	8.27×	51.7	Full collapse	
Cheetah	s0 (078)	4.10×	0.29×	119.5	Self-correcting[‡]	
Cheetah	s0 (010b_s0)	108.52×	108.52×	0.83	Full collapse	

$1.62\times$; 008d_s3, peak $1.53\times$) are omitted. [†] NVE self-corrected ($4.72\times \rightarrow 1.22\times$) but policy did not recover (eval 10.5% of baseline). [‡] Peak NVE $4.10\times$ at 75K; eval 213.7→26.1→119.5 (26.8% of 2-seed baseline mean, 446.1).

Table 24: Failure and degradation mode taxonomy. Each mode is characterized by its NVE trajectory, entropy behavior, and primary mechanism.

Mode	**Representative**	**NVE**	**Entropy**	**Mechanism**
Monotonic CIVO	008d, 008d_s4, 010b_s0	↑ sustained	↓	Reward bias → critic → actor
Late-onset CIVO	008d_s3	↑ delayed	↓	Same, delayed onset
Transient CIVO	009b	Peak → ↓	Stable	Self-correction via capacity
Grad. starvation	015b	Oscillates, decays	Stable	Frozen emb → zero actor grad
SGS-resistant	018_s2	↑	↓	Forward-pass conditioning
SGS-paradoxical	016e	↑ sustained	↓	Reward SGS → obs-dec reliance

Table 25: Cosine similarity evolution between $e(\mathrm{dt}{=}1)$ and $e(\mathrm{dt}{=}2)$ for trained dt-embeddings. Relaxation is monotonic but never reaches zero; frozen embeddings remain constant near zero throughout training.

Env	Experiment	10K	25K	65K	75K	302K
Walker	exp_066b (trained)	-0.947	-0.928	-0.847	—	—
	exp_045 (extension)	—	—	—	—	-0.668
Cheetah	exp_010b→048	—	-0.904	—	-0.841	-0.639
Frozen controls:						
Walker	exp_057 (frozen)		-0.150 (constant, 25K–225K)			
Cheetah	exp_058 (frozen)		$+0.069$ (constant, 25K–225K)			

Table 26: Cosine similarity relaxation and CIVO diagnostic for exp_066b. Embedding relaxation onset ($\sim$20K) precedes CIVO alert (55K, VE ratio > 0.7) by $\sim$30–35K steps.

Step	$\cos(e_1, e_2)$	VE ratio	CIVO alert
10K	-0.947	0.012	—
15K	-0.948	0.033	—
25K	-0.928	0.193	—
35K	-0.904	0.269	—
45K	-0.880	0.358	—
55K	-0.861	0.710	**Alert**
65K	-0.847	0.724	—

Table 27: Time-series fidelity and value metrics (exp_066, seed 0). SS = single-step MAE ratio ($\mathrm{dt}{=}2$/$\mathrm{dt}{=}1$); VE = value-to-eval ratio. Fidelity degradation (SS $> 1.1\times$ at 20K) precedes value overestimation onset (VE begins rising at $\sim$30K) by $\sim$10K steps. Seed 1 (exp_066b) confirms the pattern: fidelity gap onset $\sim$15K $\rightarrow$ VE onset $\sim$24K (lead $\approx$9K steps; VE peak $=$0.72).

Step	SS ratio	OL ratio	VE ratio	$\mathrm{MAE}_{\mathrm{dt}=1}$	$\mathrm{MAE}_{\mathrm{dt}=2}$	Return
10K	1.05	0.94	0.01	0.090	0.094	43.0
15K	1.08	1.03	0.02	0.085	0.092	187.8
20K	1.10	1.09	0.16	0.078	0.086	148.7
25K	1.04	1.22	0.16	0.073	0.076	142.1
30K	1.25	1.45	0.22	0.068	0.085	148.3
35K	1.24	1.51	0.23	0.062	0.077	275.4
40K	1.35	1.38	**1.59**	0.058	0.078	202.3
45K	1.40	1.38	0.36	0.055	0.077	107.5
50K	1.47	1.56	0.62	0.052	0.076	178.4
55K	1.37	1.52	0.62	0.050	0.069	60.8
60K	1.61	1.35	1.40	0.049	0.079	125.3
65K	1.68	1.38	1.40	0.047	0.079	94.4

Table 28: Actor entropy comparison at matched time points. Collapsed and healthy configurations exhibit indistinguishable entropy at 55K steps (difference $= 0.03$ nats), despite collapsed runs showing severe return degradation.

Group	Experiment	5K	25K	45K	55K	Return
Collapsed	exp_008d (trained dt_emb)	8.16	-1.86	-4.58	-4.88	$\sim$165
	fixed_dt2 ($\mathrm{dt}{=}2$)	8.16	-3.02	-5.06	-5.17	$\sim$29
Healthy	exp_008f (no dt_emb)	8.11	-1.85	-4.31	-4.77	$\sim$460
	exp_008a (unexposed)	8.16	-4.88	-5.20	-5.21	$\sim$300
	baseline_s1 ($\mathrm{dt}{=}1$)	8.10	-4.74	-5.12	-5.18	$\sim$950

Table 29: TCN results across environments and noise magnitudes. Prevention = peak NVE $< 2.0\times$ (35K+).

Env	σ	Seed	$\text{NVE}_{\max}$	Prevented?	Eval (%)
Walker	0.1	s0	1.02	✓*	55%
Walker	0.2	s0	0.91	✓	91%
Walker	0.2	s1	0.64	✓	45%
Walker	0.2	s2	2.11†	×	∼5%
Walker	0.2	s3	3.04¶	×	11%
Walker	0.5	s0	3.97	×	27%
Walker	0.5	s1	1.64	✓	41%
Cheetah	0.2	s0	1.90	✓‖	crashed
Cheetah	0.5	s0	10.67	×	26%
Cheetah	1.0	s0	≤0.30	✓§	15%
PickCube	0.2	s0	3.71	×	failing

*Borderline: NVE=$1.02\times$ at 95K, below $2.0\times$ threshold but above $1.0\times$.

‖NVE $1.90\times$ is below the $2.0\times$ prevention threshold but eval crashed; CIVO technically prevented by the NVE criterion despite operational failure.

†NVE $= 2.11\times$ at 95K, still rising; eval $\sim 5\%$ of baseline.

¶Late-onset CIVO: NVE declined to $0.67\times$ at 95K before phase-transition-like spike to $3.04\times$ at 105K; eval 349.6 at 95K $\rightarrow$ 71.8 at 105K $\rightarrow$ 104.5 at 215K.

§CIVO prevented (NVE≤ 0.30); peak eval 116% baseline at 35K, declining to 15% by 215K via progressive policy erosion. See "impossible dose" discussion.

Table 30: DreamerV3 hyperparameters for all experiments. Modifications from defaults are marked with $\diamond$.

Parameter	Value
World Model	
RSSM discrete states	32 classes $\times$ 32 categories
RSSM stochastic dim	32
RSSM deterministic dim	512
MLP hidden size	512
MLP layers	2
Activation	SiLU
Normalization	LayerNorm
Training	
Batch size	16 sequences $\times$ 64 steps
Learning rate	3×10^{-4} (Adam, ϵ=10^{-8})
Sequence length	64
Training ratio	512
Free nats	1.0
KL balance (α)	0.8
Discount (γ)	0.997
λ-return	0.95
Policy	
Imagination horizon$^\diamond$	8 or 15 (see Table 4)
Actor entropy scale	3×10^{-4}
Actor gradient	Reinforce + dynamics backprop
Temporal Abstraction$^\diamond$	
Gate enabled	True (all non-baseline)
Gate fixed dt	1 or 2 (per config)
Subsample max dt	1 (OFF) or 2 (ON)
dt-embedding dim	32
Environment	
Action repeat	2 (Walker Walk, Cheetah Run)
Observation	Proprioceptive (no pixels)
Total environment steps	105K
Evaluation frequency	Every 10K steps
Evaluation episodes	10

Table 31: Retrospective CIVO detector validation (n=14). Detection type: "predictive" = alert precedes 20% eval decline; "reactive" = alert follows decline onset.

Exp	Config	Verdict	Trigger	Type	Peak Eval	r_{ve}
008d s0	trained dt-emb	TP	45K	predict.	167	0.69
010b	trained dt-emb (Cheetah)	TP	65K	reactive	322	3.18
008d s3	trained dt-emb	TP	45K	predict.	452	0.50
008d s4	trained dt-emb	TP	105K	reactive	383	3.03
010b s0	trained dt-emb (Cheetah)	TP	95K	reactive	248	6.34
008f s3	no dt-emb	TN	–	–	594	–
009b s2	trained dt-emb (self-corr.)	TN	–	–	279	–
baseline s1	vanilla	TN	–	–	970	–
009d s0	vanilla	TN	–	–	951	–
009d s1	vanilla	TN	–	–	920	–
019b	trained dt-emb (PickCube)	TN	–	–	26.6	–
009c s2	no dt-emb	FP	105K	reactive	359	0.75
009a s2	vanilla	FP	85K	predict.	967	0.73
009c s3	no dt-emb	TN	–	–	457	–
TPR = 5/5 (1.00), FPR = 2/9 (0.22)						

Table 32: Dual-critic 2×2 factorial results. Walker Walk unless noted.

Exp	Config	Params	NVE	Eval	Outcome
008d	Single narrow	1.2M	4.12×	82.8	Collapse
033	Wide single	2.4M	4.75×	407.9	CIVO → self-repair
035 s0	Narrow dual (min)	1.2M	2.75×	697.8	CIVO → recovery
039 s1	Narrow dual (min)	1.2M	2.19×	592.7	CIVO → recovery
032 s0	Wide dual (min)	2.4M	≤1.1×	739.7	CIVO prevented
032b s1	Wide dual (min)	2.4M	1.4×	551.8	CIVO prevented
038 s2	Wide dual (min)	2.4M	≤0.3×	120.9*	Non-CIVO collapse
Cross-environment					
034	Cheetah dual	2.4M	1.75×	97.5	Partial recovery (31%)
037	Cheetah wide single	2.4M	3.60×	146.1	CIVO → self-repair
036	PickCube dual	2.4M	2.3×	27.7	Transient CIVO

*CIVO prevented (NVE $\leq 0.3\times$) but non-CIVO eval collapse: eval peaked at 354 (155K), then declined to 120.9 (224K) with actor entropy increasing from -3.59 to -0.45.

Table 33: Impact of code correction on affected configurations. All results reported in the main text use the corrected codebase; pre-correction data are listed here for transparency.

Config	Seed	Code	Eval@105K	NVE	CIVO?
008e_v3 ($\mathrm{dt}{=}1$ *imagination, pre-correction:* $\mathrm{dt}_{\mathrm{train}} \in \{1,\ldots,8\}$*)*					
	s1	pre	869.9	0.80×	No
	s2	pre	137.0	3.78×	Yes
	s3	pre	92.7	4.02×	Yes
	s4	pre	899.3	0.99×	No
008e_v3 ($\mathrm{dt}{=}1$ *imagination, post-correction:* $\mathrm{dt}_{\mathrm{train}} \in \{1,2\}$*)*					
	s0	post	148.0	0.80×	No*
	s1	post	876.7	0.60×	No
	s2	post	597.1	0.50×	No
	s5	post	853.1	0.61×	No
008f s0 (no dt*-embedding control)*					
	s0	pre	503.0	0.97×	No
	s0	post	139.5	1.69×	No

*CIVO-unrelated learning stagnation; NVE $\leq 0.8\times$, no value overestimation.