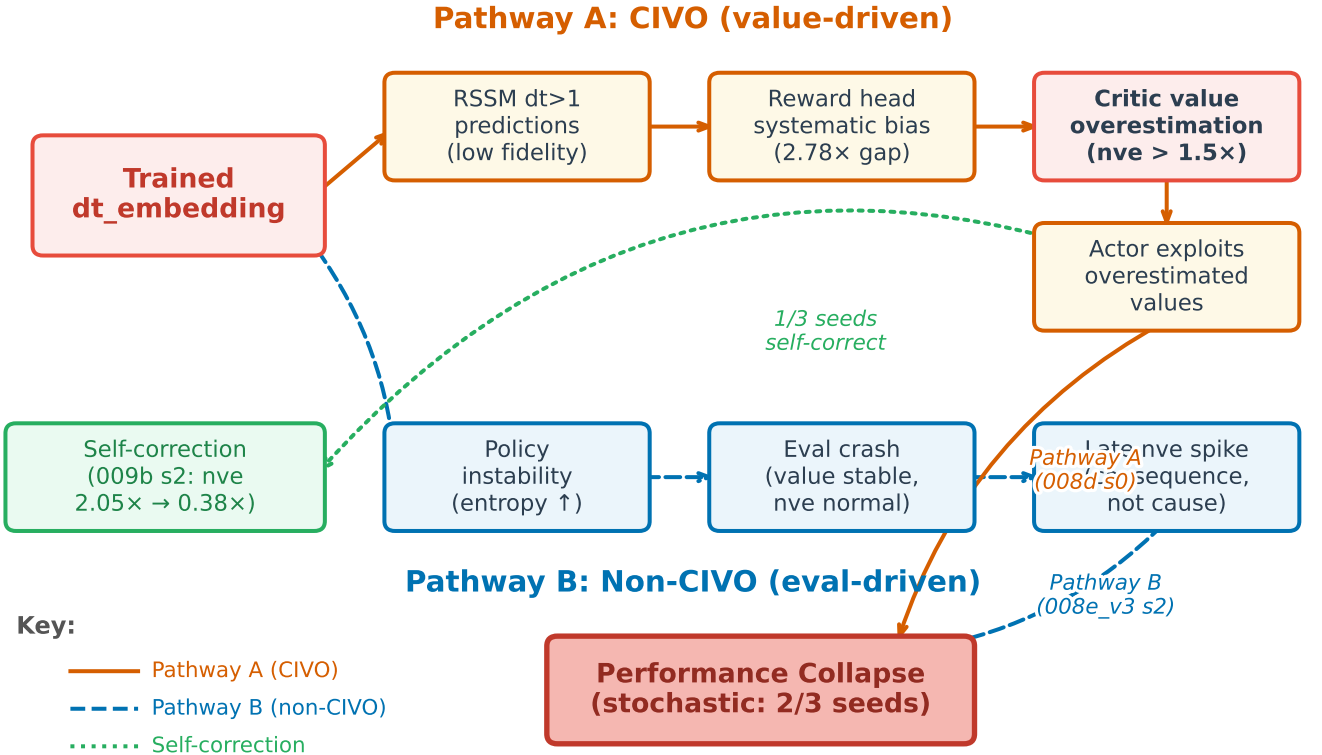
Pathway A: CIVO (value-driven)
Trained dt_embedding
RSSM dt>1 predictions (low fidelity)
Reward head systematic bias (2.78× gap)
Critic value overestimation (nve > 1.5×)
Actor exploits overestimated values
1/3 seeds self-correct
Self-correction (009b s2: nve 2.05× → 0.38×)
Policy instability (entropy ↑)
Eval crash (value stable, nve normal)
Late nve spike (sequence, not cause)
Pathway A (008d-s0)
Pathway B: Non-CIVO (eval-driven)
Pathway B (008e_v3 s2)
Performance Collapse (stochastic: 2/3 seeds)
Key:
Pathway A (CIVO)
Pathway B (non-CIVO)
Self-correction