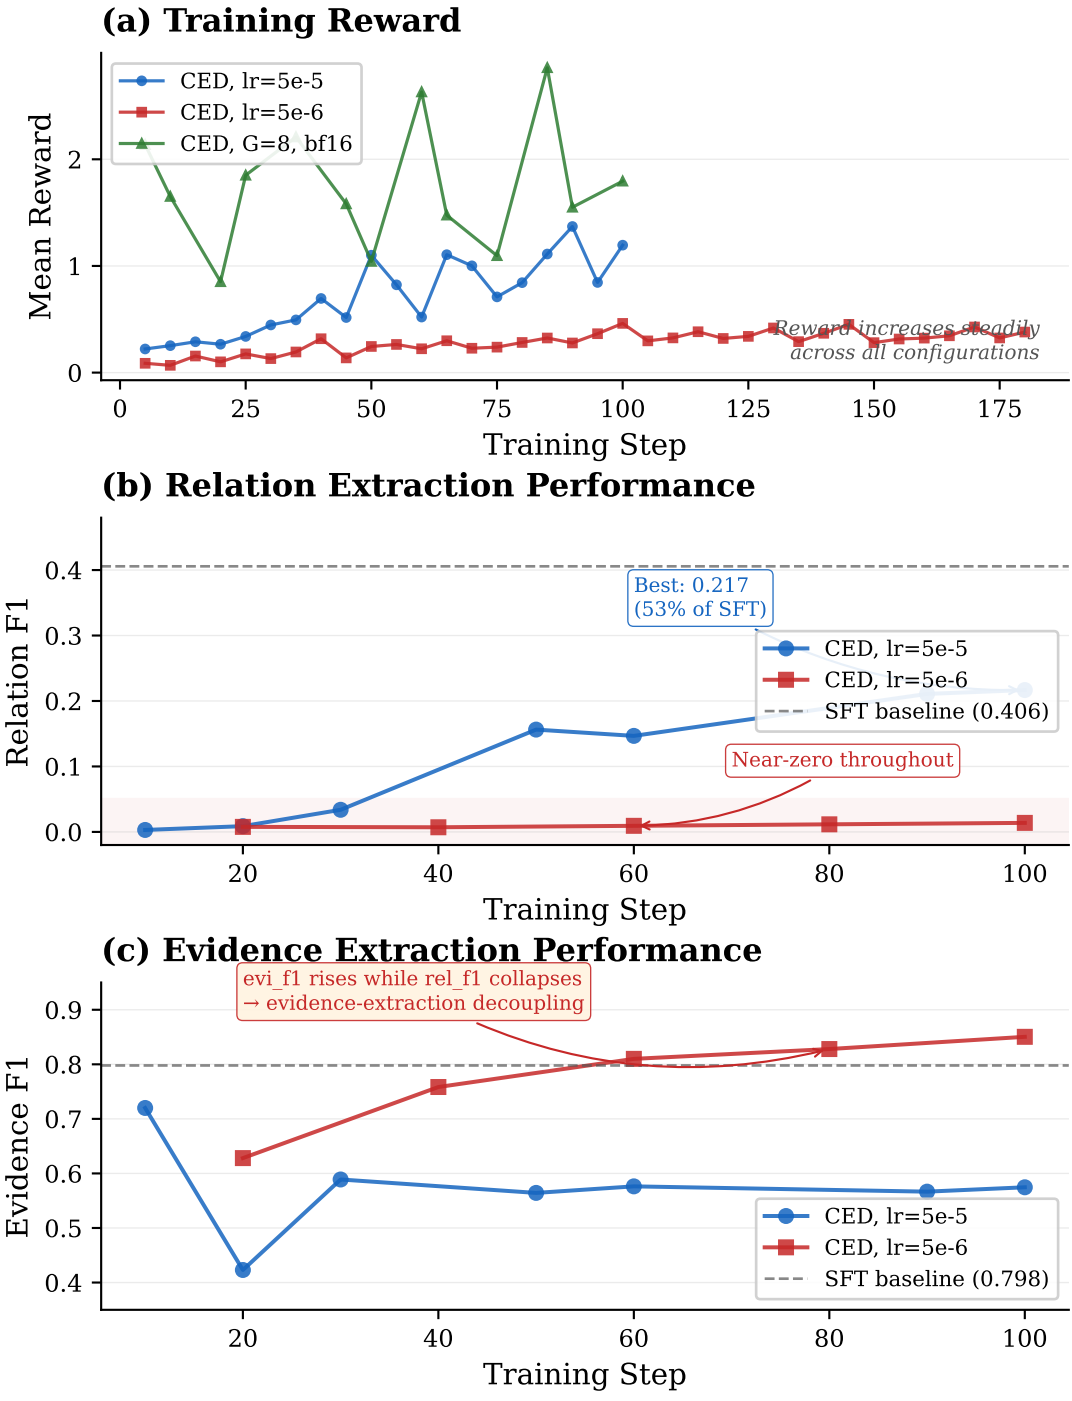
(a) Training Reward
CED, lr=5e-5
CED, lr=5e-6
CED, G=8, bf16
Mean Reward
Reward increases steadily
across all configurations
Training Step
(b) Relation Extraction Performance
Best: 0.217
(53% of SFT)
CED, lr=5e-5
CED, lr=5e-6
SFT baseline (0.406)
Near-zero throughout
Relation F1
Training Step
(c) Evidence Extraction Performance
evi_f1 rises while rel_f1 collapses
→ evidence-extraction decoupling
CED, lr=5e-5
CED, lr=5e-6
SFT baseline (0.798)
Evidence F1
Training Step