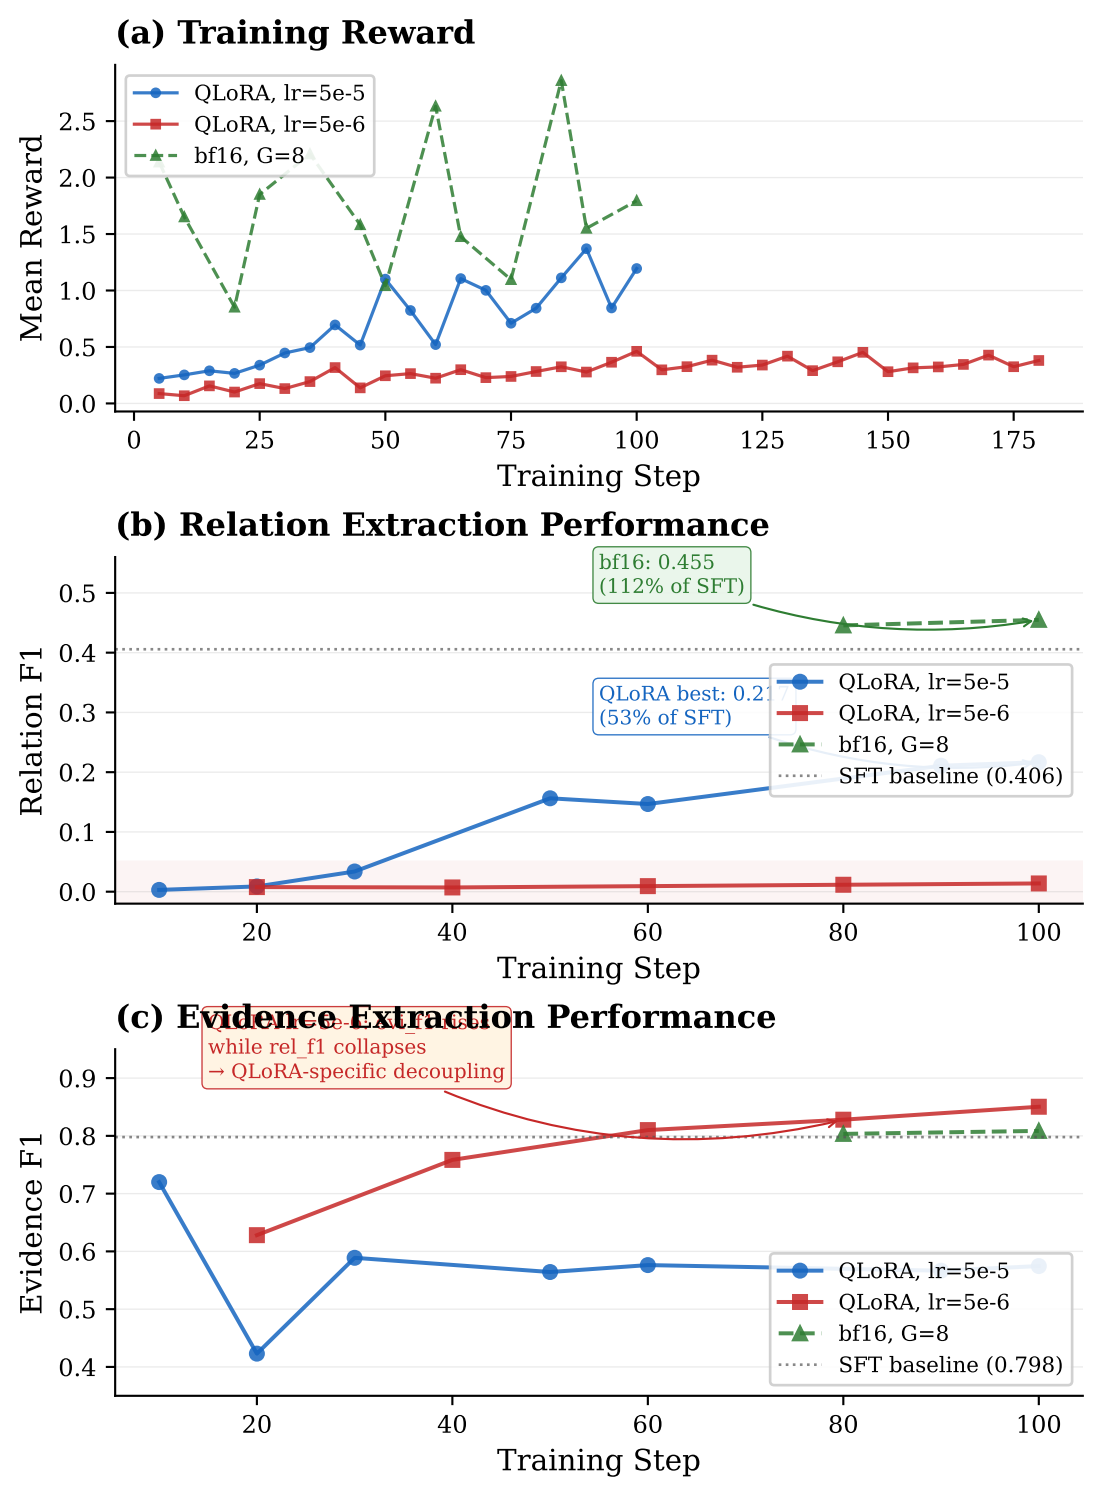
(a) Training Reward
QLoRA, lr=5e-5
QLoRA, lr=5e-6
bf16, G=8
Mean Reward
Training Step
(b) Relation Extraction Performance
bf16: 0.455
(112% of SFT)
QLoRA best: 0.21
(53% of SFT)
QLoRA, lr=5e-5
QLoRA, lr=5e-6
bf16, G=8
SFT baseline (0.406)
Relation F1
Training Step
(c) Evidence Extraction Performance
while rel_f1 collapses
→ QLoRA-specific decoupling
QLoRA, lr=5e-5
QLoRA, lr=5e-6
bf16, G=8
SFT baseline (0.798)
Evidence F1
Training Step