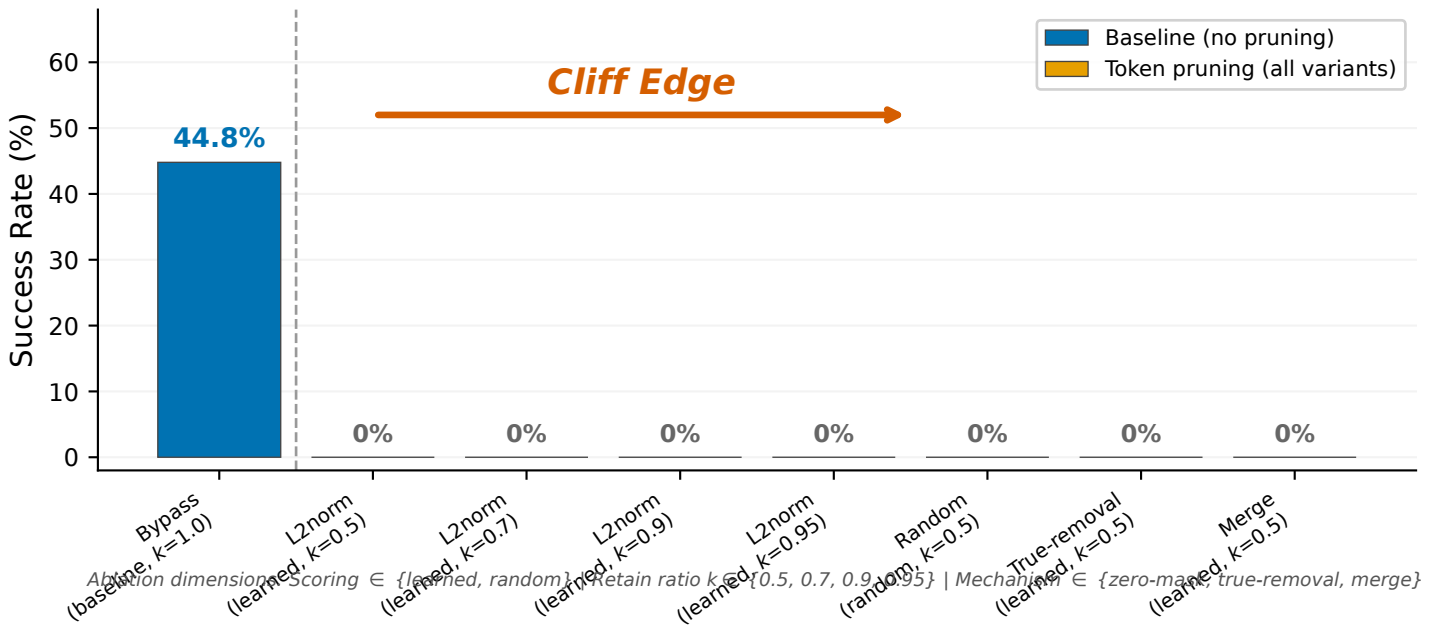
Baseline (no pruning)
Token pruning (all variants)
Cliff Edge
Success Rate (%)
0
10
20
30
40
50
60
44.8%
0%
0%
0%
0%
0%
0%
0%
Bypass (baseline, k=1.0)
L2norm (learned, k=0.5)
L2norm (learned, k=0.7)
L2norm (learned, k=0.9)
L2norm (learned, k=0.95)
Random (random, k=0.5)
True-removal (learned, k=0.5)
Merge (learned, k=0.5)
Ablation dimensions: Scoring ∈ {learned, random} | Retain ratio k ∈ {0.5, 0.7, 0.9, 0.95} | Mechanism ∈ {zero-masked, true-removal, merge}