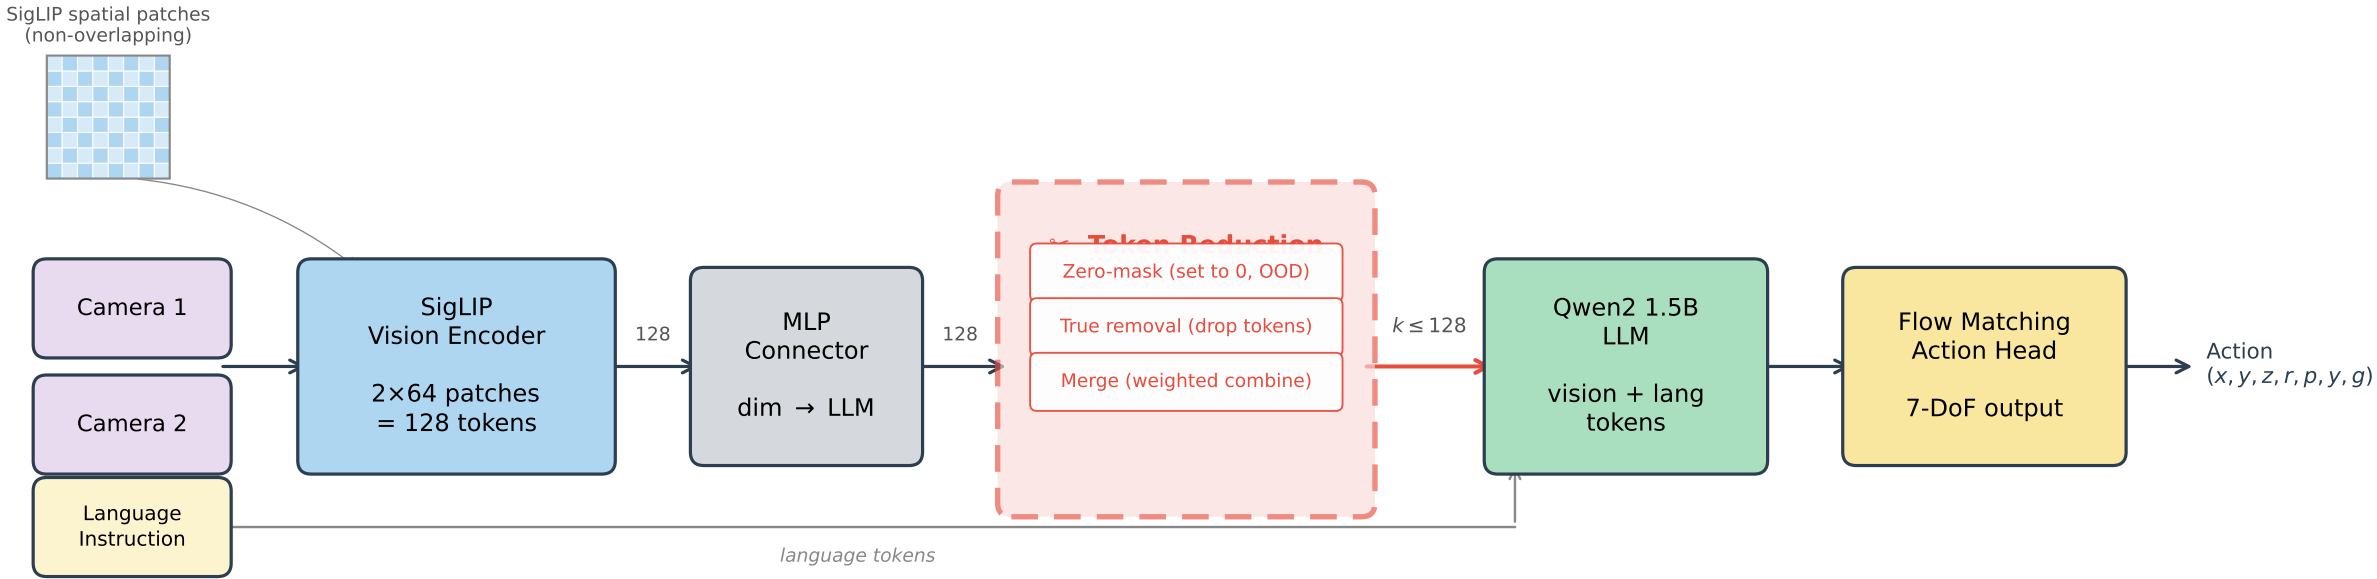
SigLIP spatial patches
(non-overlapping)
Camera 1
Camera 2
Language Instruction
SigLIP
Vision Encoder
2×64 patches
= 128 tokens
128
MLP
Connector
dim → LLM
128
Token Reduction
Zero-mask (set to 0, OOD)
True removal (drop tokens)
Merge (weighted combine)
k ≤ 128
Qwen2 1.5B
LLM
vision + lang
tokens
Flow Matching
Action Head
7-DoF output
Action
(x, y, z, r, p, y, g)
language tokens