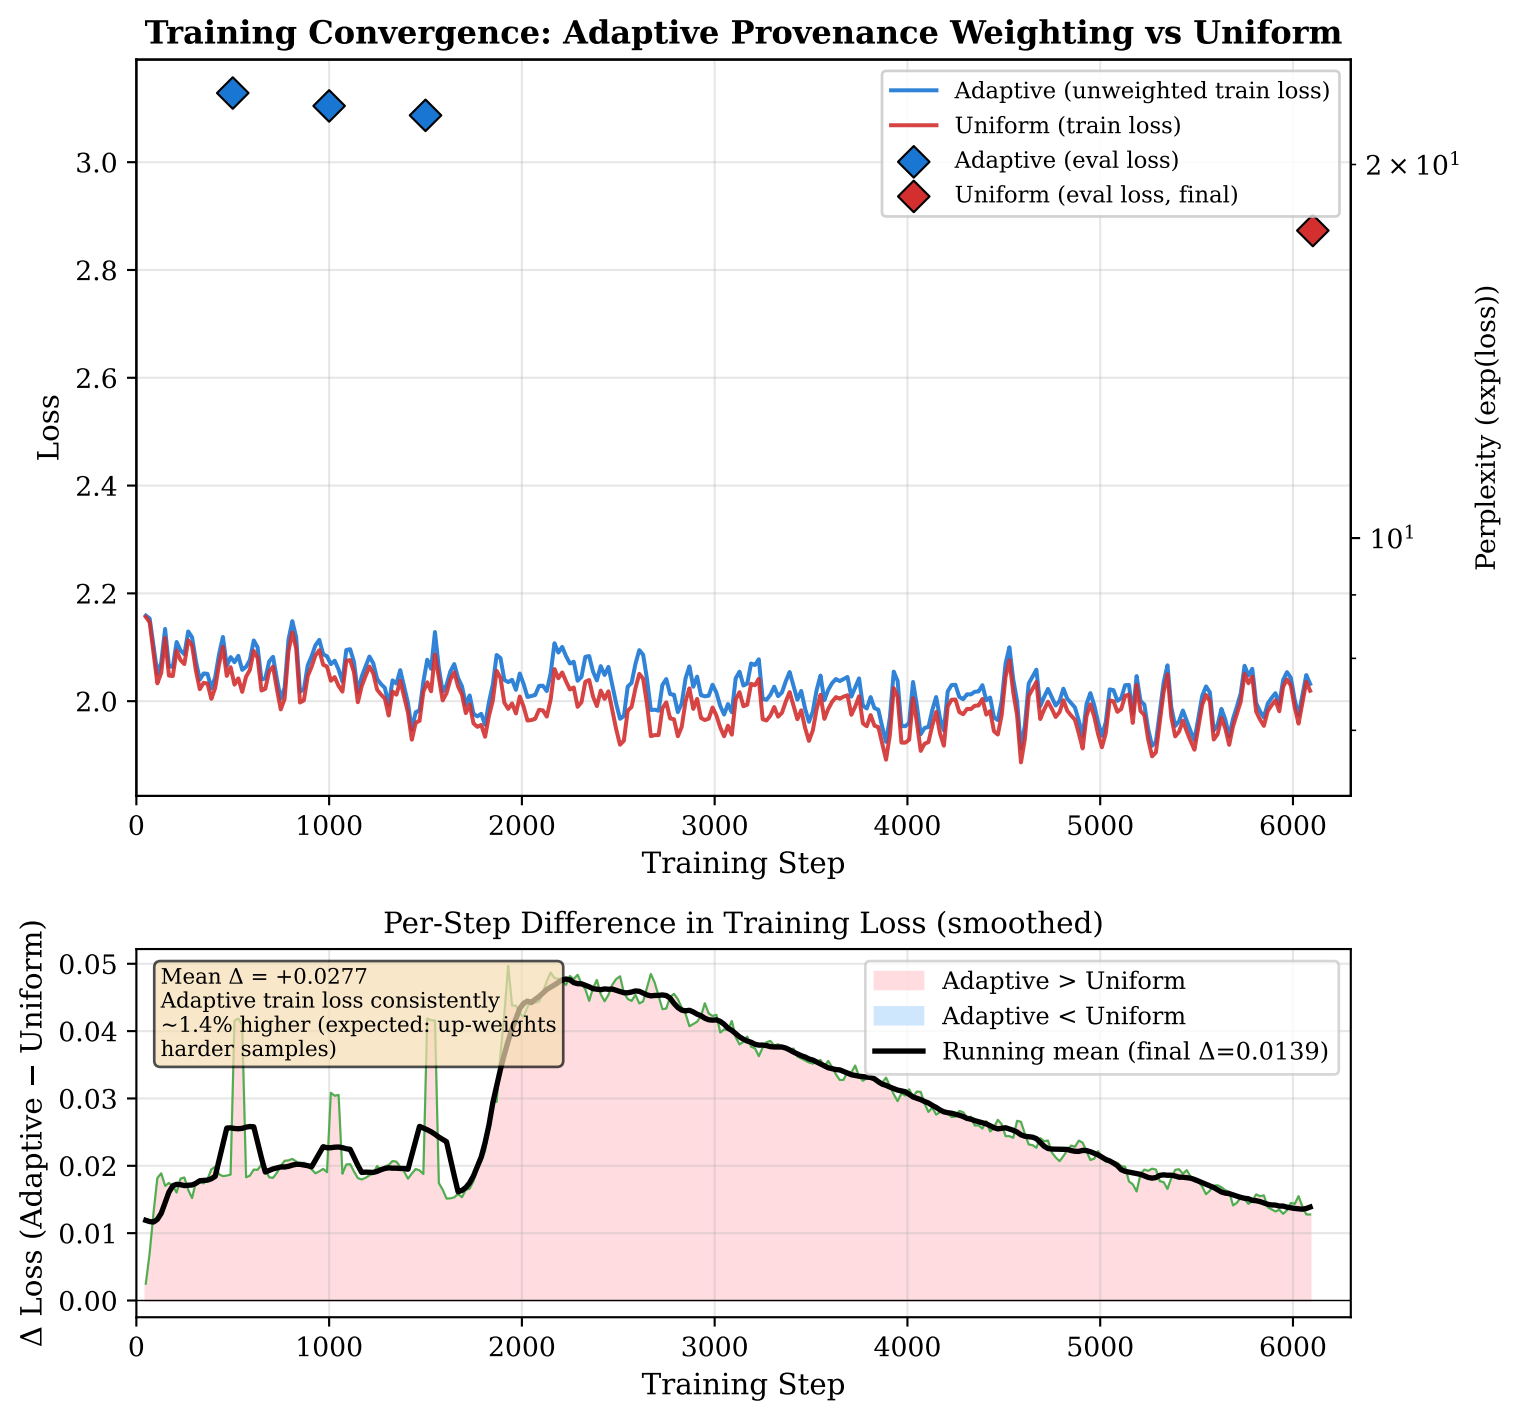
Training Convergence: Adaptive Provenance Weighting vs Uniform
Adaptive (unweighted train loss)
Uniform (train loss)
Adaptive (eval loss)
Uniform (eval loss, final)
Loss
Perplexity (exp(loss))
Training Step
Per-Step Difference in Training Loss (smoothed)
Mean Δ = +0.0277
Adaptive train loss consistently ~1.4% higher (expected: up-weights harder samples)
Adaptive > Uniform
Adaptive < Uniform
Running mean (final Δ=0.0139)
Δ Loss (Adaptive − Uniform)
Training Step