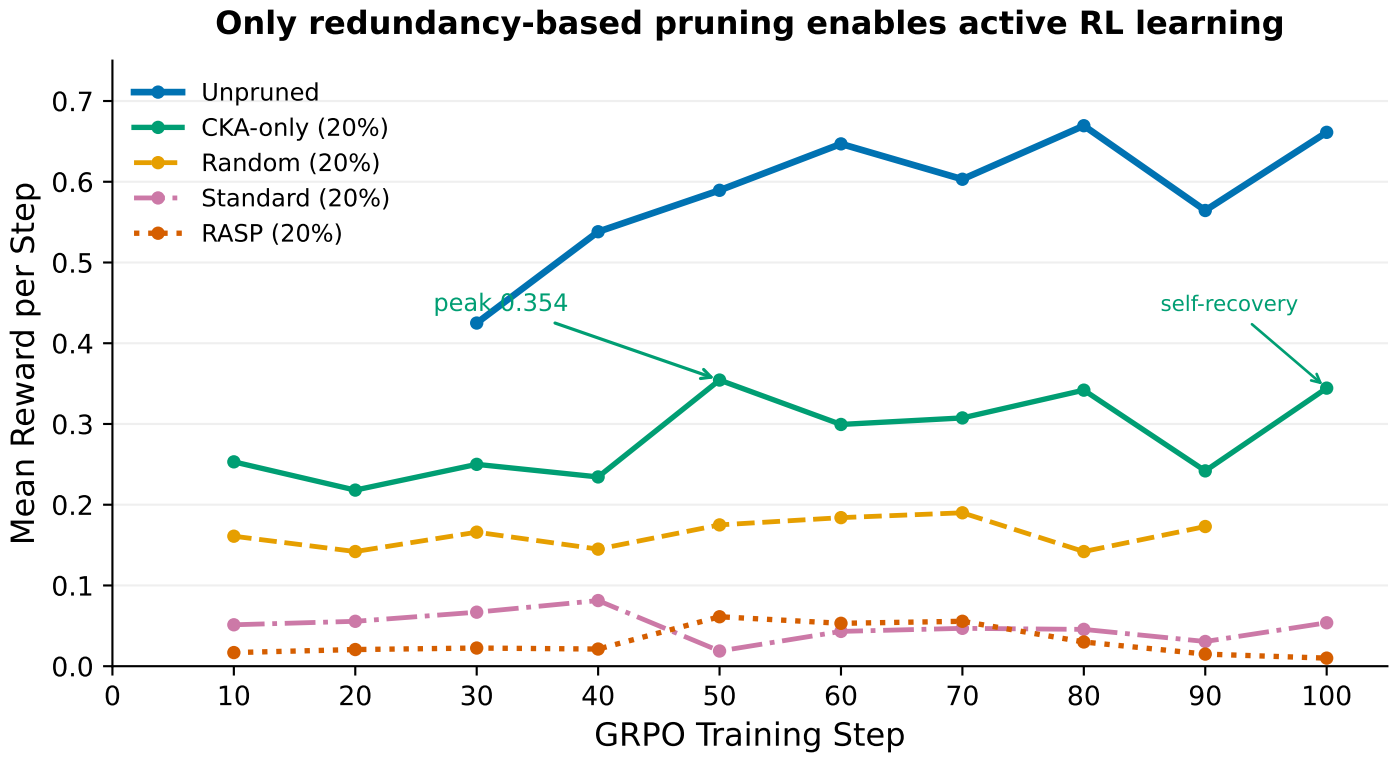
Only redundancy-based pruning enables active RL learning
Unpruned
CKA-only (20%)
Random (20%)
Standard (20%)
RASP (20%)
peak 0.354
self-recovery
Mean Reward per Step
GRPO Training Step
0.0
0.1
0.2
0.3
0.4
0.5
0.6
0.7
0
10
20
30
40
50
60
70
80
90
100