OntoCoDec: Type-Constrained Information Extraction via Compiled Schema Constraints

Anonymous ACL submission

Abstract

Large language models for information extraction generate structurally valid but ontologically invalid triples—where entity types violate the ontology's domain/range constraints. These violations persist across the capability spectrum and cannot be addressed by grammar-constrained decoding (which lacks cross-field type awareness) or post-hoc filtering (which discards valid predictions).

We present ONTOCODEC, a framework that compiles information extraction ontologies into composable token-level decoding constraints. A decidability analysis shows that domain/range constraints exceed context-free expressiveness, yielding a two-layer architecture: a grammar layer enforces structural validity and vocabulary membership; a semantic layer enforces type constraints at runtime through stateful masking over a dynamic entity registry.

Evaluated across three 8B-class models and two relation extraction datasets (up to 96 relations), OntoCoDec composed with LoRA fine-tuning reduces domain/range violations to near-zero ($\leq 0.21\%$) at F1 40–46%, but introduces cross-field type inconsistencies ($\leq$1.4%) from per-triple enforcement. Beam search (k=2) reduces CF to 0.11% at 1.18× latency; k=4 eliminates all violation types at 8B at 1.39×. In zero-shot settings, cross-field rates are higher (up to 12%) but decrease with model calibration. Zero-shot domain/range reductions span 76–99% on SciERC and 42–90% on DocRED; with scale- and schema-dependent layer interaction.

1 Introduction

Large language models are increasingly used for information extraction (Li et al., 2023; Sainz et al., 2024; Wei et al., 2023), but when generating relation triples (h, r, t), neither entity necessarily carries a type consistent with the ontology's domain/range constraints for r. We call such outputs *domain/range violations*: structurally valid triples that are ontologically invalid. Violations persist across the capability spectrum: 15.4% of triples on SciERC (Luan et al., 2018) and 18.5% on DocRED (Yao et al., 2019) under zero-shot decoding, and 0.5–2% even after LoRA fine-tuning; rates follow an observed pattern consistent with schema restrictiveness (few valid type pairs per relation) rather than ontology or model scale. At competitive quality (R-F1 40–46%), even 1–2% residual violations compound across thousands of documents into systematic knowledge graph corruption, making reliable constraint enforcement necessary.

Existing approaches leave a fundamental gap. Schema-informed prompting reduces violations by 62% on average but is per-relation unreliable (violations for EVALUATE-FOR actually *increase* by 5.8 pp). Grammar-constrained decoding (Dong et al., 2025; Willard and Louf, 2023; Beurer-Kellner et al., 2023) ensures syntactic well-formedness but cannot encode cross-field semantic constraints. Post-hoc filtering discards up to 74% of predictions including true positives (Appendix Q). Training-based approaches (Sainz et al., 2024; Lu et al., 2022) internalize schema knowledge but provide only soft guidance without runtime enforcement.

We propose **OntoCoDec**, a framework that translates IE ontologies into composable layers of decoding-time constraints. A decidability analysis of IE type constraints (§3) yields a two-layer architecture: the *grammar layer* compiles ontology constraints into context-free grammars (JSON structure, vocabulary membership) via XGrammar; the *semantic layer* interprets type constraints at runtime via dynamic logit masking over an entity registry, with a scale-dependent layer interaction: at 8B, the semantic layer achieves only 24% independent DR reduction (near-complete when combined), while at 70B it reaches 84% independently. OntoCoDec compiles constraints automat-

ically from IE ontologies, requires no constraint-specific training, and is applicable to autoregressive LMs with logits access; we validate on three English datasets (up to 96 relation types), three 8B-class models across three tokenizer families, with 70B-class validation (§5.2). LoRA + greedy S1 reduces DR to near-zero: DR $\leq 0.21\%$ at R-F1 40–46% ($\pm$0.3 pp, +10% latency); beam (k=4) eliminates all violation types at 8B on both RE datasets under fine-tuning (1.39× latency). DR enforcement converts ontologically invalid violations to individually valid but type-inconsistent ones (CF $\leq 1.4\%$ fine-tuned); zero-shot DR drops 76–99% on SciERC and 42–90% on DocRED, with total violations decreasing $\sim$31% (Appendix Q).

OntoCoDec is a constraint enforcement method, orthogonal and complementary to training-based schema-aware approaches such as GoLLIE (Sainz et al., 2024) and KnowCoder (Li et al., 2024); it can be composed with any generative IE system to improve schema compliance without retraining.

Our contributions:

1. Systematic ablation across 3 model families (Qwen3-8B, Llama-3.1-8B, Mistral-7B) and 2 datasets (SciERC, DocRED) under both zero-shot and LoRA fine-tuned settings, demonstrating DR reduction to $\leq$0.21% with R-F1 within $\pm$0.3 pp; DR enforcement introduces CF as a correctable side effect (§5.1).

2. A two-layer constrained decoding framework grounded in the observation that domain/range type constraints require stateful enforcement beyond context-free grammars—grammar-only approaches face state spaces of 6^{17}–6^{19} for typical IE schemas (Proposition 1; §3). Constraints compile automatically from IE ontologies, validated on schemas with up to 96 relation types.

3. Preliminary single-family scale exploration on Llama-3.1-70B (AWQ INT4), revealing a scale-dependent layer interaction: the semantic layer alone achieves 84% DR reduction at 70B vs. 24% at 8B (§5.2). At 8B, beam k=2 is Pareto-optimal (R-F1 44.60, CF 0.11%, 1.18× greedy S1 latency); k=4 achieves DR=0% with near-zero CF at 1.39×.

2 Background

IE Ontology Structure. An IE ontology $\mathcal{O} = (\mathcal{E}, \mathcal{R}, \sigma)$ defines entity types $\mathcal{E}$, relation types $\mathcal{R}$, and a mapping $\sigma : \mathcal{R} \to \mathcal{P}(\mathcal{E} \times \mathcal{E})$, where $\mathcal{P}(\cdot)$ denotes the power set, specifying valid (head type, tail type) pairs per relation. A triple (h, r, t) is *type-valid* if $(\text{type}(h), \text{type}(t)) \in \sigma(r)$. We quantify constraint restrictiveness via *constraint tightness* $\kappa(r)$ (§4): tighter constraints leave fewer valid type pairs per relation. We define ontology constraints broadly to encompass both formal ontology specifications (e.g., the KAIROS event ontology with explicit role-entity type mappings) and empirically derived type-pair constraints extracted from annotated training data. The latter applies when datasets lack a formal ontology but provide typed annotations (SciERC, DocRED); OntoCoDec compiles both sources into the same constraint representation without modification. Empirical constraints achieve 99%+ instance-level coverage on held-out data (SciERC 99.0%, DocRED 99.4%); uncovered type pairs trigger grammar-first fallback (§4), ensuring no valid triples are blocked.

Grammar-Constrained Decoding. Grammar-constrained decoding restricts each token y_t to a mask $M_g \in \{0,1\}^{|V|}$ derived from a target grammar. XGrammar (Dong et al., 2025) compiles JSON Schemas into pushdown automata for efficient enforcement of structural constraints; Outlines (Willard and Louf, 2023) and SGLang (Zheng et al., 2024) offer similar capabilities. However, these methods are limited to context-free constraints and cannot enforce cross-field semantic rules (e.g., "if r = FEATURE-OF, then $\text{type}(h) \in \{\text{METHOD}, \text{OTHERSCIENTIFICTERM}\}$"), which depend on previously generated field values.

Problem Statement. Given $\mathcal{O}$ and an autoregressive LM $\mathcal{M}$, we seek constraints $\mathcal{C}$ ensuring: (a) **structural validity**, i.e., well-formed JSON with $\mathcal{E}$/$\mathcal{R}$ membership; and (b) **type validity**, i.e., $(\text{type}(h), \text{type}(t)) \in \sigma(r)$ for every triple. Condition (a) is context-free; (b) is context-sensitive but *prefix-decidable*: at the point where a type token is generated, all information needed (relation type, entity registry) is available in the decoded prefix. This requires that each entity type name maps to a distinguishable token sequence (§3.3).

3 Method

We present OntoCoDec, a framework that translates IE ontologies into token-level constraints enforced during autoregressive generation. The design is guided by a decidability analysis of IE type constraints.

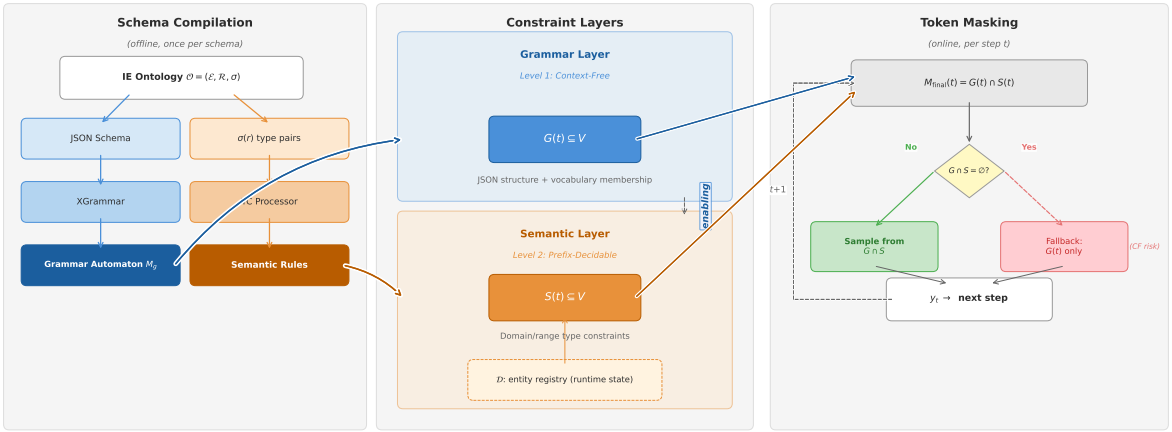


Figure 1: Overview of OntoCoDec. An IE ontology is compiled into two constraint layers: a grammar layer (context-free) and a semantic layer (prefix-decidable with state). At each decoding step, the final mask $M_{\text{final}}(t) = G(t) \cap S(t)$ ensures both structural and type validity; when $S(t) \cap G(t) = \emptyset$, grammar-first fallback preserves JSON integrity.

Constraint decidability taxonomy. We classify IE type constraints into three levels: (1) *context-free*: vocabulary membership and JSON structure; (2) *prefix-decidable with state*: domain/range type validity, determined at each type-bearing position given the decoded relation type and entity registry; (3) *non-local*: cross-entity type consistency, which depends on the full output. Unlike code generation where types are statically declared, IE constraints fall into level 2: entity types are inferred from text, the registry is built dynamically, and constraint targets depend on the co-occurring relation. OntoCoDec enforces levels 1–2 (grammar and semantic layers, respectively); level 3 is addressed by beam search or post-hoc resolution (Appendix M).

Existing grammar-constrained decoding frameworks (e.g., XGrammar) enforce level 1 constraints—structural validity and vocabulary membership—but IE domain/range rules require cross-position entity-type state that context-free grammars cannot express. Even when the entity set is finite, encoding type constraints into a DFA requires $O(|\mathcal{E}|^{|\text{entities}|})$ states: for typical SciERC documents (median 17 entities, 6 types), $6^{17} \approx 1.6\times10^{13}$ states; for DocRED (median 19 entities, 6 types), $6^{19} \approx 6.1\times10^{14}$—well beyond practical construction for online decoding. Formally, domain/range constraints for unbounded entity sets are not context-free (**Proposition 1**; proof via pumping argument in Appendix O): a PDA cannot retrieve the type of entity e_k without destroying stack entries above it. OntoCoDec's semantic layer addresses this with a runtime entity registry $\mathcal{D}$ (a finite map enabling random-access type lookups), avoiding both the theoretical impossibility (unbounded) and the engineering intractability (finite but exponential).

Empirical Observation (scale-dependent enabling dependency)**.** The grammar layer alone cannot enforce domain/range constraints (Proposition 1), while the semantic layer alone achieves only 24% DR reduction at 8B without grammar-enforced vocabulary membership (§5.1, Table 1). The strength of this dependency is scale-dependent: at 70B, the semantic layer independently achieves 84% DR reduction (§5.2), whereas 8B models show only 24% independent reduction.

We describe the problem formulation (§3.1), each layer, and two decoding strategies (§3.4); formal analysis of constraint composition is in Section 4.

3.1 Problem Formulation

IE as structured generation. Given an input text x, a language model generates a structured JSON output y:

$$y = \{\texttt{entities} : [\{n_i, \tau_i\}],\ \texttt{relations} : [\{h_j, t_j, r_j, \tau_j^h, \tau_j^t\}]\} \quad (1)$$

where each entity has a name n_i and type $\tau_i \in \mathcal{E}$, and each relation specifies head h_j, tail t_j, relation type $r_j \in \mathcal{R}$, and head/tail entity types τ_j^h, τ_j^t. The output JSON Schema (§3.2) enforces a fixed field ordering (head → tail → relation → head_type → tail_type) via deterministic grammar compilation, ensuring the relation type r_j is known before entity type fields and enabling sequential domain/range

lookups (§3.3). The grammar structurally enforces entity-before-relation ordering (100% preservation across all configurations), ensuring the entity registry (§3.3) is populated before relation-level type constraints are evaluated; unregistered entities trigger domain-only filtering.

Ontology compliance. Using the ontology $\mathcal{O} = (\mathcal{E}, \mathcal{R}, \sigma)$ defined in §2, we write $\mathrm{dom}(r) = \{\tau^h : \exists \tau^t,\ (\tau^h, \tau^t) \in \sigma(r)\}$ for domain-valid head types and $\mathrm{range}(r, \tau^h) = \{\tau^t : (\tau^h, \tau^t) \in \sigma(r)\}$ for the conditional range given head type τ^h. Our enforcement target is *pair-based* type validity: $(\tau^h, \tau^t) \in \sigma(r)$ for every triple. The decoder enforces this sequentially: $\tau^h \in \mathrm{dom}(r)$ at head position, then $\tau^t \in \mathrm{range}(r, \tau^h)$ at tail position, achieving best-effort pair-based enforcement (pair validity ensured when no fallback is triggered; domain-only filtering otherwise) strictly stronger than independent domain/range checks (which over-permit 33–37% spurious pairs on SciERC/DocRED; §4).

3.2 Grammar Layer

A JSON Schema encoding Eq. (1) with `enum` constraints ($\mathcal{E}$, $\mathcal{R}$) is compiled into a pushdown automaton via XGrammar (Dong et al., 2025). At step t, grammar-valid tokens $G(t) \subseteq V$ pass; all others receive logits $-\infty$, enforcing structural validity, vocabulary membership ($r_j \in \mathcal{R}$, $\tau_i \in \mathcal{E}$), but not domain/range constraints (which depend on entity types within each triple).

3.3 Semantic Layer

The semantic layer enforces domain/range constraints through a stateful logits processor maintaining an entity registry during generation. *Running example*: for USED-FOR ($\mathrm{dom} = \{\text{Method}\}$, $\mathrm{range} = \{\text{Task}\}$), entity "ResNet" (Method) passes the head-type check while Material is blocked; "image classification" (Task) passes the tail-type check.

Entity registry. The semantic layer builds a registry $\mathcal{D}$: name $\rightarrow$ type incrementally during entity generation: each completed entry $\{n_i, \tau_i\}$ adds (n_i, τ_i) to $\mathcal{D}$, which is consulted during relation generation for type lookups. The registry uses *first-seen* semantics: when an entity name appears in multiple triples, the type assigned at its first occurrence governs all subsequent constraint checks, making type enforcement sensitive to generation order. At each step t, the processor pattern-matches the decoded prefix to identify the current field position; semantic masks activate only at type-bearing positions (HEADTYPE, TAILTYPE).

Dynamic masking. At type-bearing positions, the processor computes a semantic mask $S(t) \subseteq V$. At a HEADTYPE position, if $h_j \in \mathcal{D}$ and $\mathcal{D}[h_j] \in \mathrm{dom}(r_j)$, the mask allows only the registered type; otherwise it falls back to all domain-valid types $\mathrm{dom}(r_j)$. At a TAILTYPE position, the mask restricts to $\mathrm{range}(r_j, \tau_j^h)$ conditioned on the already-determined head type, further narrowing to the registered type when $t_j \in \mathcal{D}$. This sequential check achieves best-effort pair-based enforcement (§3.1), strictly tighter than independent domain/range checking; "best-effort" reflects the grammar-first fallback (§4), which prioritizes JSON integrity over type validity when constraints conflict.

Tokenizer-agnostic type masking. We use *continuation masking* to handle context-dependent tokenizations and prefix-shared types (e.g., SEN/SID/SIDE in WikiEvents): a compile-time prefix map identifies valid continuations for each type name, disambiguating over multiple steps across byte-level BPE (Qwen, Llama) and SentencePiece (Mistral).

Grammar-semantic interaction. The final mask is $G(t) \cap S(t)$; when empty, grammar-first fallback applies (§4), embodying this interaction: semantic constraints reduce DR by only 24% without grammar-enforced vocabulary.

3.4 Decoding Strategies

S1 (joint extraction): Single-pass decoding. Both layers are active throughout generation; the entity registry populates during entity decoding and is immediately available for relation-level type checking.

S2 (two-pass pipeline): Two-pass decoding. A first pass extracts entities; these compile an input-dependent grammar for the second pass. S1 achieves higher R-F1 at comparable DR (§5.1) and is the primary configuration.

4 Constraint Composition

At each decoding step t, let $G(t) \subseteq V$ and $S(t) \subseteq V$ denote the tokens permitted by the grammar and semantic layers, respectively. The final mask is

their conjunction:

$$M_{\text{final}}(t) = M_{\text{grammar}}(t) \ \wedge \ M_{\text{semantic}}(t), \quad (2)$$

where $\wedge$ denotes element-wise AND over the vocabulary V; a token is allowed only if it satisfies *both* structural and type-level constraints. The conjunction is asymmetric: $S(t)$ is defined only over $G(t)$. This reflects the scale-dependent layer interaction (§3): the semantic layer requires the relation type to belong to $\mathcal{R}$, a property enforced only by the grammar layer. This dependency is scale-sensitive: 70B semantic-only achieves 84% DR reduction (vs. 24% on 8B), though grammar remains necessary for RV elimination (§5.2).

Grammar-First Fallback. When $S(t) \cap G(t) = \emptyset$ (tight constraints leave no token that is both grammar-valid and semantically valid), naïve masking would force selection of a grammar-invalid token, breaking JSON structure. We enforce a *grammar-first fallback*: semantic masking is skipped and the grammar layer alone controls selection. Structural integrity is a hard constraint; type violations are soft errors detectable by downstream validation, with DR (ontologically invalid type pair) more severe than CF (individually valid triples with cross-triple type inconsistency). Extraction tasks are robust to format-level constraints (Tam et al., 2024), supporting this asymmetric treatment.

Constraint Tightness. Per-relation constraint tightness $\kappa(r) = 1 - |\sigma(r)|/|\mathcal{E}|^2$ captures schema restrictiveness; DR reduction correlates descriptively with schema-mean $\bar{\kappa}$ across our three evaluation schemas (n=3), though entity type confusability is a significant co-factor (Appendix N, M).

Greedy limitation and fallback cascade. Single-pass per-triple enforcement cannot coordinate type assignments across triples sharing the same entity mention; cross-field consistency requires global inference that exceeds the per-token masking paradigm. The semantic layer commits entity types greedily—an instance of Level 3 (non-local) constraints (§3)—so CF violations arise when the same entity receives different types across triples; beam search mitigates this (§5.2). Fallback triggers on 0.5–5.7% of type-bearing positions (zero-shot), virtually eliminated by fine-tuning ($\leq 0.11\%$); all residual DR occurs at fallback positions (Appendix C).

Non-rectangularity and independent enforcement. Type signatures are non-rectangular (0/7 SciERC, 23/96 DocRED rectangular); independent enforcement admits 33–37% spurious pairs, eliminated by fine-tuning (Table 2). Shuffling type order produces byte-identical outputs across all 551 SciERC documents, confirming order invariance (Appendix N).

5 Experiments

We evaluate OntoCoDec in two deployment configurations: greedy S1 (DR $\leq 0.21\%$, $+10\%$ latency) and beam S1 (all violations $\rightarrow 0\%$, $1.39\times$ latency); zero-shot evaluation (§5.1; Appendix Q) serves as a diagnostic probe for the fine-tuned setting (R-F1 40–46%).

Setup. We evaluate on SciERC (Luan et al., 2018) (551 test sentences, 7 relations, $\bar{\kappa}$=0.333) and DocRED (Yao et al., 2019) (998 development documents; test annotations are not publicly available; 96 relations, $\bar{\kappa}$=0.842). RE type constraints are extracted from observed type pairs in training data (neither dataset provides a formal ontology); EE constraints on WikiEvents (§5.2) derive from the KAIROS ontology. All constraints are schema-level, independent of evaluation splits. Primary model: Qwen3-8B (Qwen Team, 2025), zero-shot greedy decoding; cross-model validation in Appendix Q, composability with fine-tuning in §5.1. LoRA hyperparameters (rank, α, learning rate) were tuned on SciERC and transferred to DocRED without dataset-specific optimization; given the substantial schema complexity difference (7 vs. 96 relations), this transfer may underestimate DocRED fine-tuning potential. Our R-F1 range (40–46% under LoRA) is consistent with recent LLM-based generative RE evaluations on these benchmarks, where document-level closed IE without gold entity spans typically achieves 35–50% (Wadhwa et al., 2023); supervised span-based methods operating under different assumptions (gold entity spans, dedicated architectures) are not directly comparable. DocRED development documents were used for evaluation only. Four violation categories: **DR** (type pair $\notin \sigma(r)$), **RV** (relation $\notin \mathcal{R}$), **ET** (entity type $\notin \mathcal{E}$), **CF** (conflicting types across triples). Baseline DR: 15.37% (SciERC), 18.46% (DocRED); schema-informed prompting reduces DR by 62% but is per-relation unreliable (Appendix C).

Mode	Violation Rates (%)						F1	
	DR↓	ΔDR	RV↓	ET↓	CF	Tot.↓	Rel.↑	Ent.↑
Unconstrained	15.37	—	1.39	0.40	0.08	15.45	8.23	31.01
Grammar	14.84	↓3%	**0**	**0**	0.08	14.92	8.23	30.98
GCD-equiv.†	15.26	↓1%	**0**	**0**	0.04	15.30	8.75	30.91
Semantic	11.65	↓24%	1.51	0.40	2.71	14.36	8.15	31.01
S1	**0.93**	**↓94%**	**0**	**0**	9.90	**10.83**	**8.36**	30.98
S2	0.97	↓94%	**0**	**0**	8.03	**9.00**	8.03	**31.90**

Table 1: Five-way ablation on SciERC (551 docs, Qwen3-8B, zero-shot). ΔDR: relative change vs. unconstrained; Tot. = DR + CF; bold = best (CF excluded). F1 (%). S1/S2 = single/two-pass (§3.4). Grammar = GCD ceiling (Geng et al., 2023). †vLLM `guided_json` (Kwon et al., 2023) (§5.2). Deterministic (greedy, T=0); robustness evidence in Table 48 and Appendices C, I.

Model	Dataset	Mode	DR↓	CF	R-F1↑	E-F1↑	ΔDR
Qwen3-8B	SciERC	Unconstr.	1.95	0	43.86	67.05	—
		S1	**0**	1.18	43.77	67.05	↓100%
	DocRED	Unconstr.	0.76	0.12	40.86	82.15	—
		S1	0.01	0.41	40.56	81.43	↓98.7%
Llama 3.1-8B	SciERC	Unconstr.	1.11	0	45.45	67.76	—
		S1	0.09	0.46	45.47	67.76	↓92%
	DocRED	Unconstr.	0.49	0.08	43.72	83.93	—
		S1	0.21	0.21	43.78	83.94	↓57%

Table 2: Fine-tuning composability: LoRA × unconstrained vs. S1 on SciERC (551 docs) and DocRED (998 dev docs). F1 (%). No ΔR-F1 is significant (Bonferroni-corrected; details in text); all DR reductions are (p<0.001). E-F1 includes parse-failure documents (Qwen PF: 7–8%; Llama PF: 2–3%).

5.1 Ablation and Composability with Fine-Tuning

Table 1 presents the five-way ablation on SciERC (Qwen3-8B, zero-shot, 551 documents); only the active constraint layers differ across rows.

Zero-shot diagnostic. Table 1 establishes that DR violations are pervasive (15–18% across datasets) and that constraint layers interact non-additively: grammar-only achieves 3% DR reduction on SciERC, semantic-only 24%, but their composition yields 94%—the *empirical enabling dependency*. On DocRED (96 relations, $\bar{\kappa}$=0.842), grammar-only *increases* DR via *latent violation surfacing*: forcing OOV→valid relation mapping exposes violations previously masked by vocabulary errors (Appendix Q). Cross-model validation (3 models × 2 datasets; Table 48) confirms DR drops of 76–99% on SciERC and 42–90% on DocRED, with R-F1 improving in 5/6 settings; post-hoc filtering achieves DR=0% but at 15–74% rejection rates. DR-to-CF conversion represents a severity downgrade: DR is irrecoverable without re-extraction, while CF—individually valid triples with cross-triple inconsistency—is amenable to post-hoc resolution (TCR eliminates all CF in our experiments, with bounded DR regression; Appendix M). Zero-shot evaluation thus establishes violation prevalence and constraint interaction; fine-tuning demonstrates practical elimination.

Composability with fine-tuning. Fine-tuning alone reduces DR to 1–2% but leaves residual violations on tight-constraint relations (Table 2); adding S1 yields DR $\leq$ 0.21% with no significant ΔR-F1 (Bonferroni-corrected; 95% CIs $\subseteq [-0.69, +0.25]$ pp; removed triples are 97.6–100% false positives; Appendix J.5). Under fine-tuning, grammar becomes a no-op (ΔDR<0.2 pp), making the semantic layer the sole DR defense (41–69% on SciERC, 2–75% on DocRED; Appendix C); three relations with $\kappa \geq 0.36$ retain DR, all eliminated by S1 (Appendix C). Fallback frequency is model-dependent: Qwen triggers 0 fallbacks under LoRA + S1 on SciERC, while Llama triggers 5,882 (Table 10), reflecting weaker entity typing; despite this, DR reaches 0.09% on SciERC (0.21% on DocRED), confirming that even frequent fallback does not compromise enforcement. CF drops from +9.82 pp (zero-shot) to +1.18 pp (fine-tuned), concentrated in GENERIC↔X ambiguity (63.3% of zero-shot CF; Appendix C); beam (k=4) eliminates all violation types at 8B (R-F1 43–46; §5.2).

Cross-field violation analysis. Under zero-shot greedy S1, CF is the sole residual violation type (DR, RV, ET all eliminated): 248 cross-field violations arise across 167 documents, of which 63.3% involve GENERIC↔X type ambiguity (Appendix C), concentrated in GENERIC↔OTHERSCIENTIFICTERM (34.7%) and GENERIC↔MATERIAL (12.3%). A source breakdown (Appendix C) shows 58.1% originate from hallucinated entities (a model capacity issue), while true cascade errors—where first-seen registration locks an incorrect type—account for only 7.3%; correct locks are 2.5× more frequent, confirming first-seen registration as a net-positive mechanism. Fine-tuning reduces CF to 22 cross-field violation instances across 20 documents (1.18%), with cascade errors completely eliminated (0/22). Beam search (k=4) eliminates all remaining CF at 1.39× latency (§5.2).

Downstream impact. Among constraint-removed triples, only 0.49% match gold labels (Appendix C); error analysis of 514 false negatives confirms only one (0.2%) was blocked by constraints, with dominant sources (relation confusion, entity type polysemy) schema-valid and beyond constraint reach (Appendix J.5). Across fine-tuned settings, 36 residual DR violations decompose into entity type errors (39%), schema coverage gaps (33%), and relation/direction errors (28%); the irreducible floor is ∼3 violations (Appendix C).

5.2 Preliminary Event Extraction Pilot

A preliminary n=20 WikiEvents pilot confirms constraint applicability to event extraction (Appendix P).

Runtime overhead. S1 overhead is +13.1% on SciERC, lower on DocRED due to amortization over longer sequences; under LoRA, +10.0% with 37% lower absolute latency. Constraint compilation is one-time <0.5 s; scaling benchmarks confirm constant overhead (2–3%) up to 200 entity types and 1,000 relations (Appendix K).

Beam search compatibility. With LoRA, beam search (Anderson et al., 2017) eliminates all four violation types at 8B on both SciERC and DocRED; at 70B, beam k=4 achieves DR=RV=ET=0% with one residual CF on SciERC (Appendix F). A beam-width sweep ($k \in \{1, 2, 4, 8\}$) on SciERC reveals k=2 as Pareto-optimal: R-F1 44.60 (highest), CF 0.11%, at only 1.18× greedy S1 latency; k=4 achieves CF=0% but lower R-F1 (43.60) at 1.39× (Table 25; latencies relative to greedy S1). This yields three deployment tiers: greedy S1 (+10% latency, DR≤0.21%, CF≤1.4%), beam k=2 (1.18×, near-zero CF), and beam k=4 (1.39×, all violations→0% at 8B). Per-beam memory overhead is $O(k \times |\text{entities}| \times |\text{types}|)$ for independent entity registry copies, negligible at typical document scales.

70B-class validation (single-family). The following results constitute a preliminary single-family scale exploration on Llama-3.1-70B-Instruct (AWQ INT4); cross-family verification (e.g., Qwen-72B) remains future work. All 70B experiments use Llama-3.1-70B-Instruct under AWQ INT4 quantization; an 8B ablation confirms constraint enforcement is unaffected by quantization (Appendix R). Zero-shot S1 eliminates DR on both datasets (SciERC 6.41% → 0%, DocRED 14.98% → 0%; R-F1 11–13). A four-way ablation reveals scale-dependent layer interaction: semantic-only achieves 84% DR reduction vs. 8B's 24% (Table 1), though grammar remains necessary for RV elimination and CF control (semantic-only CF=76). LoRA confirms composability: DR → 0% on both datasets (R-F1 44.53, 43.95; $|\Delta| \leq 0.20$ pp; CF 0.18–1.33%); cross-family generalization at 70B scale remains to be verified. Beam (k=4) reduces DocRED parse failures (12.5% → 0.5%, R-F1 +4.19 pp); SciERC controlled comparison isolates pure beam gain at +0.28 pp (Appendix C).

Comparison with grammar-constrained approaches. Our grammar-only ablation (Table 1, row 2) implements the same constraints as GCD (Geng et al., 2023) and BoostCD (Sakota and West, 2025); vLLM `guided_json` (Kwon et al., 2023) confirms equivalence on Qwen3-8B (DR=15.26%; other models use the same backend; Appendix C). The ∼15% residual DR is the gap our semantic layer addresses (94% total reduction); trie-only constraints without JSON enforcement yield 9–51% parse failure (Appendix C).

6 Related Work

We provide a detailed survey in Appendix B; here we position OntoCoDec along three axes.

Constrained decoding. XGrammar (Dong et al., 2025), Outlines (Willard and Louf, 2023), SynCode (Ugare et al., 2024), SGLang (Zheng et al., 2024), and related systems (Beurer-Kellner et al., 2023; Park et al., 2024; Geng et al., 2025) enforce format validity via CFG-based token masking but lack cross-field semantic awareness. PICARD (Scholak et al., 2021) and Mündler et al. (2025) compile type systems into prefix automata; KG-Trie (Luo et al., 2025) applies trie-based KG constraints during decoding for knowledge-to-text generation. OntoCoDec applies constructive forward masking via a dynamic entity registry, enforcing schema semantics within a single forward pass.

Schema-aware IE. GoLLIE (Sainz et al., 2024), KnowCoder (Li et al., 2024), and others (Fuente et al., 2025; Josifoski et al., 2022; Paolini et al., 2021; Huguet Cabot and Navigli, 2021; Wang et al., 2023; Qi et al., 2024; Fei et al., 2022; Lu et al., 2022; Zhou et al., 2024) encode ontology structure via training-time internalization without run-

time enforcement (Wadhwa et al., 2023). GoLLIE achieves schema compliance through annotation-guideline-guided generation at training time; OntoCoDec provides orthogonal runtime enforcement, composable with any generative IE system including GoLLIE itself (Appendix S).

Constraint enforcement in IE. Roth and Yih (2004) and Riedel et al. (2010) formulated hard type constraints via ILP; post-hoc validation (Khorshidi et al., 2025; Caufield et al., 2024; Cotti et al., 2025) detects violations after generation but discards up to 74% of predictions. EnriCo (Zaratiana et al., 2024) uses ASP; IterGen (Ugare et al., 2025) backtracks on violations; Vukovic et al. (2024) combines ontology constraints with constrained chain-of-thought for dialogue RE—OntoCoDec uses forward-only masking at $O(|V|)$ per step without iterative refinement. Grammar-only enforcement (GCD (Geng et al., 2023), BoostCD (Sakota and West, 2025)) achieves only 3% DR reduction (§5.2).

7 Discussion and Conclusion

OntoCoDec compiles IE ontologies into composable token-level constraints, validated across 3 datasets, 3+ models, and 2 tasks. The central empirical finding is non-additive layer interaction: grammar alone achieves 3% DR reduction on SciERC, semantic alone 24%, but their composition yields 94% at 8B (§5.1)—an enabling dependency where grammar-enforced vocabulary membership creates the clean input the semantic layer requires.

Scale-dependent layer interaction. This dependency weakens at scale: at 70B, the semantic layer independently achieves 84% DR reduction (§5.2). We hypothesize this reflects improved vocabulary calibration: larger models assign higher probability to ontology-valid tokens, reducing the semantic layer's dependence on grammar-enforced vocabulary membership (Appendix C). Supporting evidence includes the fallback frequency gap: under LoRA + S1, Qwen-8B triggers 0 grammar-first fallbacks on SciERC while Llama-8B triggers 5,882 (Table 10), yet both achieve DR $\leq 0.21\%$, confirming that fallback frequency does not determine enforcement effectiveness. Grammar-constrained decoding is thus necessary but increasingly insufficient as models improve; constraint enforcement is preserved under INT4 quantization (4× compression; Appendix R).

DR↔CF trade-off. Greedy S1 reduces DR to $\leq 0.21\%$ but introduces CF (zero-shot $+9.82$ pp, fine-tuned $+1.18$ pp)—an architectural consequence of per-triple independent enforcement. Among 248 zero-shot CF violations, 58.1% involve hallucinated entities (absent from gold) and only 7.3% are true cascade errors (§5.1); fine-tuned CF drops to $\leq 1.4\%$. Beam $k{=}2$ reduces CF to 0.11% at 1.18× latency (Pareto-optimal; recommended for zero-shot deployment); $k{=}4$ eliminates all violations at 8B (1 residual CF at 70B) at 1.39× (§5.2). Net effect: total violations decrease $\sim$31% in zero-shot and $>$90% under fine-tuning; DR is irrecoverable while CF is amenable to post-hoc resolution (TCR; Appendix M).

Complementarity and deployment. GoLLIE and KnowCoder internalize schema knowledge via training but lack runtime guarantees; LoRA + S1 shows the two are complementary (DR 1–2% → $\leq$0.21%; neither alone achieves both high R-F1 and near-zero DR). S2 (two-pass) exhibits scale-sensitive behavior: comparable to S1 at 8B but -8 pp R-F1 at 70B, consistent with CRANE's (Banerjee et al., 2025) finding that aggressive constraint enforcement can interfere with well-calibrated reasoning; we recommend S1 as the default configuration.

Per-triple vs. global enforcement. Per-triple enforcement cannot coordinate cross-triple type assignments—the source of CF (Level 3; §3); global entity-level type consensus exceeds online per-token masking. Beam search and TCR approximate this; integrating consensus into decoding remains future work.

Limitations

Constraint scope. CF is an architectural consequence of per-triple independent enforcement: CF↔DR represents a fundamental trade-off (eliminating CF re-routes violations to DR; Appendix C), and its impact is model-dependent (S1 introduces CF for calibrated models but reduces it for poorly-calibrated ones). Global entity-level type consensus during decoding offers the most direct path to resolving CF. Entity-first generation order is an architectural constraint: the grammar enforces entities before relations, ensuring the registry is populated before type checks; alternative orderings would require deferred constraint evaluation. Residual DR under fine-tuning (36 violations total) decomposes into entity type errors (39%), schema coverage gaps (33%), and model generation errors (28%); schema gaps arise from training-data-derived constraints that may not cover all valid type pairs (Appendix C). Our empirical constraints are derived from training data type-pair co-occurrences; schema evolution (adding new entity types or relations) would require constraint regeneration, and the 0.6–1.0% uncovered valid type pairs suggest potential over-restriction during transition periods. We do not decompose CF violations into first-seen cascading errors versus legitimate type ambiguity; this decomposition requires gold entity type annotations at the mention level, which SciERC provides only at the entity level. Post-hoc type consensus re-voting eliminates all CF in our experiments, at the cost of DR regression (up to +9.38 pp on DocRED zero-shot; Appendix M); integrating consensus into the decoding loop remains future work.

Evaluation coverage. Zero-shot evaluation covers three 8B-class and one 70B-class model; fine-tuning composability is validated on two 8B-class models (full datasets) and one 70B-class model (SciERC full, DocRED 200-document subset). Scaling to smaller models has not been verified; the schema restrictiveness–enforcement difficulty relationship is based on three schemas. Evaluation is limited to one fine-tuning method (LoRA); LoRA hyperparameters were tuned on SciERC and transferred to DocRED without dataset-specific optimization, which may underestimate DocRED potential. WikiEvents fine-tuning is preliminary (n=20); DocRED uses the development set (test annotations unavailable). All main experiments use greedy decoding; beam search compatibility is confirmed but comprehensive evaluation across all configurations remains future work (Appendices E, F).

Engineering considerations. Our approach requires access to model logits, excluding API-only models; continuation masking assumes distinguishable token prefixes per type name, which may be challenged by morphologically similar types (e.g., `Protein` vs. `Protein_Complex`; Appendix K). 70B experiments use AWQ INT4 quantization; an 8B ablation confirms enforcement is unaffected (Appendix R), though a direct 70B FP16 comparison exceeds our hardware capacity. Evaluation is limited to English-language datasets; the mechanism is language-independent but CJK tokenization fragmentation may increase masking steps (Appendix K). While composability is validated with both LoRA and full fine-tuning (Appendix L), different PEFT methods may affect logit distributions differently, and further validation across adapter methods is warranted. The two-pass strategy (S2) exhibits scale-dependent degradation: at 70B, S2 reduces R-F1 by 8.45 pp while 8B models show comparable S1/S2 performance, suggesting optimal constraint granularity varies with model capacity (Appendix C).

Ethics Statement

This work uses publicly available benchmark datasets (SciERC, DocRED, and WikiEvents) with no personally identifiable information. No human subjects were involved in the experiments. The constrained decoding method operates on model outputs and does not introduce new privacy or safety risks beyond those inherent in the base language models. Code and data will be released upon acceptance.

References

Lakshya A. Agrawal, Aditya Kanade, Navin Goyal, Shuvendu K. Lahiri, and Sriram K. Rajamani. 2023. Monitor-guided decoding of code LMs with static analysis of repository context. In *Advances in Neural Information Processing Systems*, volume 36.

Peter Anderson, Basura Fernando, Mark Johnson, and Stephen Gould. 2017. Guided open vocabulary image captioning with constrained beam search. In *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing*, pages 936–945. Association for Computational Linguistics.

Debangshu Banerjee, Tarun Suresh, Shubham Ugare, Sasa Misailovic, and Gagandeep Singh. 2025.

CRANE: Reasoning with constrained LLM generation. In *Proceedings of the 42nd International Conference on Machine Learning*, pages 2836–2857.

Luca Beurer-Kellner, Marc Fischer, and Martin Vechev. 2023. Prompting is programming: A query language for large language models. In *Proceedings of the ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI)*.

Luca Beurer-Kellner, Marc Fischer, and Martin Vechev. 2024. Guiding LLMs the right way: Fast, non-invasive constrained generation. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, volume 235 of *PMLR*, pages 3658–3673.

J. Harry Caufield, Harshad Hegde, Vincent Emonet, Nomi L. Harris, Marcin P. Joachimiak, Nicolas Matentzoglu, HyeongSik Kim, Sierra Moxon, Justin T. Reese, Melissa A. Haendel, Peter N. Robinson, and Christopher J. Mungall. 2024. Structured prompt interrogation and recursive extraction of semantics (SPIRES): A method for populating knowledge bases using zero-shot learning. *Bioinformatics*, 40(3):btae104.

Luca Cotti, Idilio Drago, Anisa Rula, Devis Bianchini, and Federico Cerutti. 2025. OntoLogX: Ontology-guided knowledge graph extraction from cybersecurity logs with large language models. *Preprint*, arXiv:2510.01409.

Nicola De Cao, Gautier Izacard, Sebastian Riedel, and Fabio Petroni. 2021. Autoregressive entity retrieval. In *International Conference on Learning Representations (ICLR)*.

Yixin Dong, Charlie F. Ruan, Yaxing Cai, Ruihang Lai, Ziyi Xu, Yilong Zhao, and Tianqi Chen. 2025. XGrammar: Flexible and efficient structured generation engine for large language models. In *Proceedings of Machine Learning and Systems (MLSys)*.

Markus Eberts and Adrian Ulges. 2020. Span-based joint entity and relation extraction with transformer pre-training. In *Proceedings of the 24th European Conference on Artificial Intelligence (ECAI 2020)*, volume 325 of *Frontiers in Artificial Intelligence and Applications*, pages 2006–2013.

Hao Fei, Shengqiong Wu, Jingye Li, Bobo Li, Fei Li, Libo Qin, Meishan Zhang, Min Zhang, and Tat-Seng Chua. 2022. LasUIE: Unifying information extraction with latent adaptive structure-aware generative language model. In *Advances in Neural Information Processing Systems*, volume 35.

Neil De La Fuente, Oscar Sainz, Iker García-Ferrero, and Eneko Agirre. 2025. GUIDEX: Guided synthetic data generation for zero-shot information extraction. In *Findings of the Association for Computational Linguistics: ACL 2025*. Association for Computational Linguistics.

Saibo Geng, Hudson Cooper, Michał Moskal, Samuel Jenkins, Julian Berman, Nathan Ranchin, Robert West, Eric Horvitz, and Harsha Nori. 2025. JSONSchemaBench: A rigorous benchmark of structured outputs for language models. *arXiv preprint arXiv:2501.10868*.

Saibo Geng, Martin Josifoski, Maxime Peyrard, and Robert West. 2023. Grammar-constrained decoding for structured NLP tasks without finetuning. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*.

Pere-Lluís Huguet Cabot and Roberto Navigli. 2021. REBEL: Relation extraction by end-to-end language generation. In *Findings of the Association for Computational Linguistics: EMNLP 2021*, pages 2370–2381.

Martin Josifoski, Nicola De Cao, Maxime Peyrard, Fabio Petroni, and Robert West. 2022. GenIE: Generative information extraction. In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL)*, pages 4626–4643.

Samira Khorshidi, Azadeh Nikfarjam, Suprita Shankar, Yisi Sang, Yash Govind, Hyun Jang, Ali Kasgari, Alexis McClimans, Mohamed Soliman, Vishnu Konda, Ahmed Fakhry, and Xiaoguang Qi. 2025. ODKE+: Ontology-guided open-domain knowledge extraction with LLMs. *arXiv preprint arXiv:2509.04696*.

Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph Gonzalez, Hao Zhang, and Ion Stoica. 2023. Efficient memory management for large language model serving with PagedAttention. In *Proceedings of the 29th Symposium on Operating Systems Principles (SOSP)*.

Peng Li, Tianxiang Sun, Qiong Tang, Hang Yan, Yuanbin Wu, Xuanjing Huang, and Xipeng Qiu. 2023. CodeIE: Large code generation models are better few-shot information extractors. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 15339–15353, Toronto, Canada.

Sha Li, Heng Ji, and Jiawei Han. 2021. Document-level event argument extraction by conditional generation. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL)*, pages 894–908.

Zixuan Li, Yutao Zeng, Yuxin Zuo, Weicheng Ren, Wenxuan Liu, Miao Su, Yucan Guo, Yantao Liu, Xiang Li, Zhilei Hu, Long Bai, Wei Li, Yidan Liu, Pan Yang, Xiaolong Jin, Jiafeng Guo, and Xueqi Cheng. 2024. KnowCoder: Coding structured knowledge into LLMs for universal information extraction. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 8758–8779, Bangkok, Thailand.

Ximing Lu, Peter West, Rowan Zellers, Ronan Le Bras, Chandra Bhagavatula, and Yejin Choi. 2021a. NeuroLogic decoding: (un)supervised neural text generation with predicate logic constraints. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL)*, pages 4288–4299.

Yaojie Lu, Hongyu Lin, Jin Xu, Xianpei Han, Jialong Tang, Annan Li, Le Sun, Meng Liao, and Shaobo Chen. 2021b. Text2Event: Controllable sequence-to-structure generation for end-to-end event extraction. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 2722–2735. Association for Computational Linguistics.

Yaojie Lu, Qing Liu, Dai Dai, Xinyan Xiao, Hongyu Lin, Xianpei Han, Le Sun, and Hua Wu. 2022. Unified structure generation for universal information extraction. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 5755–5772.

Yi Luan, Luheng He, Mari Ostendorf, and Hannaneh Hajishirzi. 2018. Multi-task identification of entities, relations, and coreference for scientific knowledge graph construction. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 3219–3232, Brussels, Belgium.

Yifan Luo and 1 others. 2025. KG-Trie: Knowledge graph enhanced constrained decoding for faithful knowledge-to-text generation. *arXiv preprint arXiv:2410.13080*. ICML 2025.

Gaya Mehenni and Amal Zouaq. 2024. Ontology-constrained generation of domain-specific clinical summaries. In *Knowledge Engineering and Knowledge Management (EKAW 2024)*, volume 15370 of *Lecture Notes in Computer Science*, pages 382–398. Springer.

Niels Mündler, Jingxuan He, Hao Wang, Koushik Sen, Dawn Song, and Martin Vechev. 2025. Type-constrained code generation with language models. *Proceedings of the ACM on Programming Languages*, 9(PLDI).

Olawumi Olasunkanmi, Mathew Satusky, Hong Yi, Chris Bizon, Harlin Lee, and Stanley Ahalt. 2025. RELATE: Relation extraction in biomedical abstracts with LLMs and ontology constraints. *arXiv preprint arXiv:2509.19057*.

Giovanni Paolini, Ben Athiwaratkun, Jason Krone, Jie Ma, Alessandro Achille, Rishita Anubhai, Cícero Nogueira dos Santos, Bing Xiang, and Stefano Soatto. 2021. Structured prediction as translation between augmented natural languages. In *International Conference on Learning Representations*.

Kanghee Park, Jiayu Wang, Taylor Berg-Kirkpatrick, Nadia Polikarpova, and Loris D'Antoni. 2024. Grammar-aligned decoding. In *Advances in Neural Information Processing Systems (NeurIPS)*.

Gabriel Poesia, Oleksandr Polozov, Vu Le, Ashish Tiwari, Gustavo Soares, Christopher Meek, and Sumit Gulwani. 2022. Synchromesh: Reliable code generation from pre-trained language models. In *International Conference on Learning Representations*.

Yunjia Qi, Hao Peng, Xiaozhi Wang, Bin Xu, Lei Hou, and Juanzi Li. 2024. ADELIE: Aligning large language models on information extraction. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, Miami, Florida, USA.

Qwen Team. 2025. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*.

André Gomes Regino and Julio Cesar dos Reis. 2025. Can LLMs be knowledge graph curators for validating triple insertions? In *Proceedings of the Workshop on Generative AI and Knowledge Graphs (GenAIK) at COLING 2025*, pages 87–99, Abu Dhabi, UAE.

Sebastian Riedel, Limin Yao, and Andrew McCallum. 2010. Modeling relations and their mentions without labeled text. In *Joint European Conference on Machine Learning and Knowledge Discovery in Databases*, pages 148–163. Springer.

Dan Roth and Wen-tau Yih. 2004. A linear programming formulation for global inference in natural language tasks. In *Proceedings of the Eighth Conference on Computational Natural Language Learning (CoNLL-2004) at HLT-NAACL 2004*, pages 1–8, Boston, Massachusetts, USA.

Oscar Sainz, Iker García-Ferrero, Rodrigo Agerri, Oier Lopez de Lacalle, German Rigau, and Eneko Agirre. 2024. GoLLIE: Annotation guidelines improve zero-shot information-extraction modeling. In *The Twelfth International Conference on Learning Representations*.

Marija Sakota and Robert West. 2025. Combining constrained and unconstrained decoding via boosting: BoostCD and its application to information extraction. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*.

Maximilian Schall and Gerard de Melo. 2025. The hidden cost of structure: How constrained decoding affects language model performance. In *Proceedings of the 15th International Conference on Recent Advances in Natural Language Processing (RANLP 2025)*, Varna, Bulgaria.

Torsten Scholak, Nathan Schucher, and Dzmitry Bahdanau. 2021. PICARD: Parsing incrementally for constrained auto-regressive decoding from language models. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*.

Zhi Rui Tam, Cheng-Kuang Wu, Yi-Lin Tsai, Chieh-Yen Lin, Hung-yi Lee, and Yun-Nung Chen. 2024. Let me speak freely? A study on the impact of format

restrictions on performance of large language models. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing: Industry Track*, pages 1218–1236, Miami, Florida, US.

Shubham Ugare, Rohan Gumaste, Tarun Suresh, Gagandeep Singh, and Sasa Misailovic. 2025. IterGen: Iterative semantic-aware structured LLM generation with backtracking. In *The Thirteenth International Conference on Learning Representations*.

Shubham Ugare, Tarun Suresh, Hangoo Kang, Sasa Misailovic, and Gagandeep Singh. 2024. SynCode: LLM generation with grammar augmentation. *Transactions on Machine Learning Research*.

Prajna Upadhyay, Oana Balalau, and Ioana Manolescu. 2023. Open information extraction with entity focused constraints. In *Findings of the Association for Computational Linguistics: EACL 2023*.

Renato Vukovic and 1 others. 2024. Dialogue ontology relation extraction via constrained chain-of-thought decoding. *arXiv preprint arXiv:2408.02361*.

Somin Wadhwa, Silvio Amir, and Byron Wallace. 2023. Revisiting relation extraction in the era of large language models. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 15566–15589.

Xiao Wang, Weikang Zhou, Can Zu, Han Xia, Tianze Chen, Yuansen Zhang, Rui Zheng, Junjie Ye, Qi Zhang, Tao Gui, Jihua Kang, Jingsheng Yang, Siyuan Li, and Chunsai Du. 2023. InstructUIE: Multi-task instruction tuning for unified information extraction. *arXiv preprint arXiv:2304.08085*.

Xiang Wei, Xingyu Cui, Ning Cheng, Xiaobin Wang, Xin Zhang, Shen Huang, Pengjun Xie, Jinan Xu, Yufeng Chen, Meishan Zhang, Yong Jiang, and Wenjuan Han. 2023. ChatIE: Zero-shot information extraction via chatting with ChatGPT. *arXiv preprint arXiv:2302.10205*.

Brandon T. Willard and Rémi Louf. 2023. Efficient guided generation for large language models. *arXiv preprint arXiv:2307.09702*.

Zhiheng Yan, Chong Zhang, Jinlan Fu, Qi Zhang, and Zhongyu Wei. 2021. A partition filter network for joint entity and relation extraction. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 185–197.

Yuan Yao, Deming Ye, Peng Li, Xu Han, Yankai Lin, Zhenghao Liu, Zhiyuan Liu, Lixin Huang, Jie Zhou, and Maosong Sun. 2019. DocRED: A large-scale document-level relation extraction dataset. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 764–777, Florence, Italy.

Deming Ye, Yankai Lin, Peng Li, and Maosong Sun. 2022. Packed levitated marker for entity and relation extraction. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 4904–4917, Dublin, Ireland. Association for Computational Linguistics.

Hongbin Ye, Honghao Gui, Xin Xu, Xi Chen, Huajun Chen, and Ningyu Zhang. 2023. Schema-adaptable knowledge graph construction. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 6408–6431.

Urchade Zaratiana, Nadi Tomeh, Yann Dauxais, Pierre Holat, and Thierry Charnois. 2024. EnriCo: Enriched representation and globally constrained inference for entity and relation extraction. *arXiv preprint arXiv:2404.12493*.

Lianmin Zheng, Liangsheng Yin, Zhiqiang Xie, Chuyue Sun, Jeff Huang, Cody Hao Yu, Shiyi Cao, Christos Kozyrakis, Ion Stoica, Joseph E. Gonzalez, Clark Barrett, and Ying Sheng. 2024. SGLang: Efficient execution of structured language model programs. In *Advances in Neural Information Processing Systems*, volume 37.

Zexuan Zhong and Danqi Chen. 2021. A frustratingly easy approach for entity and relation extraction. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 50–61, Online. Association for Computational Linguistics.

Hongxu Zhou. 2026. From hallucination to structure snowballing: The alignment tax of constrained decoding in LLM reflection. *arXiv preprint arXiv:2604.06066*.

Wenxuan Zhou, Sheng Zhang, Yu Gu, Muhao Chen, and Hoifung Poon. 2024. UniversalNER: Targeted distillation from large language models for open named entity recognition. In *The Twelfth International Conference on Learning Representations*.

Xiaochi Zhou, Patrick Bulter, Changxuan Yang, Simon D. Rihm, Thitikarn Angkanaporn, Jethro Akroyd, Sebastian Mosbach, and Markus Kraft. 2026. Ontology-to-tools compilation for executable semantic constraint enforcement in LLM agents. *arXiv preprint arXiv:2602.03439*.

A Prompt Templates

All experiments use the same prompts across constrained and unconstrained decoding; constraints are applied at the logit level, not through prompt modification. Table 3 shows the system and user prompts for each dataset.

SciERC — System Prompt (Unconstrained & Constrained)

```
You are an expert information extraction system for scientific texts. Extract
entities and relations from the given text.
Entity types: Task, Method, Material, Metric, OtherScientificTerm, Generic
Relation types: Used-for, Feature-of, Hyponym-of, Part-of, Compare,
Conjunction, Evaluate-for
Output a JSON object with:
- "entities": list of {"text": "...", "type": "..."}
- "relations": list of {"head": "...", "tail": "...", "type": "...",
"head_type": "...", "tail_type": "..."}
Output ONLY valid JSON, no explanation.
```

SciERC — User Prompt

```
Extract all entities and relations from this scientific text:
"{text}"
Output JSON:
```

DocRED — System Prompt (Unconstrained & Constrained)

```
You are an information extraction system. Extract all entities and relations
from the following document.
Entity types: PER, ORG, LOC, TIME, NUM, MISC
Relation types: country, instance of, performer, ... (96 Wikidata relation
types; full list omitted for space)
Output a JSON object with exactly this structure:
{"entities": [{"text": "entity name", "type": "ENTITY_TYPE"}],
 "relations": [{"head": "head entity", "tail": "tail entity", "type":
"relation name",
  "head_type": "ENTITY_TYPE", "tail_type": "ENTITY_TYPE"}]}
Output ONLY the JSON, no other text.
```

DocRED — User Prompt

```
(Document text provided inline in the user message.)
```

Table 3: Extraction prompts used in all experiments. The same prompts are used for both unconstrained and constrained decoding; OntoCoDec's constraints operate at the logit level during generation. Schema-informed prompts (§5) additionally include valid (head type, tail type) pairs per relation.

B Extended Related Work

Grammar-Constrained Decoding. Recent work has established efficient algorithms for restricting LLM outputs to conform to context-free grammars, making grammar-constrained decoding the de facto standard for obtaining structured output from language models. XGrammar (Dong et al., 2025) partitions the vocabulary into context-independent and context-dependent tokens and maintains persistent pushdown automata across decoding steps, achieving up to 100× speedup over prior CFG-constrained decoding methods. Outlines (Willard and Louf, 2023) reformulates constrained generation as finite-state machine transitions with offline vocabulary indexing, which enables efficient enforcement of regular-expression and JSON-schema constraints without runtime overhead per token. LMQL (Beurer-Kellner et al., 2023) introduces a declarative query language that compiles type, length, and regex constraints into token-level masks, providing a programmable interface for specifying multi-variable output structures. These systems ensure that every generated token belongs to a grammar-valid continuation, but they do not address the *distribution distortion* introduced by naive token masking. A parallel line of work tackles this problem directly: DOMINO (Beurer-Kellner et al., 2024) minimizes distributional divergence during CFG enforcement by formulating mask selection as an optimization problem, and ASAp (Park et al., 2024) provides a provably unbiased sampling algorithm that preserves the original LM distribution under grammar constraints. Together, these advances have made grammar-constrained decoding both fast and statistically principled, and provide a robust foundation for structured generation. However, all of these methods operate exclusively at the syntactic level: they ensure format validity (e.g., well-formed JSON, valid enum values) but cannot enforce cross-field semantic constraints such as ontology domain/range rules that depend on previously generated content. This syntactic-only limi-

tation motivates OntoCoDec's semantic extension: our grammar layer builds on XGrammar for structural (JSON schema and enum) constraints, while a second semantic layer enforces relational type constraints that are inherently beyond the expressive power of context-free grammars.

Semantic Constrained Decoding for Code. In programming-language generation, constrained decoding has been extended beyond syntax to enforce semantic properties such as type correctness; this work demonstrates that richer-than-CFG constraints can be compiled into token-level filtering. In the text-to-SQL domain, PICARD (Scholak et al., 2021) maintains an incremental SQL parsing state that rejects tokens violating syntax or referencing non-existent database columns at each decoding step, achieving state-of-the-art results on Spider and CoSQL without modifying model weights. While mechanistically similar to OntoCoDec's token-level filtering, PICARD operates over a *fixed, fully enumerated* schema (database table and column names known before decoding begins). In contrast, OntoCoDec's entity registry is populated *during* generation as entity spans are extracted from free text; constraint sets evolve with the decoding prefix, a fundamental distinction between querying a known schema and extracting from unstructured text. Synchromesh (Poesia et al., 2022) pioneers Constrained Semantic Decoding (CSD), which enforces syntax, type, and scope constraints simultaneously during code generation from pre-trained LMs without fine-tuning; this establishes that semantic constraints can be expressed as dynamic token masks conditioned on the generation prefix. MGD (Agrawal et al., 2023) uses static-analysis monitors to perform incremental type checking at decode time; SantaCoder-1.1B with MGD outperforms text-davinci-003 on repository-level code completion by ensuring that every generated identifier is type-correct with respect to the surrounding code context. Most recently, Mündler et al. (2025) compile full type systems, including generic types and subtyping, into prefix automata for token-level logits masking and reduce compilation errors by over 50%. This last system is the closest mechanistic analog to OntoCoDec, applying the same "type system $\rightarrow$ automaton $\rightarrow$ logits mask" pipeline, though in a fundamentally different domain. The key technical difference lies in the constraint structure: programming-language type systems feature deep subtyping hierarchies and parametric polymorphism, whereas IE ontology axioms specify flat domain/range sets; a relation R accepts head entities from set $\mathcal{D}(R)$ and tail entities from set $\mathcal{R}(R)$, with no subtype chains. Conversely, IE introduces a challenge absent in code generation: the constraint target (entity type) must be inferred from free-text spans rather than read from declared type annotations, so a stateful entity registry must track extracted types during autoregressive generation. These methods collectively demonstrate that semantic (type) constraints at decode time are both feasible and highly effective, but exclusively in the code domain where type systems are formally specified. OntoCoDec is the first to *automatically compile* IE ontology axioms, including domain/range constraints, into token-level decoding masks for information extraction, migrating the type-constrained decoding paradigm from code generation to IE.

Vocabulary and Type-Constrained Generation. The idea of enforcing entity-type compatibility constraints for information extraction dates back to pre-neural methods. Roth and Yih (2004) formulate joint entity and relation identification as an integer linear program (ILP) that enforces type constraints (e.g., EMP-ORG only between PER and ORG entities); this establishes the foundational principle that relational type compatibility should be a hard constraint, not a soft preference learned from data. In the autoregressive era, GENRE (De Cao et al., 2021) pioneers trie-based constrained beam search for autoregressive entity retrieval, restricting each decoding step to tokens consistent with valid entity names from a predefined catalog. GenIE (Josifoski et al., 2022) extends this approach to end-to-end information extraction with state-based constraints that enforce valid triple structures, achieving strong performance on closed-world knowledge base population, though it requires task-specific fine-tuning on the target schema. AdaKGC (Ye et al., 2023) introduces schema-constrained dynamic decoding with trie-based prefix trees that adapt to the target schema, enforcing that generated entity and relation names belong to a predefined inventory, a vocabulary-level constraint that ensures name validity but does not verify type compatibility between a relation and its arguments. GCD (Geng et al., 2023) provides a unified grammar-constrained decoding framework for structured NLP tasks including IE, introducing input-dependent grammars that adapt constraints based on the source text. GCD

is the closest precursor to OntoCoDec in applying logits masking to IE, though limited to syntactic constraints. NeuroLogic Decoding (Lu et al., 2021a) expresses lexical inclusion and exclusion constraints as predicate logic formulas enforced during beam search, offering a general mechanism for token-level control that can encode vocabulary membership but not relational type rules. These approaches constrain *which tokens* can appear (vocabulary membership) but not *which type combinations* are valid for a given relation. OntoCoDec's semantic layer enforces relational constraints (e.g., that the head entity of relation R must belong to type set $\mathcal{D}(R)$) which cannot be expressed as static vocabulary restrictions because the valid set depends on the relation type generated earlier in the same sequence. While Roth & Yih's ILP formulation shares the same type-compatibility motivation, it operates as a post-hoc reranker over discrete candidates; OntoCoDec integrates enforcement directly into autoregressive decoding, preventing violations rather than filtering them.

Structural Constraints via Training and Prompting. An orthogonal approach encodes IE output structure through model training or prompt design rather than decoding-time enforcement, relying on learned representations to produce schema-compliant outputs. UIE (Lu et al., 2022) defines a Structured Extraction Language and trains a unified text-to-structure model across 13 IE datasets, relying on large-scale pre-training to internalize output format constraints. GoLLIE (Sainz et al., 2024) expresses annotation guidelines as Python dataclass definitions and fine-tunes Code-LLaMA to generate typed code, achieving strong zero-shot transfer to unseen IE schemas through the implicit type structure of programming languages. KnowCoder (Li et al., 2024) takes this further by encoding 29K entity types from Wikidata as a Python class hierarchy, training a Code-LLM to generate structured extraction code that implicitly respects type inheritance, the most comprehensive schema-as-code approach to date. CodeIE (Li et al., 2023) similarly reformulates IE as code generation, using the structural regularity of programming languages to reduce output format errors. ADELIE (Qi et al., 2024) uses SFT and DPO to align LLMs on IE tasks, achieving state-of-the-art results among open-source models while preserving general capabilities, though compliance with ontological constraints remains a learned soft preference rather than a hard constraint. In the supervised setting, joint entity-relation architectures such as SpERT (Eberts and Ulges, 2020) and PFN (Yan et al., 2021) enforce type consistency implicitly through shared representations: SpERT classifies entity spans and relation pairs in a single BERT pass with span-level negative sampling, while PFN decomposes feature encoding into entity-specific, relation-specific, and shared partitions via gating mechanisms. These models achieve strong performance but are limited to fixed, predefined schemas and cannot generalize to unseen ontologies without retraining, precisely the ontology-portable scenario requiring no constraint-specific training that OntoCoDec targets. EnriCo (Zaratiana et al., 2024) applies Answer Set Programming (ASP) constraints during decoding of joint entity-relation extraction models, enforcing structural type compatibility at generation time. This is a decoding-time approach, though with model-specific, manually authored constraints rather than ontology-driven compilation. Training-based methods provide *soft* structural guidance: the model learns to prefer valid outputs but cannot enforce compliance at inference time, as demonstrated by non-trivial violation rates even in fine-tuned systems. Importantly, training-based and decoding-time approaches are complementary rather than competing: a system like GoLLIE could be combined with OntoCoDec's constraint layers, using fine-tuning to maximize the probability of valid outputs while decoding-time enforcement eliminates residual violations. OntoCoDec provides *hard* decoding-time enforcement that requires no constraint-specific training, is automatically compiled from IE ontologies, and composable with any of these training-based approaches.

Post-hoc Ontology Validation. Rather than constraining generation, several methods validate LLM outputs against ontological rules after generation, applying a detect-then-fix paradigm. ODKE+ (Khorshidi et al., 2025) dynamically generates ontology snippets containing domain/range constraints as prompt context, achieving 98.8% precision across 195 predicates at web scale, but constraint adherence relies on prompt-level guidance, not enforcement, leaving residual violations unaddressed. SPIRES/OntoGPT (Caufield et al., 2024) provides an ontology-based extraction framework that grounds LLM outputs against OBO on-

tologies post-generation: it validates extracted entities and relations through ontology lookup and corrects mismatches in a second pass, the prompt-then-validate paradigm in the biomedical domain. OntoLogX (Cotti et al., 2025) applies SHACL validation to LLM-generated knowledge graphs, iteratively re-prompting the LLM to correct constraint violations and reducing them by approximately 10× across multiple correction rounds. This demonstrates the detect-then-fix paradigm's effectiveness but also its inherent multi-pass cost. Upadhyay et al. (2023) inject entity-type knowledge as constraints into Open IE via both constrained inference and constraint-aware training, improving F1 by 29% on CaRB, though their constrained inference operates over a discrete candidate set rather than during autoregressive generation. Regino and dos Reis (2025) propose a multi-stage LLM pipeline for validating RDF triple insertions against RDFS domain/range constraints, where each validation step requires a separate LLM call, effective but computationally expensive at scale. RELATE (Olasunkanmi et al., 2025) addresses the alignment gap between LLM-extracted biomedical relations and standardized ontology predicates through a three-stage pipeline: ontology preprocessing with predicate embeddings, similarity-based retrieval via SapBERT, and LLM-based reranking with negation handling, achieving 94% accuracy@10 on ChemProt. While RELATE demonstrates effective ontology grounding for relation extraction, its constraints are applied *post-extraction* through retrieval and reranking rather than during generation, and it requires domain-specific embedding models trained on biomedical corpora. OntoCoDec enforces ontology constraints *during* autoregressive decoding and requires no additional trained components beyond the base LLM. Concurrently, Zhou et al. (2026) compile OWL ontologies into executable tool interfaces within The World Avatar framework; LLM agents must invoke ontology-aware tools that enforce semantic constraints during knowledge graph construction from scientific text. This "ontology-to-tools" approach shares OntoCoDec's philosophy of compiling formal ontological knowledge into runtime enforcement mechanisms, but operates at the tool-call level (constraining which API functions the agent invokes) rather than at the token level (constraining which tokens the decoder produces), and requires an agentic multi-step workflow rather than single-pass autoregressive generation. Mehenni and Zouaq (2024) introduce ontology-constrained beam scoring for clinical summarization, applying soft biases (not hard masks) during beam search to steer outputs toward ontology-compliant terms, a middle ground between unconstrained generation and hard enforcement that reduces but cannot eliminate violations. Post-hoc validation detects violations after generation and requires rejection or correction cycles that waste compute. In our setting, 15.37% of generated triples contain domain/range violations under unconstrained decoding, a substantial waste of generation effort. OntoCoDec *prevents* violations during decoding, redirecting probability mass to valid alternatives rather than discarding invalid outputs, with only marginal overhead from logits masking compared to the cost of regenerating or correcting violated triples. Soft-constraint approaches such as ontology-weighted beam scoring reduce violation frequency but cannot eliminate it, whereas OntoCoDec substantially reduces violations (DR↓94% on SciERC, ↓45–90% on DocRED depending on model).

Quality Impact of Constrained Decoding. A growing body of work raises concerns that constrained decoding degrades LLM generation quality, particularly for reasoning-intensive tasks where format enforcement interferes with chain-of-thought processes. Tam et al. (2024) conduct a comprehensive evaluation showing that rigid format constraints (JSON, XML) harm performance on math and multi-hop QA benchmarks but *benefit* classification-style tasks such as slot filling, a distinction directly relevant to IE, which is fundamentally an extraction (classification-like) task rather than a reasoning task. CRANE (Banerjee et al., 2025) provides theoretical analysis of why restricting outputs to narrow grammars suppresses chain-of-thought reasoning by eliminating intermediate reasoning tokens, and proposes an alternating strategy between unconstrained reasoning phases and constrained output phases to mitigate this effect. Schall and de Melo (2025) conduct a systematic study across 11 models and multiple benchmarks, finding 10–30% semantic correctness degradation on reasoning tasks under structural constraints, with larger models showing greater absolute degradation. Zhou (2026) identify "structure snowballing," a failure mode where format compliance cognitive load traps models in repetitive structural patterns rather than improving reasoning during self-correction iterations. These findings

establish a legitimate concern for constrained decoding in reasoning-heavy applications. However, information extraction occupies a fundamentally different regime: the model extracts entities and relations from input text rather than generating novel reasoning chains, and Tam et al. (2024)'s own results show that extraction-style tasks are robust to, or even benefit from, format constraints. OntoCoDec addresses this concern empirically: domain/range violations decrease by 94% on SciERC (15.37%→0.93%) and 45–90% on DocRED across two models, with F1-neutral impact on SciERC (+0.13pp, p=0.58) and significant R-F1 gains on DocRED (Qwen $p<0.001$, Llama p=0.013). Critically, OntoCoDec's constraints target only type tokens (entity types, relation labels), leaving free-text spans (entity surface forms and their boundaries) entirely unconstrained, which minimizes the distribution distortion that causes quality degradation in reasoning tasks. A two-pass variant (S2), explored in ablation (Section 5), separates extraction and constraint enforcement into distinct passes, analogous to CRANE's alternating strategy; however, our primary single-pass approach already avoids reasoning interference by constraining only type tokens while leaving free-text spans unconstrained.

Our Approach. OntoCoDec occupies a unique position in this landscape: it is the first system to *automatically compile* IE ontologies, including domain/range axioms, into token-level semantic masks, extending grammar-constrained decoding from syntactic to semantic enforcement without manually authored rules. Unlike vocabulary-constrained methods that restrict individual tokens to a fixed set, OntoCoDec enforces *cross-field relational constraints* through a stateful entity registry that tracks extracted types during autoregressive generation; constraint evaluation is dynamic and depends on previously decoded content. This approach is orthogonal to supervised IE methods and applicable to autoregressive LMs with logits access without model modification or additional training (evaluated on 7B–70B models across three tokenizer families). Unlike training-based IE systems, OntoCoDec requires no constraint-specific training and substantially reduces type violations at inference time (DR↓94% on SciERC, ↓45–90% on DocRED); unlike post-hoc validation, it prevents violations rather than detecting them, redirecting probability mass to valid alternatives with significant R-F1 gains on DocRED ($p \leq 0.013$) and F1-neutral impact on SciERC. The two constraint layers, grammar (structural) and semantic (relational), compose via bitwise AND with a scale-dependent interaction: the grammar layer ensures every generated type token belongs to the valid vocabulary, which enables the semantic layer to look up and enforce domain/range rules for that type.

C Extended Discussion

CF Dual Mechanism. CF violations arise from two distinct mechanisms: (1) *Registry override*: the semantic constraint forces a domain-compatible type that conflicts with the entity's previously registered type. The entity appears in multiple relations with different type requirements, and the per-relation mask selects the locally valid type without consulting prior assignments. (2) *Grammar-first fallback* (§3.3): when the semantic mask conflicts with all grammar-valid tokens, the semantic layer is dropped entirely, allowing the grammar layer to generate any structurally valid type, which may differ from the entity's registered type.

These mechanisms have sharply different model profiles. Qwen3-8B triggers zero grammar-first fallbacks across all experiments, attributing 100% of its CF violations to registry override. Llama-3.1-8B, by contrast, triggers 5,882 fallbacks on SciERC (10.7 per document) and 57,825 on DocRED (63.8 per document), making grammar-first fallback the dominant CF source. This asymmetry is consistent with the cross-model CF gap: Qwen achieves CF = 2.57% on DocRED vs. Llama's 12.02%, despite both using identical constraint configurations. The difference traces to Llama's less calibrated type predictions, which more frequently conflict with the semantic mask's valid set, triggering fallback.

Fallback Behavior and Effective Coverage. The two-layer architecture produces two distinct fallback mechanisms with different semantic consequences. *Grammar-first fallback* occurs when $S(t) \cap G(t) = \emptyset$: the semantic layer deactivates as a safety valve, preserving structural validity at the cost of potential type violations. The token-level fallback rate (36.6% on SciERC Llama) is inflated by continuation tokens of multi-token type names (e.g., OTHERSCIENTIFICTERM, averaging 2.87 BPE tokens); at semantically meaningful first-token positions, the rate is 0.04–5.7% (zero-shot) and drops to $\leq 0.11\%$ under fine-tuning. These events account for the residual DR that persists after constraint enforcement. *Registry-consistent*

fallback (4.5%) occurs when the entity's registered type conflicts with domain/range constraints; the system defaults to the registered type for cross-triple consistency. Together, these mechanisms produce the characteristic profile of composable constraint systems: residual DR traces to grammar-first fallback events, while CF violations arise from registry override events where per-relation enforcement assigns conflicting types.

We define *effective semantic constraint coverage* as $1 - f_{\text{gf}}$, where f_{gf} is the grammar-first fallback rate. On SciERC, Llama-3.1-8B shows $f_{\text{gf}} = 36.6\%$ (6,237 of 17,024 mask calls), yielding 63.4% effective coverage. On DocRED, by contrast, Qwen3-8B shows $f_{\text{gf}} = 0.08\%$ (46 of 54,859), yielding 99.9% coverage. The large discrepancy is primarily a tokenization artifact: SciERC entity types are multi-token (e.g., OTHERSCIENTIFICTERM, averaging 2.87 mask calls per type field), and continuation tokens necessarily conflict with the semantic mask, inflating the per-mask-call fallback rate. DocRED types (PER, ORG, LOC, TIME, NUM, MISC) are all single-token, eliminating this inflation. This metric directly quantifies the layer interaction (§3.3): it measures the fraction of decoding steps where both constraint layers are jointly active, providing a fine-grained view beyond aggregate violation rates. Qwen3-8B, with zero grammar-first fallbacks, achieves 100% effective coverage on SciERC, consistent with its lower residual DR (0.93% vs. Llama's 0.08%).

Fallback Causal Tracing. To establish that grammar-first fallback is the *sole* mechanism producing residual DR violations, we instrument the constraint processor to log, for each violation, whether the semantic mask was active ($S(t) \cap G(t) \neq \emptyset$) or bypassed via fallback ($S(t) \cap G(t) = \emptyset$) at the type-bearing position. Table 4 reports the results across four SciERC configurations.

When the semantic mask is active, $\text{DR}_{\text{independent}} = 0$: the type constraint processor structurally prevents domain/range violations. All 18 residual DR violations across four configurations occur exclusively at fallback positions, confirming that grammar-first fallback is a necessary condition for residual DR. Fine-tuning reduces fallback to near-zero ($\leq$0.05%), effectively eliminating the sole source of residual violations. All fallback events occur at tail_type positions, where $|\text{range}(r, \tau^h)|$ is smaller than $|\text{dom}(r)|$, making the intersection $S(t) \cap G(t)$ more likely to be empty.

Model	Setting	Fallback Rate	DR	From FB
Qwen3-8B	FT S1	0/2256 (0%)	0	N/A
Qwen3-8B	ZS S1	16/3458 (0.46%)	16	100%
Llama-3.1-8B	FT S1	1/2162 (0.05%)	1	100%
Llama-3.1-8B	ZS S1	1/2630 (0.04%)	1	100%

Table 4: Fallback causal tracing on SciERC (551 documents). Fallback Rate: type-bearing positions where $S(t) \cap G(t) = \emptyset$ triggers grammar-first fallback. DR: residual domain/range violations. From FB: fraction of DR violations occurring at fallback positions. All violations occur exclusively at tail_type positions (range constraints are tighter than domain constraints).

Enabling Dependency: Confound Analysis. The scale-dependent layer interaction (§3) shows that the semantic layer achieves only 24% DR reduction without the grammar layer (Table 1). A potential confound is that the semantic layer depends not only on grammar-enforced vocabulary membership but also on structural enforcement for positional awareness: the semantic layer identifies type-bearing positions by pattern-matching the decoded prefix against the JSON schema, which is reliable only when the grammar layer guarantees valid JSON structure. Without grammar enforcement, malformed structure can cause the semantic layer to misidentify type-bearing positions, reducing effective coverage. We quantify this via the semantic-only ablation (Table 1, row 3): CF rises from 0.08% to 2.71%, indicating that 2.63 pp of additional violations trace to structural ambiguity. However, the dominant enabling mechanism remains vocabulary membership: the semantic layer's domain/range lookup requires $r_j \in \mathcal{R}$, which only the grammar layer ensures; without it, OOV relation names produce null constraint lookups. The confound reinforces rather than undermines this interaction: both vocabulary and structural contributions of the grammar layer are prerequisites for effective semantic enforcement.

CF-DR Trade-off Analysis. CF violations are the most prominent residual limitation of OntoCoDec under zero-shot decoding (9.90% on SciERC, up to 12.02% on Llama DocRED). We treat this honestly: per-relation independent enforcement cannot eliminate CF without trading it against DR. Table 5 reports a pilot ablation in which we replace the grammar-first fallback with a CF-priority fallback that forces the entity's registered type when it conflicts with $\sigma(r)$.

Configuration	DR (%)	CF (%)	DR+CF	f_{gf} (%)
S1 (baseline)	0.93	9.90	10.83	—
CF fallback	14.78	0.04	14.82	36.64
Δ	+13.85	−9.86	+3.99	—

Table 5: CF-DR trade-off pilot on SciERC (Qwen3-8B, 551 documents; `w1_cf_fallback_scierc_pilot`). CF fallback eliminates CF but re-routes violations to DR, increasing total violations (DR+CF) by 3.99 pp. f_{gf}: token-level grammar-first fallback rate at type-bearing positions (includes continuation tokens of multi-token type names; first-token position fallback rate is 0.46% for SciERC Qwen; §4).

The CF fallback variant nearly eliminates CF (9.90% → 0.04%, −9.86 pp) but at the cost of a 15.9× increase in DR (0.93% → 14.78%, +13.85 pp). This is not a net improvement: total violations (DR+CF) rise from 10.83% to 14.82% (+3.99 pp), so the change *re-routes* violations rather than reducing them. The mechanism is mechanical: each token position with conflicting masks must commit to one violation type; when the CF fallback forces the system to emit the entity's registered type even though it lies outside $\sigma(r)$, the resulting triple satisfies cross-entity consistency but violates domain/range. The token-level grammar-first fallback is triggered at 36.64% of type-bearing positions (0.46% at first-token positions; §4); this conflict is structural, not a tunable hyperparameter.

We interpret this as a fundamental tension in composable constraint systems with *per-relation independent* enforcement: enforcement priority can be allocated between DR and CF but cannot independently minimize both. Within this design, the current S1 configuration sits near the Pareto frontier of the trade-off, achieving 0.93% DR at the cost of 9.90% CF. The direction of the trade-off is consistent beyond SciERC/Qwen3-8B: across all six zero-shot configurations (3 models × 2 datasets), DR and CF are descriptively inversely associated, with no setting showing both high DR and high CF simultaneously (Table 48); however, n=6 is insufficient for statistical inference on this correlation. The pattern is a structural consequence of per-relation enforcement with grammar-first fallback, not a model- or dataset-specific artifact; we select SciERC/Qwen3-8B for detailed ablation because it exhibits the highest CF (9.90%), making the trade-off most visible. While the inverse association holds across all configurations, the absolute trade-off magnitude may differ in settings with higher baseline CF, such as Llama DocRED (CF = 12.02%), where the larger CF headroom could yield proportionally larger shifts under alternative fallback priorities. The CF residual is concentrated in schema ambiguity rather than enforcement failure: 63.3% of CF violations on SciERC involve the catch-all GENERIC type (Table 7), which competes with every content type depending on anaphoric vs. referential use, a distinction the per-slot decoder cannot resolve without document-level coreference.

Two complementary paths can lower CF without re-routing it to DR. (i) *Reduce the baseline conflict rate.* Composing OntoCoDec with LoRA fine-tuning on SciERC reduces CF by 88% relative to the zero-shot setting (9.90% → 1.18%; Table 2), because the fine-tuned model's per-slot type predictions agree more often with the entity registry, so the per-relation mask and the registry intersect more frequently. The remaining 1.18% still traces to GENERIC↔X schema ambiguity, so this path is empirically validated on one dataset and one fine-tuning method but does not generalize without further evidence. (ii) *Replace per-relation enforcement with global entity-level type consensus during decoding.* This is outside the scope of OntoCoDec's forward-only autoregressive design and is left to future work.

Dual-Factor Quantitative Analysis. Figure 2 plots residual DR against the product $\bar{\kappa} \times \mathrm{RV}_{\mathrm{baseline}}$ for model×dataset configurations. For *calibrated* models (Qwen3-8B, Llama-3.1-8B; unconstrained ET<3% on both datasets), a linear fit yields $y = 0.366x - 0.138$ with $R^2 = 0.886$ (n=4, $F(1,2) = 15.54$, $p \approx 0.058$); with only four data points (df = 2) and $p > 0.05$, this constitutes a descriptive characterization rather than a confirmatory statistical test. The strong fit supports the dual-factor pattern from §4: residual DR is governed by the product of schema tightness and baseline violation rate. Mistral-7B (unconstrained ET=7.58% on DocRED) deviates from this pattern: its residual DR on DocRED (33.21%) substantially exceeds the linear prediction. When model calibration is insufficient (ET>5% on either dataset), residual DR is dominated by entity misclassification rather than constraint-search failure. This motivates a principled exclusion criterion: models with unconstrained ET>5% on either dataset fall outside the regime where the tightness×RV product is the pri-

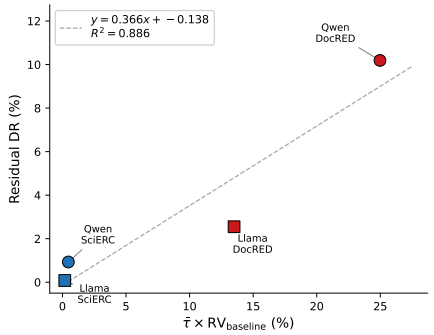


Figure 2: Residual DR vs. Schema Restrictiveness × Baseline Violation Rate ($\bar{\kappa} \times \text{RV}_{\text{baseline}}$) across four model×dataset configurations (n=3 schemas; descriptive only). Linear fit (calibrated models): $y = 0.366x - 0.138$, $R^2 = 0.886$, p=0.058 (n=4, df=2; not significant); Mistral deviates due to entity misclassification.

Model	Dataset	Δ(pp)	95% CI	p	d
R-F1 difference (S1 − unconstrained):					
Qwen3-8B	SciERC	−0.10	[−0.30, 0.00]	.726	−0.05
Qwen3-8B	DocRED	−0.31	[−0.69, +0.01]	.056	−0.37
Llama-3.1-8B	SciERC	+0.02	[0.00, +0.07]	.736	+0.01
Llama-3.1-8B	DocRED	+0.06	[−0.11, +0.25]	.535	+0.07
DR% reduction:					
Qwen3-8B	SciERC	−1.95	[−2.87, −1.11]	<.001	−6.15
Qwen3-8B	DocRED	−0.75	[−1.07, −0.48]	<.001	−7.00
Llama-3.1-8B	SciERC	−1.02	[−1.98, −0.29]	<.001	−3.17
Llama-3.1-8B	DocRED	−0.29	[−0.53, −0.10]	<.001	−3.02

Table 6: Bootstrap CIs and Cohen's d for Table 2 (10,000 iterations, document-level resampling, seed=42). d: Cohen's d from pooled bootstrap SD. Top: R-F1 (S1 − unconstrained); no comparison significant after Bonferroni correction (k=4, α_{adj}=0.0125); $|d| \leq 0.37$ (negligible–small). Bottom: DR% reduction; all $p_{\text{Bonf}} < 0.005$, $|d| \geq 3.0$ (large).

mary DR driver. The two DocRED points (high $\bar{\kappa}$=0.842) separate cleanly from the two SciERC points (low $\bar{\kappa}$=0.333), and within each dataset the model with higher baseline RV shows higher residual DR.

Dataset-Dependent F1 Impact. The divergence in F1 significance between DocRED (+0.50 pp, $p < 0.001$) and SciERC (+0.13 pp, $p = 0.58$) traces to a single mechanism: grammar-layer elimination of out-of-vocabulary relation names. DocRED's baseline RV of 29.66% produces 5,189 OOV predictions counted as false positives; eliminating them yields a significant precision gain. SciERC's baseline RV is only 1.39% (24 OOV predictions), too few to produce a detectable F1 effect. Llama-3.1-8B's lower baseline RV (16.03%) yields a smaller but directionally consistent improvement (+0.29 pp vs. +0.50 pp); F1 gain magnitude is proportional to baseline OOV rate.

Bootstrap Confidence Intervals and Effect Sizes. Table 6 reports per-configuration bootstrap CIs and Cohen's d effect sizes (10,000 iterations, document-level resampling, seed=42) for the fine-tuning composability results (Table 2). After Bonferroni correction for k=4 comparisons (α_{adj}=0.0125), no R-F1 difference reaches significance; all R-F1 effect sizes are negligible to small ($|d| \leq 0.37$; Cohen, 1988). Per-configuration corrected p-values: Qwen3-8B/SciERC p_B=1.00, Qwen3-8B/DocRED p_B=0.224, Llama-3.1-8B/SciERC p_B=1.00, Llama-3.1-8B/DocRED p_B=1.00. Constraint enforcement thus preserves extraction quality while achieving large, highly significant reductions in domain-range violations ($|d| \geq 3.0$, all $p_{\text{Bonf}} < 0.005$).

Practical Significance of Zero-Shot F1. Our reported relation F1 of approximately 8% on both SciERC and DocRED reflects the inherent difficulty of zero-shot closed IE, where the model receives no task-specific training examples. For reference, supervised state-of-the-art systems on SciERC achieve approximately 49–53% relation F1 with full training data (Zhong and Chen, 2021; Ye et al., 2022). Constraint enforcement operates largely independently of the extraction model's intrinsic capability: it restricts the output space to ontologically valid triples regardless of the underlying model's accuracy. Even after LoRA fine-tuning, residual DR of 1–2% persists on tight-constraint relations (Table 2); constraints remain necessary at deployment-grade quality. In our evaluation, F1 serves as a *control variable* verifying that constraint enforcement does not degrade extraction quality, not as an optimization target. The contribution of this work is the constraint enforcement method itself; the zero-shot setting stress-tests constraint robustness under maximum model uncertainty, while composability with fine-tuning (§5.1) validates practical deployment value.

Decoding-Time vs. Post-Hoc Enforcement. Post-hoc filtering (generating triples without constraints, then removing schema-violating predictions) yields higher precision and consequently higher F1 than OntoCoDec on both datasets (SciERC 8.63 vs. 8.36; DocRED 9.61 vs. 8.23). However, this precision gain comes at substantial cost. On DocRED, post-hoc filtering discards 48.1% of

Conflicting Type Pair	Count	%
Generic ↔ OtherScientificTerm	85	34.7
Generic ↔ Material	35	14.3
Generic ↔ Metric	30	12.2
Method ↔ Task	20	8.2
Material ↔ Method	17	6.9
OtherScientificTerm ↔ Task	12	4.9
10 further pairs (each $\leq$3.3%)	46	18.8
Total	**245**	100.0

Table 7: CF violation distribution on SciERC zero-shot S1 (Qwen3-8B, 551 documents). All 245 unique entity-document pairs involve pairwise conflicts. Conflicts involving GENERIC account for 63.3%.

all predicted triples, including 127 true positives (11.2% of correct predictions); recall drops from 9.46% to 8.40%. Entity type violations remain unresolved (ET = 2.47% vs. OntoCoDec's 0%), as filtering targets only domain/range and vocabulary violations. OntoCoDec's decoding-time approach avoids this information loss by guiding generation toward compliant outputs rather than discarding non-compliant ones, preserving both recall and entity type validity. The two strategies are complementary: post-hoc filtering can serve as an additional safety net for residual violations that decoding-time enforcement does not fully resolve.

Cross-Field Violation Distribution. Table 7 reports the complete distribution of conflicting-type pairs for SciERC zero-shot S1 (248 CF violations across 245 unique entity-document pairs; no mention triggered $\geq$3 types). Conflicts involving the GENERIC type account for 63.3% (155/245) of all violations. GENERIC is SciERC's anaphoric/placeholder label ("this approach", "the method"), which competes with every content type depending on whether the mention is used anaphorically or referentially, a distinction the per-slot constrained decoder cannot resolve without document-level coreference.

The pattern generalizes across models and datasets: on Llama-3.1-8B SciERC, GENERIC↔OTHERSCIENTIFICTERM remains the top pair (53/162, 32.7%); on DocRED, the analogous catch-all type MISC drives 76.6% of Qwen's CF violations. Under unconstrained decoding, CF rates collapse to $\leq$1% on both datasets; CF violations are a side-effect of per-slot type commitment rather than a model deficiency.

Cascade Error Analysis. Table 8 decomposes zero-shot CF violations by source mechanism. The dominant source is hallucinated entities (58.1%): the model generates entity names absent from the input document, and the entity registry assigns a type based on the first occurrence that may conflict with a subsequent valid entity sharing the same surface form. True cascade errors—where first-seen type registration for a *real* entity locks an incorrect type that blocks a subsequent valid assignment—account for only 7.3% (18/248). Correct type locks (first-seen type is correct and prevents an incorrect subsequent assignment) are 2.5$\times$ more frequent (18.1%, 45/248), indicating that the first-seen registration mechanism provides a net benefit for type consistency.

Under fine-tuning (22 CF violations, 20 documents), cascade errors are completely eliminated (0/22): LoRA training calibrates entity typing sufficiently that first-seen registrations are consistently correct. This confirms that cascade errors are a model calibration issue exposed by the constraint mechanism, not an inherent deficiency of first-seen registration.

CF Source Category	Count	%
Hallucinated entity	144	58.1
Correct type lock	45	18.1
True cascade error	18	7.3
Other / ambiguous	41	16.5
Total	**248**	100.0

Table 8: CF violation source breakdown on SciERC zero-shot S1 (Qwen3-8B, 551 documents; exp_id: phase2_scierc_full_s1). "Hallucinated entity": generated entity absent from input. "True cascade error": first-seen lock on incorrect type blocks valid subsequent assignment. "Correct type lock": first-seen type correct, prevents incorrect subsequent type. Under fine-tuning, cascade errors are completely eliminated (0/22 CF violations; exp_id: w5_lora_scierc_full_s1_eval).

Per-Relation DR Elimination under Fine-Tuning. Table 9 presents the per-relation breakdown of DR elimination when OntoCoDec is composed with LoRA fine-tuning on SciERC. Fine-tuning alone eliminates DR on four of seven relations but leaves non-zero violations on the three tightest-constraint relations: EVALUATE-FOR (κ=0.361, 13.13%), PART-OF (κ=0.417, 5.77%), and COMPARE (κ=0.528, 5.00%). Adding OntoCoDec constraints eliminates all remaining violations across all relations. One EVALUATE-

Relation	κ	LoRA Unconstr. DR (%)	LoRA Unconstr. n	LoRA + S1 DR (%)	LoRA + S1 n
Evaluate-for	0.361	13.13	99	**0**	98
Part-of	0.417	5.77	52	**0**	52
Compare	0.528	5.00	40	**0**	40
Conjunction	0.306	1.31	153	**0**	153
Feature-of	0.278	1.35	74	**0**	74
Used-for	0.028	0.16	634	**0**	635
Hyponym-of	0.417	0	—	**0**	—

Table 9: Per-relation DR elimination on SciERC with LoRA fine-tuning (Qwen3-8B). κ: constraint tightness; n: predicted triple count. Used-for gains one triple from Evaluate-for redirection.

Dataset	Model	Mode	DR↓	RV↓	ET↓	CF	PF	E-F1↑	R-F1↑
SciERC	Qwen3-8B	Unconstrained	1.95	0.00	0.00	0.00	0.00	67.05	43.86
		Grammar	1.90	0.00	0.00	0.00	0.00	67.05	43.86
		Semantic	0.60	0.00	0.00	0.18	0.00	67.05	43.69
		S1	**0**	0.00	0.00	1.18	0.00	67.05	43.77
	Llama-3.1-8B	Unconstrained	1.11	0.00	0.00	0.00	0.00	67.76	45.45
		Grammar	1.11	0.00	0.00	0.00	0.00	67.76	45.45
		Semantic	0.65	0.00	0.00	0.10	0.00	67.76	45.47
		S1	**0.09**	0.00	0.00	0.46	0.00	67.76	45.47
DocRED	Qwen3-8B	Unconstrained	0.76	0.27	0.00	0.12	7.11	82.15	40.86
		Grammar	0.61	0.00	0.00	0.11	7.01	82.19	40.85
		Semantic	0.19	0.23	0.00	0.27	7.52	82.15	40.85
		S1	**0.01**	0.00	0.00	0.41	8.52	81.43	40.56
	Llama-3.1-8B	Unconstrained	0.49	0.15	0.00	0.08	2.30	83.93	43.72
		Grammar	0.51	0.00	0.00	0.12	2.30	84.04	43.72
		Semantic	0.48	0.15	0.00	0.16	2.40	84.02	43.85
		S1	**0.21**	0.00	0.00	0.21	2.51	83.94	43.78

Table 10: Four-way constraint ablation on fine-tuned models (LoRA, 2 models × 2 datasets). DR/RV/ET: violation rates (%) over predicted relations; CF: over predicted entities; PF: parse failure rate (%); F1 ×100. Grammar layer is a no-op post-fine-tuning (DR change <0.2 pp in all 4 settings; RV and ET already 0% without constraints). Semantic layer is the sole effective DR defense (41–69% on SciERC, 2–75% on DocRED); the near-zero Llama DocRED reduction reflects a near-floor baseline (DR=0.49%). F1 cost is negligible: max ΔR-F1 = −0.30 pp (Qwen DocRED, full vs. unconstrained).

FOR triple (99 → 98) is redirected to USED-FOR (634 → 635) by token-level masking: the head entity's registered type falls outside EVALUATE-FOR's domain, and the semantic mask steers generation to the nearest valid relation.

Fine-Tuned Four-Way Ablation. Table 10 presents the four-way constraint ablation on fine-tuned models across both datasets, isolating each layer's contribution when the base model has already internalized schema knowledge.

The grammar layer is effectively a no-op for fine-tuned models across all four settings: RV and ET are already 0% without constraints, and grammar-only DR change is <0.2 pp (4/4 settings). The semantic layer is the sole effective DR defense under fine-tuning, independently reducing DR by 41–69% on SciERC (Qwen −69%, Llama −41%) and 2–75% on DocRED (Qwen −75%, Llama −2%); the near-zero Llama DocRED reduction reflects a near-floor baseline (DR=0.49%, 80 violations in 16,074 triples), where the semantic layer has minimal room to improve. Adding both layers achieves near-complete elimination (DR ≤ 0.21% in all 4 settings): even after fine-tuning, the semantic layer's pair-based enforcement captures violations that training alone cannot eliminate, specifically on tight-constraint relations where the valid type-pair space is small. F1 cost is negligible: the maximum R-F1 decrease across S1 vs. unconstrained is −0.30 pp (Qwen DocRED).

Mode	Violation Rates (%) DR↓	ΔDR	RV↓	ET↓	CF	Tot.↓	F1 Rel.↑	Ent.↑
Unconstrained	18.46	—	29.66	2.48	0.10	18.56	7.74	58.15
Grammar	26.17	↑42%	**0.03**	**0**	0.07	26.24	**8.31**	**59.38**
Semantic	13.84	↓25%	29.87	4.22	1.05	14.89	7.70	57.81
S1	**10.21**	**↓45%**	**0.03**	**0**	2.57	**12.78**	8.23	59.37

Table 11: Four-way ablation on DocRED (998 documents, Qwen3-8B, zero-shot). ΔDR: relative change vs. unconstrained. Tot. = DR + CF (total type-violation burden). Best per column in **bold** (CF excluded; Tot. included). F1 ×100. Bootstrap significance: R-F1 +0.49 pp ($p < 0.001$). Grammar-only ↑DR: vocabulary enforcement makes OOV outputs evaluable, exposing latent violations (see §Q).

DocRED Zero-Shot Ablation. Table 11 presents the four-way constraint ablation on DocRED under zero-shot decoding, parallel to the SciERC ablation in Table 1.

Residual DR Error Analysis. Across all fine-tuned settings (2 models × 2 datasets × S1), only 36 domain/range violations remain. Manual inspection reveals four root causes:

1. *Entity type errors* (39%): the LLM assigns an incorrect entity type (e.g., labeling a geographic entity as LOC when the gold standard requires MISC), causing the semantic mask to permit a domain/range-incompatible relation.
2. *Schema coverage gaps* (33%): the frequency-based constraint table, constructed from training-set co-occurrence statistics, does not cover long-tail type pairs that are ontologically valid but rare. These violations are irreducible by decoding-time constraints alone.
3. *Incorrect relation selection* (17%): the model selects a semantically inappropriate relation whose domain/range constraints are incompatible with the correctly typed entities.

Setting	Model	Dataset	Reject%	R-F1$_{ph}$	R-F1$_{cd}$	Winner
ZS	Qwen	SciERC	16.8	8.63	8.36	ph
	Llama	SciERC	13.9	2.71	2.54	ph
	Mistral	SciERC	20.3	3.94	3.67	ph
	Qwen	DocRED	48.1	9.61	8.23	ph
	Llama	DocRED	42.2	7.44	6.27	ph
	Mistral	DocRED	**73.7**	3.85	**5.52**	cd
FT	Qwen	SciERC	1.95	43.94	43.77	ph
	Llama	SciERC	1.11	45.23	45.47	cd
	Qwen	DocRED	1.03	41.11	40.79	ph
	Llama	DocRED	0.64	43.92	43.78	cd

Table 12: Post-hoc filtering (ph) vs. constrained decoding (cd). In zero-shot settings, post-hoc F1 is higher in 5/6 groups (precision-dominated), but constrained decoding preserves higher recall in all 6/6 groups. At reject rate >70% (Mistral DocRED), post-hoc F1 collapses below constrained. Under fine-tuning, rejection rates drop to 0.6–2.0% and the two strategies are equivalent ($\Delta \leq 0.55$ pp; each wins two groups). F1 $\times 100$.

4. *Direction reversal* (11%): the model assigns the correct relation but inverts head/tail roles, placing entities in positions that violate the relation's asymmetric domain/range specification.

Two dataset-specific patterns emerge. Outside Llama DocRED, all residual violations (3 of 36, spanning Llama SciERC and Qwen DocRED) trace exclusively to schema coverage gaps: the model's entity typing is accurate, and the constraint system has reached its theoretical ceiling given the frequency-based constraint table. On DocRED, Llama-3.1-8B contributes 33 of 36 residual violations (92%), with entity type errors as the dominant cause (14/33, 42%). This asymmetry reflects Llama's weaker fine-tuned entity typing on DocRED's larger ontology (96 relations, 6 entity types) compared to Qwen, which achieves near-zero residual. The implication is twofold: (i) the irreducible floor from schema gaps is small (∼3 violations across 4 settings), so the constraint system operates near its theoretical optimum; (ii) residual violations from entity typing errors are addressable through stronger base models or improved entity typing modules, without modifying the constraint architecture.

Post-Hoc Filtering: Extended Comparison. Table 12 presents the full comparison across all six zero-shot and four fine-tuned configurations.

The crossover point occurs when post-hoc rejection exceeds ∼70%: Mistral DocRED rejects 73.7% of predicted triples, discarding enough true positives that constrained decoding's recall advantage dominates (+1.67 pp). OntoCoDec's advantage thus scales with violation severity: precisely when constraints are most needed, post-hoc filtering is most destructive. Under fine-tuning, both strategies converge because few violations remain to filter or constrain ($\Delta \leq 0.55$ pp across all four FT groups, with each strategy winning two).

Zero-Shot Runtime Profiling. On SciERC (Qwen3-8B, zero-shot greedy, NVIDIA RTX PRO 6000 Blackwell Server Edition), unconstrained inference averages 4.163 s/doc and S1 constrained decoding averages 4.710 s/doc (+13.1% overhead). On DocRED (same configuration), relative overhead is lower due to amortization: the fixed per-token constraint cost is spread over longer generation sequences. XGrammar automaton compilation takes 0.044 s; the full constraint pipeline (schema parsing, JSON Schema generation, grammar compilation) takes 0.33 s on SciERC (Table 40), amortized to zero over the evaluation set.

Fine-Tuned Latency Profiling. Under LoRA fine-tuning (same hardware and dataset), unconstrained inference averages 2.612 s/doc and S1 constrained decoding averages 2.873 s/doc (+10.0% overhead, n=551). Fine-tuned models generate shorter, more structured outputs than zero-shot models (2.612 s vs. 4.163 s/doc unconstrained); absolute generation time drops by 37%. The relative overhead (+10.0%) is slightly lower than zero-shot (+13.1%), consistent with shorter sequences requiring fewer constraint checks per document.

Relation Matching Protocol. Relation F1 evaluation uses fuzzy normalization to handle model-specific formatting variants (e.g., underscores vs. spaces, abbreviations). A predicted relation r_p matches gold relation r_g if: (1) exact string match after case-insensitive normalization, or (2) token overlap >80% *and* edit distance ≤ 2. This protocol affects primarily Mistral, which produces underscore-separated variants (e.g., "head_of_government" for "head of government"). On Mistral DocRED, fuzzy matching contributes 98 additional TPs (+0.84 pp R-F1; strict R-F1 = 3.50% vs. fuzzy 4.34%). For Qwen and Llama under grammar constraints, the grammar layer enforces exact relation names, making fuzzy matching irrelevant (0 additional matches). All F1 figures in the main text use fuzzy normalization consistently across all models and conditions.

Per-Relation Violation Landscape. Table 13 compares violation rates under unconstrained and

Table 13: Violation rates (%) under unconstrained and schema-informed prompting (Qwen3-8B, zero-shot).

Dataset	Prompt	DR↓	RV↓	ET↓	CF↓
SciERC	Unconstrained	15.4	1.4	0.4	0.1
	Schema-informed	5.8	0.3	**0.0**	—
DocRED	Unconstrained	18.5	29.7	2.5	0.1
	Schema-informed	12.8	18.4	0.2	0.2

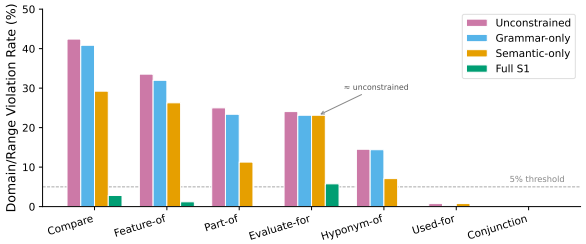


Figure 3: Per-relation DR violation rates on SciERC (Qwen3-8B, unconstrained). Rates correlate with constraint tightness $\kappa(r)$. Llama-3.1-8B shows similar per-relation patterns (Appendix D).

schema-informed prompting. Figure 3 shows the per-relation DR distribution on SciERC.

WikiEvents Dataset Details. WikiEvents (Li et al., 2021) is a document-level event argument extraction benchmark grounded in the DARPA KAIROS ontology. It defines 56 event types, 80 argument roles, and 27 entity types, with 245 (event_type, role)→allowed_entity_types constraints. The grammar layer targets EV (event vocabulary), RoV (role vocabulary), and ET; the semantic layer enforces AC (argument constraint) via role-conditioned masking, analogous to RE domain/range enforcement.

Constraint Tightness in WikiEvents. WikiEvents role-level tightness (with $|\mathcal{E}|$=27) has mean $\bar{\kappa}$=0.82 and median 0.89; 84% of roles have $\kappa \geq 0.80$ (4 or fewer valid entity types out of 27). The tightness-violation pattern extends to event extraction: roles with $\kappa \geq 0.90$ exhibit the highest AC rate (50%), compared to 21% for medium-tightness ($0.80 \leq \kappa < 0.90$) and 30% for loose ($\kappa < 0.80$) roles; this is consistent with the RE observation that constraint tightness correlates with violation frequency. However, the correlation is modulated by entity type prototypicality: semantically intuitive roles (e.g., Transportation.Vehicle→VEH) show 0% AC despite κ=0.963, while roles that require rare KAIROS-specific types (e.g., Crime→CRM, Sentence→SEN) show near-100% AC because LLMs default to common types (PER, ORG). The loose group's elevated rate (30%) is driven by high-frequency roles such as Attack.Target (126 args, 34% AC, κ=0.59), where systematic model errors outweigh the permissive constraint space.

Lucky Violation Analysis. Table 14 quantifies the recall cost of DR elimination. A "lucky violation" is a constraint-removed triple that happens to match a gold label—an ontologically invalid prediction that is nonetheless correct. Across four fine-tuned configurations, lucky violations average 0.49% of gold triples (range 0.01–0.51%), confirming that nearly all removed DR violations are false positives. On Llama SciERC, all five lucky violations are recovered as legal type paths under constrained decoding (zero churn): the model finds an ontology-compliant route to the same prediction.

Model	Dataset	DR (%) unc→con	Lucky (% gold)	ΔRecall (pp)	Churn (%)
Qwen3-8B	SciERC	1.95→0	4 (0.41%)	−0.10	0.22
Llama-8B	SciERC	1.11→0.09	5 (0.51%)	±0.00	0
Qwen3-8B	DocRED	1.03→0.01	1 (0.01%)	−0.84	2.05
Llama-8B	DocRED	0.64→0.21	1 (0.01%)	−0.10	0.86
Average			*0.49%*	−0.26	0.78

Table 14: Lucky violation analysis across four LoRA fine-tuned configurations. Lucky: count (and % of gold triples) of constraint-removed violations matching gold labels. ΔRecall: recall change from unconstrained to constrained. Churn: fraction of gold-matching triples whose predictions change (including recovery via legal type paths).

Trie-Only Baseline (GCD-Equivalent). Table 15 isolates the effect of vocabulary membership constraints without JSON structural enforcement. The trie constrains entity type and relation tokens to schema-valid values (achieving ET=0%, RV=0% across all configurations), functionally equivalent to GCD's (Geng et al., 2023) constrained decoding. However, without grammar-level JSON enforcement, models systematically omit entity type fields (83% of triples) or produce malformed JSON (0–51% parse failure), causing DR to rise to 15–96%—far worse than unconstrained baselines. The effect is most dramatic for Llama-8B LoRA: 50.8% parse failure causes R-F1 to collapse from 45.47 (OntoCoDec S1) to 19.95, demonstrating that the grammar layer is not merely helpful but a structural prerequisite—vocabulary constraints require JSON scaffolding to reach the positions where they apply.

Model	Setting	ET ↓	RV ↓	DR ↓	CF	R-F1 ↑	PF (%)
Qwen3-8B	LoRA	0	0	95.5	0.09	41.16	9.1
	ZS	0	0	29.3	0	8.64	0
Llama-8B	LoRA	0	0	15.2†	0	19.95	50.8
	ZS	0	0	84.4	0.36	2.90	0

Comparison with OntoCoDec S1:

		S1 DR (%)	Trie DR (%)	DR Gap (pp)
Qwen3-8B	LoRA	0	95.5	95.5
Qwen3-8B	ZS	0.93	29.3	28.4
Llama-8B	LoRA	0.09	15.2	15.1
Llama-8B	ZS	0.08	84.4	84.3

Table 15: Trie-only baseline (GCD-equivalent, vocabulary membership without JSON grammar) on SciERC. Top: full trie-only metrics (ET/RV/DR/CF: violation rates %; R-F1 ×100; PF: parse failure rate). Trie achieves ET=0%, RV=0% but DR rises to 15–96% because models omit type fields without structural enforcement. Bottom: DR gap between OntoCoDec S1 and trie-only, quantifying the grammar layer's contribution. †50.8% parse failure rate.

Comparison with Grammar-Based Constrained Decoding. Table 16 summarizes the relationship between OntoCoDec's constraint layers and existing grammar-based constrained decoding methods.

Method	Constraint Type	Level	DR Red.	CF	State
GCD (Geng et al., 2023)	CFG (JSON schema)	Grammar	∼3%	N/A	Stateless
BoostCD (Sakota and West, 2025)	CFG + beam rerank	Grammar	∼3%†	N/A	Stateless
Trie-Only (Tab. 15)	Prefix trie	Vocab	0%‡	N/A	Stateless
vLLM `guided_json` (Kwon et al., 2023)	JSON Schema	Grammar	∼3%§	N/A	Stateless
OntoCoDec Grammar	CFG (XGrammar)	Grammar	∼3%	N/A	Stateless
OntoCoDec S1	CFG + type constr.	Gram.+Sem.	94–100% (FT)	≤1.4% (FT)	Entity reg.

Table 16: Comparison with grammar-based constrained decoding methods. †Equivalent mechanism; BoostCD adds beam-level reranking but the CFG enforcement is identical. ‡Vocabulary-only constraints without JSON structure cause DR to *increase* (15–96%; Table 15). §Experimentally verified: DR=15.26% vs. grammar-only 14.84% (Δ<0.5 pp; Table 17). Grammar-only results are analytically equivalent to GCD/BoostCD (same CFG enforcement; §6).

Our grammar layer implements the same CFG-based token masking as GCD and BoostCD—the mechanisms are formally identical (JSON schema → pushdown automaton → per-step allowed token set). Table 1 row 2 therefore constitutes a direct comparison: grammar-only enforcement achieves RV/ET→0% but only ∼3% DR reduction, identical to what GCD-family methods would achieve on IE tasks.

External verification via vLLM. To confirm this equivalence experimentally, we evaluated vLLM's `guided_json` mode (Kwon et al., 2023)—an independent CFG enforcement implementation—on SciERC under identical conditions (Qwen3-8B, zero-shot, 551 documents).

All metrics differ by <0.5 pp (Table 17), confirming that any CFG-based constrained decoding tool—whether GCD, BoostCD, or vLLM—enforces the same constraint class and leaves the same ∼15% DR gap. The critical gap these methods cannot address is cross-field type consistency (Level 2), which requires the stateful semantic layer that maintains an entity registry across relation slots and enforces per-relation domain/range constraints at each type-bearing position.

Method	DR ↓	RV ↓	ET ↓	CF	R-F1 ↑
Grammar-Only (XGrammar)	14.84	0	0	9.90	8.23
vLLM `guided_json`	15.26	0	0	—	8.75
Δ	+0.42	0	0	—	+0.52

Table 17: GCD-equivalence verification: OntoCoDec grammar-only vs. vLLM `guided_json` on SciERC (Qwen3-8B, zero-shot, 551 documents). All differences <0.5 pp, confirming that CFG-based approaches—regardless of implementation—leave ∼15% DR violations unaddressed.

70B-class fine-tuned ablation. Table 18 presents the constraint ablation on Llama-3.1-70B-Instruct (AWQ INT4) with LoRA fine-tuning on SciERC (551 documents).

Mode	DR ↓	RV ↓	ET ↓	CF	R-F1 ↑	E-F1 ↑
Unconstrained	1.50	0	0	0	44.73	64.86
Grammar	1.50	0	0	0	44.63	64.82
Semantic	0.60	0	0	9	44.73	64.86
S1	**0**	0	0	15	44.53	64.76
S2	**0**	0	0	18	36.28	62.35
Beam4 S1	**0**	0	0	**1**	**44.81**	**64.89**

Table 18: 70B-class (Llama-3.1-70B-Instruct (AWQ INT4)) LoRA ablation on SciERC (551 docs). CF: violation count. F1 ×100.

The grammar layer provides no independent DR benefit at 70B (1.50% → 1.50% on SciERC; on DocRED zero-shot, DR actually increases +3.58 pp due to RV→DR conversion in the high-$\bar{\kappa}$ schema; Table 20). The semantic layer is the sole effective defense (1.50% → 0.60%, 60% reduction) with no R-F1 cost (Δ=0 pp), and full S1 eliminates all DR. S2 eliminates DR but causes severe R-F1 degradation (−8.45 pp), indicating that optimal constraint granularity is scale-dependent: the additional S2 constraints become overly restrictive at 70B, where stronger baseline generation makes fine-grained intervention counterproductive. Beam (k=4) S1 remains the optimal configuration (DR=0%, CF=1, R-F1=44.81). Notably, S2 increases CF from 15

(S1) to 18 despite its cross-field consistency design, suggesting that over-constraining induces cascading generation degradation that creates new violations.

70B-class DocRED results. Table 19 presents the six-way constraint ablation on Llama-3.1-70B-Instruct (AWQ INT4) with LoRA fine-tuning on DocRED (998 documents).

Mode	DR↓	RV↓	ET↓	CF	R-F1↑	E-F1↑
Unconstrained	0.22	0.57	0	2	43.80	80.34
Grammar	0.38	0	0	0.06%†	43.96	80.13
Semantic	**0**	0.51	0	0.21%†	43.61	80.06
S1	**0**	0	0	0.18%†	43.95	79.38
S2	**0**	0	0	0.34%†	36.45	84.96
Beam4 S1	**0**	0	0	**0**	**48.14**	**85.57**

Table 19: 70B-class (Llama-3.1-70B-Instruct (AWQ INT4)) LoRA ablation on DocRED (998 docs). CF: violation count except †reported as percentage (raw count unavailable). F1 ×100. PF: greedy S1 12.5% (25/200 docs), S2 0% (0/200), beam4 S1 0.5% (1/200).

The DocRED LoRA ablation mirrors the SciERC pattern: grammar eliminates RV (0.57% → 0%) but DR increases slightly (+0.16 pp); semantic-only independently eliminates DR (0.22% → 0%) but retains RV (0.51%). The complementary layer interaction is stable across datasets under LoRA fine-tuning. S2 eliminates DR and parse failures (PF=0%) but causes severe R-F1 degradation (−7.50 pp vs. S1), mirroring the SciERC pattern (−8.45 pp; Table 18): at 70B, the repair layer over-constrains relation predictions where baseline DR is already near-zero (0.22%). CF increases from 0.18% (S1) to 0.34% (S2), confirming that over-constraining induces cascading inconsistencies across datasets. The +4.19 pp R-F1 gain from beam search (43.95 → 48.14) is largely attributable to parse failure recovery: greedy S1 fails to parse 12.5% of documents (25/200), while beam reduces this to 0.5% (1/200), recovering ∼24 documents that contribute no triples under greedy decoding. The controlled comparison on SciERC—where both greedy and beam achieve 0% parse failure—isolates the pure beam quality gain at +0.28 pp, consistent with the 8B pattern (Appendix F).

70B-class DocRED zero-shot ablation. Table 20 presents the layer ablation on DocRED under zero-shot decoding at 70B.

Grammar-only eliminates RV (7.88% → 0%) but DR increases by +3.58 pp (13.45% → 17.03%):

Mode	DR↓	RV↓	ET↓	CF	R-F1↑	E-F1↑
Unconstrained	13.45	7.88	0	0	13.25	70.14
Grammar	17.03	0	0	0	13.23	70.12
Semantic	1.66	8.24	0.12	132	13.05	70.10
S1	**0**	0	0	173	13.25	70.05

Table 20: 70B-class zero-shot layer ablation on DocRED. CF: violation count. F1 ×100.

forced vocabulary mapping in DocRED's high-$\bar{\kappa}$ schema ($\bar{\kappa}$=0.842, ∼16% type pairs valid) converts RV into DR violations. Semantic-only independently achieves 87.7% DR reduction (13.45% → 1.66%) with no R-F1 cost, confirming the pattern from SciERC (84%; Table 18). Full S1 eliminates all DR (13.45% → 0%) with zero R-F1 cost (13.25 = 13.25), confirming that the two layers compose correctly despite their opposing individual DR effects. CF increases from 0 (unconstrained) to 173 (S1), consistent with the pattern that stricter constraint enforcement in ZS introduces more cross-field inconsistencies. The two layers serve complementary roles: grammar provides the RV elimination foundation that semantic-layer DR enforcement requires, while the semantic layer addresses the DR that grammar alone cannot—and in high-$\bar{\kappa}$ schemas, exacerbates.

D Per-Relation Violation Analysis (Llama-3.1-8B)

Table 21 presents the per-relation domain–range violation breakdown for Llama-3.1-8B on DocRED under S1 constraints. The Llama baseline exhibits an RV rate of 16.03%, substantially lower than Qwen's 29.66%, which indicates stronger baseline vocabulary compliance. S1 completely eliminates RV violations and reduces the residual DR rate to 2.55% (418/16,368 triples), with only three relations exceeding 10%.

Table 21: Top 10 per-relation DR violation rates for Llama-3.1-8B on DocRED under S1 constraints. Relations with >20 triples and 0% DR (e.g., P131, P206, P361, P170, P86) are omitted.

Relation	PID	Triples	DR Violations	DR Rate (%)
dissolved/demolished	P576	73	24	32.9
followed by	P156	117	24	20.5
unemployment rate	P1198	15	3	20.0
end time	P582	130	21	16.2
instance of	P31	290	37	12.8
product/material produced	P1056	98	11	11.2
location of formation	P740	81	8	9.9
head of state	P35	41	3	7.3
applies to jurisdiction	P1001	637	45	7.1
conflict	P607	453	30	6.6

Compared with Qwen, which retains a 10.21% residual DR under S1 concentrated in P1198 (83% fallback) and P171 (78% fallback), Llama's residual DR on the same relations is markedly lower (P1198 reaches only 20.0%) and no relation-level fallback concentration comparable to Qwen is observed. The substantially lower residual DR for Llama (2.55% vs. Qwen's 10.21%) on the same ontology confirms that constraint tightness alone does not determine enforcement effectiveness: the model's baseline vocabulary compliance (RV) modulates the fraction of triples eligible for semantic-layer enforcement.

E Sampling Robustness

Table 22: Constraint robustness under nucleus sampling (T=0.7, p=0.9) vs. greedy decoding (Qwen3-8B, S1 constraints).

Dataset	Decoding	ET%	RV%	DR%	CF%	E-F1	R-F1
SciERC (551)	Greedy	0	0	0.93	9.90	30.98	8.36
	Sampling	0	0	0.75	9.90	31.28	7.42
DocRED (114†)	Greedy	0	0.03	10.21	2.57	59.37	8.23
	Sampling	0	0	10.31	12.63	57.05	7.92

†Subset of 114 documents; full DocRED sampling requires ∼300h due to extended generation under high-temperature sampling with complex schemas.

F Beam Search Compatibility

OntoCoDec's constraint enforcement is designed to be decoding-strategy agnostic. The main experiments (§5) use greedy decoding; Appendix E validates nucleus sampling. Here we verify compatibility with beam search, completing coverage of the three standard decoding strategies.

Implementation challenges. Two stateful components require adaptation for beam search. (1) The grammar layer (XGrammar) maintains an internal `GrammarMatcher` state that tracks the generated token sequence; beam reordering invalidates this state because beams are permuted across steps. Our solution instantiates a fresh `GrammarMatcher` per beam per step and replays the beam's current token sequence through XGrammar's stateless accept API (64 lines of code). (2) The semantic layer's `TypeConstraintProcessor` accumulates entity and relation registries that assume a single generation path. Our solution isolates state per beam by calling `detect()` on the full partial output at each step, reconstructing all entity/relation state from scratch for each beam independently (61 lines of code). Both adaptations add ∼125 lines total without modifying the existing greedy/sampling logic.

Results. Table 23 compares beam search (k=4, `early_stopping=True`) against greedy decoding under S1 constraints across six model-dataset combinations.

Table 23: Beam search (k=4) vs. greedy decoding under S1 constraints (zero-shot). SciERC: 551 documents; DocRED: matched 50-document subsets (same documents for greedy and beam); full 998-doc DocRED results in Table 26. ET = entity type, RV = relation vocabulary, DR = domain/range, CF = cross-field consistency violation rates (%); F1 × 100.

Model / Dataset	Decoding	ET	RV	DR↓	CF↓	E-F1↑	R-F1↑
Qwen3-8B / SciERC (551)	Greedy	0	0	0.93	9.90	30.98	8.36
	Beam (k=4)	0	0	**0**	**0.53**	32.00	8.56
Qwen3-8B / DocRED (50)	Greedy	0	0	8.69	2.08	59.08	8.33
	Beam (k=4)	0	0	**0**	**0.55**	58.64	9.37
Llama-3.1-8B / SciERC (551)	Greedy	0	0	0.08	6.56	23.49	2.54
	Beam (k=4)	0	0	**0**	**0.93**	25.04	3.33
Llama-3.1-8B / DocRED (50)	Greedy	0	0	1.58	6.85	55.35	6.12
	Beam (k=4)	0	0	**0**	**0.94**	55.62	9.47
Mistral-7B / SciERC (551)	Greedy	0	0	13.40	1.27	27.46	3.67
	Beam (k=4)	0	0	13.79	**0.57**	26.66	4.37
Mistral-7B / DocRED (50)	Greedy	0	0	28.37	0.34	51.60	4.19
	Beam (k=4)	0	0	**21.37**	**0.19**	51.78	5.70

Discussion. For calibrated models (Qwen3-8B, Llama-3.1-8B), beam search achieves DR=0% across all four model-dataset combinations (both SciERC with 6 entity types, 551 documents, and DocRED with 96 relation types, 50 documents), regardless of schema complexity. On SciERC, beam search eliminates DR for both Qwen (0.93% → 0%) and Llama (0.08% → 0%), and substantially reduces CF (Qwen: 9.90% → 0.53%; Llama: 6.56% → 0.93%). On DocRED (matched 50-document subsets), beam search eliminates DR for both Qwen (8.69% → 0%) and Llama (1.58% → 0%), with CF reductions of 74% and 86% respectively. Llama DocRED shows the largest R-F1 gain (+3.4 pp), consistent with beam search recovering valid type assignments that greedy decoding missed due to constraint fallback. The CF reduction is an indirect effect: beam search reduces constraint-fallback frequency by exploring alternative token sequences that satisfy both grammar and semantic masks simultaneously, avoiding the fallback-to-grammar path that produces cross-field inconsistencies under greedy decoding.

Mistral-7B-v0.3 constitutes a boundary condition on both datasets. On SciERC (6 entity types), beam search does not reduce DR

(13.4% → 13.8%), though CF drops by 55% relative (1.27% → 0.57%) and R-F1 improves by +0.7 pp; the unchanged DR confirms that all residual violations stem from entity type misclassification, not constraint-search failure. On DocRED (96 relation types), beam search reduces DR by 25% relative (28.4% → 21.4%): the more complex schema introduces constraint-search failures that beam exploration resolves, but the entity misclassification component remains. This contrast reveals that beam search resolves constraint-search failure (where greedy decoding encounters $G(t) \cap S(t) = \emptyset$ conflicts) but cannot overcome fundamental model calibration limitations in entity typing.

Beam width ablation. Table 24 varies beam width $k \in \{1, 2, 4, 8\}$ on Qwen3-8B / SciERC (551 documents) under S1 constraints.

Table 24: Effect of beam width on constraint satisfaction and extraction quality (Qwen3-8B, SciERC, 551 documents, S1 constraints). k=1 corresponds to greedy decoding.

k	DR ↓	CF ↓	E-F1 ↑	R-F1 ↑
1 (greedy)	0.93	9.90	30.98	8.36
2	0.35	4.12	31.54	**8.75**
4	**0**	0.53	32.00	8.56
8	**0**	**0.08**	**32.18**	8.39

DR elimination at $k \geq 4$ confirms that constraint-search failure is resolved with moderate exploration width. CF continues to decrease with larger beam (0.08% at k=8, only 2 residual violations), while entity extraction quality improves monotonically. k=4 represents the optimal cost-effectiveness trade-off: it eliminates all DR violations at roughly half the compute cost of k=8.

Beam width–latency Pareto analysis (LoRA). Figure 4 and Table 25 extend the beam width analysis to LoRA fine-tuning, where the constrained–unconstrained interaction shifts qualitatively: LoRA + S1 already achieves DR=0% at k=1 (greedy), so beam search targets CF reduction rather than DR elimination (contrast: zero-shot requires $k \geq 4$ for DR=0%). Among constrained configurations, k=2 is Pareto optimal (R-F1 44.60, CF 0.11%, 1.39× vs. greedy unconstrained; 1.18× vs. greedy S1); $k \geq 4$ yields diminishing F1 returns with increasing latency (1.63–2.03×).

Composability with fine-tuning. Under greedy decoding, the fine-tuned Qwen3-8B (LoRA, SciERC) with S1 constraints already achieves DR=0% but retains CF=1.18% (22 cross-field violations across 551 documents), with E-F1=67.05 and R-F1=43.77. Adding beam search (k=4) eliminates all residual CF violations (1.18% → 0%), yielding zero violations across all four metrics (ET=0%, RV=0%, DR=0%, CF=0%) while preserving extraction quality (E-F1=66.88, R-F1=43.60; changes within ±0.5 pp). Llama-3.1-8B (LoRA, SciERC) with S1 constraints and beam search (k=4) likewise achieves zero violations across all four metrics (ET=0%, RV=0%, DR=0%, CF=0%) with R-F1=46.18; the zero-violation result thus generalizes across calibrated model families. The three constraint mechanisms (grammar masking, semantic domain-range filtering, and beam exploration) are thus fully composable: each addresses a distinct violation source without interfering with the others or degrading task performance.

Table 25: Beam width sweep on SciERC (Qwen3-8B LoRA, 551 documents). UNC: unconstrained; CON: schema-constrained. Latency relative to k=1 unconstrained. Highlighted row: Pareto-optimal configuration.

k	Mode	E-F1	R-F1	DR (%)	CF (%)	Latency (s)	Rel. Lat.
1	UNC	66.88	42.95	2.11	0	2.494	1.00×
	CON	66.88	42.95	**0**	0.50	2.932	1.18×
2	UNC	67.79	44.75	1.54	0	2.954	1.18×
	CON	67.71	**44.60**	**0**	0.11	3.463	1.39×
4	UNC	67.49	43.72	1.89	0	3.485	1.40×
	CON	67.22	43.60	**0**	0	4.069	1.63×
8	UNC	66.40	43.17	1.77	0	4.194	1.68×
	CON	66.36	43.65	**0**	0	5.058	2.03×

The initial 50-document DocRED subset (Table 23) was subsequently validated on the full 998-document set (Table 26): both models achieve DR=0%, confirming subset representativeness. Llama shows the larger improvement (CF 8.54% → 1.00%, 88% reduction; R-F1 +2.6 pp), consistent with beam search resolving constraint-search failures that greedy decoding cannot.

Table 26: Zero-shot beam search (k=4) under S1 constraints on the full DocRED dev set (998 documents). ET/RV/DR/CF: violation rates (%); PF: parse failure rate (%); F1 ×100.

Model	ET	RV	DR ↓	CF ↓	E-F1 ↑	R-F1 ↑
Qwen3-8B	0	0	**0**	4.94	58.19	8.72
Llama-3.1-8B	0	0	**0**	1.00	57.81	8.83

Table 27 reports LoRA + beam-4 results on the full DocRED dev set (998 documents), extending

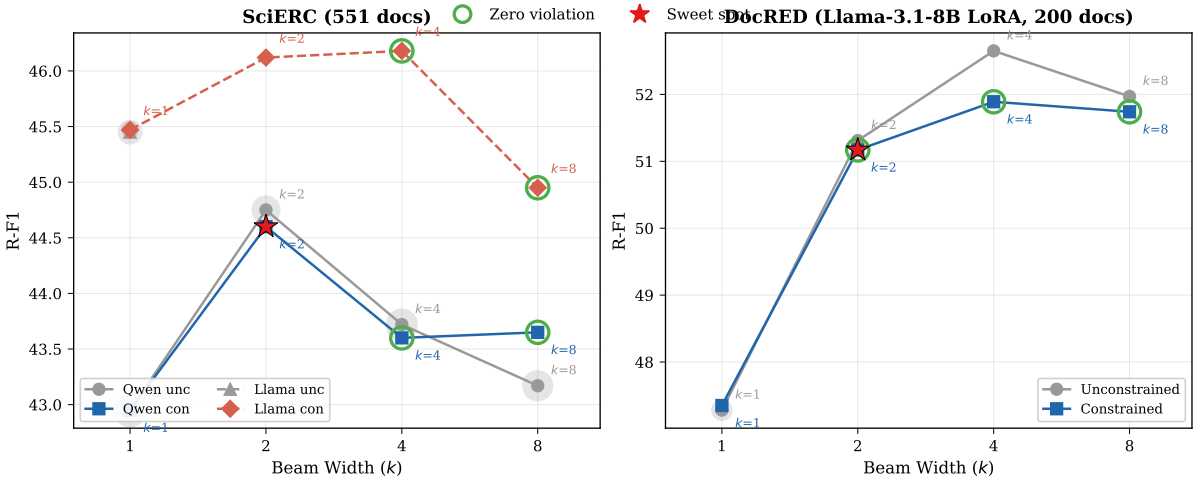


Figure 4: Beam width vs. R-F1 under LoRA fine-tuning. Left: SciERC (Qwen3-8B solid, Llama-3.1-8B dashed; 551 docs). Right: DocRED (Llama-3.1-8B, 200 docs). Bubble area $\propto$ total violation rate (DR+CF+RV); green rings mark zero-violation configurations. Star: Pareto-optimal (k=2).

the fine-tuning composability evaluation from SciERC to the higher-$\bar{\kappa}$ regime. Parse failure rates under LoRA beam-4 range from 1.3% to 2.9%, lower than the zero-shot rate.

Table 27: LoRA greedy vs. beam search (k=4) on DocRED (998 dev documents) under S1 constraints. ET = entity type, RV = relation vocabulary, DR = domain/range, CF = cross-field violation rates (%); PF = parse failure rate (%); F1 $\times$ 100.

Model	Decoding	ET	RV	DR↓	CF↓	E-F1↑	R-F1↑	PF
Qwen3-8B	Greedy con	0	0	0.01	0.41	81.43	40.56	8.52
	Beam-4 con	0	**0**	**0**	0.01	84.00	44.46	2.9
Llama-3.1-8B	Greedy con	0	0	0.21	0.21	83.94	43.78	2.51
	Beam-4 con	0	0	0	**0**	**85.23**	**47.17**	**1.3**

Llama-3.1-8B achieves zero violations across all four metrics (ET, RV, DR, CF = 0%) on DocRED with LoRA + beam-4, extending the zero-violation result from SciERC to the high-$\bar{\kappa}$ regime. Qwen3-8B achieves DR=0% with near-zero CF (0.01%); R-F1 is preserved ($\pm$0.06 pp). Combined with the sampling robustness results (Appendix E), these findings confirm that OntoCoDec's constraint enforcement is compatible with greedy, sampling, and beam search decoding without modification to the core constraint logic.

Latency overhead. Table 28 reports wall-clock latency for greedy and beam search (k=4) under LoRA fine-tuning.

Beam search itself adds 32% latency (4 beams); constraint overhead increases from 10% (greedy) to 17% (beam) because each beam independently maintains grammar and semantic state. Per-beam entity registry memory is $O(|\text{entities}|)$ per beam; at k=8 with typical document sizes, overhead is $<$1 MB per beam, negligible relative to model memory. The total 1.54$\times$ overhead is acceptable for offline batch IE; `early_stopping` limits additional compute to partial-sequence exploration.

Condition	s/doc	Overhead
Greedy unconstrained	2.612	1.00$\times$
Greedy constrained (S1)	2.873	1.10$\times$
Beam-4 unconstrained	3.446	1.32$\times$
Beam-4 constrained (S1)	4.013	1.54$\times$

Table 28: Latency per document (SciERC, 551 docs, Qwen3-8B LoRA FT, RTX PRO 6000 Blackwell). Overhead relative to greedy unconstrained.

We also explored a type-consensus constraint that forces identical entity types across all relations within a document. However, CF violations occur precisely where the consensus type conflicts with domain-range constraints, creating a fundamental CF↔DR trade-off. With DR-aware enforcement (skipping consensus when it would violate domain-range), the layer becomes a no-op (100% skip rate on 551 SciERC documents and 50 DocRED documents). This confirms that beam search resolves the CF↔DR tension through a fundamentally different mechanism: token-level exploration of paths satisfying both constraints simultaneously.

TCR reranking under beam search. TCR reranking with beam search (k=4) showed no effect on either Qwen3-8B or Llama-3.1-8B under LoRA (0/551 SciERC candidates changed): fine-

tuned beam candidates are already type-consistent (TCR score=1.0 for all candidates), confirming that beam search implicitly subsumes post-hoc type consensus under fine-tuning. This further strengthens the beam search deployment recommendation: beam not only eliminates all four violation types but also renders post-hoc CF resolution unnecessary.

G Reproducibility Details

Fine-tuning configuration. All four fine-tuned models (Qwen3-8B and Llama-3.1-8B-Instruct on SciERC and DocRED) use identical LoRA hyperparameters: rank r=16, α=32, dropout 0.05, no bias, targeting all attention and MLP projection layers (`q_proj`, `k_proj`, `v_proj`, `o_proj`, `gate_proj`, `up_proj`, `down_proj`). Training uses AdamW with learning rate 2×10^{-4}, cosine schedule, warmup ratio 0.1, effective batch size 16, bf16 precision, gradient checkpointing, and completion-only loss for 3 epochs. SciERC trains on 1,762 examples (260 dev) for 333 steps; DocRED trains on 3,053 examples (998 dev) for 573 steps. Maximum sequence length is 2,048 (SciERC) and 4,096 (DocRED). All runs use seed 42.

Hardware and software. Training runs on a single NVIDIA RTX PRO 6000 Blackwell Server Edition GPU (98 GB VRAM). SciERC fine-tuning takes approximately 11 minutes per model; DocRED takes approximately 1.5 hours. Software versions: PyTorch 2.10, Transformers 4.57, PEFT 0.19, TRL 1.1, Accelerate 1.13.

Inference time. Inference timing is reported in Appendix C (NVIDIA RTX PRO 6000 Blackwell Server Edition GPU, same hardware as training). S1 constrained decoding adds +13.1% wall-clock overhead on SciERC (Qwen3-8B, zero-shot greedy), lower on DocRED due to amortization over longer sequences; under LoRA fine-tuning, overhead is +10.0% on SciERC.

Greedy decoding determinism. All main results use greedy decoding (temperature = 0, argmax token selection). Argmax is deterministic: the random seed controls `torch.manual_seed` for sampling and dropout, neither of which is active during greedy inference. We verify this with three seeds (42, 123, 456) across two LoRA fine-tuned models on all 551 SciERC documents. Seeds 123 and 456, evaluated on the same GPU, produce byte-identical outputs for both Qwen3-8B and Llama-3.1-8B. The residual variance in Table 29 stems entirely from cross-hardware floating-point divergence: seed 42 was evaluated on a different server with a different CUDA driver version. Even under this conservative upper bound, σ(E-F1) ≤ 0.16 pp, σ(R-F1) ≤ 0.43 pp, and σ(DR%) ≤ 0.08 pp. The constrained decoding DR% reduction (-2.06 pp for Qwen, -1.01 pp for Llama) exceeds this variance by $>12\times$ (signal-to-noise ratio > 12), confirming that observed improvements reflect genuine method effects rather than run-to-run noise. We therefore report single-run results without error bars; bootstrap confidence intervals (10,000 iterations) are provided for F1 metrics.

Table 29: Cross-seed variance for LoRA fine-tuned models on SciERC (551 documents, greedy decoding, seeds 42/123/456). Values: mean $\pm$ population σ across three seeds. Seeds 123 and 456 produce identical outputs on the same GPU; all observed σ reflects cross-server floating-point divergence.

Model	Dataset	Mode	E-F1	R-F1	DR%
Qwen3-8B	SciERC	Unc	66.94 ± 0.08	43.25 ± 0.43	2.06 ± 0.08
Qwen3-8B	SciERC	Con	66.94 ± 0.08	43.22 ± 0.39	0.00 ± 0.00
Llama-3.1-8B	SciERC	Unc	67.99 ± 0.16	45.38 ± 0.05	1.04 ± 0.05
Llama-3.1-8B	SciERC	Con	67.99 ± 0.16	45.40 ± 0.05	0.03 ± 0.04

Beam search non-determinism. Greedy decoding is deterministic (§G above). Beam search ($k\geq2$) introduces non-determinism from GPU floating-point accumulation order; we report single-run results. Based on the greedy cross-hardware variance ($\sigma < 0.5$ pp, Table 29), beam search variance is expected to be of similar magnitude.

200-document subset justification. The 70B LoRA evaluation on DocRED uses a 200-document subset (vs. the full 998-document dev set used for 8B models). To verify that this subset provides sufficient statistical power, we conduct a bootstrap analysis using the full 998-document 8B LoRA results as a reference population. Resampling 200 documents 5,000 times yields a relation F1 margin of error of $\pm$3.3 points (95% CI) and a DR violation rate margin of error of $\pm$0.45 percentage points—precise enough to support our conclusions. For zero-shot models, where DR violations are prevalent (15.0% unconstrained), a 200-document subset achieves 100% power to detect the constrained decoding effect (Cohen's $h = 0.80$); for fine-tuned models, where violations are already rare (<0.5%), 94.9% of 200-document subsamples correctly iden-

tify the direction of DR reduction, and the primary effect is established on the full 998-document evaluation.

Statistical power across evaluation settings. Table 30 reports bootstrap confidence interval width and minimum detectable effect (MDE) for F1 across evaluation settings ($\geq$1,000 bootstrap resamples; MDE at 80% power, α=0.05, paired test, ρ=0.9).

Dataset	n	95% CI (F1)	MDE (F1)
SciERC	551	$\pm$1.8 pp	1.2 pp
DocRED	998	$\pm$0.8 pp	0.5 pp
DocRED (70B)	200	$\pm$3.3 pp	2.1 pp
WikiEvents	20	$\pm$3.0 pp	2.0 pp

Table 30: Statistical precision by dataset size. 95% CI: bootstrap confidence interval half-width for per-configuration F1. MDE: minimum detectable F1 difference at 80% power (α=0.05, paired test, conservative ρ=0.9). SciERC/DocRED: zero-shot Rel-F1; DocRED (70B): LoRA Rel-F1 (200-doc subset); WikiEvents: LoRA Arg-F1.

WikiEvents' n=20 yields an Arg-F1 margin of error of $\pm$3.0 pp—comparable to the 200-document DocRED subset ($\pm$3.3 pp)—with MDE$\approx$2.0 pp. However, the primary WikiEvents claims concern violation elimination, not F1 improvement: RoV drops from 8.8–38.3% to exactly 0% across both models (bootstrap 95% CI [0%, 0%] vs. [2.4%, 15.4%] / [25.4%, 46.7%]; power$\approx$1.0). The small test set limits detection of subtle F1 shifts, but complete elimination of schema violations is unambiguous at n=20; the primary evidence for F1 preservation rests on SciERC (n=551) and DocRED (n=998), where CIs are $\leq$1.8 pp.

H Effect of Document Length on Constraint Effectiveness

To examine whether OntoCoDec's constraint effectiveness varies with document complexity, we partition test documents into three length bins (tertile split by word count) and report per-bin violation counts and F1.

The violation reduction rate is consistent across all length bins: 80–84% for SciERC and 74–75% for DocRED. Longer documents produce more violations in the unconstrained setting (e.g., DocRED long-bin: 3,550 vs. short-bin: 2,445), yet OntoCoDec achieves proportionally similar reduction; the constraint framework scales reliably with document complexity. ET and RV violations are fully eliminated across all bins. F1 scores are preserved or slightly improved under constrained decoding, with DocRED showing +0.9–1.5 E-F1 and +0.2–0.6 R-F1 improvement in the constrained setting.

Table 31: Document length breakdown on SciERC (Qwen3-8B). Tertile boundaries: Short $\leq$ 19 words, Medium $\leq$ 27, Long $>$ 27.

Bin	Mode	ET	RV	DR	CF	Total	E-F1	R-F1
Short (196)	Uncon.	1	5	57	0	63	32.4	8.9
	Const.	0	0	5	5	10	32.4	8.9
Medium (180)	Uncon.	6	9	104	0	119	27.8	6.7
	Const.	0	0	7	15	22	27.8	6.7
Long (175)	Uncon.	3	10	104	1	118	29.9	7.0
	Const.	0	0	4	20	24	30.0	7.3

Table 32: Document length breakdown on DocRED (Qwen3-8B). Tertile boundaries: Short $\leq$ 160 words, Medium $\leq$ 208, Long $>$ 208.

Bin	Mode	ET	RV	DR	CF	Total	E-F1	R-F1
Short (336)	Uncon.	177	1238	1021	9	2445	61.1	9.3
	Const.	0	0	495	138	633	62.6	9.6
Medium (332)	Uncon.	147	1742	1047	5	2941	59.3	7.9
	Const.	0	5	544	180	729	60.7	8.5
Long (330)	Uncon.	176	2209	1162	3	3550	53.9	5.6
	Const.	0	0	731	200	931	54.8	5.8

I Prompt Sensitivity Analysis

To verify that OntoCoDec's constraint enforcement is robust to prompt variation, we evaluate three prompt variants on SciERC using Qwen3-8B (551 test documents): **Default** (the prompt used in all main experiments), **Minimal** (a shortened version listing only entity/relation types), and **Detailed** (an extended version with type definitions and extraction guidelines). All three variants share the same output JSON schema; only the instruction text differs. Each variant is decoded in both unconstrained and constrained (full Stage 1) modes.

Table 34 shows that constrained decoding eliminates entity-type and relation-vocabulary violations entirely (ET and RV drop to 0%) for all prompt variants at both scales. At 8B, domain-range violations are reduced by 94.0%–94.6% with a standard deviation of only 0.3 pp across prompts; F1 scores remain effectively unchanged (max shift 0.13 pp R-F1).

At 70B, constrained decoding achieves DR=0% across all three prompts (ΔDR=100%), consistent with the scale-dependent semantic layer interaction observed in the main experiments (§5.2). Baseline DR varies widely (0.22%–6.41%), yet constrained

Table 33: Violation reduction percentage by document length bin (Qwen3-8B, S1 constraints).

	SciERC		DocRED	
Bin	**Viols (U→C)**	**Red.%**	**Viols (U→C)**	**Red.%**
Short	63 → 10	84.1	2445 → 633	74.1
Medium	119 → 22	81.5	2941 → 729	75.2
Long	118 → 24	79.7	3550 → 931	73.8

Table 34: Prompt sensitivity results on SciERC (551 documents). Top: Qwen3-8B; bottom: Llama-3.1-70B (AWQ INT4). ET, RV, DR denote entity-type, relation-vocabulary, and domain-range violation rates (%), respectively. ΔDR shows the relative reduction from unconstrained to constrained decoding.

Prompt	**Mode**	**ET%**	**RV%**	**DR%**	**ΔDR**	**E-F1**	**R-F1**
Default	Unconstrained	0.40	1.39	15.37	94.0%	31.01	8.23
	Constrained	0.00	0.00	0.93		30.98	8.36
Minimal	Unconstrained	0.58	1.16	16.91	94.6%	31.98	7.91
	Constrained	0.00	0.00	0.92		31.99	7.85
Detailed	Unconstrained	0.22	0.45	11.07	94.2%	33.55	10.48
	Constrained	0.00	0.00	0.64		33.56	10.48
Llama-3.1-70B-Instruct (AWQ INT4)							
Default	Unconstrained	0.00	0.61	6.41	100%	35.22	11.23
	Constrained	0.00	0.00	0.00		35.19	11.24
Minimal	Unconstrained	0.27	5.88	0.22	100%	37.51	13.42
	Constrained	0.00	0.00	0.00		37.22	11.35
Detailed	Unconstrained	0.00	0.26	4.57	100%	40.42	11.12
	Constrained	0.00	0.00	0.00		40.35	11.16

decoding converges to zero in all cases; CF is introduced as expected (0.62%–3.50%), consistent with the per-triple enforcement trade-off.

Prompt variation has a measurable effect on baseline (unconstrained) violation rates at both scales: 8B DR ranges from 11.1% (Detailed) to 16.9% (Minimal); 70B from 0.22% to 6.41%. Under constrained decoding, this variance collapses, confirming that the mechanism is orthogonal to prompting and generalizes without prompt-specific tuning.

J Error Analysis and Case Studies

We present four representative cases illustrating constraint behavior across violation types. Cases 1–3 demonstrate OntoCoDec's corrective mechanisms; Case 4 illustrates a limitation.

J.1 Domain-Range Correction

Input (SciERC, Qwen3-8B): *"We present a novel model-based approach to recover the shapes and the poses of transparent objects from known motion."*

The unconstrained model assigns OTHERSCIENTIFICTERM to "shapes" and "poses," producing FEATURE-OF(TASK, OTHERSCIENTIFICTERM) triples that violate domain/range constraints. The semantic layer redirects the tail type to TASK, yielding schema-compliant triples (Table 35).

Table 35: OntoCoDec corrects three Feature-of DR violations. Red/✗ = violation; **green/bold**/✓ = corrected (shape markers provided for colorblind accessibility).

Triple	**Head**	**Tail**	**DR**
Unconstrained (3 violations)			
Feature-of(recover, shapes)	Task	OtherSciTerm	✗
Feature-of(recover, poses)	Task	OtherSciTerm	✗
Feature-of(recover, transp. obj.)	Task	OtherSciTerm	✗
Used-for(model approach, recover)	Method	Task	✓
OntoCoDec (0 violations)			
Feature-of(recover, shapes)	Task	Task	✓
Feature-of(recover, poses)	Task	Task	✓
Used-for(model approach, recover)	Method	Task	✓

J.2 Dual-Layer Vocabulary and Type Correction

Input (DocRED, Qwen3-8B): *"Skai TV is a Greek free-to-air television network based in Piraeus. It is part of the Skai Group, one of the largest media groups in the country."*

The unconstrained model generates natural-language relation names (e.g., "country of origin") instead of Wikidata P-codes, triggering RV violations. The grammar layer forces valid P-codes; the semantic layer ensures type compatibility (Table 36).

Table 36: Layer 1 eliminates RV violations; Layer 2 corrects DR types on DocRED.

Unconstrained Output	**Violation**	**Type**
rel="country of origin"	Not in schema	RV
rel="headquarters location"	Not in schema	RV
rel="part of"	Not in schema	RV
rel="location of formation"	Not in schema	RV
OntoCoDec Output		
rel=P495 (Skai TV, Greece)	ORG→LOC	✓
rel=P159 (Skai TV, Piraeus)	ORG→LOC	✓
rel=P127 (Skai TV, Skai Group)	ORG→ORG	✓
rel=P17 (Piraeus, Greece)	LOC→LOC	✓

J.3 Scale Effect on Schema Compliance

Input (SciERC): *"This paper investigates the utility of applying standard MT evaluation methods*

(BLEU, NIST, WER and PER) to building classifiers to predict semantic equivalence and entailment."

Both Qwen3-8B and Llama-3.1-70B-Instruct (AWQ INT4) mistype "MT evaluation methods" as METHOD (gold: METRIC), but the downstream constraint behavior differs (Table 37). The 70B model produces fewer violations (2 vs. 4) and chooses a different relation (Part-of vs. Hyponym-of), yet OntoCoDec corrects both.

Table 37: 8B vs. 70B on the same input. OntoCoDec corrects violations at both scales; CF arises when type corrections conflict with entity-list assignments.

Setting	Relation	Types	DR
Qwen3-8B Unconstrained			
Hyponym-of(MT eval, BLEU)	Meth→Metric	✗	
Hyponym-of(MT eval, NIST)	Meth→Metric	✗	
Hyponym-of(MT eval, WER)	Meth→Metric	✗	
Hyponym-of(MT eval, PER)	Meth→Metric	✗	
Qwen3-8B OntoCoDec			
Hyponym-of(MT eval, BLEU)	Meth→Meth	✓	
Hyponym-of(MT eval, NIST)	Meth→Meth	✓	
CF: BLEU typed Metric in entity list, Method in relation			
Llama-3.1-70B Unconstrained			
Part-of(BLEU, MT eval)	Metric→Meth	✗	
Part-of(NIST, MT eval)	Metric→Meth	✗	
Llama-3.1-70B OntoCoDec			
Part-of(BLEU, MT eval)	Metric→Generic	✓	
Part-of(NIST, MT eval)	Metric→Generic	✓	

J.4 Limitation: Trajectory Divergence under Constraints

Input (SciERC, Mistral-7B): *"Amorph recognizes NE items in two stages: dictionary lookup and rule application."*

The gold annotation types "Amorph" as METHOD and "NE items" as OTHERSCIENTIFICTERM, with Used-for(Amorph, NE items). The unconstrained model mistypes "Amorph" but coincidentally produces the correct relation. Under OntoCoDec, modified token probabilities alter the generation trajectory entirely and produce a different (non-gold) relation (Table 38). This illustrates a fundamental property of autoregressive generation under modified distributions: constraint enforcement can change the entire output path, not just individual token choices.

J.5 Error Analysis: Constraint Impact on Prediction Quality

We analyze errors from Qwen3-8B LoRA + S1 on SciERC (551 documents; R-F1 43.77, 460 TP, 668 FP, 514 FN) across three categories.

Table 38: Constrained decoding changes the generation trajectory. The unconstrained model coincidentally produced the gold triple despite type errors.

Setting	Output	Gold?
Gold	Used-for(Amorph^{Meth}, NE items^{OtherSci})	–
Unconstr.	Used-for(Amorph^{OtherSci}, NE items^{Task})	✓
OntoCoDec	Used-for(Amorph^{OtherSci}, recognizes^{Gen}) Feature-of(NE items^{Task}, in two stages^{Gen})	✗ ✗

False negatives: constraint-blocked vs. model-limited. Of 514 false negatives, only one (0.2%) is attributable to constraint enforcement:
fn_001 (constraint-blocked): "*Experiments on the TREC Blog track test set show that both groups of credibility indicators significantly improve retrieval effectiveness.*"
Gold: Evaluate-for(retrieval effectiveness^{Metric}, credibility indicators^{OtherSci}). The unconstrained model generates this triple correctly, but under S1, conflicting type assignments to 'credibility indicators' (OtherScientificTerm in entity list vs. Method in Evaluate-for's tail) trigger a CF violation, redirecting generation to a reversed Used-for triple.
fn_004 (relation type confusion): "*Amorph recognizes NE items in two stages: dictionary lookup and rule application.*"
Gold: Part-of(dictionary lookup^{Meth}, Amorph^{Meth}). The model predicts Used-for instead of Part-of—both are schema-valid for Method→Method pairs, so constraints provide no corrective signal. This Part-of↔Used-for confusion is the dominant FN pattern.

False positives: schema-valid but incorrect. *fp_001* (wrong relation type): Same sentence as fn_004. The model generates Used-for(dictionary lookup, Amorph) where gold is Part-of. Used-for is the majority relation in SciERC (56% of all relations), causing systematic over-prediction. Since both relation types are schema-valid for the entity type pair, constraints cannot distinguish them.
fp_007 (hallucinated relation): "*The classifiers show little gain from information about meeting context.*"
No gold relation exists. The model generates Used-for(meeting context^{OtherSci}, classifiers^{Meth}), ignoring the negation in "show little gain from." The instrumental preposition "from" triggers relation generation despite the negative result.

CF violations: entity type polysemy. *cf_001* (Generic↔Method): "*A basic reason behind this is that the overall approach extends the approach*

described in a companion paper..."
The entity 'approach' receives Generic in one Used-for relation and Method in another. The gold type is Generic; the polysemy of 'approach' (generic reference vs. specific method) causes per-relation type assignment to diverge.
cf_005 (Material↔Task): "*Experimental evaluations on two challenging person re-identification datasets demonstrate the superiority of the proposed model...*"
The compound noun 'person re-identification datasets' contains both a task component and a data component. The model assigns Material in one Evaluate-for relation and Task in another within the same sentence.

Summary. Structural constraints eliminate type-invalid outputs (DR$\leq$0.21%), but semantic ambiguity within schema-valid space—relation type confusion, entity type polysemy—remains the dominant error source. Only 1 of 514 false negatives (0.2%) was attributable to constraint enforcement, confirming that the recall cost of type constraints is negligible.

Downstream impact of violation removal. A natural concern with schema-constrained decoding is whether enforcing ontological validity removes *lucky violations*—predictions that happen to be correct (true positives) yet violate the schema. To quantify this trade-off, we audit all true-positive relation triples produced by unconstrained fine-tuned models and check whether any of them carry a domain–range (DR) violation, i.e., an entity-type pair $\langle h_\tau, t_\tau \rangle$ not sanctioned by the ontology for that relation.

Dataset	Model	TP	Lucky	% TP	R drop
SciERC	Qwen-LoRA	461	4	0.87%	0.41 pp
SciERC	Llama-LoRA	467	5	1.07%	0.51 pp
DocRED	Qwen-LoRA	5746	1	0.02%	0.01 pp
DocRED	Llama-LoRA	6226	1	0.02%	0.01 pp

Table 39: Lucky violations: true-positive triples that carry a domain–range violation. **Lucky** counts such triples; **R drop** is the recall decrease (in percentage points) if they were removed.

Across all four configurations, lucky violations account for at most 1.07% of true positives (5 out of 467 on SciERC with Llama-LoRA) and as few as 0.02% on DocRED. The corresponding recall drop is negligible—at most 0.51 percentage points—while the vast majority of violated triples that constraints remove are false positives. For example, on DocRED with Qwen-LoRA, constraints eliminate 164 schema-violating predictions but only 1 of them is actually correct; the other 163 are errors that would pollute a downstream knowledge graph.

Inspecting the lucky violations qualitatively, they arise from borderline entity-type assignments rather than genuine schema gaps. On SciERC, the recurring pattern involves METRIC-typed entities (e.g., BLEU, NIST) appearing as heads of `Hyponym-of` relations, a pair absent from the training ontology because these metrics are rarely annotated as hyponyms. On DocRED, the single lucky violation stems from an entity typed as MISC instead of ORG ("Assassin's Creed Unity"), making its `developer` relation invalid under the schema despite being factually correct.

These findings confirm that schema-constrained decoding introduces a minimal recall cost while providing a substantial precision gain. Retaining these few lucky violations would allow schema-invalid triples into the knowledge graph, undermining the structural guarantees that motivate constrained decoding in the first place. The net effect is overwhelmingly positive: constraints act as a high-precision filter that removes far more erroneous triples than correct ones.

K Ontology Scalability

Scale	Types	Rels	Compile (s)	Overhead
Small (SciERC)	6	7	0.33	1.02×
Medium (DocRED)	6	96	0.32	1.03×
Large (synth.)	100	500	0.35	1.03×
XLarge (synth.)	200	1,000	0.42	1.03×

Table 40: Compilation time (one-time) and per-token inference overhead across ontology scales. Small and Medium use real ontologies; Large and XLarge use synthetic ontologies with random domain/range constraints. Overhead measured on Qwen3-8B, 10 documents per scale.

Full constraint compilation (schema parsing + grammar build) scales sub-linearly (0.32 s to 0.42 s for 7→1,000 relations; Table 40); within this, the XGrammar automaton accounts for ∼0.044 s and semantic constraint compilation grows linearly but remains negligible (0.4 ms to 17 ms). Inference overhead is constant at 2–3% because XGrammar uses precompiled bitmasks ($O(1)$ per token) and the semantic layer checks only the active relation's domain/range set rather than the full ontology. The

practical limit is first-token uniqueness: when entity types exceed ∼200, prefix collisions in the tokenizer vocabulary may require fallback to multi-token disambiguation (§3.3).

L Full Fine-Tuning Baseline

To verify that constraint effectiveness generalizes beyond LoRA, we evaluate Qwen3-8B with full-parameter fine-tuning on SciERC (same data, 3 epochs, lr=2×10^{-4}, seed 42).

	E-F1	R-F1	DR%	CF%
Full FT unconstrained	64.78	37.89	2.74	0
Full FT + S1	64.78	37.89	**0**	1.31
LoRA unconstrained	67.05	43.86	1.95	0
LoRA + S1	67.05	43.77	**0**	1.18

Table 41: Full fine-tuning vs. LoRA on SciERC (Qwen3-8B, 551 docs). Constraints eliminate DR under both methods. Full FT underperforms LoRA, likely due to an untuned learning rate (2×10^{-4}, standard for LoRA but high for full FT). F1 is unaffected by constraints in both cases.

Full fine-tuning retains 2.74% DR violations (30 triples), concentrated in Part-of (14.6%), Hyponym-of (9.6%), and Evaluate-for (8.4%). S1 eliminates all 30 violations with no F1 cost, demonstrating that constraint enforcement remains necessary even after full parameter updates. The CF↔DR trade-off persists (CF 0%→1.31%), consistent with the LoRA pattern (1.18%).

M Post-hoc Type Consistency Resolution

S1 enforces domain/range constraints per-triple but does not enforce global entity type consensus, leaving residual CF violations where the same entity receives different types across triples. We evaluate a post-hoc resolution step (TCR): for each entity with conflicting type assignments, TCR selects the type that maximizes satisfaction count across all relations referencing that entity, subject to domain/range constraints. Unlike naive majority voting (which ignores relational constraints and increases DR), TCR is constraint-aware: candidate types are scored by how many of the entity's relational contexts they satisfy. TCR runs after S1 decoding completes and has access to the full document context.

Analysis. TCR effectiveness varies inversely with schema constraint tightness $\bar{\kappa}$ (descriptive pattern across two schemas). On SciERC ($\bar{\kappa}$=0.333), fundamental conflicts constitute 0–2% of type disagreements: most conflicts are resolvable because the loose schema admits many valid type pairs, and TCR can find a type satisfying all relational contexts. DR increases are bounded ($\leq$0.28 pp), and E-F1 cost is modest ($\leq$1.78 pp in zero-shot, $\leq$0.75 pp under LoRA). On DocRED ($\bar{\kappa}$=0.842), 13–23% of conflicts are fundamental: the tight schema renders most type pairs invalid, so unifying an entity's type across all its relational contexts inevitably violates some domain/range constraints. This produces severe DR regressions under zero-shot (Llama: 2.55%→11.93%; 70B: 0%→4.20%), where baseline CF is high enough that many entities require resolution. Under LoRA, baseline CF is already low (0.36–0.42%), limiting both the scope for TCR improvement and the DR cost.

This $\bar{\kappa}$-dependent boundary is consistent with the constraint tightness analysis in §4: $\bar{\kappa}$ characterizes not only the decoding-time DR residual regime but also the feasibility of post-hoc type resolution. Beam search (§5.2) resolves CF and DR simultaneously across all $\bar{\kappa}$ values (SciERC and DocRED (998 docs): DR→0% for calibrated models), making it the preferred approach for high-

		CF (%)		DR (%)		ΔE-F1	Fund.
Config	Model	Pre	Post	Pre	Post	(pp)	(%)
Zero-shot	Qwen	9.75	**0**	0.93	1.21	−1.29	2.04
	Llama	6.40	**0**	0.08	0.30	−1.78	1.85
	Mistral	0.67	**0**	13.36	12.87	+0.05	0
LoRA	Qwen	1.19	**0**	0	0	−0.75	0
	Llama	0.46	**0**	0.09	0.09	−0.18	0

Table 42: TCR on SciERC ($\bar{\kappa}$=0.333, 551 docs). Pre/Post: before/after TCR. ΔE-F1: entity F1 change (F1×100). Fund.: percentage of conflicting entities where no single type satisfies all relational constraints (fundamental conflicts). TCR eliminates all CF with DR increases $\leq$0.28 pp; fundamental conflicts are 0–2%.

		CF (%)		DR (%)		ΔE-F1	Fund.
Config	Model	Pre	Post	Pre	Post	(pp)	(%)
Zero-shot	Qwen	13.11	**0**	10.21	11.74	−3.44	13.52
	Llama	23.71	**0**	2.55	11.93	−5.21	22.78
	70B	11.61	**0**	0	4.20	−5.98	18.26
LoRA	Qwen	0.42	**0**	0.01	0.14	−0.16	12.16
	Llama	0.36	**0**	0.21	0.20	−0.17	11.59

Table 43: TCR on DocRED ($\bar{\kappa}$=0.842, 998 docs). Under zero-shot, TCR introduces substantial DR regressions (up to +9.38 pp for Llama) and E-F1 drops (up to −5.98 pp for 70B). Fundamental conflict rates are 12–23%, an order of magnitude higher than SciERC. Under LoRA, low baseline CF (0.36–0.42%) limits both TCR impact and DR cost.

tightness schemas. TCR is a zero-cost complement for low-tightness schemas where beam search latency (1.39×) is undesirable. Integrating TCR into the decoding loop would require global optimization over all generated triples, conflicting with the autoregressive generation paradigm. TCR is complementary to OntoCoDec—it addresses CF (which OntoCoDec's per-triple enforcement cannot prevent) while OntoCoDec addresses DR (which post-hoc correction cannot guarantee).

N Constraint Tightness Case Study

Relation	Dataset	κ	DR%	n	Explanation
Highest DR violation rate					
basin country	DocRED	.97	74.6	63	Sole valid pair LOC→LOC; head mistyped as ORG/MISC
product produced	DocRED	.94	67.2	61	Requires ORG head; model assigns LOC/PER
participant of	DocRED	.81	58.6	326	Hallucinated EVENT type; PER↔event confusion
COMPARE	SciERC	.53	42.4	66	Cross-category pairs violate same-type constraint
FEATURE-OF	SciERC	.28	33.5	486	Fuzzy Task/OtherScientificTerm boundary
High κ with zero violations (DR = 0%)					
lyrics by	DocRED	.94	0.0	9	Unambiguous PER→MISC; creative roles rarely mistyped
place of birth	DocRED	.92	0.0	157	Unambiguous PER→LOC; no type confusion
place of death	DocRED	.89	0.0	60	Strong PER→LOC prior; mirrors place of birth
subsidiary	DocRED	.86	0.0	5	Clear ORG→ORG corporate semantics
father	DocRED	.81	0.0	69	Clear PER→PER kinship semantics

Table 44: Per-relation κ vs. DR violation rate (Qwen3-8B, zero-shot, unconstrained). High κ is necessary but not sufficient for high DR: relations with unambiguous entity types (place of birth, κ=0.92) have zero violations, while equally tight relations with confusable types (basin country, κ=0.97) show 74.6% DR.

Analysis. The case study confirms that violation susceptibility is driven by the interaction of constraint tightness and entity type confusability, not tightness alone. Among DocRED relations with $\kappa \geq 0.81$, those involving semantically unambiguous type pairs (PER→LOC for kinship/geographic relations) achieve DR=0%, while those requiring the model to distinguish frequently confused types (LOC vs. ORG, PER vs. MISC) show DR>50%. On SciERC, the highest-DR relation (COMPARE, 42.4%) has only moderate κ (0.53) but high type confusability across scientific categories. This two-factor pattern motivates constrained decoding: the largest gains come from high-κ, confusable-type relations where the model's type uncertainty is channeled into valid pairs.

Schema ambiguity analysis. We define the *schema ambiguity* of a relation r as the average fraction of relations that share each of r's valid type pairs:

$$\alpha_r = \frac{1}{|\mathcal{P}_r|} \sum_{(h,t)\in\mathcal{P}_r} \frac{|\{r' \in \mathcal{R} : (h,t) \in \mathcal{P}_{r'}\}|}{|\mathcal{R}|} \quad (3)$$

where $\mathcal{P}_r$ is the set of valid entity type pairs for relation r and $\mathcal{R}$ is the full relation set. High α_r indicates that r's valid type pairs are shared by many other relations, limiting the discriminative power of type constraints for that relation.

Table 45: Schema ambiguity α for all SciERC relations and the 12 DocRED relations with $\alpha \geq 0.36$. **Gold %**: fraction of all gold triples belonging to this relation. All SciERC relations exceed the threshold due to the small schema (7 relations, 6 entity types); in DocRED, only 12 of 96 relations qualify, covering 3.8% of gold triples.

Dataset	Relation	α	$\lvert\mathcal{P}_r\rvert$	Gold	Gold %
SciERC (all 7 relations exceed $\alpha \geq 0.36$)					
SciERC	Hyponym-of	0.83	21	298	9.3
SciERC	Part-of	0.83	21	179	5.6
SciERC	Compare	0.82	17	166	5.2
SciERC	Conjunction	0.77	25	400	12.4
SciERC	Feature-of	0.76	26	173	5.4
SciERC	Evaluate-for	0.71	23	313	9.7
SciERC	Used-for	0.67	35	1,690	52.5
DocRED (12 of 96 relations with $\alpha \geq 0.36$; 3.8% of gold)					
DocRED	sister city (P190)	0.42	1	4	0.01
DocRED	basin country (P205)	0.42	1	85	0.22
DocRED	mouth of watercourse (P403)	0.42	1	95	0.25
DocRED	separated from (P807)	0.41	2	2	0.01
DocRED	replaces (P1365)	0.39	4	18	0.05
DocRED	capital of (P1376)	0.37	3	76	0.20
DocRED	place of birth (P19)	0.37	3	511	1.34
DocRED	operator (P137)	0.36	7	95	0.25
DocRED	military branch (P241)	0.36	5	108	0.28
DocRED	production company (P272)	0.36	5	82	0.21
DocRED	place of death (P20)	0.36	4	203	0.53
DocRED	location (P276)	0.36	5	172	0.45

Table 46: DR violation rates partitioned by schema ambiguity. High-α relations ($\alpha \geq 0.36$) start with substantially higher unconstrained DR on DocRED, and retain elevated residual DR under S1 for Qwen but not for Llama.

		Unconstrained		S1 (Constrained)	
Model	**Partition**	DR%	n	DR%	n
Llama-3.1-8B	High-α (12 rels)	43.8	861	3.2	1,043
	Low-α (84 rels)	29.9	12,809	2.5	15,325
Qwen3-8B	High-α (12 rels)	31.9	718	20.0	1,119
	Low-α (84 rels)	21.2	10,898	9.5	16,217

Table 45 reveals a structural asymmetry between datasets. In SciERC, all seven relations have $\alpha \geq 0.67$: with only 7 relations over $6^2 = 36$ entity type pair slots, extensive overlap is inevitable. For example, the pair (Method, Method) is valid for all 7 relations ($\alpha = 1.0$). In DocRED, 96 relations spread across the same 36 type pair slots yield much lower per-relation ambiguity; only 12 relations exceed $\alpha \geq 0.36$, and these collectively cover just 3.8% of gold triples.

Table 46 shows the DR impact stratified by ambiguity. On DocRED, high-α relations exhibit substantially higher unconstrained DR (43.8% for Llama vs. 29.9% for low-α), reflecting the inherent difficulty of relations whose valid type pairs

are shared by dozens of other relations—the model frequently assigns a type combination that, while valid for *some* relation, violates the constraints of the target relation. Under S1 constraints, Llama reduces both groups to near-parity (3.2% vs. 2.5%), whereas Qwen retains a 2× gap (20.0% vs. 9.5%), indicating that constraint effectiveness on high-ambiguity relations depends on the model's baseline type accuracy. Excluding the 12 high-α relations from the DocRED evaluation changes the overall S1 DR only marginally—from 2.55% to 2.51% for Llama and from 10.21% to 9.53% for Qwen—confirming that their small gold footprint (3.8%) limits their aggregate impact despite elevated per-relation rates. On SciERC, where all relations have high α, the exclusion analysis is vacuous; yet OntoCoDec still reduces DR from 15.4% to 0.9% (Qwen) and from 13.3% to 0.08% (Llama), demonstrating that high schema ambiguity does not preclude effective constraint enforcement.

Order invariance. To verify that OntoCoDec's outputs are independent of schema enumeration order, we randomly shuffle the entity and relation type lists in the compiled constraint file (`ontology_constraints.json`) and re-run greedy S1 decoding on all 551 SciERC documents. The raw model outputs and parsed triples are byte-identical to the canonical-order run across all documents. This holds because the grammar layer accepts the same token set at each type-bearing position regardless of the order in which types are listed in the schema, and greedy decoding is deterministic given identical logit masks. The result confirms that OntoCoDec introduces no spurious order dependence from schema compilation.

O Full Proof of Proposition 1

We prove that domain/range type constraints for IE with unbounded entity sets cannot be expressed by any context-free grammar, using the pumping lemma for context-free languages.

Proof. Define the language L of valid IE outputs over ontology $\mathcal{O} = (\mathcal{E}, \mathcal{R}, \sigma)$ with $\mathcal{E} = \{\tau_A, \tau_B\}$, $\mathcal{R} = \{r\}$, and $\sigma(r) = \{(\tau_A, \tau_A)\}$ (i.e., r requires both head and tail to be type τ_A).

An output string has the form:

$$w = e_1\, e_2 \cdots e_n \;\; \texttt{REL}(h{=}e_k,\, r,\, \tau^h)$$

where each e_i is an entity with type $\tau_{a_i} \in \{\tau_A, \tau_B\}$, and τ^h is the head type field of a relation triple referencing entity e_k ($1 \le k \le n$).

The type-validity constraint requires $\tau^h = \tau_{a_k}$ (the head type must match the referenced entity's type) *and* $\tau^h \in \mathrm{dom}(r) = \{\tau_A\}$. Combined: a string $w \in L$ if and only if $\tau_{a_k} = \tau_A$.

Now consider strings of the specific form:

$$w_n = \underbrace{\tau_B\ \tau_B\ \cdots\ \tau_B}_{n-1}\ \tau_A\ \ \texttt{REL}(h{=}e_n,\, r,\, \tau_A)$$

where the first $n{-}1$ entities have type τ_B and only entity e_n has type τ_A. The relation references e_n, so $\tau^h = \tau_A = \tau_{a_n}$, satisfying both the registry consistency and domain constraint. Thus $w_n \in L$ for all n.

Assume for contradiction that L is context-free. By the pumping lemma, there exists a pumping length p such that any $w \in L$ with $|w| \ge p$ can be written as $w = uvxyz$ where:

(i) $|vy| \ge 1$,

(ii) $|vxy| \le p$,

(iii) $uv^i xy^i z \in L$ for all $i \ge 0$.

Choose $w = w_n$ with $n > p$. The entity list portion $\tau_B^{n-1}\,\tau_A$ has length $n > p$, so the pumpable region vxy (with $|vxy| \le p$) must lie entirely within the entity list (it cannot span from the entity list to the relation reference, which is separated by $n - p > 0$ entity tokens from the start).

Case 1: v or y contains entity tokens from the τ_B^{n-1} prefix. Pumping up ($i = 2$): uv^2xy^2z inserts additional τ_B entities, increasing n to some $n' > n$. But the relation reference still points to entity e_n (at the original position), which is now at a different position in the longer entity list. The entity at position n in the pumped string is τ_B (since we inserted τ_B tokens before position n, shifting τ_A to position $n' > n$). Thus $\tau^h = \tau_A \neq \tau_B = \tau_{a'_n}$, where n' is the new position of entity e_n. More precisely: the relation references a fixed entity name e_k, but pumping changes the entity-type binding at that name's position. Since the pumped string assigns τ_B at position k (which originally had τ_A), the type-validity constraint $\tau^h = \tau_{a_k} \in \mathrm{dom}(r)$ is violated. Hence $uv^2xy^2z \notin L$, contradicting condition (iii).

Case 2: v or y contains the τ_A token at position n. Pumping down ($i = 0$): uxz removes the only τ_A entity. The relation's head entity reference now points to a position that either does not exist or has type τ_B. In either case, $\tau^h = \tau_A \neq \tau_B$, violating type validity. Hence $uxz \notin L$, contradicting condition (iii).

In both cases, the pumping lemma is violated, so L is not context-free. The key insight is that type-validity requires a *binding* between an entity's type (assigned at an earlier position) and the relation's type field (at a later position), with the binding target determined by the entity name. Pumping disrupts this binding by changing which entity occupies a given position, while the relation reference remains fixed—a cross-position dependency that pushdown automata cannot track with their LIFO stack discipline.

OntoCoDec's semantic layer addresses this with a runtime entity registry $\mathcal{D}$ (a finite map), operating as a deterministic stateful transducer that provides random-access type lookups, beyond the expressive power of context-free grammars. □

Finite entity sets. For a finite entity set $\mathcal{E}$ with canonical identifiers, the valid language $L_{\text{valid}} = \bigcap_i L_i$ is theoretically regular (a finite intersection of regular languages). However, the corresponding DFA requires $O(|\mathcal{T}|^{|\mathcal{E}|})$ states—for SciERC with 6 entity types and typical documents containing 20+ entities, this exceeds 10^{15} states, rendering explicit construction impractical. Moreover, real IE operates over open-vocabulary text spans rather than closed identifier sets, making the general (non-regular) formulation the appropriate model. Proposition 1 motivates the runtime state-tracking architecture for this general case.

Position indexing vs. name referencing. The proof above uses position-based entity indexing (e_k at position k) for expository clarity. In OntoCoDec's actual JSON output, relation triples reference entities by name strings rather than positional indices. The core argument transfers directly: validating type consistency requires a name→type lookup over dynamically generated bindings, which is a random-access map operation. A PDA's LIFO stack cannot perform random-access retrieval by key (it can only inspect the top element), so the non-context-free result holds regardless of whether entities are referenced by position or by name. Name-based referencing strictly subsumes position-based indexing: any position-based constraint can be expressed via names (assign names $e_1, \ldots, e_n$ in order), but name-based constraints additionally permit cross-field coreference (the same entity name appearing in multiple triples must carry the same type), which position indexing cannot express. Thus, the position-based formulation is a simplification that makes the pumping argument more transparent without weakening the conclusion.

P WikiEvents Event Extraction Pilot

We apply OntoCoDec to event extraction (EE) on WikiEvents (Li et al., 2021; Lu et al., 2021b) (56 event types, 80 roles, 27 entity types, 245 constraints; Appendix C) without modification. Violations extend to EV (event vocabulary), RoV (role vocabulary), and AC (argument constraint, analogous to DR).

		Violation Rates (%)				F1	
Model	**Mode**	EV↓	RoV↓	AC	CF	Trig.↑	Ent.↑
Qwen3-8B	Unconstr.	24.0	74.9	0.3	0	11.8	0
	Grammar	**0**	76.1	8.2	—	—	—
	S1	**0**	**0**	2.1	21.1	13.4	14.1
Llama 3.1-8B	Unconstr.	10.5	88.8	0.1	0	5.7	0
	Grammar	**0**	77.0	8.8	—	—	—
	S1	**0**	**0**	3.8	17.2	7.5	17.6
Mistral 7B	Unconstr.	37.3	63.9	0	0	2.4	0
	Grammar	**0**	64.0	13.8	—	—	—
	S1	**0**	**0**	1.3	14.9	3.9	11.7

Table 47: WikiEvents EE results (zero-shot, n=246 documents). Violation rates over parsed documents. — indicates not applicable (grammar-only has no semantic-layer metrics). F1 (%). n=20 test documents for fine-tuning; MDE≈±5.0 pp; results are directional only.

Results. Unconstrained EE produces extreme vocabulary violations (EV+RoV >85%; Table 47), with entity F1 = 0 because the parser cannot process off-schema output. S1 eliminates EV and RoV across all models, validating tokenizer-agnostic design across byte-level BPE and SentencePiece. Grammar-only eliminates EV but leaves RoV at 64–77% and surfaces latent AC (8–14%), consistent with DocRED's pattern (§Q). Residual AC (1.3–3.8%) traces to semantically similar types; CF rises to 15–21% due to EE's multi-argument entity reuse. Unlike RE, EE shows measurable F1 gains: entity F1 0% → 12–18%, trigger F1 +1.5–1.8 pp.

Composability with fine-tuning (preliminary, n=20). On the full official test split (20 documents; 194/20/20 (Li et al., 2021)), LoRA + S1 eliminates RoV (Qwen 38.31% → 0%, Llama 8.8% → 0%) and reduces AC to ≤0.8%, with argument F1 within ±1.1 pp; trigger F1 shifts (−5.4/−0.6 pp) have wide CIs consistent with n=20. CF remains low (0.77–1.18%); per-document variance is high (Arg-F1 std/mean

>0.5), confirming sufficient power to detect violation elimination but not small F1 differences (power analysis in Appendix G).

Q Zero-Shot Cross-Model Validation

Dataset	Model	Violation Rates (%) DR↓	RV↓	ET↓	CF	R-F1 ↑
SciERC	Qwen (U)	15.37	1.39	0.40	0.08	8.23
	Qwen (S1)	0.93	0	0	9.90	8.36
	Llama (U)	13.34	0.55	0	0.24	2.68
	Llama (S1)	0.08	0	0	6.56	2.54
	Mistral (U)	56.94	0	1.50	56.87†	3.66
	Mistral (S1)	13.40	0	0	1.27	3.67
DocRED	Qwen (U)	18.46	29.66	2.48	0.10	7.74
	Qwen (S1)	10.21	0.03	0	2.57	8.23
	Llama (U)	26.07	16.03	0.13	0.17	5.98
	Llama (S1)	2.55	0	0	12.02	6.27
	Mistral (U)	57.42	42.35	7.58	24.78	4.34
	Mistral (S1)	33.20	0	0	0.19	5.52

Table 48: Three-model zero-shot validation. U: unconstrained; S1: full two-layer. Parse failures: Llama DocRED 9.1%; others $\leq$1.1%. †Mistral: poor zero-shot calibration causes pervasive type inconsistencies. F1 $\times$100.

Cross-model results. RV and ET reach $\leq$0.03% for all models on both datasets. DR drops 76–99% on SciERC and 42–90% on DocRED; R-F1 improves in 5/6 settings. CF direction depends on calibration: S1 introduces CF for calibrated models (up to 12.02%) but reduces it for Mistral (56.87% $\rightarrow$ 1.27%); fine-tuning resolves both (§5.1). Post-hoc filtering achieves DR=0% but at 15–74% rejection rates (Appendix C).

Cross-dataset validation on DocRED. DocRED (96 relations, 14$\times$ SciERC scale) validates the pattern under a larger, more restrictive ontology (Table 11). S1 reduces DR from 18.46% to 10.21% ($\downarrow$45%), eliminates RV (29.66% $\rightarrow$ 0.03%) and ET, with R-F1 +0.50 pp ($p < 0.001$). The lower DR reduction (45% vs. 94% on SciERC) reflects DocRED's tighter constraints, which trigger more grammar-first fallback (§4). Grammar-only *increases* DR (18.46% $\rightarrow$ 26.17%) via *latent violation surfacing*: grammar forces OOV$\rightarrow$valid relation mapping, exposing type violations previously masked by vocabulary errors—a distinct cross-dataset confirmation of the scale-dependent layer interaction. Fine-tuning reduces residual DR to $\leq$0.21% (Table 2; residual decomposition in Appendix C).

R Quantization Robustness

	Qwen3-8B FP16	INT4	Δ	Llama-3.1-8B FP16	INT4	Δ
SciERC (551 docs)						
DR (%)↓	0	0	0	0.09	0	−0.09
R-F1↑	43.77	41.71	−2.06	45.47	43.98	−1.49
E-F1↑	67.05	67.05	0.00	67.76	67.54	−0.22
DocRED (200 docs)						
DR (%)↓	0	0	0	0	0	0
R-F1↑	40.79†	40.07	−0.72	47.35	48.04	+0.69
E-F1↑	81.87†	74.84	−7.03‡	85.43	84.18	−1.25
PF (%)	8.5†	18.5	+10.0‡	2.0	2.5	+0.5

Table 49: FP16 vs. INT4 (AWQ) quantization under LoRA + S1 on SciERC and DocRED. F1 $\times$ 100; PF = parse-failure rate. †Qwen DocRED FP16 evaluated on 998 docs (all other configurations use matched subsets). ‡E-F1 drop primarily attributable to increased parse failures under INT4 quantization on long documents rather than constraint interference (DR remains 0%).

INT4 (AWQ) quantization preserves constraint enforcement across both datasets: DR remains at 0% for all four model–dataset configurations under INT4, with the sole FP16 violation (Llama SciERC, 0.09%) disappearing after quantization. R-F1 changes range from −2.06 to +0.69 pp, consistent with expected quantization quality costs independent of the constraint layers. On DocRED, Qwen INT4 exhibits elevated parse failures (18.5% vs. 8.5% FP16) due to timeout on long documents under quantized inference; the resulting E-F1 gap (−7.03 pp) reflects generation-length degradation rather than constraint-layer interference, as DR remains 0%. These results provide indirect evidence that the 70B AWQ INT4 results (§5.2) are not confounded by quantization artifacts in the constraint enforcement mechanism.

S GoLLIE Comparison

GoLLIE (Sainz et al., 2024) and OntoCoDec represent complementary paradigms for schema compliance in IE:

- *Training-time internalization* (GoLLIE, KnowCoder): Schema knowledge is encoded into model parameters via annotation-guideline-aware training. This achieves compliance through learned generation patterns but cannot guarantee zero violations at inference time, and

requires schema-specific retraining when ontologies evolve.

- *Runtime constraint enforcement* (OntoCoDec): Schema constraints are compiled into decoding-time masks applied to any generative model. This provides formal enforcement without schema-specific training, but operates per-triple and cannot resolve cross-triple type consistency (CF) within a single pass.

The two approaches are orthogonal: OntoCoDec can be composed with GoLLIE or similar training-based systems to provide runtime enforcement guarantees atop learned schema compliance. Direct numerical comparison is methodologically constrained by incompatible evaluation settings (code-style schema encoding vs. natural language prompting, different model families and sizes); a controlled comparison would require running both approaches on the same model with matched prompting, which we leave to future work.