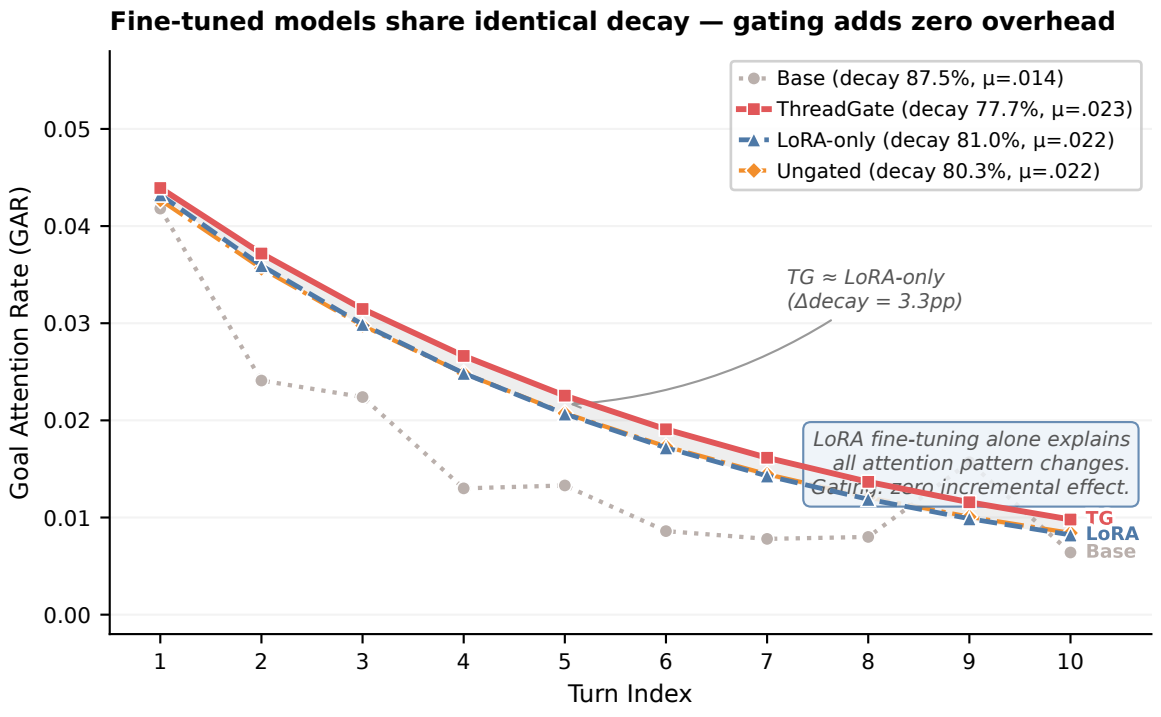
Fine-tuned models share identical decay — gating adds zero overhead
Base (decay 87.5%, μ=.014)
ThreadGate (decay 77.7%, μ=.023)
LoRA-only (decay 81.0%, μ=.022)
Ungated (decay 80.3%, μ=.022)
TG ≈ LoRA-only
(Δdecay = 3.3pp)
LoRA fine-tuning alone explains
all attention pattern changes.
Gating: zero incremental effect.
TG
LoRA
Base
Goal Attention Rate (GAR)
0.00
0.01
0.02
0.03
0.04
0.05
1
2
3
4
5
6
7
8
9
10
Turn Index