Death Avoidance as the Dominant Routing Signal: A Counterfactual Analysis on ScienceWorld

Anonymous ACL submission

Abstract

Adaptive computation routing, deciding when an LLM agent should rely on internal reasoning versus external verification, is a growing focus in agent design, yet the field lacks a principled way to diagnose what signal actually drives routing preferences. Without such diagnosis, practitioners cannot distinguish environments where learned routing adds genuine value from those where a trivial baseline suffices. We introduce an episode-level counterfactual framework that decomposes routing signal into mechanistic components by comparing pure-strategy outcomes before any router is built. Applied to ScienceWorld, the framework reveals that 78.9% of non-tie routing labels trace to a single mechanism: *death asymmetry*, where internal reasoning triggers catastrophic failure while the deterministic strategy survives. This dominance replicates across three model families (7B–70B); at 70B scale, death avoidance accounts for the entire performance gap. A trivial always-deterministic policy matches 96.3% of oracle decisions, leaving negligible room for learned routing. Boundary probes on MATH and ALFWorld show the framework distinguishes qualitatively different signal sources, including environments where no routing signal exists[1].

1 Introduction

Large language models deployed as interactive agents must repeatedly decide *how* to compute: reason internally, or delegate to an external tool. Existing routing methods, whether inserting API calls by perplexity reduction (Schick et al., 2023), decomposing tasks into tool-augmented programs (Paranjape et al., 2023; Lu et al., 2023), or adaptively selecting reasoning depths (Yang et al., 2026b,a) and model tiers (Zhang et al., 2026c,a), share an implicit assumption: that routing signal stems from diverse sources and a learned router can exploit it. But before training a router, we should understand what mechanistic signal drives the routing preference. If signal collapses to a single dominant mechanism, a learned step-level router may not be the right solution.

We propose an *episode-level counterfactual diagnostic framework* to address this prior question. The framework compares pure-strategy outcomes (internal-only vs. deterministic-only) and introduces the Death Avoidance Factor (DAF) to decompose the routing score gap into death-avoidance and non-death components (§3.2–3.4). Operating at episode granularity is essential: per-step oracle labeling degenerates under sparse terminal rewards (97.4% tie rate in ScienceWorld; §3.2).

Applied to ScienceWorld (Wang et al., 2022), the framework reveals that **78.9% of non-tie routing labels arise from death asymmetry**, where the internal strategy triggers a catastrophic failure while the deterministic strategy survives. This pattern replicates across three architecturally distinct model families (7–70B parameters) and intensifies at larger scale, where death avoidance accounts for the entire performance gap (§4.2–4.3). Always-deterministic matches 96.3% of oracle routing decisions; a trained step-level router fails to exceed this trivial baseline (§4.5–4.6). Boundary probes on MATH and ALFWorld show that the framework distinguishes qualitatively different mechanisms, including environments with no routing signal and those driven by information asymmetry (§4.4).

Our contributions: (1) An **episode-level counterfactual diagnostic framework** and the DAF metric that decompose routing signal into mechanistic components before any router is built, overcoming per-step oracle degeneracy in sparse-reward environments. (2) Evidence that **routing signal in ScienceWorld is death-dominated**: 78.9% of non-tie labels trace to death asymmetry, replicated

[1] Code and data are available at https://anonymous.4open.science/r/simgate-modular-routing-ED12/.

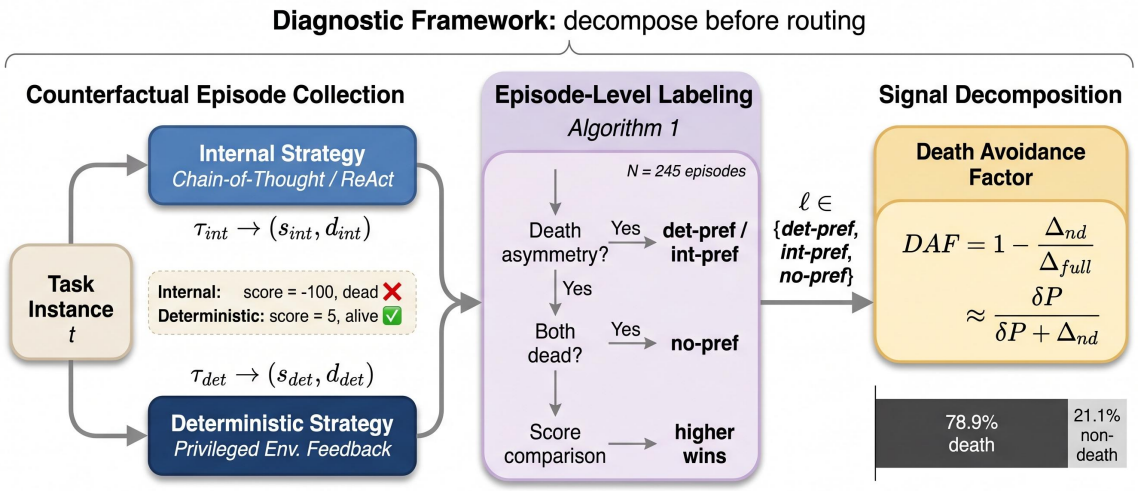


Figure 1: Counterfactual episode collection and analysis pipeline. Each task instance is run under both pure strategies independently, producing episode-level labels via Algorithm 1.

across three model families and two scales (7B and 70B). A ∼50-episode pilot suffices for practitioners to assess whether learned routing adds value (Appendix R). (3) **Boundary probes** on MATH and ALFWorld showing the framework can diagnose qualitatively different routing mechanisms, while clarifying that full quantitative validation remains specific to ScienceWorld.

2 Related Work

Tool-Augmented Agents and Adaptive Routing. Methods for augmenting LLMs with external tools span self-supervised invocation (Schick et al., 2023; Liu et al., 2025), task-level decomposition (Paranjape et al., 2023; Lu et al., 2023), and RL-based tool selection (Feng et al., 2025; Qu et al., 2024). Concurrent work examines *when* to invoke tools: Sun et al. (2026) analyze accuracy–latency tradeoffs, Zhang et al. (2026b) quantify invocation overhead, and Wu et al. (2026); Zeng et al. (2026) investigate overuse patterns. On the routing side, adaptive methods range from learned halting (Graves, 2016) and token-level recursion (Bae et al., 2025) to per-step depth adjustment (Yang et al., 2026b,a) and multi-turn model selection (Zhang et al., 2026c). Model-selection routing (Ong et al., 2024; Chen et al., 2023; Aggarwal et al., 2024) optimizes *which* model to invoke based on query difficulty; Zhang et al. (2026a) train episode-level routing under sparse rewards via counterfactual rollouts, the closest setting to ours, though they focus on learning the router rather than diagnosing what signal it captures. All these approaches presuppose that diverse routing signals exist; our decomposition reveals that in ScienceWorld, a single mechanism accounts for 78.9% of routing labels.

Counterfactual Methods and Agent Environments. Prior work generates counterfactual explanations by perturbing individual actions (Olson et al., 2021; Gajcin, 2023; Amitai et al., 2024); recent extensions target LLM agents via abstract state comparison (Pona et al., 2025), axis-aligned perturbations (Gyevnár et al., 2025), and causal structure discovery (Li et al., 2026b). These approaches explain *why* an agent succeeded or failed on a single trajectory. We instead execute complete episodes under two strategies to decompose routing signal *before* training. Counterfactual data augmentation (Kaushik et al., 2020) edits inputs to improve robustness; we execute full counterfactual *episodes* under controlled strategy variation, answering what *type* of signal drives strategy preferences. ScienceWorld (Wang et al., 2022) serves as our primary testbed; ALFWorld (Shridhar et al., 2021) and AgentBench (Liu et al., 2024) provide complementary evaluation. ReAct (Yao et al., 2023) is our internal baseline; Reflexion (Shinn et al., 2023) and LATS (Zhou et al., 2024) extend it. The concurrent SALT (Li et al., 2026a) assigns step-level credit on ALFWorld, complementing our episode-level decomposition.

Safe RL, Agent Safety, and Causal Inference. Death avoidance connects to safe RL (García and Fernández, 2015; Achiam et al., 2017), shielding (Alshiekh et al., 2018), and the growing literature on LLM agent safety evaluation (Ruan et al., 2024; Lee et al., 2025; Zhang et al., 2024). Classical safe RL methods require known transition

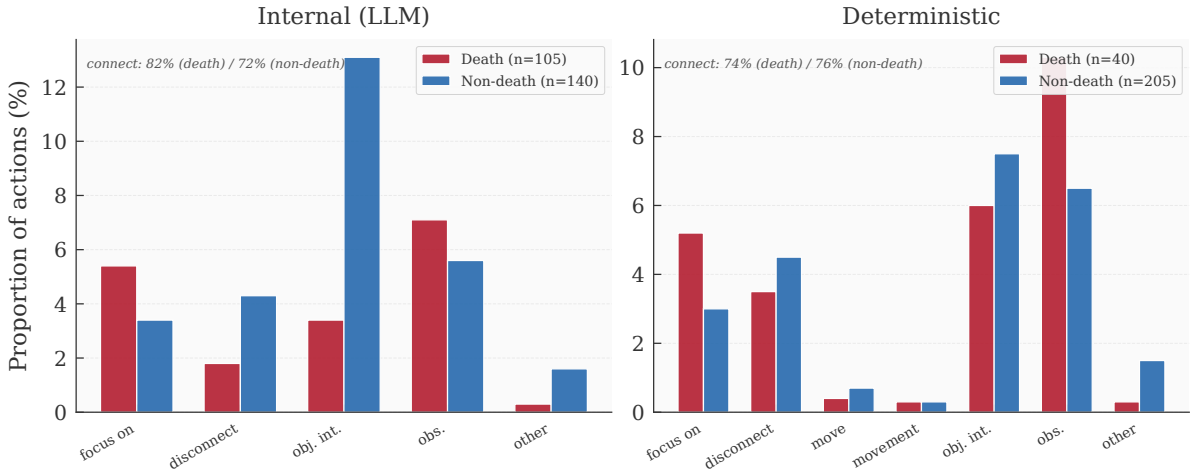


Figure 2: Non-`connect` action distribution comparison between internal (left) and deterministic (right) strategies within death episodes. Both strategies exhibit comparable `focus on` rates (∼5%); the internal strategy's higher death count (105 vs. 40) reflects more frequent entry into catastrophic states, not different action selection within those states.

dynamics or a differentiable safety model, neither available for black-box LLM agents. Unlike token-level adaptive computation (Raposo et al., 2024; Xin et al., 2020) and mixture-of-experts routing (Fedus et al., 2022), which optimize compute allocation at inference time, our framework diagnoses aggregate routing signal at episode granularity. Methodologically, our pure-strategy comparison is a simplified potential outcomes framework (Rubin, 1974): executing both strategies under identical initial conditions eliminates confounding without structural causal models (Pearl, 2009), at the cost of requiring paired episode collection.

3 Method

Our diagnostic framework (Figure 1) quantifies what drives routing signal in interactive environments through two components: (i) episode-level counterfactual labeling that generates routing preferences from paired pure-strategy outcomes (§3.2–3.3); and (ii) the Death Avoidance Factor (DAF), which decomposes the routing score gap into death-avoidance and non-death components (§3.4). A supplementary Signal Concentration Index (SCI; Appendix A) measures whether signal is concentrated in a few task types. Episode-level bootstrap resampling provides 95% confidence intervals for all metrics.

3.1 Problem Setup

Consider an interactive environment with observation space $\mathcal{O}$ and action space $\mathcal{A}$. An agent policy $\pi : \mathcal{O}^* \to \mathcal{A}$ maps observation histories to actions. We define two computation modes $\mathcal{M} = \{\texttt{int}, \texttt{det}\}$: **internal** (`int`), where the agent predicts intermediate observations via chain-of-thought reasoning (ReAct (Yao et al., 2023)), with prediction errors compounding across steps; and **deterministic** (`det`), which uses the same LLM to generate candidate actions but evaluates them through privileged environment feedback (ground-truth simulator state in ScienceWorld, code interpreter in MATH, admissible-action enumeration in ALFWorld; see §4.4 for environment-specific operationalizations). Crucially, the deterministic strategy does not provide optimal actions; it reduces observation uncertainty by grounding the LLM's choices in verified environment state rather than predicted state. These qualitative differences affect what routing signals the framework can detect. A *routing policy* $\rho : \mathcal{O}^* \to \mathcal{M}$ selects which computation mode to use at each step; our analysis investigates what signal is available to learn ρ.

3.2 Counterfactual Episode Collection

Per-step oracle labeling degenerates under sparse terminal rewards: on ScienceWorld (Wang et al., 2022), 97.4% of steps (5,935 of 6,094 across 11 prescreened task types) produce identical intermediate scores, yielding no per-step routing signal. Episode-level comparison (directly measuring terminal outcomes) is therefore the principled granularity for sparse-reward environments.

Pure-strategy trajectories. For a task instance t and computation mode $m \in \mathcal{M}$, a *pure-strategy trajectory* is the complete episode obtained by applying mode m at every step:

$$\tau_m^{(t)} = (o_1, a_1, o_2, a_2, \ldots, o_T, a_T) \quad (1)$$

where $T \leq T_{\max}$ is the episode length (capped at $T_{\max} = 30$ steps). Both trajectories $\tau_{\texttt{int}}^{(t)}$ and $\tau_{\texttt{det}}^{(t)}$ are executed under identical initial conditions (same random seed, same task variation), so performance differences arise solely from the computation mode choice.

Each trajectory τ yields a terminal score $s(\tau) \in \mathbb{Z}$ and a death indicator $d(\tau) \in \{0, 1\}$. When the agent dies ($d = 1$), the environment imposes a penalty: $s(\tau) = -P$ where $P > 0$ ($P = 100$ in ScienceWorld). The episode pair produces a tuple:

$$\big(s^{\texttt{int}}, s^{\texttt{det}}, d^{\texttt{int}}, d^{\texttt{det}}\big) \triangleq \big(s(\tau_{\texttt{int}}),\, s(\tau_{\texttt{det}}),\, d(\tau_{\texttt{int}}),\, d(\tau_{\texttt{det}})\big) \quad (2)$$

We use Qwen2.5-7B-Instruct (Yang et al., 2024) (vLLM v0.21.0) on ScienceWorld (Wang et al., 2022), a text-based environment spanning 30 task types with a -100 death penalty; complete prompt templates are provided in Appendix T. After prescreening, we collect 245 episode pairs across 11 task types (20–40 per type). The framework is applicable to any environment where paired episodes under two computation modes can be collected with binary success/failure indicators. We apply it unchanged to ALFWorld and MATH in §4.4–4.4.

3.3 Episode-Level Labeling

We define a labeling function L that maps a counterfactual pair to a routing preference:

$$L\big(\tau_{\texttt{int}}^{(t)}, \tau_{\texttt{det}}^{(t)}\big) \rightarrow \ell \in \{\texttt{det-pref}, \texttt{int-pref}, \texttt{no-pref}\} \quad (3)$$

The labeling function is *death-aware*: it prioritizes survival over score, then compares scores among same-status episodes (Algorithm 1). Both-dead episodes receive `no_preference` because ScienceWorld's fixed -100 penalty makes death timing uninformative. ScienceWorld scores are integer-valued, so we use exact (in)equality with no threshold. A quality audit confirming zero tie-breaking artifacts is in Appendix B. Label distributions are reported in §4.1.

Algorithm 1 Episode-Level Counterfactual Labeling

Require: Task instance t, computation modes $\mathcal{M} = \{\texttt{int}, \texttt{det}\}$
Ensure: Label $\ell \in \{\texttt{det-pref}, \texttt{int-pref}, \texttt{no-pref}\}$

```
for each m ∈ M do
    Run episode τ_m with pure-strategy m
    Record score s_m and death indicator d_m
end for
if d_int = 1 and d_det = 0 then
    ℓ ← det-pref {death asymmetry}
else if d_int = 0 and d_det = 1 then
    ℓ ← int-pref {death asymmetry}
else if d_int = d_det = 1 then
    ℓ ← no-pref {both dead}
else if s_det > s_int then
    ℓ ← det-pref {score comparison}
else if s_int > s_det then
    ℓ ← int-pref {score comparison}
else
    ℓ ← no-pref {tied alive}
end if
return ℓ
```

3.4 Death Avoidance Factor

To quantify the contribution of death avoidance to routing signal, we decompose the mean score gap between strategies. Let $\delta = \mathbb{E}[d^{\texttt{int}}] - \mathbb{E}[d^{\texttt{det}}]$ denote the death-rate differential and Δ_{nd} the mean score gap restricted to episodes where neither strategy dies. Partitioning episodes by death status, death-asymmetric episodes (where only one strategy dies) contribute a per-episode gap of $P + \bar{s}_{\text{surv}}$, where $\bar{s}_{\text{surv}}$ is the surviving strategy's mean score; when $\bar{s}_{\text{surv}} \ll P$ (in our data, $\bar{s}_{\text{surv}} = 0.87$, or 0.9% of P), these episodes contribute $\approx \delta P$ in aggregate. The full gap thus decomposes as:

$$\Delta_{\text{full}} \approx \underbrace{\delta \cdot P}_{\text{death contribution}} + \Delta_{\text{nd}} \quad (4)$$

The approximation residual is $29.60 - 29.37 = 0.23$ (0.8% of Δ_{full}).

We define the Death Avoidance Factor (DAF) as the fraction of the full score gap attributable to death avoidance:

$$\text{DAF} = 1 - \frac{\Delta_{\text{nd}}}{\Delta_{\text{full}}} \approx \frac{\delta P}{\delta P + \Delta_{\text{nd}}} \quad (5)$$

[2]

[2] When $\Delta_{\text{nd}} < 0$ (deterministic strategy scores lower than internal on non-death episodes), DAF exceeds 1; death avoidance then overcompensates for a non-death disadvantage.

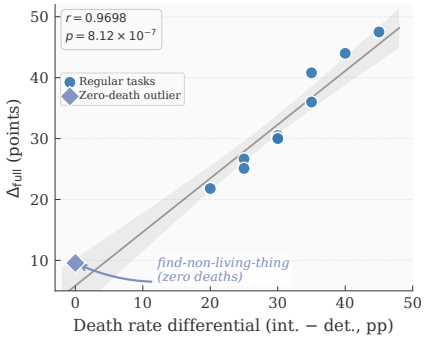


Figure 3: Per-task death rate differential vs. performance gap (n=11 task types). The near-perfect correlation ($r = 0.97$) is a structural consequence of death dominance (DAF $\to$ 1 $\Rightarrow$ $r \to 1$), not an independent finding. The sole outlier (`find-non-living-thing`, diamond marker) has zero deaths; its $\Delta_{\text{full}} = +9.58$ arises entirely from non-death score differences.

Condition	Death%	N
Single-trial baseline	42.9	245
Reflexion (3×)	31.4	220
Qwen2.5-72B (10×)	30.6	255
Deterministic baseline	16.3	245
10 death-prone tasks only:		
Death-prone baseline	48.0	200
Retry (5×)	43.5	200

Table 1: Internal-strategy death rates under improvement methods. Top: all 11 tasks. Bottom: 10 death-prone tasks (`find-non-living-thing` excluded). No intervention approaches deterministic-level safety.

DAF increases monotonically with P (Table 14, Appendix K); we treat the label fraction (78.9%) as the primary, penalty-invariant diagnostic. In ScienceWorld, P affects only post-hoc score calculation; death rates and Δ_{nd} are P-invariant by design (Ng et al., 1999). Uncertainty is quantified via episode-level bootstrap (B=10,000, percentile 95% CIs).

4 Experiments

We analyze 11 of ScienceWorld's 30 task types (245 episodes); the 19 excluded types lack actionable routing signal due to development contamination (5), taxonomic redundancy (7), or agent-floor performance (7). An anti-cherry-picking analysis confirms death dominance holds in 18 of these 19 types ($\overline{\Delta\text{DR}} = +38.2$ pp $>$ included tasks' $+28.7$ pp; Appendix C), with extended validation on representative excluded tasks in Appendix S.

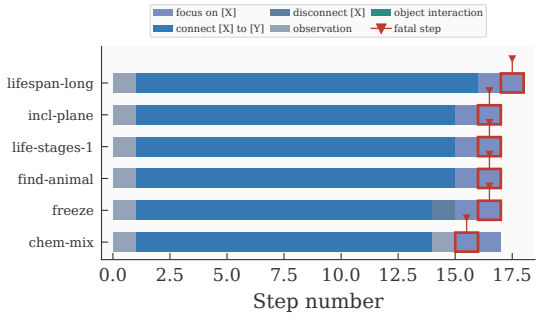


Figure 4: Action sequences for 6 representative death episodes. The fatal `focus on` (arrow) appears in final steps after long `connect` sequences ("connect loop $\to$ fatal focus").

Task	n	Δ_{full}	D%$_{\text{int}}$	D%$_{\text{det}}$	D / I / N
id-life-stages-1	20	+47.50	55	10	9 / 0 / 11
lifespan-long	25	+44.00	56	16	11 / 1 / 13
lifespan-l-then-s	20	+40.80	45	10	8 / 1 / 11
inclined-plane	20	+36.00	55	20	7 / 0 / 13
freeze	20	+30.50	55	25	6 / 0 / 14
meas-melt-point	20	+30.05	55	25	6 / 0 / 14
id-life-stages-2	20	+30.00	50	20	6 / 0 / 14
find-animal	20	+26.65	50	25	5 / 0 / 15
grow-fruit	20	+25.10	45	20	6 / 1 / 13
chem-mix	20	+21.80	45	25	6 / 2 / 12
find-non-living†	40	+9.58	0	0	16 / 4 / 20
All (11 tasks)	**245**	**+29.60**	**42.9**	**16.3**	**86 / 9 / 150**

Table 2: Per-task decomposition. D%: death rate for internal/deterministic. D/I/N: det-preferred / int-preferred / no-preference labels. † Only task with zero deaths; n=40 (two batches).

4.1 Death Avoidance Decomposition

Of the 95 non-tie routing labels, 75 (78.9%; 95% CI: [70.7%, 87.0%], episode-level bootstrap) arise from death asymmetry: the internal strategy triggers a catastrophic -100 penalty while the deterministic strategy survives. This penalty-independent label fraction is the primary diagnostic metric. Of these, 54 are sole-signal labels where death was the only discriminator; at P=0 these reduce to ties, but the labeling function itself is penalty-independent (Algorithm 1).

DAF = 90.4% at P=100 ($\Delta_{\text{full}} = +29.60$, $\Delta_{\text{nd}} = +2.84$; 95% CI$_{\text{DAF}}$: [81.3%, 99.1%], episode-level bootstrap). The 78.9% label fraction provides the more robust, penalty-invariant characterization (§3.4).

Per-task structure. The internal strategy dies in 42.9% of episodes versus 16.3% for deterministic;

Model	Env	n	Int DR	Det DR	ΔDR	DAF [95% CI]	Signal
Qwen2.5-7B	SW	245	42.9	16.3	+26.5	90.4 [81.3, 99.1]	death
Llama-3.1-8B	SW	220	41.4	11.4	+30.0	88.2 [80.2, 95.1]	death
Gemma-2-9B	SW	220	36.4	14.1	+22.3	89.7 [79.5, 98.4]	death
Qwen2.5-72B	SW	255	30.6	2.8	+27.8	92.4 [86.1, 97.9]	death
Llama-70B	SW	220	48.6	20.0	+28.6	101.0 [95.3, 107.5]	death
Qwen2.5-7B	ALF	135	0	0	0	0	info-asym
Llama-3.1-8B	ALF	135	0	0	0	0	info-asym
Qwen2.5-7B	MATH	795	—	—	—	—	none

Table 3: Cross-model and cross-environment summary. SW = ScienceWorld, ALF = ALFWorld. Death dominance replicates across three 7–9B families and intensifies under 10× scaling; ALFWorld shows information-asymmetry signal; MATH shows none. DR = death rate (%).

all 105 internal deaths are triggered by the `focus on` action (Appendix N). Per-task ΔDR correlates with Δ_{full} at $r = 0.97$ (Figure 3; leave-one-out and bootstrap robustness in Appendix D), a structural consequence of DAF$\rightarrow$ 1 rather than an independent finding. Figure 4 illustrates the typical death trajectory.

Binary structure and robustness. All 10 death-prone task types produce 100% death-asymmetry labels; only `find-non-living-thing` (zero deaths) produces non-death signal (SCI = 0.21; Appendix A). Leave-one-task-out (Table 4) confirms robustness: removing any single task changes DAF by ≤3.0 pp. Bootstrap 95% CIs: DAF $[81.3\%, 99.1\%]$; cluster-robust $[73.3\%, 100.0\%]$ (K=11; Appendix E). Two replications (Gemma-2-9B-it; seed=123) converge within 0.7 pp.

4.2 Multi-Model Validation (7–9B Tier)

Llama-3.1-8B (Grattafiori et al., 2024) (N=220) and Gemma-2-9B-it (Gemma Team, 2024) (N=220) independently replicate death dominance: DAF = 88.2% and 89.7% respectively, with 10/10 directional consistency on death-prone tasks (Appendix M; per-task CIs in Appendix F). Three architecturally distinct families (Qwen, Llama, Gemma) converge on DAF 88–90% (spread 2.2 pp), suggesting death dominance is an environment-level property. Scaling 10× to two 70B-class models (§4.3) yields DAF=92.4–101.0%, indicating the pattern persists at larger scale.

4.3 Agent Improvement Baselines

We test whether agent-level improvement reduces death dominance (Table 1). Retry (re-executing the episode without explicit reflection; N=5) yields marginal reduction (43.5% vs. 48.0%). Reflexion (Shinn et al., 2023) reduces death to 31.4% but the gap to deterministic remains 15.1 pp; under matched-episode analysis (N=3), $\Delta_{\text{nd}} = -0.47$, yielding DAF > 100%. Scaling to Qwen2.5-72B (N=255) reduces death rate to 30.6% with DAF = 92.4% (95% CI: $[86.1\%, 97.9\%]$). A second 70B model, Llama-3.1-70B (N=220), yields DAF = 101.0% ($[95.3\%, 107.5\%]$): death avoidance more than fully accounts for the gap. Both 70B-class models suggest that scaling does not attenuate death dominance.

Removed Task	N_{rem}	DAF (%)	Δ (pp)
(none, full)	245	78.9	—
Find Non-Living†	205	100.0	+21.1
Lifespan Longest	220	75.9	−3.0
Life Stages 1	225	76.7	−2.2
Lifespan Long+Short	225	76.7	−2.2
Chemistry Mix	225	77.0	−1.9
Grow Fruit	225	77.3	−1.7
Inclined Plane	225	77.3	−1.7
Freeze	225	77.5	−1.4
Life Stages 2	225	77.5	−1.4
Melt Pt Unknown	225	77.5	−1.4
Find Animal	225	77.8	−1.2

Table 4: Leave-one-task-out DAF sensitivity (label fraction). Removing any death-prone task changes DAF by ≤3.0 pp. † Sole death-free task; removal raises DAF to 100% by construction.

Dose-response analysis. We extend both interventions to N=5 trials on death-prone tasks (n=200; Figure 5). Retry is nearly flat (48.0% → 43.5%). Reflexion drops steeply to 34.0% at N=3 ($\Delta = -13.5$ pp, $p = 0.003$) then stabilizes (34–36.5% for $N \geq 3$, pairwise $p > 0.5$); both the decline and plateau are well-powered (minimum detectable effect ∼8 pp at 80% power; Appendix H).

4.4 Cross-Environment Boundary Probes

We probe two environments as boundary cases to test the framework's specificity (Figure 6). On MATH Level 4–5 (Hendrycks et al., 2021) (N=795), internal reasoning and code-interpreter accuracy are near-identical (73.5% vs. 73.0%, $p = 0.81$): without a catastrophic failure mechanism, no routing signal emerges despite substantial task difficulty. ALFWorld (Shridhar et al., 2021) (135 games, no death mechanism, DAF = 0%) shows 94.6% of non-tie labels favoring the information-advantaged strategy (which enumerates admissible

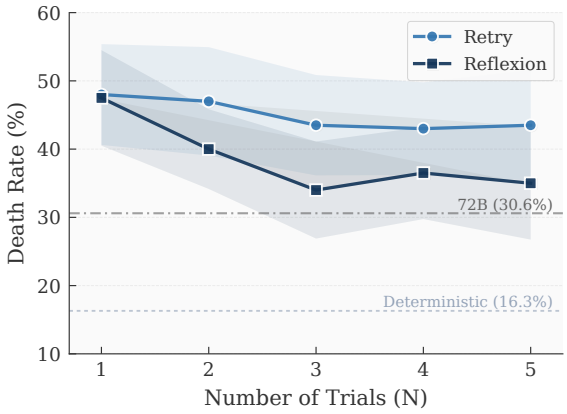


Figure 5: Death rate vs. trial count for Retry and Reflexion on 10 death-prone tasks (n=200). Reflexion drops steeply through N=3 then plateaus; Retry is nearly flat. Horizontal lines: 72B (30.6%) and deterministic (16.3%) baselines. Error bands: Wilson 95% CIs.

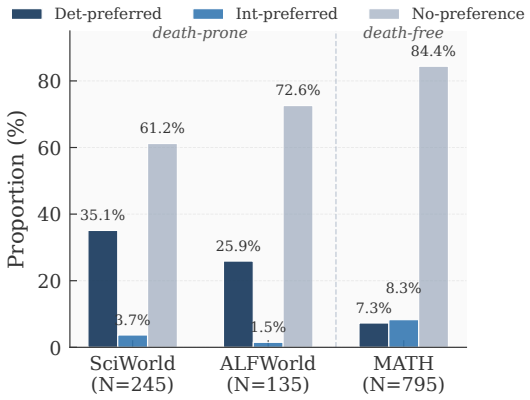


Figure 6: Cross-environment label distributions. ScienceWorld and ALFWorld both favor deterministic via different mechanisms (death avoidance vs. information asymmetry); MATH shows no routing signal.

Strategy	Score	95% CI	Death %
Always-Internal	−39.07	[−46.40, −32.27]	42.9
Always-Deterministic	−9.47	[−14.71, −4.29]	16.3
Random	−22.38	[−28.67, −15.75]	27.8
Oracle	−6.58	[−11.54, −1.70]	14.3
Trained Router†	−10.70	[−16.33, −5.27]	17.6

Table 5: Episode-weighted mean scores on ScienceWorld (11 task types, 245 episodes). † Single seed (42); see Appendix J for 8-seed robustness.

actions); Llama-3.1-8B replicates at 100%. The framework identifies this qualitatively different mechanism, information asymmetry rather than death avoidance, without modification, indicating that DAF correctly attributes zero signal to death when none exists (full label distributions and per-task breakdowns in Appendix L).

4.5 Baselines

Always-deterministic (−9.47) outperforms always-internal (−39.07) by 29.60 points (Table 5). Random routing (−22.38) falls between the two pure strategies, confirming that the gap is not an artifact of evaluation methodology. The oracle improves only to −6.58, with 96.3% of decisions matching always-deterministic, an *oracle degeneracy* where perfect per-episode foresight provides marginal benefit over the trivial always-deterministic baseline (Appendix Q).

4.6 Router Training: Inconclusive Results

A Qwen3 LoRA router (0.6B–4B; Appendix I) trained on counterfactual labels shows flat step-level balanced accuracy (56–59%) across all model scales; episode-level evaluation yields 47.1±2.2%, below majority-vote.[3] Leave-one-task-out degrades to chance (50.0%), suggesting signal is a per-task-type constant rather than a learnable function of step-level observations. Deployed end-to-end, the best router (4B) scores −10.70, not exceeding always-deterministic (−9.47); this negative result is consistent with the death-dominance finding, as death avoidance cannot be predicted from pre-episode features when it depends on multi-step trajectory dynamics. Additional ablations (capacity scaling, generative routing, reasoning depth) are in Appendix P.

[3]Consistent with noisy-label bounds (Natarajan et al., 2013): ∼88% of step-level labels are uninformative with 90.5%/3.7% class imbalance (Appendix J).

5 Discussion

Signal granularity and router failure. The router diagnostic (§4.6) reveals a structural mismatch: death avoidance is a task-type-level binary property, but step-level features correlate weakly because death risk depends on multi-step trajectories. Six router variants (0.6B–4B) converge on near-chance balanced accuracy (56–59%); leave-one-task-out degrades to 50.0%, suggesting the router captures only task-type identity rather than a learnable function of step-level observations. Scaling 10× to 70B-class models does not attenuate this pattern (DAF 92–101%), and cross-model death-rate stability (31–49% across five models) confirms death vulnerability is a task-environment property, better addressed through action-space constraints (Appendix O) than learned routing policies.

Condition	ΔDR	f_{death}	Prediction	Actual
Qwen-7B (SW)	+26.5	78.9%	det-fixed	det wins
Llama-8B (SW)	+30.0	>75%	det-fixed	det wins
Qwen-72B (SW)	+27.8	>75%	det-fixed	det wins
MATH	+0.5	0%	router viable	no signal

Table 6: Retrospective validation of the practitioner decision framework. SW = ScienceWorld. The three-stage diagnostic correctly classifies all four conditions (4/4).

Failure mechanism. All 105 internal-strategy deaths trace to misuse of `focus on [object]`, ScienceWorld's answer-submission command, applied to irrelevant targets (rooms, doors, inventory items). The action distribution analysis (Figure 7, Appendix N) shows that `focus on` is moderately more frequent in death episodes (5.4% vs. 3.4%), while Figure 2 reveals that within death episodes both strategies exhibit comparable rates (∼5%). The higher internal death count reflects more frequent entry into catastrophic states, not different action selection within those states. Among fatal `focus on` targets, `workshop` is the most frequent (20.0%), followed by `inventory` (17.1%) and `kitchen` (10.5%), indicating a systematic object bias rather than random exploration. Action-space filtering that blocks `focus on` during the exploration phase reduces death by 19 pp with zero leaks and zero passthroughs (Appendix O), confirming shallow behavioral origins amenable to simple interventions.

Oracle degeneracy and practical implications. Always-deterministic matches 96.3% of oracle decisions (§4.5), capping routing's marginal value at 2.89 points—an *oracle degeneracy* where $\Delta_{\text{nd}} = 2.84$ (9.6% of total gap). This degeneracy arises because the deterministic strategy already avoids the primary failure mode, leaving only a residual non-death gap for the oracle to exploit. A ∼50-episode pilot suffices to compute the death-label fraction f_{death} (the proportion of non-tie labels where exactly one strategy dies) and assess whether learned routing adds value before committing to full-scale router training. Table 6 shows the three-stage procedure (Appendix R) correctly classifies all four conditions, including ALFWorld's information-asymmetry signal where DAF = 0% correctly identifies a non-death mechanism.

Scope and applicability conditions. The decomposition requires: (i) binary catastrophic failures with fixed penalty, (ii) asymmetric failure rates, and (iii) $P \gg \Delta_{\text{nd}}$. ScienceWorld satisfies all three; MATH and ALFWorld lack the first. Graded penalties or symmetric failure rates would require DAF extensions. The penalty sensitivity analysis (Table 14, Appendix K) confirms that DAF is mechanically driven by P: the label fraction (78.9%) is P-invariant, while the score-based DAF ranges from 7.4% at P=0 to 97.9% at P=500. In death-dominated settings, targeted safety interventions may outperform general-purpose routing, shifting the problem from "which strategy to route to" to "how to prevent catastrophic actions."

6 Conclusion

Our episode-level counterfactual framework reveals that routing signal in ScienceWorld is dominated by a single mechanism: 78.9% of non-tie labels trace to death asymmetry, a finding stable across three model families and two 70B-class models (DAF 92–101%). When routing signal collapses to catastrophic-failure avoidance, always-deterministic approaches oracle performance, and learned step-level routing offers limited additional value.

The framework's diagnostic value extends beyond negative results. The practitioner decision procedure (Table 6) correctly classifies four distinct conditions, and boundary probes demonstrate that the same decomposition identifies qualitatively different signal sources without modification, suggesting utility as a general pre-routing diagnostic.

Key open directions include validating death dominance in other catastrophic-failure environments (e.g., safety-critical robotics, irreversible web actions), developing hybrid approaches that combine action-space filtering for death-prone tasks with learned routing for non-death tasks, extending DAF to graded penalties and stochastic failure severity, and investigating whether the concentration of routing signal in a single mechanism generalizes to other agent benchmarks with asymmetric failure modes. More broadly, our results suggest that diagnosing signal composition *before* training a router can save substantial compute by identifying when a simple fixed strategy already captures most of the available signal.

Limitations

Full quantitative validation (DAF, label fractions) is specific to ScienceWorld; MATH and ALFWorld serve as boundary probes, not comprehensive cross-

environment validation. DAF depends on a fixed death penalty ($P{=}100$); environments with graded or stochastic failure severity would require extensions (Appendix K). The internal baseline uses single-pass ReAct (Yao et al., 2023); more aggressive self-correction (e.g., Tree-of-Thought) may further reduce death rates and alter the signal composition. With $K{=}11$ task-type clusters, cluster bootstrap yields wider CIs; wild cluster bootstrap ($p < 0.001$; Appendix E) and 100% directional consistency across all 28 death-prone tasks provide partial mitigation. Total compute was approximately 200 GPU-hours on consumer and workstation GPUs.

References

Joshua Achiam, David Held, Aviv Tamar, and Pieter Abbeel. 2017. Constrained policy optimization. In *Proceedings of the 34th International Conference on Machine Learning (ICML)*, pages 22–31.

Pranjal Aggarwal, Aman Madaan, Ankit Anand, Srividya Pranavi Potharaju, Swaroop Mishra, Pei Zhou, and 1 others. 2024. AutoMix: Automatically mixing language models. In *Advances in Neural Information Processing Systems (NeurIPS)*.

Mohammed Alshiekh, Roderick Bloem, Rüdiger Ehlers, Bettina Könighofer, Scott Niekum, and Ufuk Topcu. 2018. Safe reinforcement learning via shielding. In *AAAI*.

Yotam Amitai, Yael Septon, and Ofra Amir. 2024. Explaining reinforcement learning agents through counterfactual action outcomes. In *Proceedings of the AAAI Conference on Artificial Intelligence*.

Sangmin Bae, Yujin Kim, Reza Bayat, Sungnyun Kim, Jiyoun Ha, and 1 others. 2025. Mixture-of-recursions: Learning dynamic recursive depths for adaptive token-level computation. In *Advances in Neural Information Processing Systems (NeurIPS)*.

A. Colin Cameron, Jonah B. Gelbach, and Douglas L. Miller. 2008. Bootstrap-based improvements for inference with clustered errors. *The Review of Economics and Statistics*, 90(3):414–427.

Lingjiao Chen, Matei Zaharia, and James Zou. 2023. FrugalGPT: How to use large language models while reducing cost and improving performance. *arXiv preprint arXiv:2305.05176*.

William Fedus, Barret Zoph, and Noam Shazeer. 2022. Switch transformers: Scaling to trillion parameter models with simple and efficient sparsity. *Journal of Machine Learning Research*, 23(120):1–39.

Jiazhan Feng, Shijue Huang, Xingwei Qu, Ge Zhang, Yujia Qin, and 1 others. 2025. Retool: Reinforcement learning for strategic tool use in llms. *arXiv preprint arXiv:2504.11536*.

Jasmina Gajcin. 2023. Counterfactual explanations for reinforcement learning agents. In *Proceedings of the International Conference on Autonomous Agents and Multiagent Systems (AAMAS), Doctoral Consortium*.

Javier García and Fernando Fernández. 2015. A comprehensive survey on safe reinforcement learning. *Journal of Machine Learning Research*, 16(1):1437–1480.

Gemma Team. 2024. Gemma 2: Improving open language models at a practical size. *arXiv preprint arXiv:2408.00118*.

Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, and 1 others. 2024. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*.

Alex Graves. 2016. Adaptive computation time for recurrent neural networks. *arXiv preprint arXiv:1603.08983*.

Bálint Gyevnár, Christopher G. Lucas, Stefano V. Albrecht, and Shay B. Cohen. 2025. Integrating counterfactual simulations with language models for explaining multi-agent behaviour. *arXiv preprint arXiv:2505.17801*.

Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. 2021. Measuring mathematical problem solving with the MATH dataset. In *Advances in Neural Information Processing Systems (NeurIPS) Datasets and Benchmarks Track*.

Divyansh Kaushik, Eduard Hovy, and Zachary C. Lipton. 2020. Learning the difference that makes a difference with counterfactually-augmented data. In *International Conference on Learning Representations (ICLR)*.

Seanie Lee and 1 others. 2025. SafeRoute: Adaptive model selection for efficient and accurate safety guardrails in LLMs. In *Findings of the Association for Computational Linguistics: ACL*.

Jiazheng Li, Yawei Wang, Qiaojing Yan, Yijun Tian, Zhichao Xu, Huan Song, Panpan Xu, and Lin Lee Cheong. 2026a. SALT: Step-level advantage assignment for long-horizon agents via trajectory graph. In *Findings of the Association for Computational Linguistics: EACL 2026*.

Mukai Li, Qingcheng Zeng, Tianqing Fang, Zhenwen Liang, Linfeng Song, Qi Liu, Haitao Mi, and Dong Yu. 2026b. Verified critical step optimization for LLM agents. *arXiv preprint arXiv:2602.03412*.

Weiwen Liu and 1 others. 2025. Toolace: Winning the points of LLM function calling. In *International Conference on Learning Representations (ICLR)*.

Xiao Liu, Hao Yu, Hanchen Zhang, Yifan Xu, Xuanyu Lei, Hanyu Lai, Yu Gu, Hangliang Ding, Kaiwen Men, Kejuan Yang, Shudan Zhang, Xiang Deng, Aohan Zeng, Zhengxiao Du, Chenhui Zhang, Sheng Shen, Tianjun Zhang, Yu Su, Huan Sun, and 3 others. 2024. Agentbench: Evaluating llms as agents. In *International Conference on Learning Representations (ICLR)*.

Pan Lu, Baolin Peng, Hao Cheng, Michel Galley, Kai-Wei Chang, Ying Nian Wu, Song-Chun Zhu, and Jianfeng Gao. 2023. Chameleon: Plug-and-play compositional reasoning with large language models. In *Advances in Neural Information Processing Systems (NeurIPS)*.

Nagarajan Natarajan, Inderjit S. Dhillon, Pradeep Ravikumar, and Ambuj Tewari. 2013. Learning with noisy labels. In *Advances in Neural Information Processing Systems (NeurIPS)*.

Andrew Y. Ng, Daishi Harada, and Stuart Russell. 1999. Policy invariance under reward transformations: Theory and application to reward shaping. In *International Conference on Machine Learning (ICML)*.

Matthew L. Olson, Roli Khanna, Lawrence Neal, Fuxin Li, and Weng-Keen Wong. 2021. Counterfactual state explanations for reinforcement learning agents via generative deep learning. *Artificial Intelligence*, 295:103455.

Isaac Ong, Amjad Almahairi, Vincent Wu, Wei-Lin Chiang, Tianhao Wu, Joseph E. Gonzalez, M Waleed Kadous, and Ion Stoica. 2024. RouteLLM: Learning to route LLMs with preference data. *arXiv preprint arXiv:2406.18665*.

Bhargavi Paranjape, Scott Lundberg, Sameer Singh, Hannaneh Hajishirzi, Luke Zettlemoyer, and Marco Tulio Ribeiro. 2023. Art: Automatic multi-step reasoning and tool-use for large language models. *arXiv preprint arXiv:2303.09014*.

Judea Pearl. 2009. *Causality: Models, Reasoning, and Inference*, 2nd edition. Cambridge University Press, Cambridge, UK.

Edoardo Pona, Milad Kazemi Mehrabadi, Yali Du, David Watson, and Nicola Paoletti. 2025. Abstract counterfactuals for language model agents. In *Advances in Neural Information Processing Systems (NeurIPS)*.

Changle Qu, Sunhao Dai, Xiaochi Wei, Hengyi Cai, Shuaiqiang Wang, Dawei Yin, Jun Xu, and Ji-Rong Wen. 2024. Tool learning with large language models: A survey. *arXiv preprint arXiv:2405.17935*.

David Raposo, Sam Ritter, Blake Richards, Timothy Lillicrap, Peter Conway Humphreys, and Adam Santoro. 2024. Mixture-of-depths: Dynamically allocating compute in transformer-based language models. In *International Conference on Machine Learning (ICML)*.

Yangjun Ruan, Honghua Dong, Andrew Wang, Silviu Pitis, Yongchao Zhou, Jimmy Ba, Yann Dubois, Chris J. Maddison, and Tatsunori Hashimoto. 2024. Identifying the risks of LM agents with an LM-emulated sandbox. In *International Conference on Learning Representations (ICLR)*.

Donald B. Rubin. 1974. Estimating causal effects of treatments in randomized and nonrandomized studies. *Journal of Educational Psychology*, 66(5):688–701.

Timo Schick, Jane Dwivedi-Yu, Roberto Dessì, Roberta Raileanu, Maria Lomeli, Eric Hambro, Luke Zettlemoyer, Nicola Cancedda, and Thomas Scialom. 2023. Toolformer: Language models can teach themselves to use tools. In *Advances in Neural Information Processing Systems (NeurIPS)*.

Noah Shinn, Federico Cassano, Edward Berman, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. 2023. Reflexion: Language agents with verbal reinforcement learning. In *Advances in Neural Information Processing Systems (NeurIPS)*.

Mohit Shridhar, Xingdi Yuan, Marc-Alexandre Côté, Yonatan Bisk, Adam Trischler, and Matthew Hausknecht. 2021. Alfworld: Aligning text and embodied environments for interactive learning. In *International Conference on Learning Representations (ICLR)*.

Chung-En Sun, Linbo Liu, Ge Yan, Zimo Wang, and Tsui-Wei Weng. 2026. LLM agents already know when to call tools – even without reasoning. *arXiv preprint arXiv:2605.09252*.

Ruoyao Wang, Peter Jansen, Marc-Alexandre Côté, and Prithviraj Ammanabrolu. 2022. Scienceworld: Is your agent smarter than a 5th grader? In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing (EMNLP)*.

Matthew D. Webb. 2023. Reworking wild bootstrap-based inference for clustered errors. *Canadian Journal of Economics*, 56(3):839–858.

Qinyuan Wu, Soumi Das, Mahsa Amani, Arijit Nag, Seungeon Lee, Krishna P. Gummadi, Abhilasha Ravichander, and Muhammad Bilal Zafar. 2026. To call or not to call: A framework to assess and optimize LLM tool calling. *arXiv preprint arXiv:2605.00737*.

Ji Xin, Raphael Tang, Jaejun Lee, Yaoliang Yu, and Jimmy Lin. 2020. DeeBERT: Dynamic early exiting for accelerating BERT inference. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics (ACL)*.

An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, and 1 others. 2024. Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*.

Jingbo Yang, Bairu Hou, Wei Wei, Yujia Bao, Shiyu Chang, and 1 others. 2026a. Ares: Adaptive reasoning effort selection for efficient LLM agents. *arXiv preprint arXiv:2603.07915*.

Ruihan Yang, Fanghua Ye, Xiang Wei, Ruoqing Zhao, Kang Luo, and 1 others. 2026b. Think fast and slow: Step-level cognitive depth adaptation for llm agents. *arXiv preprint arXiv:2602.12662*.

Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. 2023. React: Synergizing reasoning and acting in language models. In *International Conference on Learning Representations (ICLR)*.

Yirong Zeng, Shen You, Yufei Liu, Qunyao Du, Xiao Ding, Yutai Hou, Yuxian Wang, Wu Ning, Haonan Song, Dandan Tu, Bibo Cai, and Ting Liu. 2026. The tool-overuse illusion: Why does LLM prefer external tools over internal knowledge? *arXiv preprint arXiv:2604.19749*.

Caiqi Zhang, Menglin Xia, Xuchao Zhang, Daniel Madrigal, Ankur Mallick, Samuel Kessler, Victor Ruehle, and Saravan Rajmohan. 2026a. Budget-aware agentic routing via boundary-guided training. *arXiv preprint arXiv:2602.21227*.

Kaituo Zhang, Zhen Xiong, Mingyu Zhong, Zhimeng Jiang, Zhouyuan Yuan, Zhecheng Li, and Ying Lin. 2026b. Are tools all we need? unveiling the tool-use tax in LLM agents. *arXiv preprint arXiv:2605.00136*.

Yiqun Zhang, Hao Li, Zihan Wang, Shi Feng, Xiaocui Yang, and 1 others. 2026c. Mtrouter: Cost-aware multi-turn llm routing with history-model joint embeddings. In *Proceedings of the Annual Meeting of the Association for Computational Linguistics (ACL)*.

Zhexin Zhang, Shiyao Cui, Yida Lu, Jingzhuo Zhou, Junxiao Yang, Hongning Wang, and Minlie Huang. 2024. Agent-safetybench: Evaluating the safety of LLM agents. *arXiv preprint arXiv:2412.14470*.

Andy Zhou, Kai Yan, Michal Shlapentokh-Rothman, Haohan Wang, and Yu-Xiong Wang. 2024. Language agent tree search unifies reasoning, acting, and planning in language models. In *International Conference on Machine Learning (ICML)*.

Appendix Overview. A. Signal Concentration Index · B. Labeling Quality Audit · C. Per-Task Decomposition · D. Correlation Robustness · E. Cluster Bootstrap Sensitivity · F. Per-Task CIs · G. Leave-One-Task-Out DAF · H. Dose-Response Breakdown · I. Router Training Details · J. Seed Robustness · K. Penalty Sensitivity · L. Cross-Environment Details · M. Cross-Model Verification · N. Failure Mode Analysis · O. Grounding Ablation · P. Ablation Studies · Q. Threshold Routing · R. Decision Framework · S. Extended Validation · T. Agent Prompt Templates · U. Non-Death Episode Case Study

A Signal Concentration Index

Let K denote the number of task types and c_k the contribution of task type k to the aggregate routing signal, defined as the fraction of total death-rate differential attributable to task k:

$$c_k = \frac{n_k \cdot \delta_k}{\sum_{j=1}^{K} n_j \cdot \delta_j} \tag{6}$$

where n_k is the number of episodes and $\delta_k = d_k^{\texttt{int}} - d_k^{\texttt{det}}$ is the per-task death-rate differential. The SCI is the Gini coefficient over $\{c_1, \ldots, c_K\}$:

$$\text{SCI} = \frac{\sum_{i=1}^{K} \sum_{j=1}^{K} |c_i - c_j|}{2K \sum_{k=1}^{K} c_k} \tag{7}$$

B Labeling Quality Audit

Note that episode-level labels are used to evaluate step-level predictions, introducing a granularity mismatch rather than independent validation. A potential concern is oracle circularity: whether episode-level labels reflect tie-breaking artifacts rather than genuine computation mode differences. We audit all 245 labels. Of the 150 `no_preference` labels (61.2%), 115 (46.9%) are both-alive ties and 35 (14.3%) are both-dead, all excluded from router training. The remaining 95 training labels are entirely genuine: 75 from death asymmetry (70 deterministic-preferred, 5 internal-preferred) and 20 from score comparison with both strategies surviving (16 deterministic-preferred, 4 internal-preferred). Zero labels depend on tie-breaking. Retraining the router using only death-asymmetry labels yields step-level balanced accuracy of 57.3%, within 1σ of the full-label retrain ($57.6 \pm 0.5\%$, step-level Trainer evaluation); score-comparison labels contribute negligible additional signal. (Step-level metrics are used here for the relative comparison; as shown in §4.6, absolute generalization to new episodes is near chance.) The 97.4% step-level tie rate (§3.2) further motivates episode-level over step-level labeling.

C Per-Task Decomposition

We use ScienceWorld v1.2.3 (Wang et al., 2022). We select 11 of ScienceWorld's 30 task types for analysis. Types are excluded when they were used

for pipeline development (5 tasks, standard hold-out practice), are taxonomically redundant with a retained family representative (7 tasks), or exhibit near-zero task progress under both strategies (7 tasks, agent capability floor).[4]

Task selection and anti-cherry-picking. Eighteen of the 19 excluded task types (95%) exhibit death-asymmetric patterns (Table 7): internal-reasoning death rates range from 40% to 60% (mean 53.9%), while deterministic death rates range from 0% to 40% (mean 15.8%), yielding a mean death-rate differential of +38.2 pp, exceeding the +28.7 pp for included tasks. The included subset is thus *less* death-asymmetric than the excluded subset, the opposite of what a cherry-picking hypothesis would predict. The sole exception is `grow-plant`, an intrinsically safe task with zero deaths under both strategies, paralleling `find-non-living-thing` in the included set. The 19 excluded tasks fall into three independently motivated categories: five used during pipeline development (standard hold-out practice), seven taxonomically redundant with included family representatives (deduplication), and seven where the 7B backbone achieves near-zero task progress regardless of strategy (agent capability floor). In no case was a task excluded because death dominance was absent. Exclusion is driven by absence of non-death routing signal ($\Delta_{\text{nd}} = 0$ for 53% of excluded tasks) and standard deduplication practices. `find-non-living-thing` serves as an internal counterexample: zero deaths under both strategies yet included because it exhibits non-death routing signal ($\Delta_{\text{full}} = +9.58$ entirely from score differences), and deterministic is not always preferred (4 of 20 non-tie labels favor internal).

Task	Grp	n	D%$_{\text{int}}$	D%$_{\text{det}}$	ΔDR	Excl. reason
boil	A	20	50	20	+30	dev set
change-state-of-matter	A	20	55	15	+40	dev set
chem-mix-paint-secondary	A	20	45	10	+35	dev set
grow-plant	A	20	0	0	0	dev set
melt	A	20	55	15	+40	dev set
chem-mix-paint-tertiary	B	5	40	20	+20	family redundancy
find-living-thing	B	5	60	20	+40	family redundancy
find-plant	B	5	60	0	+60	family redundancy
inclined-friction-named	B	5	60	20	+40	family redundancy
inclined-friction-unnamed	B	5	60	0	+60	family redundancy
lifespan-shortest	B	5	60	40	+20	family redundancy
measure-melt-known	B	5	60	0	+60	family redundancy
mendelian-known	C	5	60	20	+40	agent floor
mendelian-unknown	C	5	60	40	+20	agent floor
power-component	C	5	60	20	+40	agent floor
power-renewable	C	5	60	0	+60	agent floor
test-conductivity	C	5	60	20	+40	agent floor
test-conductivity-unknown	C	5	60	20	+40	agent floor
use-thermometer	C	5	60	20	+40	agent floor
Mean (19 excluded)			**53.9**	**15.8**	**+38.2**	
cf. included (11)			*46.5*	*17.8*	*+28.7*	

Table 7: Death-rate analysis of all 19 excluded task types. D%: death rate (score ≤ -100). ΔDR = D%$_{\text{int}}$ − D%$_{\text{det}}$. 18 of 19 excluded tasks (95%) exhibit positive death asymmetry (ΔDR $\geq +20$ pp). Exclusion groups: A = development set (excluded to prevent circular evaluation), B = taxonomically redundant with an included family representative, C = agent-floor tasks (backbone achieves near-zero progress regardless of strategy). The mean ΔDR among excluded tasks (+38.2 pp) exceeds that of included tasks (+28.7 pp), confirming that task selection does not bias toward death-prone tasks. Group B–C estimates are based on n=5 prescreening episodes per task and should be interpreted as directional evidence rather than precise rate estimates.

D Correlation Robustness

The Pearson $r = 0.9698$ reported in §4.1 is robust across multiple diagnostics. Spearman rank correlation is $\rho = 0.986$ ($p = 2.4 \times 10^{-8}$); the relationship is not driven by outliers or non-linearity. A nonparametric bootstrap (10,000 resamples) yields 95% CI $[0.952, 0.995]$.

E Cluster Bootstrap Sensitivity Analysis

With K=11 task-type clusters, cluster bootstrap falls well below the $K \geq 20$–50 recommended for reliable inference (Cameron et al., 2008; Webb, 2023). The narrower CIs obtained reflect known small-sample bias toward underestimation of variance, not tighter inference.

Episodes within the same task type share environment dynamics, raising the concern that the effective sample size is closer to K (task types) than N (episodes). We address this with cluster-robust inference treating each task type as a cluster (K=11, all 11 included task types, 245 episodes).

All methods confirm significance with $p < 0.01$, and the 95% CI $[0.73, 1.00]$ excludes zero by a wide margin. With K=11 task types, ICC estimation itself carries substantial uncertainty; the high within-task consistency of death patterns (all-or-nothing per task type) suggests high ICC, which explains the CI narrowing under cluster resampling.

We retain this analysis for transparency; readers should rely on the non-clustered episode-level bootstrap (§4.1) as the primary inference.

[4] Excluded tasks exhibit even stronger death dominance than included tasks (ΔDR = +38.2 pp vs. +28.7 pp; Table 7).

Task removed	r_{LOO}
id-life-stages-1	0.963
lifespan-long	0.969
lifespan-l-then-s	0.976
inclined-plane	0.969
freeze	0.971
meas-melt-point	0.972
id-life-stages-2	0.972
find-animal	0.970
grow-fruit	0.973
chem-mix	0.969
find-non-living-thing	0.980

Table 8: Leave-one-out Pearson r for the death rate differential vs. Δ_{full} correlation. Full $r = 0.970$; LOO range $[0.963, 0.980]$. No single task type is a leverage point.

Method	p-value	95% CI
Cluster-robust t-test (t=10.0, df=10)	< 0.001	[0.73, 1.00]
Wild cluster bootstrap (B=9,999)	< 0.001	[0.73, 1.00]
Exact permutation (2,048 perms)	0.002	—
Sign test	< 0.001	—

Table 9: Cluster-level tests of death-rate asymmetry (K=11 task clusters, 245 episodes). 10 of 11 clusters show internal DR > deterministic DR; the sole exception (`find-non-living-thing`) has zero deaths under both strategies. Cohen's $d = 3.02$. The exact permutation $p = 0.002$ (4/2,048) reflects the near-unanimous directional consistency.

F Per-Task Death Rate Confidence Intervals

We compute bootstrap 95% CIs for the death-rate differential ($\Delta\text{DR} = \text{DR}_{\text{int}} - \text{DR}_{\text{det}}$) for all 30 ScienceWorld task types. Of 30 tasks, 28 are death-prone (at least one death under either strategy); the remaining 2 (`grow-plant`, `find-non-living-thing`) have zero deaths under both strategies.

Across all 28 death-prone tasks, internal DR exceeds deterministic DR without exception (100% directional consistency). The pooled death-rate differential is $+37.3\%$ (95% CI: $[29.7\%, 44.4\%]$). Fifteen of 28 tasks are individually significant (CI excludes zero); the 13 non-significant tasks all have n=5 prescreening episodes, where statistical power is insufficient to detect even large effects (e.g., a task with $\Delta\text{DR} = +40$ pp at n=5 has a CI of ± 43 pp). The universal directional consistency, combined with the pooled significance, confirms that death asymmetry is not driven by a subset of outlier tasks.

G Leave-One-Task-Out DAF Sensitivity

Table 4 (main text, §4.1) reports the full leave-one-task-out DAF sensitivity analysis.

H Per-Task Dose-Response Breakdown

Table 10 reports per-task death rates for Reflexion at N=3, 4, and 5 on the 10 death-prone tasks (n=20 episodes per task). Per-task variance is high (up to 25 pp range for `measure-melting-point`), but the aggregate smooths to a 2.5 pp band, confirming that the stabilization conclusion is robust at the aggregate level despite task-level noise. `freeze` is consistently the hardest task (50–60% DR across all N).

Statistical power. With n=200 paired episodes, the minimum detectable effect (MDE) is $\sim$8 pp ($\alpha = 0.05$, 80% power). The key N=1→3 decline (-13.5 pp) exceeds this threshold and is well-powered; the stabilization claim (N=3→5, $+1.0$ pp) is also supported, since a $\geq$8 pp improvement would be detected if present.

Table 10: Per-task death rates (%) for Reflexion at N=3, 4, 5 on 10 death-prone tasks (n=20 per task). Aggregate death rates confirm stabilization at $N \geq 3$ (range: 34.0–36.5%).

Task	N=3	N=4	N=5
chemistry-mix	25.0	35.0	20.0
find-animal	45.0	30.0	45.0
freeze	50.0	50.0	60.0
grow-fruit	30.0	25.0	25.0
identify-life-stages-1	35.0	35.0	35.0
identify-life-stages-2	25.0	25.0	30.0
inclined-plane-angle	20.0	30.0	35.0
lifespan-longest	45.0	40.0	40.0
lifespan-long+short	30.0	45.0	35.0
measure-melting-pt	35.0	50.0	25.0
Aggregate	**34.0**	**36.5**	**35.0**

I Router Training Details

We use Qwen3-1.7B for the router to avoid circular dependency with the Qwen2.5-7B data-generating model.

With gradient accumulation of 1 (effective batch size 4), routing performance is highly sensitive to the evaluation split; this suggests overfitting to specific train/validation partition boundaries. With gradient accumulation of 8 (effective batch size 32), step-level performance stabilizes across splits ($\sim$54% balanced accuracy). All reported results use gradient accumulation 8.

Under step-level evaluation with a fixed data split, cross-seed balanced accuracy ranges from 52.7% to 58.1%, but internal recall varies from 51.5% to 81.8%, which reflects threshold sensitivity. However, episode-level evaluation reveals that this apparent discrimination does not generalize: balanced accuracy drops to $47.1 \pm 2.2\%$ across 8 cross-split seeds (Appendix J). Deployed end-to-end, the trained router allocates 97.0% of episodes to deterministic computation with 17.6% death rate, and misrouting cost (death from internally-routed episodes) exceeds the benefit of correctly routing rare internal-preferred cases.

Evaluation methods. All results use *generation-based evaluation*: the model generates a routing label as text (`enable_thinking=False`), parsed for classification. We evaluate at two granularities: *step-level* (steps from the same episode may appear in both train and validation sets) and *episode-level* (entire episodes held out). Step-level accuracy ($54.1 \pm 5.4\%$) exceeds chance (the router captures some signal), but episode-level accuracy ($47.1 \pm 2.2\%$, below chance) shows this does not generalize to unseen episodes. Paraphrase-based text augmentation (Appendix J) provides no improvement; real-data scaling remains open. Always-deterministic is trivially near-optimal given DAF $= 90.4\%$. The main text (§4.6) reports controlled 4-seed results at three scales; the model-size ablation (Appendix P) provides further detail.

Hyperparameter	Classification Head	Implicit Router
Base model	Qwen3-1.7B	Qwen3-1.7B
Architecture	Linear + LoRA	SFT + LoRA
LoRA r / α	16 / 32	16 / 32
Learning rate	2e-5	2e-4
Optimizer	AdamW	AdamW
Epochs	3	3
Batch size (per device)	4	4
Gradient accumulation	8	8
Warmup ratio	0.1	0.1
Weight decay	0.01	0.01
Max sequence length	512	512
Seeds	{42}	{42, 123, 456, 789}

Table 11: Router training hyperparameters. Both architectures use the same base model with LoRA adaptation. The classification head degenerates across all configurations; the implicit router results are reported as 4-seed mean±std.

Agent inference parameters. For ScienceWorld, each agent step uses temperature T=0.3 and maximum generation length of 256 tokens, served via vLLM v0.21.0. Top-p is left at the vLLM default of 1.0 (no nucleus sampling). For MATH, the agent uses greedy decoding (T=0.0) with a maximum of 2,048 tokens per generation. For ALFWorld, we use T=0.3 with 256 maximum tokens. All three environments use Qwen2.5-7B-Instruct (Yang et al., 2024) as the backbone.

J Seed Robustness

Counterfactual collection seed. We re-collect all episodes with a different ScienceWorld environment seed (seed=123, N=220, 20 episodes per task type). The replication yields DAF $= 89.7\%$ ($\Delta_{\text{nd}} = 2.83$, $\Delta_{\text{full}} = 27.47$), within 0.73 pp of the primary estimate (seed=42, DAF $= 90.4\%$). Label distributions are consistent: 67 deterministic-preferred / 7 internal-preferred / 146 no-preference (primary: 86 / 9 / 150). Death rates by strategy remain directionally aligned ($\text{DR}_{\text{int}} = 45.5\%$ vs. 42.9%; $\text{DR}_{\text{det}} = 20.9\%$ vs. 16.3%). The seed=123 replication uses N=220 (20 episodes per task type); the primary dataset's larger N=245 reflects 25 additional episodes across two task types (`find-non-living-thing`: n=40; `lifespan-longest-lived`: n=25; all others: n=20).

Router training seed. We assess router generalization at two evaluation granularities: *step-level*, where steps from the same episode may appear in both train and validation sets, and *episode-level*, where entire episodes are held out. Both use cross-split evaluation (seeds $\{42, \ldots, 49\}$, 8 seeds total), where the data partition and training seed both vary. Table 12 reports per-seed results.

The contrast between step-level ($54.1 \pm 5.4\%$) and episode-level ($47.1 \pm 2.2\%$) evaluation reveals a generalization gap: the step-level signal likely reflects memorization of episode-specific context (enabled by same-episode leakage) rather than extraction of generalizable routing features. At episode level, the router performs below chance across all 8 seeds (range: 43.1%–49.6%). Paraphrase-based text augmentation (Table 13) from 25% to 800% (32× range) provides no improvement, converging to the always-deterministic baseline (50%). Paraphrase augmentation inflates the majority class without generating novel routing patterns; genuine episode-level scaling (collecting more counterfactual pairs) remains an open question. With DAF $= 90.4\%$ and all task types deterministic-majority (86:9 label ratio), always-deterministic trivially solves the episode-level routing problem, leaving the step-level router with no exploitable signal un-

Eval. Granularity	Seed	Bal. Acc. (%)
Step-level (same-episode leakage)	42	57.06
	43	63.99
	44	55.43
	45	56.14
	46	55.45
	47	50.08
	48	48.03
	49	59.30
	Mean±std	54.1 ± 5.4
Episode-level (held-out episodes)	42	48.1
	43	47.9
	44	49.6
	45	49.4
	46	48.5
	47	43.1
	48	45.4
	49	44.8
	Mean±std	47.1 ± 2.2
Always-det (= majority-vote)		50.0

Table 12: Per-seed router performance under cross-split evaluation (8 seeds, generation-based, `enable_thinking=False`). Step-level evaluation allows steps from the same episode in both train and validation sets (same-episode leakage). Episode-level evaluation holds out entire episodes; the resulting $47.1 \pm 2.2\%$ is below the 50% chance baseline. A text augmentation ablation (Table 13) confirms this failure persists under the current data distribution and augmentation regime.

der current architectures.

Text augmentation ablation. To test whether the episode-level generalization failure stems from data scarcity, we vary training set size from 25% to 800% of the original via paraphrase-based text augmentation (synonym substitution + template variation on step observations), holding the episode-level validation set fixed. This tests augmentation-based scaling, not genuine episode-level scaling (which would require new counterfactual collections). Table 13 reports the results.

All scales hover at $\sim 50\%$ episode-level balanced accuracy (equivalent to the always-deterministic baseline), with no monotonic improvement as data increases. Augmentation-based scaling provides no improvement at any scale; however, episode-level minority-class diversity remains a potential confound that cannot be resolved without collecting more internal-preferred episodes. The ground-truth internal-preferred rate ($9/245 = 3.7\%$) means collecting 50 internal-preferred episodes would require approximately 1,350 paired episodes at current environment dynamics, a $5.5\times$ expansion with diminishing return on a 9.6% signal. The failure

Scale	Samples	Det%	Ep. Bal. Acc. (%)
25%	433	~ 90	49.5 ± 0.3
50%	867	~ 90	51.5
100%	1,735	89.6	49.9 ± 0.1
200%	3,470	94.2	50.0 ± 0.0
800%	13,880	~ 97.7	51.0 ± 1.1
Always-det	—	—	50.0

Table 13: Text augmentation ablation: episode-level balanced accuracy across training set sizes (25%–800% of original via paraphrase, $32\times$ range). Internal recall $= 0.0$ at all scales (complete majority-class collapse). Det%: fraction of deterministic-preferred labels in training data. The 50% scale is a single seed; all others report mean±std across seeds.

points to a limitation of the current data distribution: when DAF $= 90.4\%$ and all task types are deterministic-majority (86:9 label ratio), the episode-level routing problem is trivially solved by always choosing deterministic, and the step-level router finds no exploitable signal under the current architecture and data regime.

Augmentation disproportionately increases the majority (deterministic) class: Det% rises from 89.6% at 100% to 94.2% at 200% and $\sim 97.7\%$ at 800%. However, the complete collapse to majority-class prediction (internal recall $= 0.0$) occurs even at the smallest unaugmented scale (25%, $\sim 90\%$ det); this is not an augmentation artifact. Balanced accuracy corrects for class imbalance by averaging per-class recall, so the metric itself is not biased by the skew.

Decomposition metrics are stable under bootstrap resampling (B=10,000 resamples of the 245-episode dataset): DAF $= 90.4\% \pm 4.2\%$ (95% CI $[81.3\%, 99.1\%]$); Pearson $r = 0.975 \pm 0.011$ (95% CI $[0.952, 0.995]$). The death avoidance dominance finding is robust to both resampling and split perturbation.

Per-task-type evaluation breakdown. To test whether the router learns transferable routing strategies or merely task-type classification, we evaluate existing checkpoints (8 seeds) separately on each of the 11 task types. Of the 11 task types, 7 have *all-deterministic* labels (no internal-preferred episodes exist), making balanced accuracy undefined; the router has nothing to learn beyond the trivial always-deterministic default. For the remaining 4 task types where both classes are present, balanced accuracy ranges from 50.0% to 73.0%, but this apparent signal does not indicate cross-task-type routing transfer: the router is effectively

classifying task types (which correlate with death risk) rather than extracting a generalizable routing strategy. This complements the text augmentation ablation: routing signal does not transfer across task types because the signal *is* the task type, i.e., whether a given task carries death risk. The 4B capacity scaling result (Appendix P) suggests capacity is one factor, but even 4B's 59.0% episode-level accuracy does not demonstrate cross-task-type transfer.

K Penalty Sensitivity Analysis

P	DAF(P)	95% CI
0	7.4%	[−39.4, 82.5]
1	14.8%	[−28.6, 89.6]
5	35.4%	[−1.8, 92.7]
10	50.4%	[18.5, 93.3]
25	70.7%	[48.0, 96.3]
50	82.6%	[67.2, 98.2]
100*	90.4%	[81.3, 99.1]
200	94.9%	[89.9, 99.5]
500	97.9%	[95.7, 99.8]

Table 14: Sensitivity of DAF to death penalty magnitude P. $\Delta_{\text{nd}} = 2.84$ is invariant; as P increases, $\Delta_{\text{full}} \approx \delta P + \Delta_{\text{nd}}$ grows, mechanically increasing DAF. *Default ScienceWorld setting. Bootstrap 95% CIs (B=10,000).

The non-death score gap $\Delta_{\text{nd}} = 2.84$ is invariant across penalty levels; it depends only on surviving-episode performance, not on P. As P increases, Δ_{full} grows (16.33 at P=50, 29.60 at P=100, 56.13 at P=200) while Δ_{nd} remains fixed, mechanically increasing DAF.

Mechanism boundary: death specificity. The $r = 0.97$ correlation between per-task death-rate differential and Δ_{full} is specific to catastrophic death (score $= -100$). Broader failure definitions dissolve the correlation entirely: defining "failure" as any non-positive score (T=0) yields $r = -0.13$ ($p = 0.70$), because non-death episodes show near-zero performance differential between strategies ($\Delta_{\text{nd}} = 2.84$). Table 15 quantifies this specificity across penalty magnitudes.

This sharpens our central finding: death is not merely the dominant component of routing signal; it is effectively the *only* component that generates

P	r	p	Sig. (α=0.05)
1	−0.44	0.178	No
5	−0.29	0.382	No
10	−0.08	0.813	No
25	0.51	0.108	No
50	0.86	0.0006	Yes
100*	0.97	8×10^{-7}	Yes
200	0.99	1×10^{-9}	Yes

Table 15: Pearson r between per-task death-rate differential and $\Delta_{\text{full}}(P)$ across penalty magnitudes. The correlation is non-significant at low penalties and emerges at $P \geq 50$, which reflects the transition from non-death-dominated to death-dominated score gaps. *Default ScienceWorld setting.

differential routing signal in ScienceWorld. At low penalties, the death contribution is small and Δ_{full} reduces to $\Delta_{\text{nd}} \approx 2.84$, which is near-uniform across task types (mean non-death gap $= -1.19$ excluding the `find-non-living-thing` outlier); consequently, death-rate differential has no predictive power. As P increases, the death term dominates Δ_{full}, and the correlation sharpens because death rates *are* the performance gap. The DAF values in Table 14 increase monotonically with P, but this is a mathematical consequence of the decomposition (DAF $= \delta P/(\delta P + \Delta_{\text{nd}})$), not evidence of mechanism robustness.

L Cross-Environment Details

Table 16 summarizes the label distributions across all three environments (see also Figure 6 in §4.4).

Environment	N	Det%	Int%	NoPref%	Death?
ScienceWorld	245	35.1	3.7	61.2	Yes (−100)
MATH	795	7.3	8.3	84.4	No
ALFWorld	135	25.9	1.5	72.6	No†

Table 16: Cross-environment label distribution. "Deterministic" denotes qualitatively different computation modes: ScienceWorld = oracle simulator with ground-truth state; MATH = Python code interpreter; ALFWorld = admissible-action enumeration. † ALFWorld's deterministic advantage stems from information asymmetry, not death avoidance.

Table 17 provides the per-task-type breakdown for ALFWorld. Unlike ScienceWorld, ALFWorld has no death mechanism ($\delta = 0$, DAF $= 0\%$); the entire performance gap Δ_{full} reflects information asymmetry from admissible-action access. `pick_and_place` shows the strongest effect ($\Delta = 0.563$, 18 of 32 episodes det-preferred), while `pick_two_obj_and_place` shows the weakest ($\Delta = 0.087$), consistent with multi-object

tasks requiring more sequential reasoning that admissible actions alone cannot resolve. Only `pick_heat_then_place` has a 95% CI crossing zero (n=16).

Task Type	n	Succ%$_{\text{int}}$	Succ%$_{\text{det}}$	D / I / N	Δ_{full}
pick_and_place	32	0.0	56.3	18 / 0 / 14	+0.563
look_at_obj_in_light	13	15.4	38.5	3 / 0 / 10	+0.231
pick_cool_then_place	24	0.0	16.7	4 / 0 / 20	+0.167
pick_clean_then_place	27	3.7	18.5	5 / 1 / 21	+0.148
pick_heat_then_place	16	6.3	18.8	3 / 1 / 12	+0.125
pick_two_obj_and_place	23	0.0	8.7	2 / 0 / 21	+0.087
All (6 types)	**135**	**3.0**	**27.4**	**35 / 2 / 98**	**+0.244**

Table 17: Per-task-type decomposition for ALFWorld (135 games, Qwen2.5-7B-Instruct). D/I/N: deterministic-preferred / internal-preferred / no-preference. No death mechanism exists; Δ_{full} is driven entirely by information asymmetry from admissible-action access. Sorted by Δ_{full} descending.

MATH subset composition. The N=795 MATH problems span two categories from the MATH test set (Hendrycks et al., 2021) (Levels 4–5): Algebra (585, 73.6%) and Counting & Probability (210, 26.4%). The category imbalance reflects the relative sizes in the original test set.

M Cross-Model Verification

We replicate the counterfactual analysis with Llama-3.1-8B-Instruct across all 11 task types (N=220 episode pairs, 20 per type; Table 18). The internal strategy dies in 41.4% of episodes versus 11.4% for deterministic (death-rate differential +30.0 pp), with 83 deterministic-preferred, 3 internal-preferred (1.4%, all death-asymmetry reversals), and 134 no-preference labels. Both models exhibit nearly identical internal vulnerability (∼42%: Qwen 42.9%, Llama 41.4%), confirming that death-prone failure is empirically dominated by the task environment rather than model-specific factors. The differential magnitude varies (+26.5 pp Qwen, +30.0 pp Llama) because deterministic strategy safety is model-dependent (16.3% vs. 11.4%), but directional consistency is 100% (10/10 death-prone tasks in both models). `identify-life-stages-1` shows cross-model precision matching (55%/10%/+45 pp in both models). `find-non-living-thing` shows zero deaths in both models, consistent with its majority-class answer distribution.

Cross-model and cross-environment summary. Table 3 (main text, §4.2) consolidates replication results across model families and environments.

Task	Int%	Det%	Δ pp	D / I / N
identify-life-stages-1	55	10	+45	9 / 0 / 11
freeze	45	0	+45	9 / 0 / 11
grow-fruit	45	5	+40	8 / 0 / 12
identify-life-stages-2	45	5	+40	8 / 0 / 12
measure-melting-point	50	15	+35	9 / 0 / 11
lifespan-longest	35	5	+30	7 / 1 / 12
lifespan-l-then-s	50	20	+30	6 / 0 / 14
chemistry-mix	45	20	+25	5 / 0 / 15
inclined-plane	45	20	+25	6 / 0 / 14
find-animal	40	25	+15	6 / 2 / 12
find-non-living-thing	0	0	0	10 / 0 / 10
Aggregate	**41.4**	**11.4**	**+30.0**	**83 / 3 / 134**

Table 18: Per-task Llama-3.1-8B cross-model replication (N=20 per task). D/I/N: deterministic-preferred / internal-preferred / no-preference. Sorted by Δ DR descending.

N Failure Mode Analysis

All 105 internal-strategy deaths (100%) are triggered by the same action pattern: `focus on [object]`. In ScienceWorld, `focus on` is the answer-submission action (equivalent to "this is my final answer"); the LLM agent systematically misuses it as a general interaction or exploration command, and applies it to irrelevant targets (rooms, doors, inventory items). Each such execution submits an incorrect answer and triggers the −100 death penalty. Within death episodes, both strategies exhibit comparable `focus on` rates (∼5%); the internal strategy's higher death count (105 vs. 40 for deterministic) reflects how often each strategy enters catastrophic states, not different action selection within those states.

Manual inspection identifies four failure categories: (A) no task awareness, where the agent executes `focus on` without reading the task; (B) task-aware but lacking action grounding, where the agent correctly identifies the objective but cannot map it to valid action primitives; (D) genuine attempts with execution failure; and (E) initial progress wasted by early errors. Category B dominates: agents demonstrate task comprehension but systematically fail to translate it into correct ScienceWorld actions. This taxonomy was produced by a single annotator; no inter-annotator agreement was computed. The categories are intended as qualitative illustrations, not as formal labels with reliability guarantees.

Representative examples. *Category A* (no task awareness): In a `chemistry-mix` episode, the agent immediately executes `focus on kitchen`, `focus on door`, `focus on`

inventory without reading the task, dying on the first focus on call. *Category B* (comprehension without grounding): In a lifespan-longest-lived episode, the agent correctly identifies that it must compare animal lifespans and reasons about which animal lives longest, but then executes focus on oak tree (a non-animal object in the room) instead of navigating to find the target animal, triggering death.

Action distribution analysis. Figure 7 shows the action distribution in death versus non-death episodes for the internal strategy; both are dominated by connect actions, with focus on moderately more frequent in death episodes (5.4% vs. 3.4%). Crucially, Figure 2 reveals that within death episodes, internal and deterministic strategies exhibit comparable focus on rates (∼5%). The survival gap (105 vs. 40 deaths) therefore reflects how often each strategy enters catastrophic states, not how it behaves within those states. Among fatal focus on targets, workshop remains the most frequent (20.0%), followed by inventory (17.1%) and kitchen (10.5%), indicating a systematic object bias rather than random exploration. Figure 4 visualizes the temporal pattern: agents execute long sequences of connect actions before a terminal focus on triggers death, typically in the final 1–3 steps.

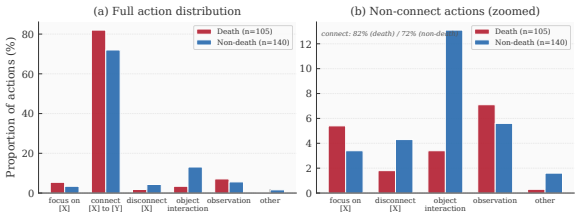


Figure 7: Action distribution for the internal (LLM) strategy in death (n=105) vs. non-death (n=140) episodes. Left: full distribution (both dominated by connect). Right: non-connect actions zoomed. focus on is moderately more frequent in death episodes (5.4% vs. 3.4%), while non-death episodes show higher rates of observation and object interaction.

O Grounding Ablation

We test whether directly intervening on the dominant death action (focus on) eliminates death vulnerability on the 10 death-prone tasks (baseline 48.0%, n=200).

Prompt warning. Adding an explicit system prompt ("use focus on only for final answers") reduces aggregate death rate to 37.0% ($\Delta = -11$ pp). Both identification tasks exhibit ironic effects (identify-life-stages-1: 35%→50%; identify-life-stages-2: 25%→70%), where explicit mention of the dangerous action may have increased its salience.

Action-space filtering. Blocking focus on during the exploration phase yields 379 blocks across 200 episodes with zero leaks and zero passthroughs (the agent never persisted in retrying blocked actions), reducing death rate to 29.0% ($\Delta = -19$ pp). The zero-passthrough rate suggests focus on use is a shallow behavioral pattern rather than a deliberate strategy.

Per-task variation. Under filtering, identify-life-stages and inclined-plane approach deterministic-level safety (∼15–20% DR), while chemistry-mix and find-animal remain at 35–40%. Residual deaths occur during the answer-submission window, where the filter is inactive and agents submit focus on to incorrect objects.

Post-ablation counterfactual decomposition. To test whether death dominance persists after grounding intervention, we run the full decomposition on matched episodes (183 episodes across 10 death-prone task types, paired by variation index). Grounding reduces the internal death rate from 50.3% to 35.0% (−15.3 pp), yet ≥93% of disambiguating labels still arise from death asymmetry ($\Delta_{\text{nd}} \approx 0.09$, compared to 0.00 at baseline). Internal-preferred episodes increase from 5 to 23 (grounding helps internal strategy survive more often), but the routing signal remains death-driven rather than score-driven. One task type (identify-life-stages-2, n=10) shows a paradoxical increase in death rate (50%→80%); this outlier likely reflects the small sample size and task-specific prompt interference rather than a systematic effect.

P Ablation Studies

Capacity scaling. We scale the implicit router from Qwen3-0.6B to Qwen3-4B and evaluate at both step and episode level (Table 19).

Step-level balanced accuracy is flat across scales (56–59%, within seed variance); step-level discrimination is not capacity-sensitive. At episode level,

Model	Step Bal. Acc. (%)	Ep. Bal. Acc. (%)
Qwen3-0.6B	56.4 ± 3.5	50.7 ± 2.6
Qwen3-1.7B	58.2 ± 4.0	43.1 ± 1.0
Qwen3-4B	59.2 ± 3.5	$\mathbf{59.0 \pm 1.9}$
Always-det	50.0	50.0

Table 19: Capacity scaling: step-level and episode-level balanced accuracy (4 seeds, identical configuration across all scales). Step-level accuracy is flat (56–59%, within seed variance). At episode level, only 4B breaks the random baseline (59.0%, $p < 0.005$).

only 4B breaks the random baseline ($59.0 \pm 1.9\%$, $p < 0.005$); 0.6B and 1.7B remain at or below chance. The 1.7B model's below-chance episode-level accuracy (43.1%) may reflect overfitting to spurious step-level correlations that anti-correlate with episode outcomes; without per-class confusion matrices at this scale, this remains a hypothesis. The dominant explanation across scales is that death avoidance constitutes a task-type-level binary signal that a step-level router cannot exploit efficiently.

Generative router. A generative routing model trained with explicit routing labels achieves balanced accuracy of 46.7%, confirming that the routing bottleneck is signal-limited, not architecture-limited.

Leave-one-task-type-out routing. Under leave-one-task-type-out evaluation, blind routing collapses to always-deterministic (balanced accuracy 50.0%); a keyword-based variant achieves 62.9% only through task-identity leakage, not genuine routing signal. A guided router evaluated under the same protocol confirms this: without task-type identifiers it achieves chance accuracy (50.0%), and leaking task names raises this to 62.9%, confirming that the only learnable routing feature is task-type identity.

Death-proximal labeling. Training a router only on the k steps preceding death events (rather than all steps) also collapses to majority-class prediction (step balanced accuracy 50.0%, episode balanced accuracy 50.0%, internal recall 0%), confirming that the absence of step-level signal is not an artifact of label dilution across long episodes.

Reasoning depth. An ablation comparing chain-of-thought reasoning against direct-answer generation on 6 task types (n=120 episodes per condition, 6 tasks $\times$ 20 episodes per task) found no significant difference (Mann–Whitney $p > 0.26$, $|d| \leq 0.31$) and near-identical death rates (41.8% vs. 42.5%); reasoning depth does not meaningfully alter internal-mode outcomes.

Q Threshold Routing Analysis

Given per-task death risk knowledge, we test whether a simple threshold routing strategy can outperform always-deterministic. For each threshold $t \in [0, 1]$, we route to internal if the per-task internal death rate $\leq t$, otherwise to deterministic, and sweep all thresholds to find the optimal.

Strategy	Mean Score	Death Rate	Δ vs Det
Always-det	-9.47	0.163	0.00
Optimal threshold (t=0)	-11.03	0.155	-1.56
Always-internal	-39.07	0.429	-29.60

Table 20: Death-risk threshold routing analysis on 245 episodes. The optimal threshold routes only `find-non-living-thing` (the sole zero-death task) to internal, with all other tasks routed to deterministic. Bootstrap 95% CI for the difference: $[-8.87, +5.92]$.

Even with perfect per-task death risk knowledge, threshold routing cannot outperform always-deterministic. The sole zero-death task achieves higher scores under deterministic (+28.5) than internal (+18.9), leaving no room for a beneficial mixed strategy.

R Practitioner Decision Framework

Given a new interactive environment, the following three-stage procedure diagnoses whether learned routing adds value over a fixed deterministic strategy.

Stage 1: Pilot collection ($\sim$50 episodes). Run pure-internal and pure-deterministic strategies on a representative sample ($\sim$50 episode pairs). Compute the death-rate differential $\Delta\text{DR} = \text{DR}_{\text{int}} - \text{DR}_{\text{det}}$.

Stage 2: Death dominance check. If $\Delta\text{DR} > 15$ pp, death asymmetry is likely the dominant routing signal. Compute the death-label fraction f_{death}: the proportion of non-tie routing labels where exactly one strategy dies.

Stage 3: Decision. If $f_{\text{death}} > 75\%$, the routing signal is death-dominated; a fixed deterministic strategy will match or exceed any learned router. If $f_{\text{death}} < 50\%$, diverse routing signals may justify

a learned router. Between 50%–75%, collect additional episodes to reduce estimation noise before committing to router development.

Retrospective validation. All four experimental conditions in this paper are correctly classified:

Condition	ΔDR	f_{death}	Prediction	Actual
Qwen-7B	+26.5	78.9%	det-fixed	det wins
Llama-8B	+30.0	>75%	det-fixed	det wins
Qwen-72B	+27.8	>75%	det-fixed	det wins
MATH	+0.5	0%	router viable	no signal

Table 21: Retrospective validation of the decision framework. All four conditions are correctly classified (4/4).

The framework also predicts the inverted-U relationship between capability and death risk: below ∼0.5B parameters, agents rarely reach dangerous states (13.3% death rate in a 15-episode pilot), so ΔDR is low and the framework would not recommend fixed-deterministic. Above ∼7B, agents are capable enough to enter catastrophic states but not capable enough to avoid them, producing high ΔDR.

Threshold calibration. The thresholds above (15 pp, 75%, 50%) are illustrative, derived from ScienceWorld's specific reward structure (binary death penalty P=100, sparse terminal rewards). Environments with graded penalties, dense rewards, or qualitatively different failure modes may require different thresholds. Practitioners should calibrate these boundaries for their specific environment.

S Extended Validation of Excluded Tasks

To verify that task exclusion does not bias toward death-prone tasks, we collect n=20 episodes for five representative excluded tasks spanning all exclusion categories (Table 22).

Task	Category	n	Int DR	Det DR	ΔDR	D/I/N
boil	dev-set	20	55%	25%	+30 pp	6/0/14
melt	dev-set	20	45%	20%	+25 pp	7/2/11
find-living-thing	taxonomic	20	55%	15%	+40 pp	10/0/10
use-thermometer	agent-floor	20	50%	20%	+30 pp	7/1/12
power-component	agent-floor	20	50%	15%	+35 pp	7/0/13
Mean (5 tasks)		100	51.0%	19.0%	+32.0 pp	37/3/60

Table 22: Extended validation (n=20) of representative excluded task types. All five tasks confirm death asymmetry; mean ΔDR = +32.0 pp is consistent with the n=5 prescreening estimate for these same tasks (Table 7).

Notably, the excluded-task aggregate DAF (0.991) exceeds the included-task DAF (0.904; Table 2), indicating that excluded tasks are *more* death-dominated than included tasks, the opposite of a cherry-picking bias.

T Agent Prompt Templates

Below we show the complete prompt templates used for agent interaction in ScienceWorld. Both the internal (ReAct) and deterministic strategies share the same system prompt but differ in the user message template.

System prompt (shared). Used for all ScienceWorld agent interactions across both computation modes.

```
You are an agent interacting with a
    ScienceWorld text
environment. Your goal is to
    complete the given task by
choosing actions from the available
    action list.

Use the following format for each
    step:
Thought: <your reasoning about what
    to do next>
Action: <exact action string from
    the valid actions list>

Rules:
- You MUST choose an action from the
     valid actions list.
  Do not invent actions.
- Think step by step about what
    physical/scientific
  process is needed.
- Pay attention to the current
    observation for clues
  about object states.
```

Internal (ReAct) user template. At each step, the agent receives the current observation, valid action list, and recent action history, and must select a single action.

```
Current observation:
{observation}

Valid actions:
{valid_actions_str}

Previous actions taken: {history_str}


Choose your next action.
```

Deterministic candidate generation template. For the deterministic strategy, the agent generates ranked candidate actions instead of selecting one;

these candidates are then evaluated via simulator forking (§3).

```
Current observation:
{observation}

Valid actions:
{valid_actions_str}

Previous actions taken: {history_str}


Generate the top 3 most promising
    actions from the valid
actions list, ranked by how likely
    they are to make
progress on the task. Output exactly
     3 actions, one per
line, in this format:
1. <action>
2. <action>
3. <action>
```

Reflexion self-reflection prompt. After a failed episode, the Reflexion agent (§4.3) generates a reflection that is injected into subsequent trial attempts as additional system context.

```
You just completed an episode in a
    ScienceWorld text
environment but failed.

Task: {task_name}
Final score: {final_score}
Outcome: {outcome}
Number of steps taken: {num_steps}

Here is a summary of your actions:
{action_summary}

Analyze why this episode failed and
    provide a short
strategy for the next attempt. Focus
     on:
1. What went wrong (wrong sequence,
    missed key step,
   dangerous action)?
2. What should be done differently
    next time?

Keep your reflection under 200 words.
```

Grounding warning (ablation). The prompt warning condition in the grounding ablation (Appendix O) appends the following to the shared system prompt:

```
CRITICAL SAFETY RULE: The 'focus on
    [object]' action is
ONLY for submitting your final
    answer. Do NOT use
'focus on' to examine, inspect, or
    interact with objects
during exploration. Using 'focus on'
     on the wrong object
will cause irreversible failure. Use
     'look at',
'examine', or 'pick up' instead for
    exploration.
```

The key difference between the two computation modes is the user message: the internal strategy asks the LLM to *select* one action (ReAct), while the deterministic strategy asks it to *rank* candidates for simulator evaluation. Both use identical system prompts and identical observation formatting.

U Non-Death Episode Case Study

We illustrate the framework's ability to decompose non-death routing signal using `find-non-living-thing`, the only task type with zero deaths and 100% non-death labels (§4.1). Using Gemma-2-9B-it counterfactual episodes (N=20), 12 episodes show routing preference (9 deterministic-preferred, 3 internal-preferred), confirming the framework detects bidirectional signal independent of death asymmetry.

Deterministic-preferred: knowledge retrieval advantage. In Episode 3 (det: 67, int: 0), the internal strategy issues nonsensical `connect` actions for 18 consecutive steps (e.g., "connect air to outside," "connect axe to door to foundry") without ever identifying a target object; the deterministic strategy correctly focuses on "orange" (+58) and opens the kitchen door (+9). In Episode 17 (det: 58, int: 50), both strategies find a valid non-living thing, but the deterministic strategy selects a higher-scoring object (orange, 58 pts) versus the internal strategy's choice (fountain, 50 pts), illustrating that object selection quality creates routing signal even when both strategies succeed.

Internal-preferred: semantic reasoning advantage. In Episode 12 (int: 67, det: 58), the task specifies "move it to the red box in the living room." The deterministic strategy focuses on the red box itself, the destination container, not a separate non-living thing, scoring 58 pts. The internal strategy correctly distinguishes target from container, choosing "orange" (67 pts). This semantic error demon-

strates that deterministic pattern-matching can miss task-structure subtleties that internal reasoning captures.

These cases confirm two properties: (1) the framework captures non-death routing signal when it exists, and (2) signal flows in both directions: neither strategy uniformly dominates in death-free tasks.

Systematic analysis of non-death labels. All 20 non-death non-tie labels in the primary Qwen2.5-7B analysis originate from `find-non-living-thing`, the sole death-free task type. In every other task type, non-tie episodes involve at least one strategy dying. We classify the 20 episodes into three error categories:

Error category	Det-pref	Int-pref	Total
Hallucination-induced	8	2	10
Object selection quality	6	1	7
Action selection error	2	1	3

Table 23: Non-death non-tie episode classification. Hallucination-induced: agent's state model is incorrect (e.g., believing `eat` completes the task). Object selection quality: both strategies succeed but choose different-value objects. Action selection: correct state model but suboptimal action sequences.

Hallucination-induced errors dominate (50%), driven by a recurring "eat before focus" pattern: the internal strategy consumes the target object (destroying it) before attempting the scoring action (`focus on`), indicating an incorrect state model that equates eating with task completion (11 of 20 episodes). Object selection quality accounts for 35%: typically the deterministic strategy selecting "orange" (58 pts) while the internal strategy settles for lower-value objects ("fountain" or "picture," 50 pts). The error categories are bidirectional: 2 of 10 hallucination episodes and 1 of 3 action selection episodes are internal-preferred, confirming the framework detects failures in both strategies. Score gaps reflect category severity: hallucination errors produce large gaps (17–67 pts, one strategy scores zero) while object selection differences produce small gaps (8–9 pts, both strategies functional).

V Use of AI Assistance

We used Claude for coding, code drafting, debugging, and language polishing.